




**POLITECNICO**  
MILANO 1863

# Switches Switching

# Switching

The core part of a router (or in general a switch) moves the packets from the input link to the output link



The switch must decide which input and output links should be matched

- packet-by-packet (nanoseconds!)
- as many couples of links as possible at the same time
- possibly multicast

Mathematically this is a bipartite matching problem

- match as many input links to as many output links as possible
- in 8 ns @40Gbit/s!

Several tricks:

- trade accuracy for time
- hardware parallelism
- randomization
- priorities

# Shared-memory Switches




Packets are read into memory as they arrive and read from memory as output links become available

Problem: memory bandwidth

- with 32 input / output 1 Gbit/s-ports, memory bandwidth must be at least 64 Gbit/s

Feasible with small number of ports (and/or slow links)

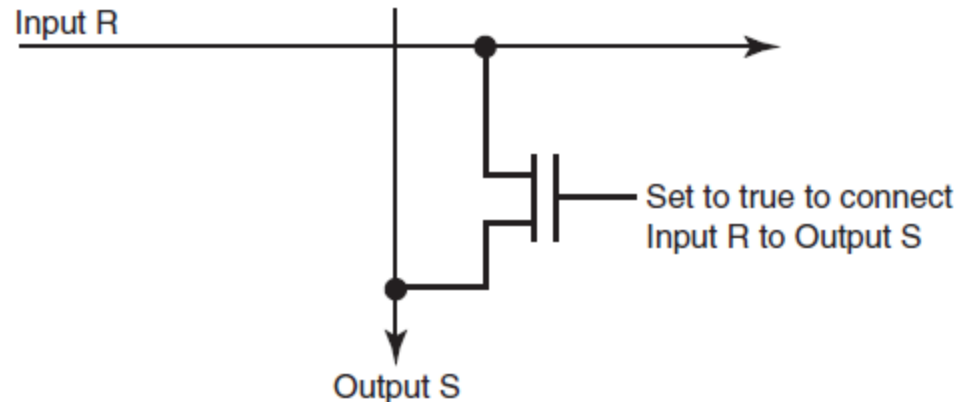

# Bus Architecture



One packet at a time on the bus. Forwarding decision by

- a single CPU
  - slow, a CPU has a lot of other things to do
  - each packet must traverse the bus twice (once to the CPU and once from the CPU)
- multiple CPUs, each CPU serves  $M/N$  line cards
  - still each packet must traverse the bus twice
  - coordination overhead
- one CPU per line card + routing CPUs on the bus
  - the bus is still a bottleneck

# The Crossbar Switch



A crossbar is a set of  $2N$  buses: one for each input, one for each output  
Possibly  $N$  times faster than a single bus if you can find  $N$  disjoint source/destination pairs

For every pair of buses there is a crosspoint ( $N^2$  total)

- conceptually a transistor, in practice more sophisticated

Variable-length packets are generally divided in fixed-size cells and time is divided in fixed-duration time slots

For each time slot the scheduler must decide the status of the  $N^2$  crosspoints.



Problems to solve:

- every output connected to at most one input
- maximise number of input/output relations

Factors limiting parallelism:

- no data at some input port
- contention at the output port

An example: take-a-ticket algorithm

- works with variable sized packets (no need to split in cells)
- similar to a real-life queue management system

# Take-a-Ticket

## Basic idea

- each output line card  $S$  maintains a distributed queue of packets waiting to be sent to  $S$
- the various pieces of the queue are stored in a single queue at the input  $R$  (input buffering)
- if line card  $R$  wants to send a packet to  $S$  asks on a control bus for a ticket number
- when  $S$  is free, it announces on the control bus the next number to be served
- $S$  sets the R/S crosspoint
- $R$  sends the packet




		NUMERO
AMB. 4	ECOGRAFIA	01
AMB. 5	PRELIEVI	01
AMB. 6	ECODOPPLER	999
AMB. 7	M.O.C.	01
AMB. 8	LABORATORIO ANALISI	999
AMB. 9	RADIOLOGIA	01



# TaT: round 1

Round 1




At the input queues each packet is labeled with its destination port


- A wants to send a packet to 1, 2, and 3
- A, B, and C ask for a ticket for exit 1
- Port 1 gives ticket #1 to A, #2 to B, and #3 to C
- A can transmit
- the other ports are blocked (Head-of-Line blocking problem!)

## TaT: round 2


### Round 2



Request



Ticket grant




Connect

- A asks for a ticket for port 2
- port 2 gives ticket #1
- port 1 announces #2
- port 1 announces #1
- A, B transmit
- C is blocked

# TaT: round 3

## Round 3



- A asks for a ticket for port 3, which gives #1
- B asks for a ticket for port 2, which gives #2
- ports 1, 2, 3 announce #3, #2, #1
- A, B, C transmit
- no port is blocked

# Head-of-Line blocking

With single input queues all the packets are blocked if the head-of-line is blocked

Rough estimate of throughput saturation:

- assume each head-of-line wants to go to one of the N output ports chosen uniformly and randomly
- Probability that a given output port is used is:

$$1 - \left(1 - \frac{1}{N}\right)^N$$

- For N large converges to  $(1 - 1/e) \sim 63\%$
- More precise calculations give efficiency  $\sim 58\%$
- Non uniform traffic distributions can give worse results

# How to deal with Head-of-Line Blocking

- Speedup the switch and queue of the output
  - to deal with the worst case the needed speedup is  $N$
  - more likely, the necessary speedup is  $k$ , where  $k$  is the expected number of cells contending for the same output
- Virtual Output Queues (VOQ)

# Parallel Iterative Matching (PIM)



Decompose the input queues in separate input queues: one for each output.

Note that PIM requires fixed size packets and that time is divided in slots.

# Parallel Iterative Matching (PIM)

## 1) Request

- Each input port communicates its status using a bitmap.
- Each bit represents an output queue.
  - 1 = virtual output queue has at least one packet
  - 2 = virtual output queue is empty

## 2) Grant

- Each output port chooses an input port.
- When an output port receives request from multiple input ports, it chooses one at random.
- Note that multiple




## 3) Accept

- If an input port was chosen by multiple output ports, it chooses one randomly.






# Parallel Iterative Matching (PIM)

Round 1



Round 2



Round 3

