


Deficit Round Robin

- **The Round-Robin scheduler**

- Round Robin:

- Packets from **different flows** are stored in **different queues**
- **Queues** are serviced in a **Round-Robin manner**

Graphically:



- **Property of the Round-Robin scheduler**

- Fairness guarantee by Round Robin:

- If flows f and g are **backlogged** over the interval $[t_1, t_2]$ then:

$$|W_f(t_1, t_2) - W_g(t_1, t_2)| \leq 1 \text{ packet}$$

- **Proof:**

- Since flows f and g are **backlogged** over the interval $[t_1, t_2]$, the **Round Robin scheduler** will **transmit one packet from flows f and g in each round**

- The **difference** $|W_f(t_1, t_2) - W_g(t_1, t_2)|$ will be **zero, except** during the time when the **Round Robin scheduler**:

- has **already** transmitted **a packet from one flow** and
- **going to** transmit **a packet from the other flow**

- **Therefore:**

$$|W_f(t_1, t_2) - W_g(t_1, t_2)| \leq 1 \text{ packet}$$

- **Conclusion:**

- The **Round Robin scheduler** can achieve **fairness** if

1. flows have **equal weights** (i.e., **unweighted**)
2. All **packets** have the **same size**

- Generalizing the Round Robin Scheduling

- Sheerhar and Varghese presented in

- "Efficient Fair Queueing using Deficit Round Robin" ([click here](#))

a **very nice generalization** of the **round robin** scheduler that can provide **fairness** when:

- flows have **different weights**
- packets** can have **different sizes**

- The *unweighted* Deficit Round Robin (DRR) Scheduler

- Information maintained by the *unweighted* DRR scheduler:

- Deficit counter** = the **number of bytes** that a **flow** is **allowed to transmit** when it is its turn.
- Quantum** = the **number of bytes** that is **added** to the **deficit counter** of flow in **each round**
(**Quantum** = amount of credit **per round**)

- Operation of the DRR scheduler:

- Packets from **flows** are **transmitted** in a **round robin** manner
- The **quantum** is **added** to the **deficit counter** of a flow *before* servicing a flow.

$$\text{deficit counter}(f) = \text{deficit counter}(f) + \text{quantum}$$

- A **packet** from a **flow** is **only transmitted** if:

$$\text{deficit counter} \geq \text{length of packet}$$

- If a **packet** is **transmitted**, the **deficit counter** of that **flow** is **updated** as follows:

$$\text{deficit counter}(f) = \text{deficit counter}(f) - (\text{\#bytes in packet})$$

 D=0
SE FLOW
VVOTO

Note:

- When the **DDR router** detects the **first packet** of a **new flow**, it **allocates** a **deficit counter** for the **new flow** and **initialize** the **deficit counter** to **zero**

- Example: *unweighted* Deficit Round Robin Scheduler

- Example:

- 3 flows
- Quantum = 500 (bytes)

Graphically:

Deficit counters	
flow 1	0
flow 2	0
flow 3	0
Quantum = 500	

- Packet arrivals:

Deficit counters				
flow 1	0	200 bytes	400 bytes	
flow 2	0	100	200	300 bytes
flow 3	0	200 bytes	200	200

Quantum = 500

- Servicing Flow 1:

- Add quantum to flow 1's deficit counter:

Deficit counter				
flow 1	500	200 bytes	400 bytes	
flow 2	0	100	200	300 bytes
flow 3	0	200 bytes	200	200

Quantum = 500

- Transmit first packet from Flow 1:

Deficit counter				
flow 1	300	400 bytes		
flow 2	0	100	200	300 bytes
flow 3	0	200 bytes	200	200

Quantum = 500

Flow 1 does **not** have enough credits for the next packet and ends its turn

Flow 1 will carry its residual "transmission" credits over to the next service round

- Servicing Flow 2:

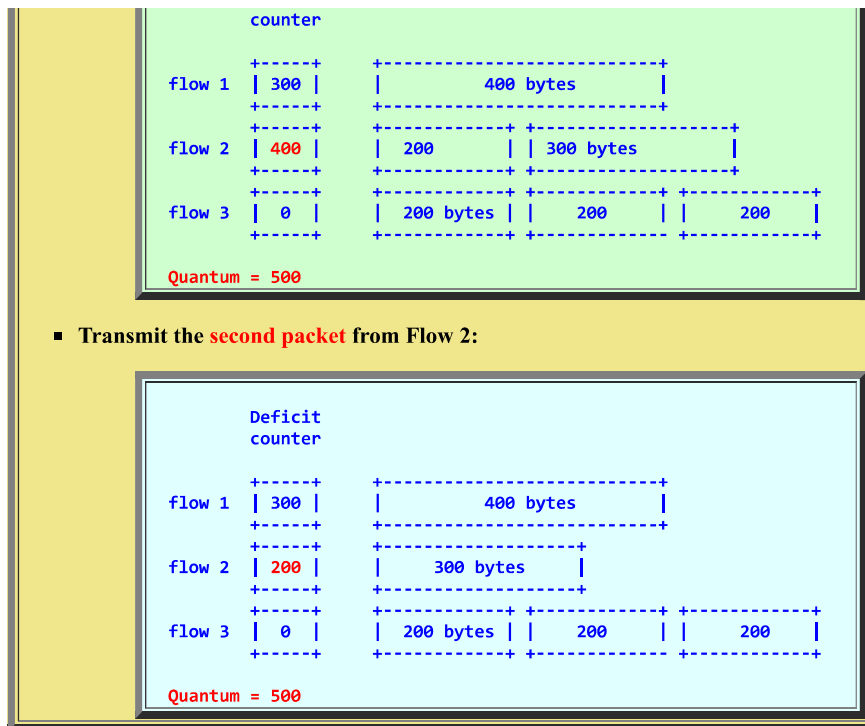
- Add quantum to flow 2's deficit counter:

Deficit counter				
flow 1	300	400 bytes		
flow 2	500	100	200	300 bytes
flow 3	0	200 bytes	200	200

Quantum = 500

- Transmit first packet from Flow 2:

Deficit				
---------	--	--	--	--

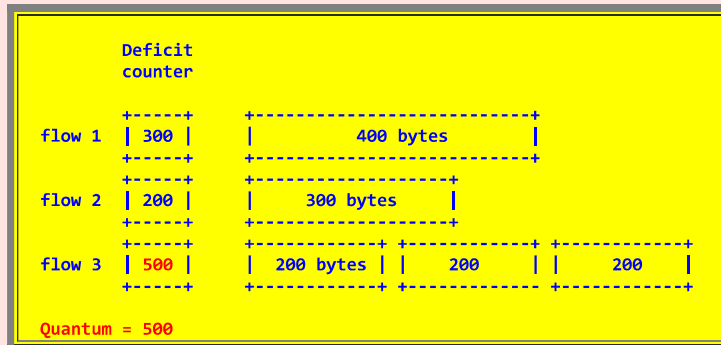


Flow 2 does **not have enough credits** for the **next packet** and **ends its turn**

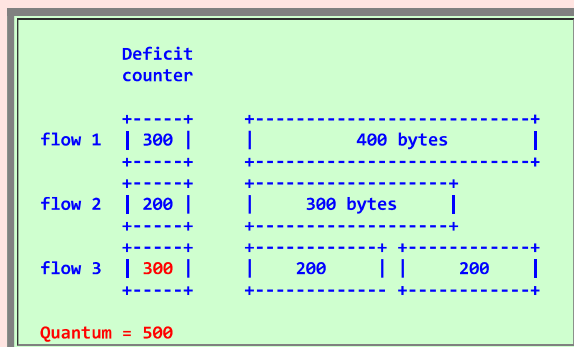
Flow 2 will **carry** its **residual "transmission" credits** over to the **next service round**

◦ Servicing Flow 3:

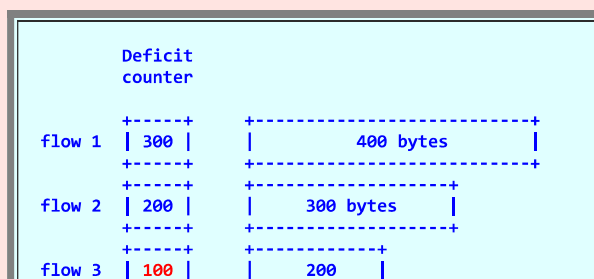
- Add **quantum** to flow 3's **deficiy counter**:




Transmit **first packet** from Flow 3:



Transmit **second packet** from Flow 3:

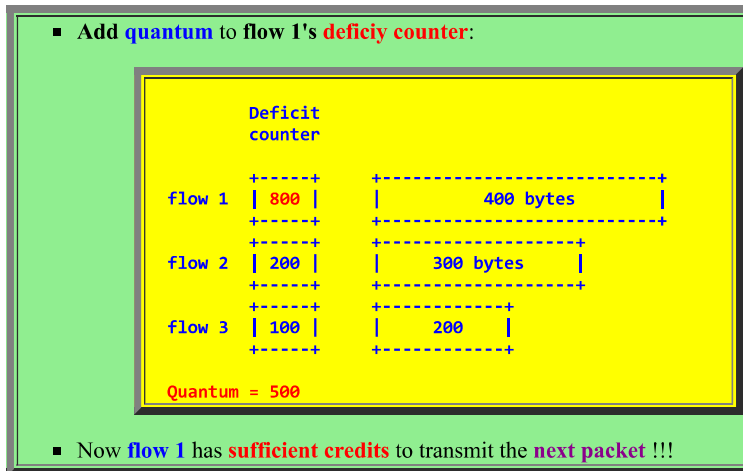




Flow 3 does **not have enough credits** for the **next packet** and **ends its turn**

Flow 3 will **carry** its **residual "transmission" credits** over to the **next service round**

- Servicing Flow 1 again:



- And so on....

- From the operation, it should be clear that:

- Flow will share the bandwidth fairly if they are **unsaturated**
- The **saturated flows** (flows that transmit at a lower rate than the fair share) will receive its complete service rate.
The left over transmission rate will be divided among the **unsaturated flows**

• Property of the Unweighted DRR scheduler

- Fairness guarantee by DRR:

▪ If the **maximum packet size** $\leq Q$ and **flows f and g** are **backlogged** over the interval $[t_1, t_2]$ then:

▪ $|W_f(t_1, t_2) - W_g(t_1, t_2)| \leq 2 \times Q$ (Quantum)

▪ **Proof:**

▪ Suppose there are **k rounds** of **transmissions** in $[t_1, t_2]$

▪ Since **flows f and g** are **continuously backlogged** over the interval $[t_1, t_2]$ and the **maximum packet size** $< Q$, then:

▪ **residual of flows f** $< Q$, and

▪ **residual of flows g** $< Q$

▪ The **worst case** is:

▪ One flow (e.g., **flow f**) has a **residual of $Q - 1$**

▪ The other flow (e.g., **flow g**) is **one round ahead** and has a **residual of 1**

▪ **Graphically:**

	transmission rounds					
	1	2	$k-2$	$k-1$	k	# bytes transmitted
flow g :	(1)		$k \times Q - 1$
flow f :	(Q-1)			$(k-2) \times Q + 1$

▪ Therefore:

$$|W_f(t_1, t_2) - W_g(t_1, t_2)| \leq 2 \times Q \text{ (bytes)}$$

• Weighted Round Robin Scheduling

- Information maintained by the *unweighted* DRR scheduler:

▪ **Deficit counter** = the **number of bytes** that a **flow** is **allowed to transmit** when it is its turn.

▪ **Quantum per flow** = the **number of bytes** that is **added** to the **deficit counter** of *that* flow in **each round**

- Operation of the DRR scheduler:

▪ Packets from **flows** are **transmitted** in a **round robin** manner

▪ The **flow's quantum** is **added** to the **flow's deficit counter** of a flow *before* servicing a flow:

$$\text{deficit counter}(f) = \text{deficit counter}(f) + \text{quantum}(f)$$

▪ A **packet** from a **flow** is **only transmitted** if:

$$\text{deficit counter} \geq \text{length of packet}$$

▪ If a **packet** is **transmitted**, the **deficit counter** of that **flow** is **updated** as follows:

$$\text{deficit counter}(f) = \text{deficit counter}(f) - (\text{\#bytes in packet})$$

- Weighted DRR:

▪ We can achieve **weighted fairness** by assigning a **different quantum** to **different flows**

- Example:

	Flow Quantum	Deficit counter			
flow 1	800	0	200 bytes	400 bytes	
flow 2	400	0	100	200	300 bytes
flow 3	600	0	200 bytes	200	200

- At the start of servicing a flow, the **flow quantum** is **added** to the **deficit counter**

• Property of the *weighted* DRR scheduler

- Fairness guarantee by *Weighted* DRR:

▪ Let Q_f = the **quantum** of flow f

▪ If the **maximum packet size of flow f** $\leq Q_f$ and **flows f and g** are **backlogged** over the interval $[t_1, t_2]$ then:

$$\left| \frac{W_f(t_1, t_2)}{Q_f} - \frac{W_g(t_1, t_2)}{Q_g} \right| \leq 2$$



(Proof is similar to the **unweighted DRR**)

• Using DRR Scheduling in NS

- The **unweighted DRR scheduler** has been implemented in NS
- To create a **DRR** queue, use:

```
$ns duplex-link $n1 $n2 0.3Mb 200ms DRR
```

- You can set the following parameters of the **DRR** queue:

- **quantum_** = size of the **quantum** that is added per round
- **buckets_** = number of buckets

Buckets and flows:

- Each **flow** is assigned a **bucket**
- A **bucket** can contain **multiple flows**

- **blimit_** = bucket limit (= the **queue size** in each bucket)

- **mask_** = how to **identify a flow** (= how to put flows into buckets).

- **mask_ = 0:**

- a **flow** is identified by (**node id, port id**)
- In this case, each flow is put in a **different bucket**

- **mask_ = 1:**

- a **flow** is identified by (**node id**) (only)
- In this case, **all flows** originating from the **same node** are put in the **same bucket**

- **Example:**


```
Queue/DRR set quantum_ 500
$ns duplex-link $n1 $n2 0.3Mb 200ms DRR
```

- **DEMO:** (2 TCP flows sharing a DRR queue)

Example

- NS Prog file: [click here](#)






- **DRR can limit unresponsive (UDP) flows**

- TCP competing with an *over-sealous* UDP flow using **Drop Tail** scheduling:

CWND of the TCP flow:



- **DEMO:** (1 UDP vs 1 TCP flows)

Example


- NS Prog - UDP and TCP flow sharing DropTail Queue: [click here](#)

The **UDP flow** makes **TCP** use **CWND = 1** all the time...

(Look at the **NAM** visualization and you will know why...)

- TCP competing with an *over-sealous* UDP flow using **DRR** scheduling:

CWND of the TCP flow:



- **Example Program:** (Demo above code)

Example

- NS Prog - **UDP and TCP flow** sharing **DRR** Queue: [click here](#)

Note:

- You can see that **DRR** can provide TCP with a **fair share** of bandwidth even when the UDP flow transmits more than the bottle neck link.