


Parallel Architectures



ESCALab@PoliMi

*Embedded Systems and Computer
Architecture Laboratory*

2011 December 1st


Outline

- **Introduction**
- **Multi-core architectures**
 - Memory hierarchy
 - Interconnect design
- **Commercial examples**
- **Parallel programming issues**
 - Forms of parallelism in software (overview)
 - Main issues in parallel programming
- **Wrap-up**

What's next

- **Introduction**
- **Multi-core architectures**
 - Memory hierarchy
 - Interconnect design
- **Commercial examples**
- **Parallel programming issues**
 - Forms of parallelism in software (overview)
 - Main issues in parallel programming
- **Wrap-up**

Moore's Law at Intel: 1965




Moore's Law turning point: 2002

- **Microarchitectural techniques for ILP**

- Pipelining, superscalar processors, VLIW, out-of-order execution, register renaming, speculative execution, branch prediction


Moore's Law ran smoothly until 2002, when the gap between performance and gate count started to appear.



-
- Moore's Gap implies that not enough ILP can be exploited from using more transistors!
 - Why?
 - Power consumption
 - Signal propagation delay $>$ transistor delay
 - Memory bottleneck
 - Memory stalls dominate

Microelectronics industry trend


- Higher performance
- Higher power
- Higher power density



Power consumption

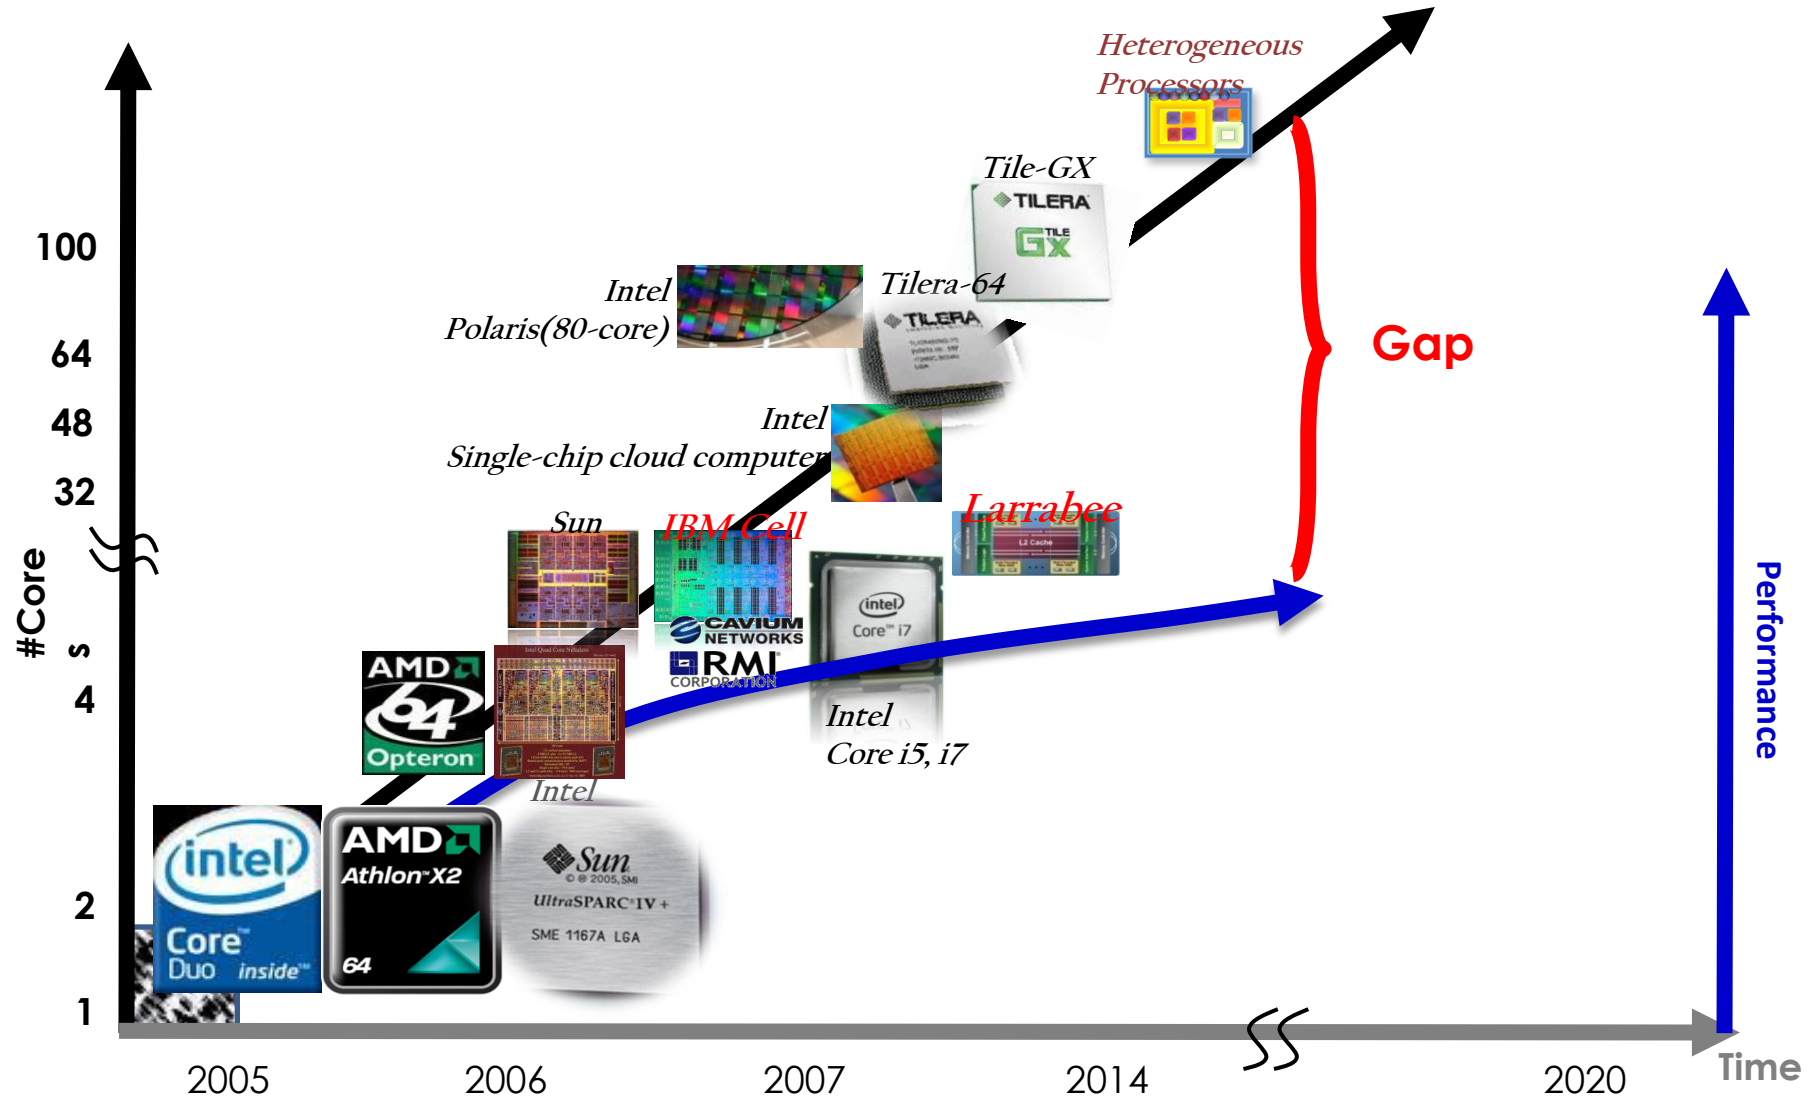
$$P_{dynamic} \propto V^2 f C_{load}$$

- Density [W/cm²]



(Courtesy of Pat Gelsinger, Intel Developer Forum, Spring 2004)

Many-core evolution




Source: Das, C. "Interconnection networks".
7th ACACES Summer School, 2011, Fuggi, July 2011

Many-core evolution

(cont'd)


Increasing number of cores per chip demands for increasing communication bandwidth and capability



Source: International Technology Roadmap for Semiconductors, Design chapter, 2008.

Energy requirements

State-of-the-Art architecture design



250 nm

Gate Delay

Global Wiring

32 nm

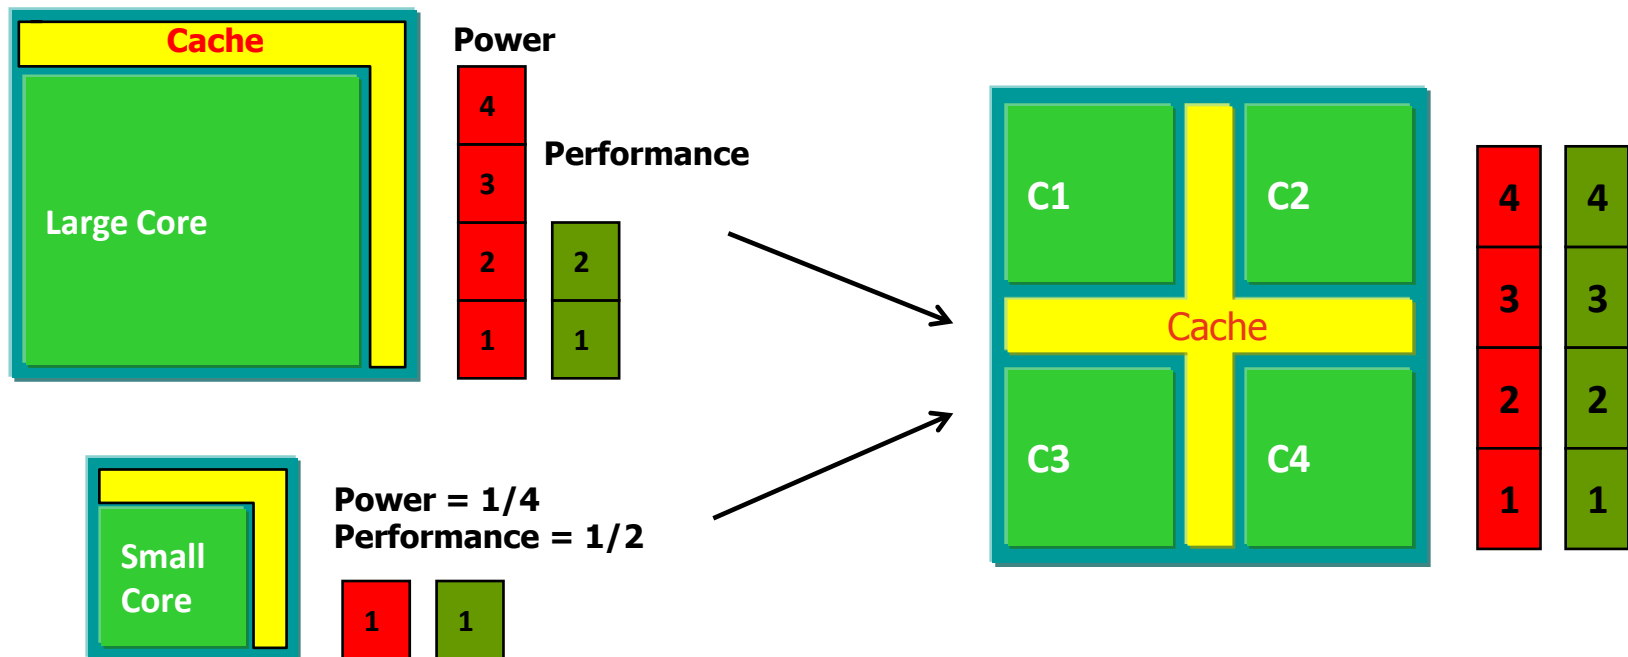
Global Wiring Delay Dominates

Operation	Energy (pJ)
<i>64bIntegerAdd</i>	1
<i>I\$ Fetch</i>	33
<i>Read64bRegister(64x32 bank)</i>	3.5
<i>Read64bRAM(64x2K)</i>	25
<i>Readtags(24x2K)</i>	8
<i>Move64b1mm</i>	6
<i>Move64b20mm</i>	120
<i>Move64boffchip</i>	256
<i>Read64bfromDRAM</i>	2000

Multi-core benefits


Multi-chip cores are not appropriate for power/performance optimization because of the off-chip energy requirements

Multi-core chips provide better power management



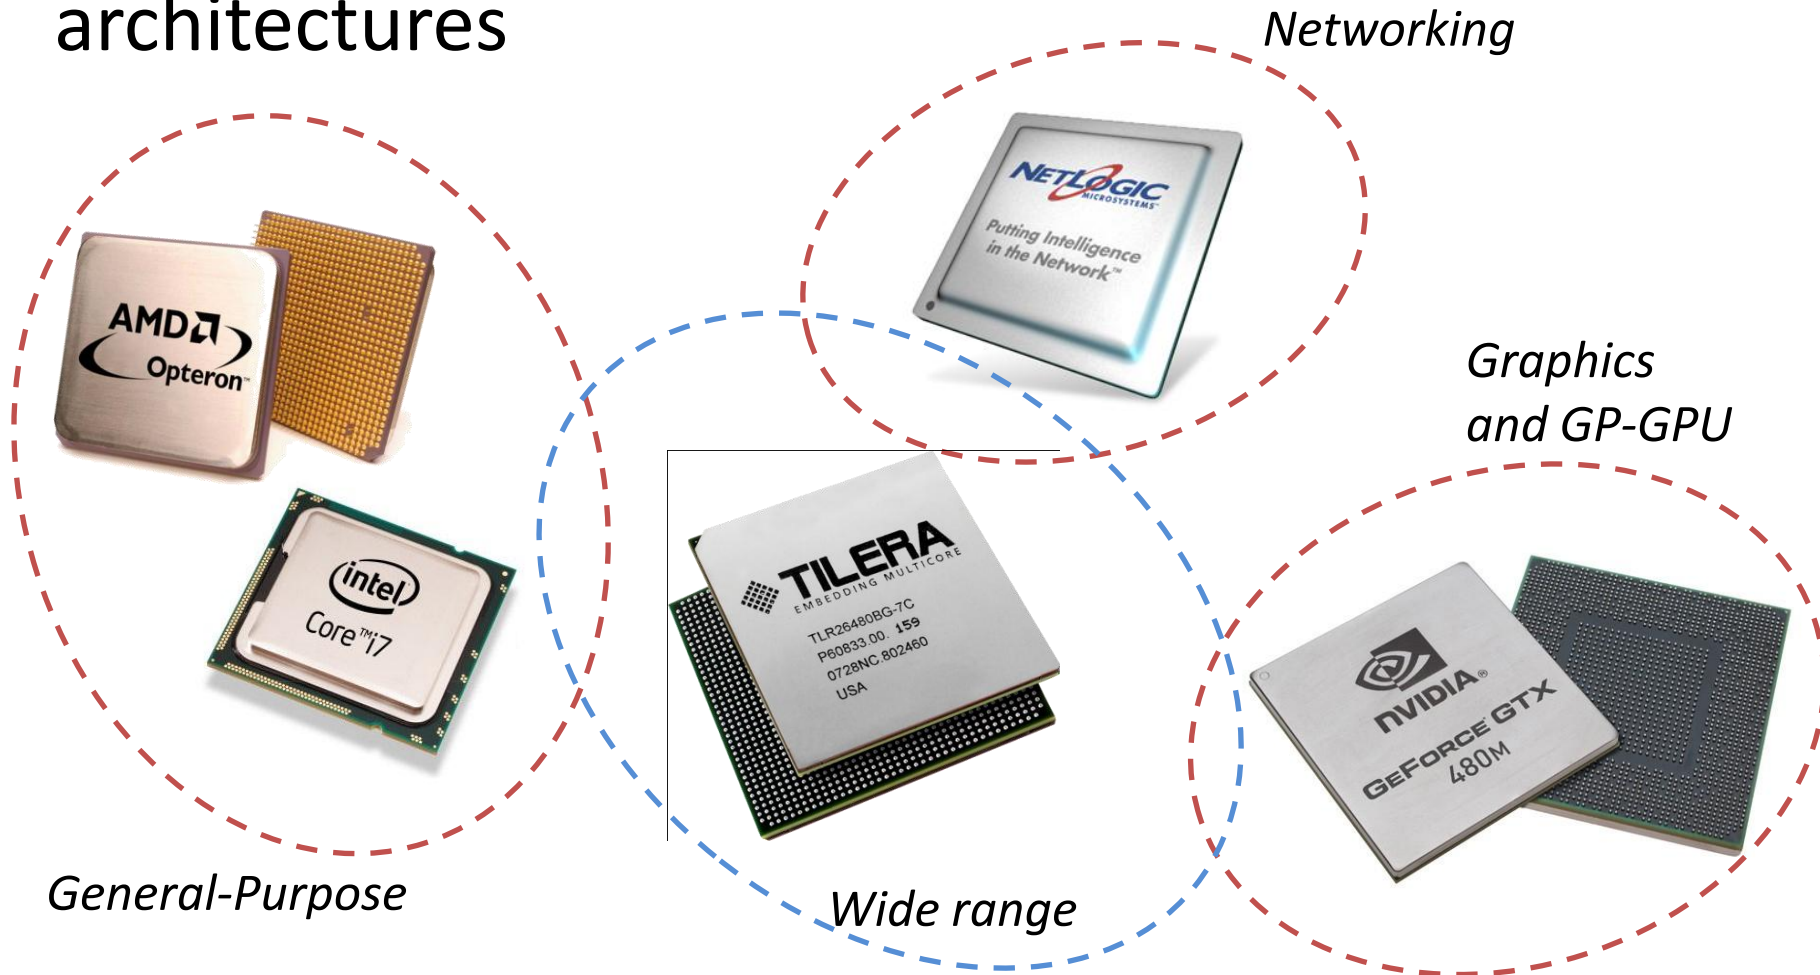
Multi-core processors adoption

- Multi-core processors are adopted world-wide

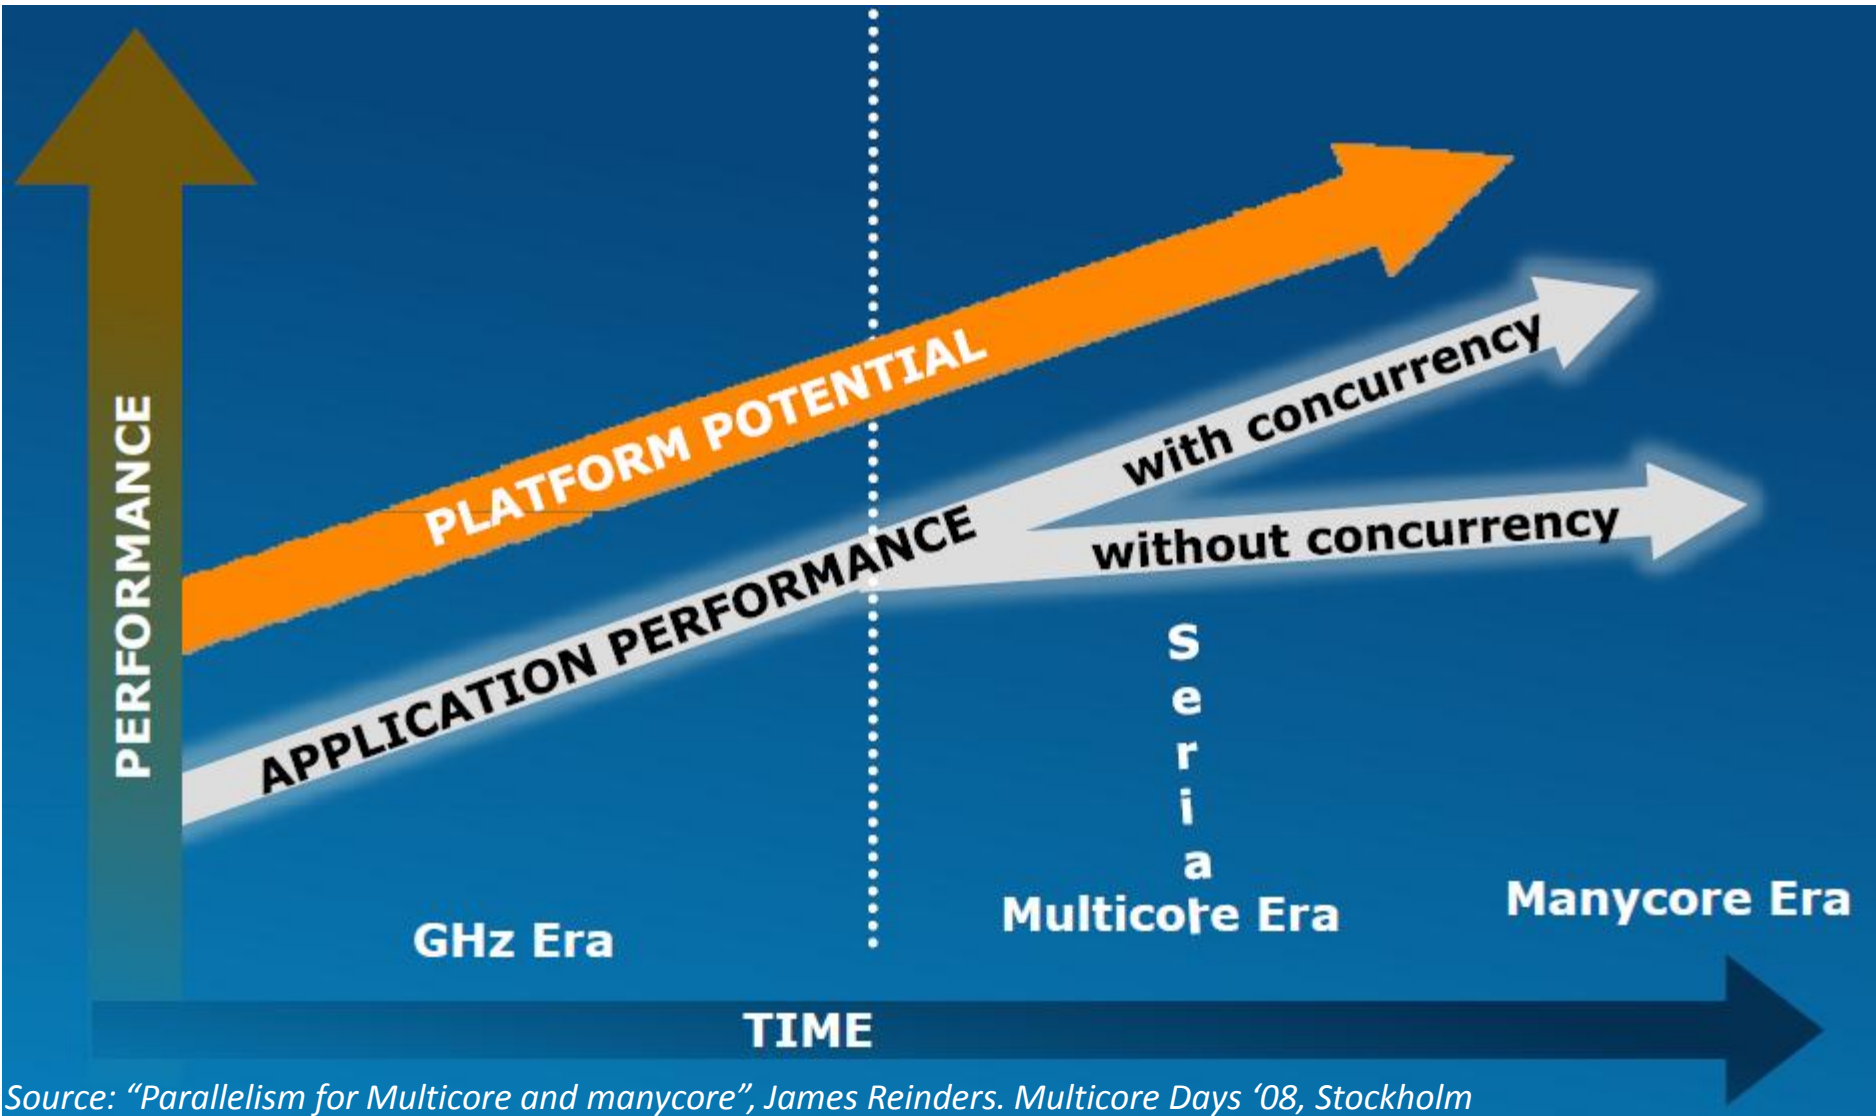


Market segments

- Different markets for multi-core and many-core architectures



The opportunity from concurrency



Parallel programming


- What should we learn from that?
 - Current parallelism support (e.g., ILP) is no longer suitable for ensuring increasing performance requirements
 - Need a programming paradigm shift
 - Parallel programming
- This comes with non-negligible cost
 - Steep (?) learning-curve

What's next

- Introduction
- **Multi-core architectures**
 - Memory hierarchy
 - Interconnect design
- Commercial examples
- Parallel programming issues
 - Forms of parallelism in software (overview)
 - Main issues in parallel programming
- Wrap-up

Old-school multiprocessors


- Processors reside on different chips
 - Off-chip memory and communication resources



Source: "The Art of Multiprocessor Programming",
Herlihy, Maurice; Shavit, Nir. Morgan Kauffman


Multicore architectures

- Processors on a single chip
 - MPSoCs




Implementation

- Multi-core architectures present a variety of shapes
 - Application-specific
 - Performance-constrained
 - Ad-hoc solutions
- Challenges
 - Interconnect
 - Memory subsystem
 - Chip layout





Shared memory

- A unique global address space




← Cache-Coherent
Uniform Memory Access
(*single physical bank*)

(Cache-Coherent)
Non-Uniform
Memory Access
(*multiple banks*)



Distributed memory

- Memory addresses in one processor do not map on another processor
- Cache coherence concept does not apply

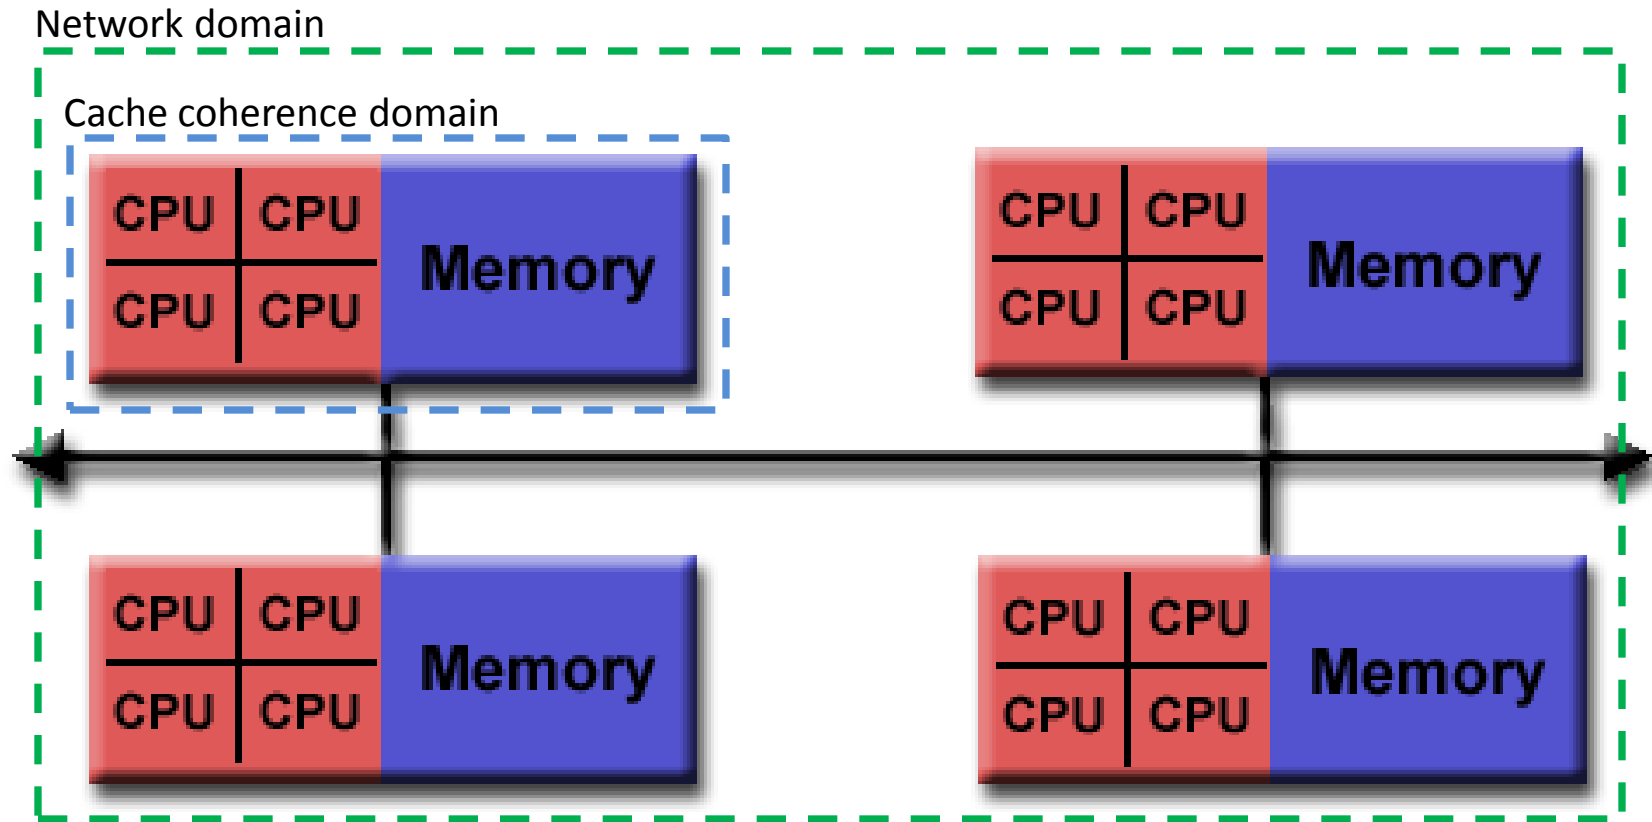


Shared versus distributed memory

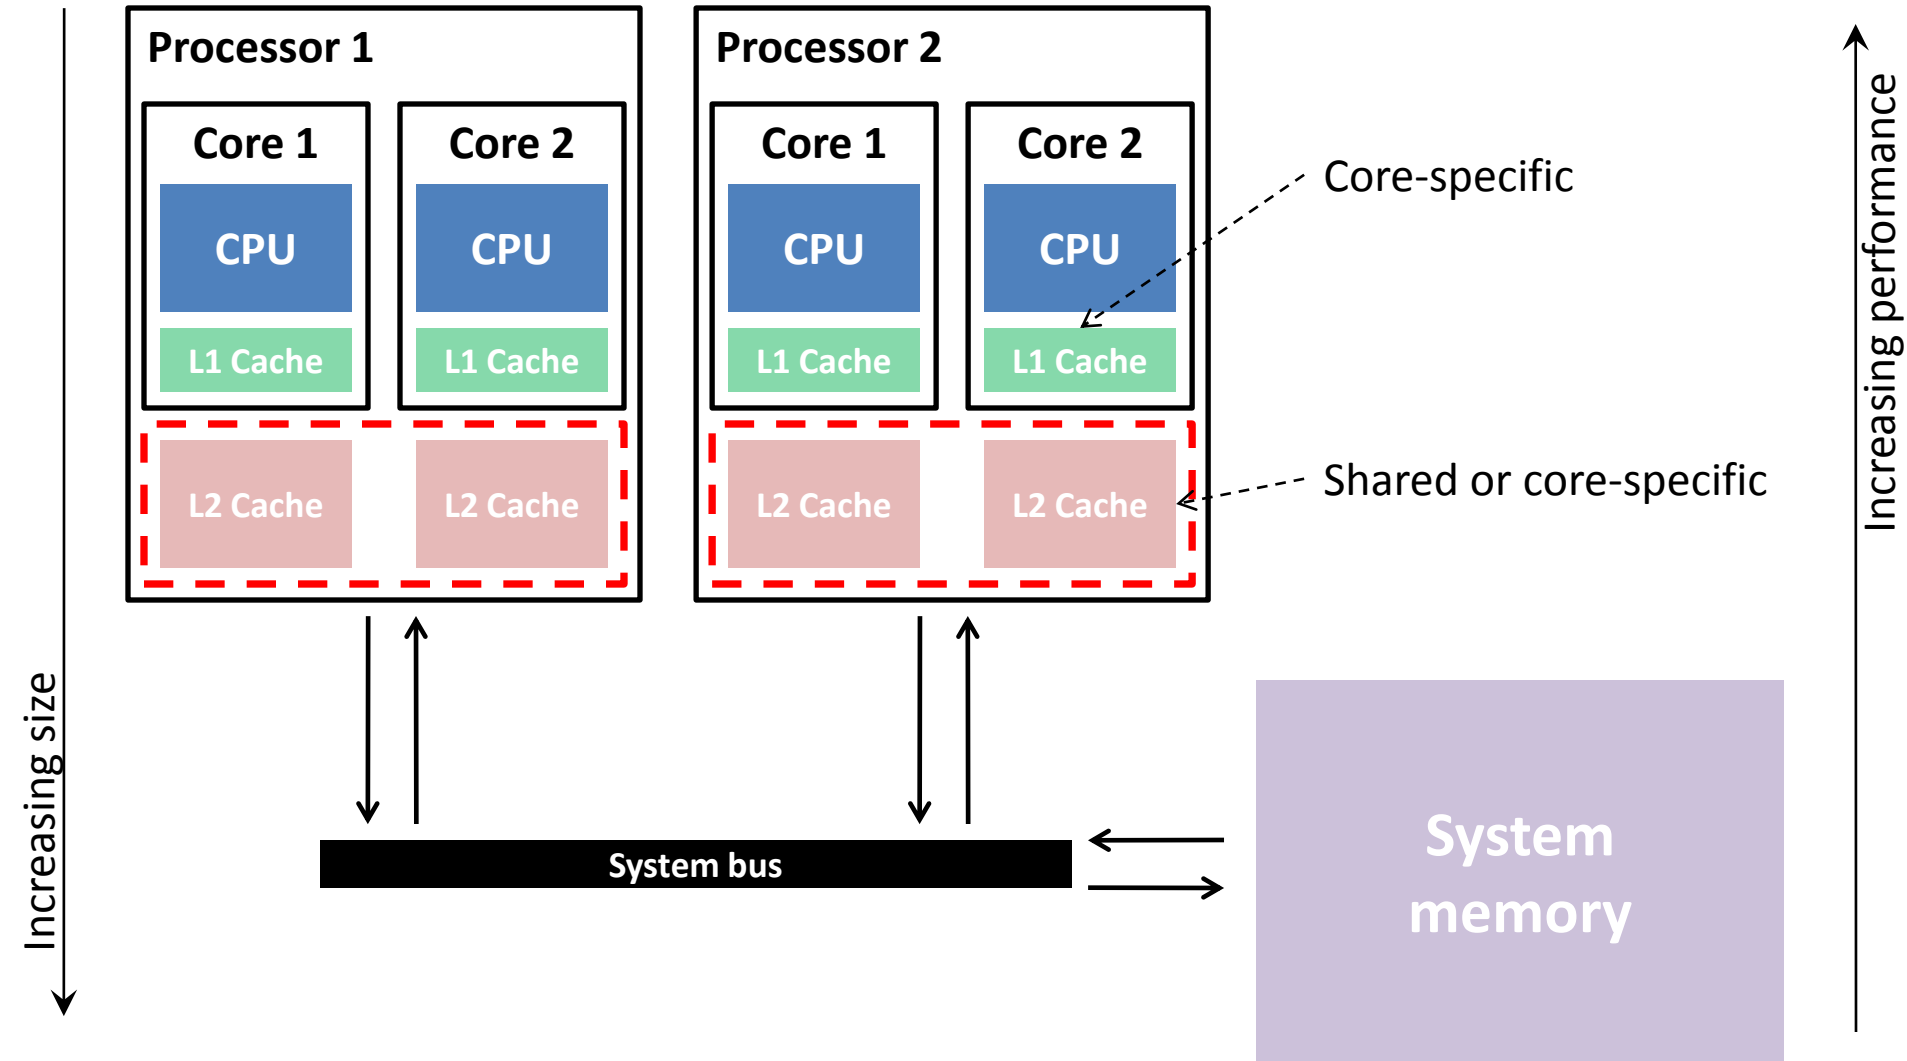
SHARED MEMORY	DISTRIBUTED MEMORY
<ul style="list-style-type: none">+ Global address space simplifies programming efforts to memory+ Fast data sharing mechanism	<ul style="list-style-type: none">+ Scalability+ Each processor can access its local memory without impacting on other processors
<ul style="list-style-type: none">- Scalability (increasing number of processors increase traffic)- It is in charge of the programmer to explicitly use synchronize constructs to access shared resources- Cost with increasing number of processors	<ul style="list-style-type: none">- Many of the communication details between processors is in charge of the programmer- NUMA in nature

An hybrid approach

- Clustered architecture

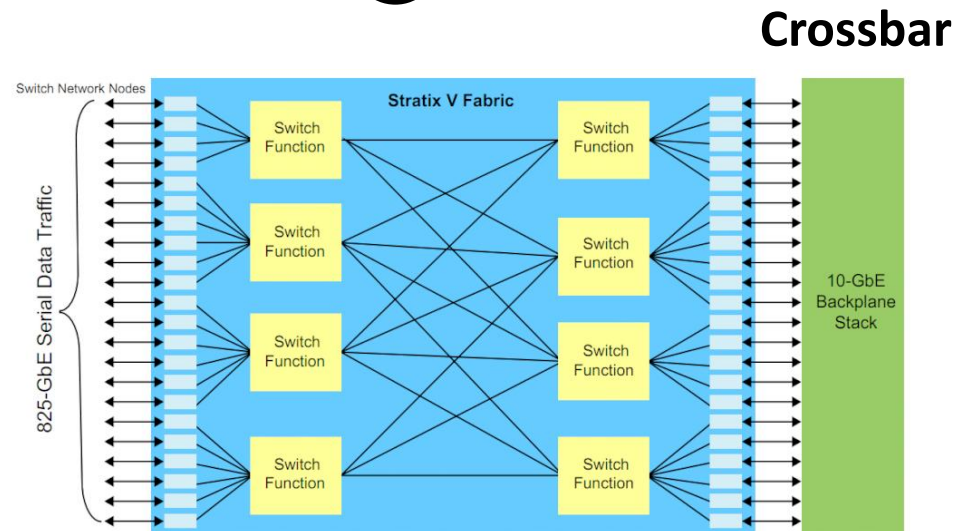

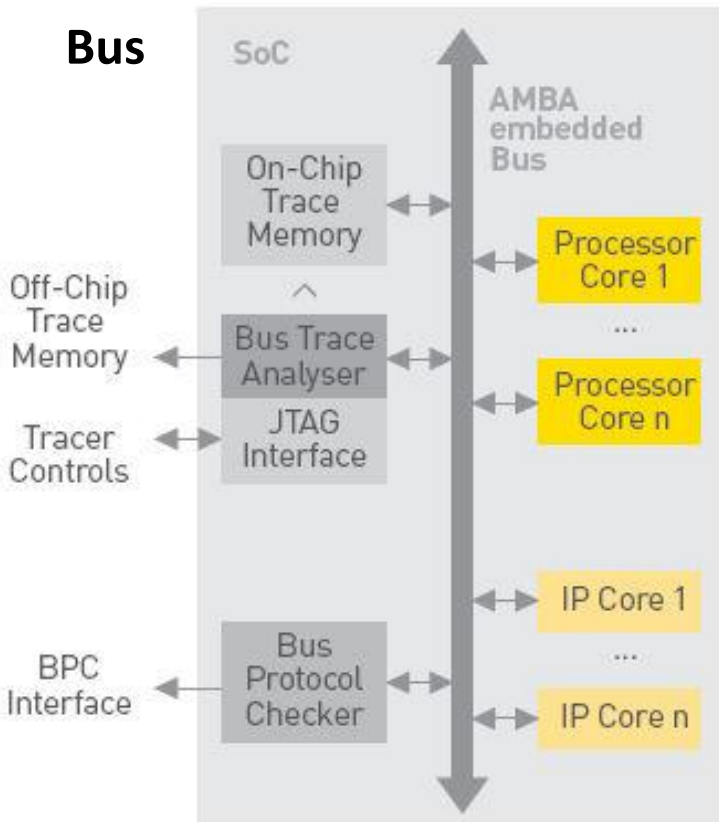


Memory hierarchy



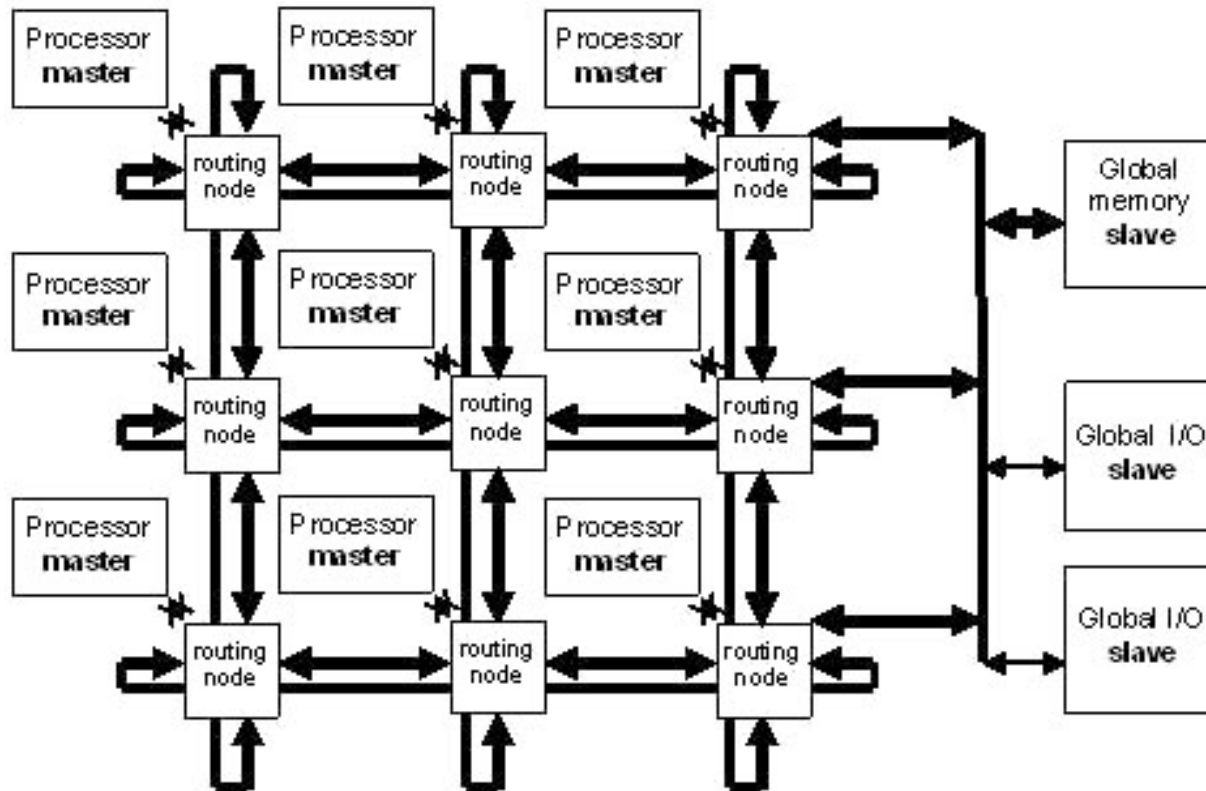
Interconnects

- Classical approaches are no longer suitable for interconnect design in MPSoC architectures



Network-on-Chip

- NoC is a good candidate for scalable communication architecture design




What's next

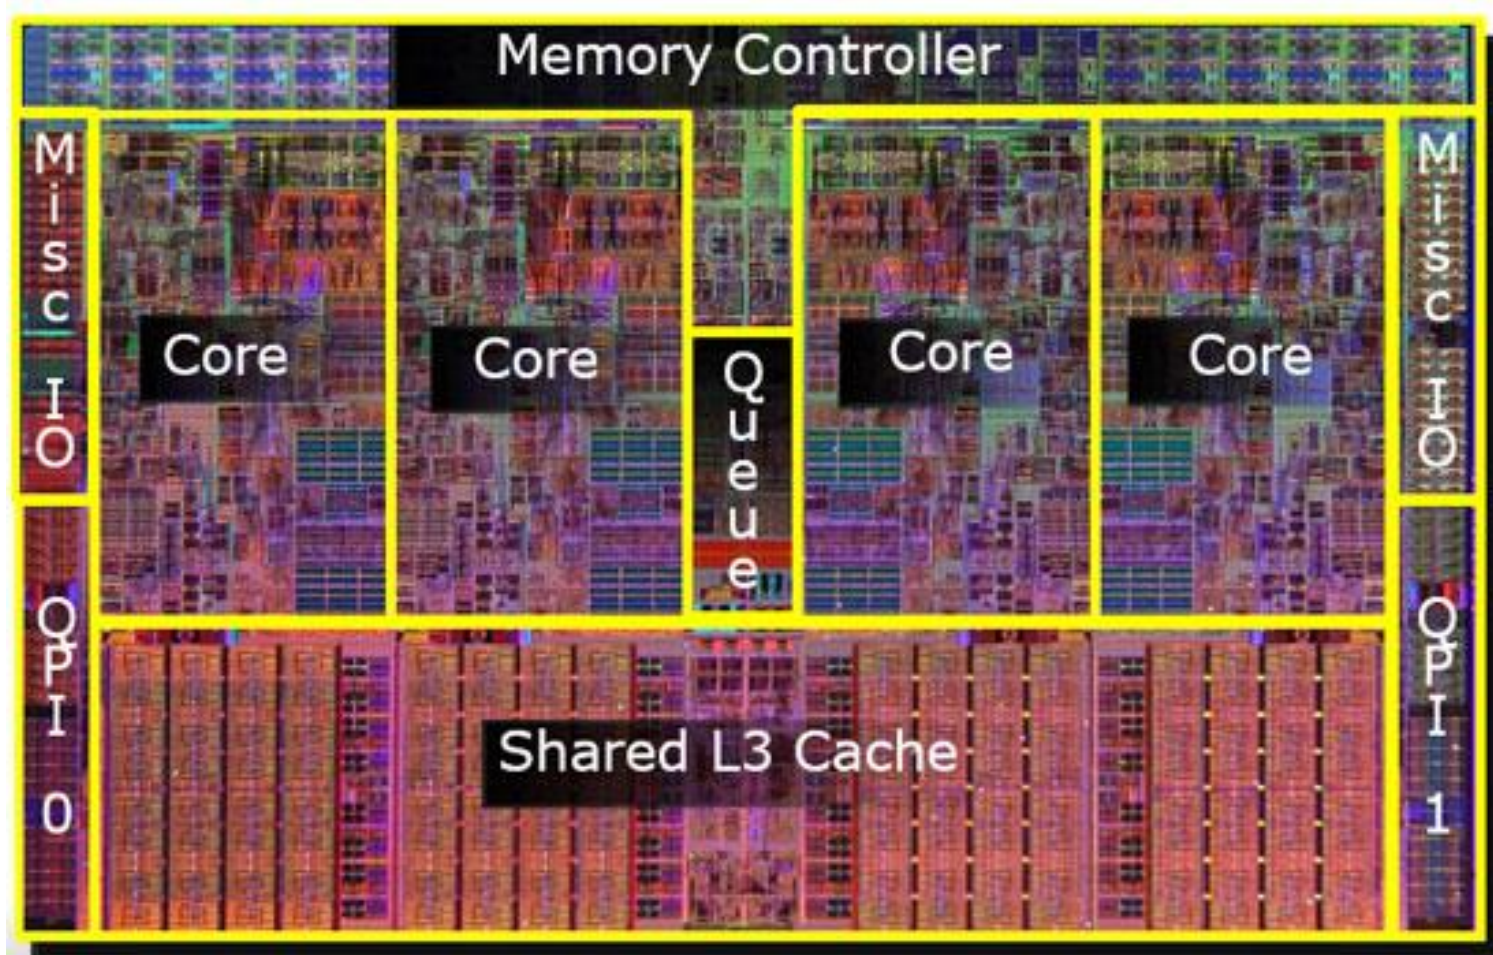
- Introduction
- Multi-core architectures
 - Memory hierarchy
 - Interconnect design
- **Commercial examples**
- Parallel programming issues
 - Forms of parallelism in software (overview)
 - Main issues in parallel programming
- Wrap-up

Commercial processors examples

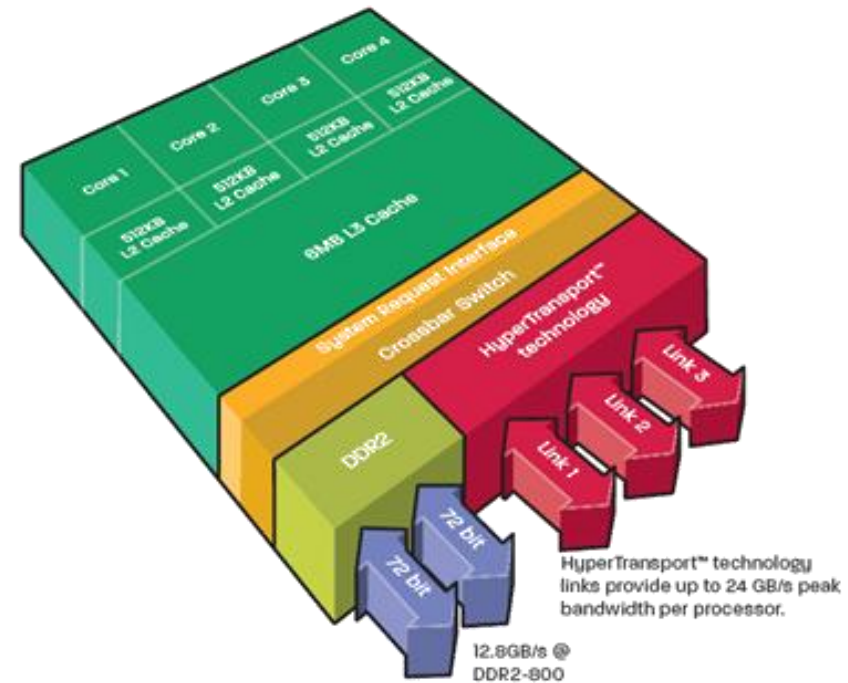

- Intel Core2-Duo Quad



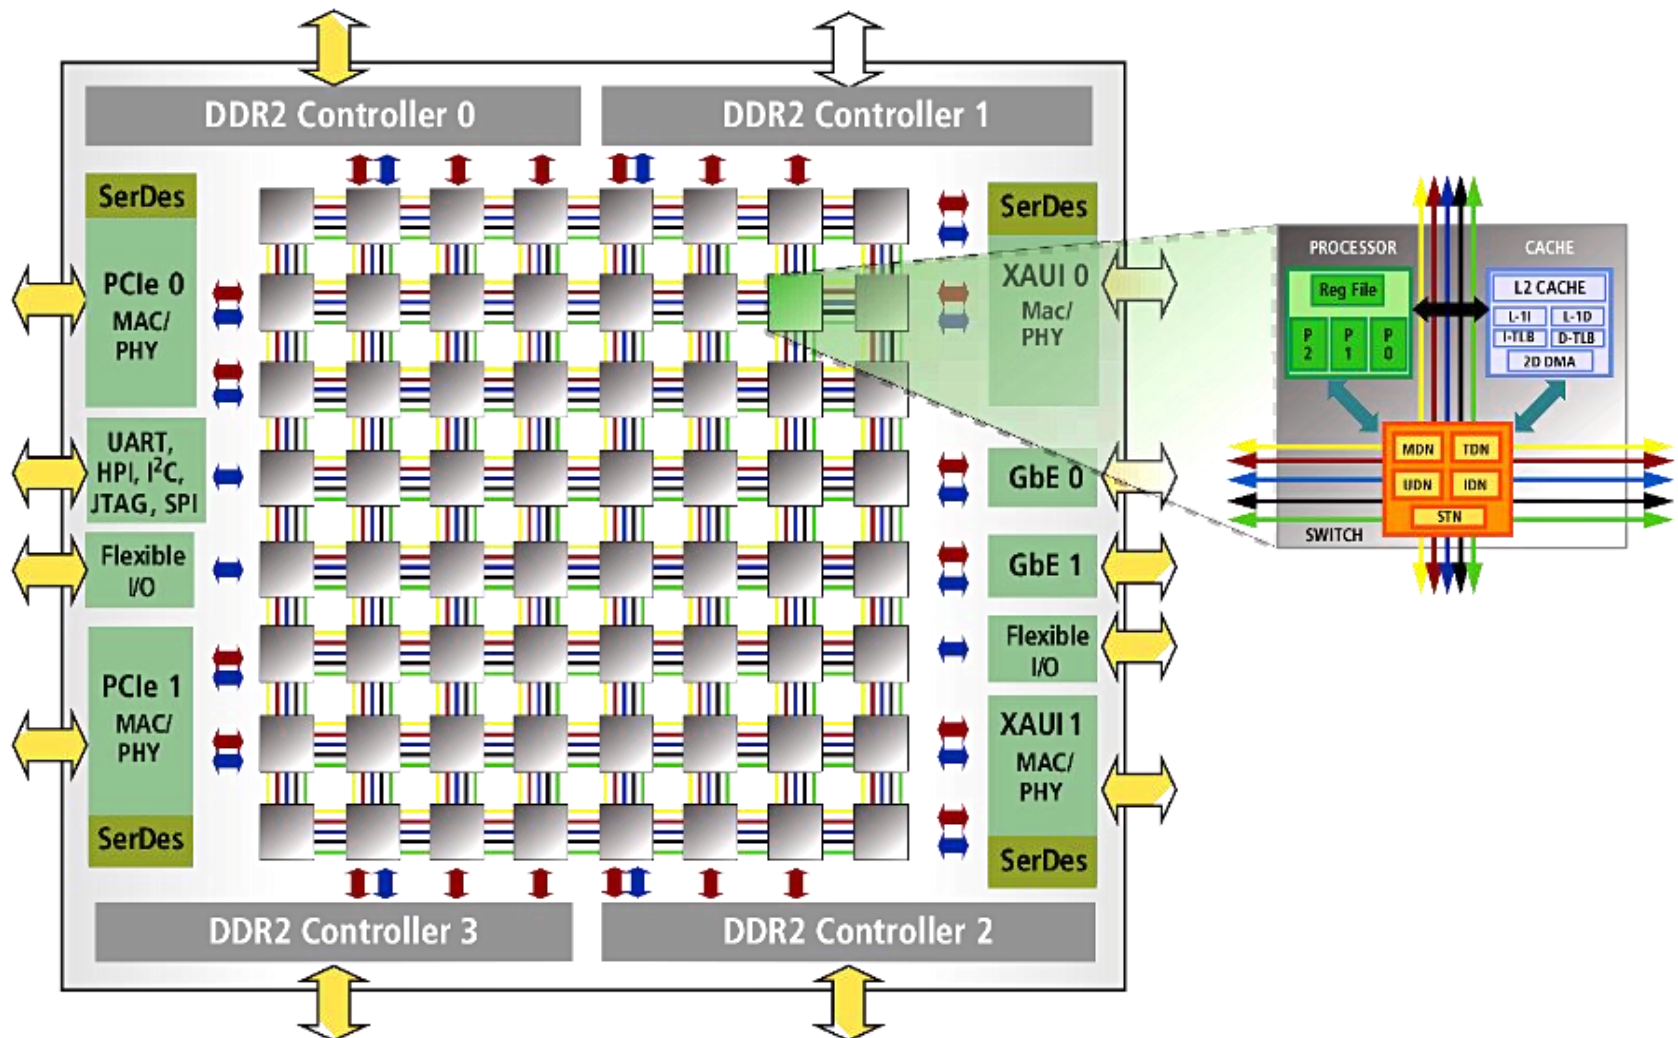
- Intel Core i7 Quad-core (Nehalem)



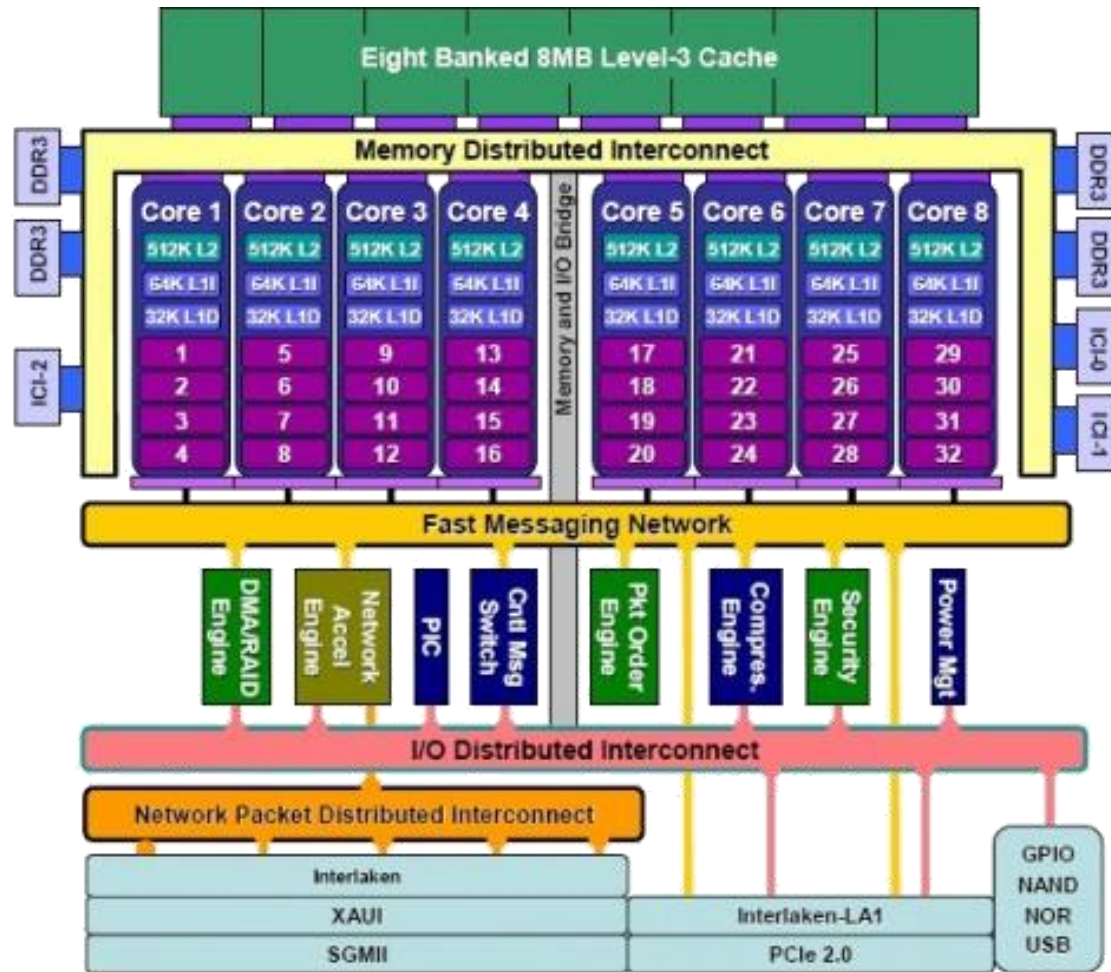
- AMD Opteron Quad-core



- Tilera TILE64



- NetLogic XLP series




XLP832 Processor Block Diagram

What's next

- Introduction
- Multi-core architectures
 - Memory hierarchy
 - Interconnect design
- Commercial examples
- **Parallel programming issues**
 - Forms of parallelism in software (overview)
 - Main issues in parallel programming
- Wrap-up

Parallel programming efforts

Source: "Parallelism for Multicore and manycore", James Reinders. Multicode Days '08, Stockholm



Forms of parallelism in software

- Instruction Level Parallelism (ILP)
 - (General term) require hardware support (e.g., instruction pipelining, VLIW processors, out-of-order execution, speculative control, ...)
- Multi-threading (MT)
 - Might be implicit (HW support) or explicit (SW structure)
- Task Level Parallelism (TLP)
- Data Level Parallelism (DLP)
- Pipelining (aka streaming)

Multi-threading

- Can be either **implicit** or **explicit**

*In charge to the software designer
(see example in next slides)*



Application is structured
as a collection of threads,
each having a specific role
(i.e., *decomposition*)

*“Hardware Threads”,
transparent to software designers*



Require HW support,
e.g. **HyperThreading**
technology from Intel

Task-Level Parallelism

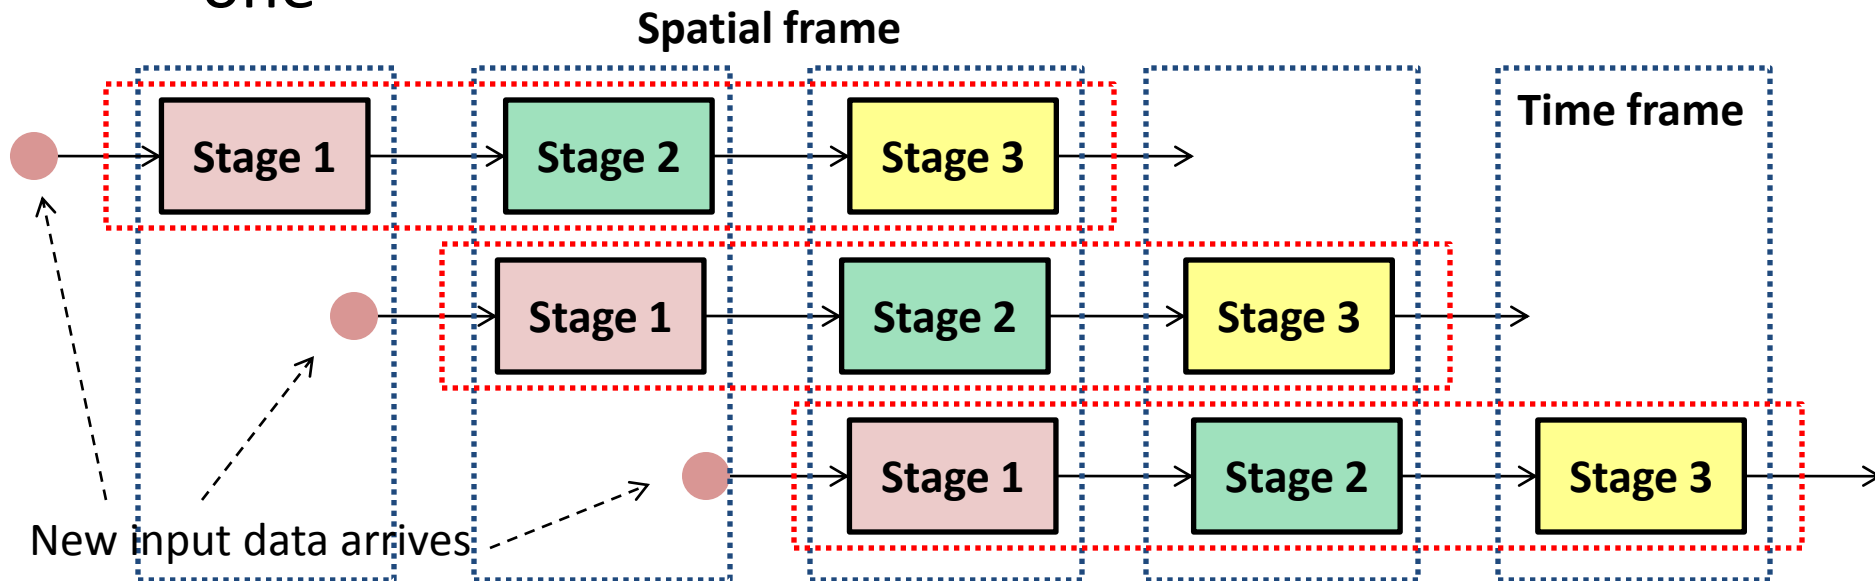
- Parallelism is exploited at process-level granularity
 - Multiple applications running concurrently on different processors
 - Processes from the same application running concurrently on different processors
- Processes execute different functionalities on different data
 - Single Program Multiple Data (SPMD) or Multiple Programs Multiple Data (MPMD)

Data-Level Parallelism

- Computational load is balanced on different processes running the same functionality on different data
 - e.g., image visualization algorithms
- Parallelism is done at the data level, not at the functionality level
 - Similar to the Single Instruction Multiple Data (SIMD) category defined by Flynn's taxonomy

Pipelining (streaming)

- An application is divided into elementary basic blocks (basic functionalities)
 - Applications is seen as a cascade of operations
 - Input in current stage is output from the previous one




Parallel programming

- Main challenges in parallel programming
 - Access to shared resources
 - Synchronization
 - Load balance
- Solutions
 - Hardware
 - Software
 - OS-supported


Shared resources

- A resource **X** can be shared among different processors
 - R/R operations aren't real issues
 - Issues when an operation is a W



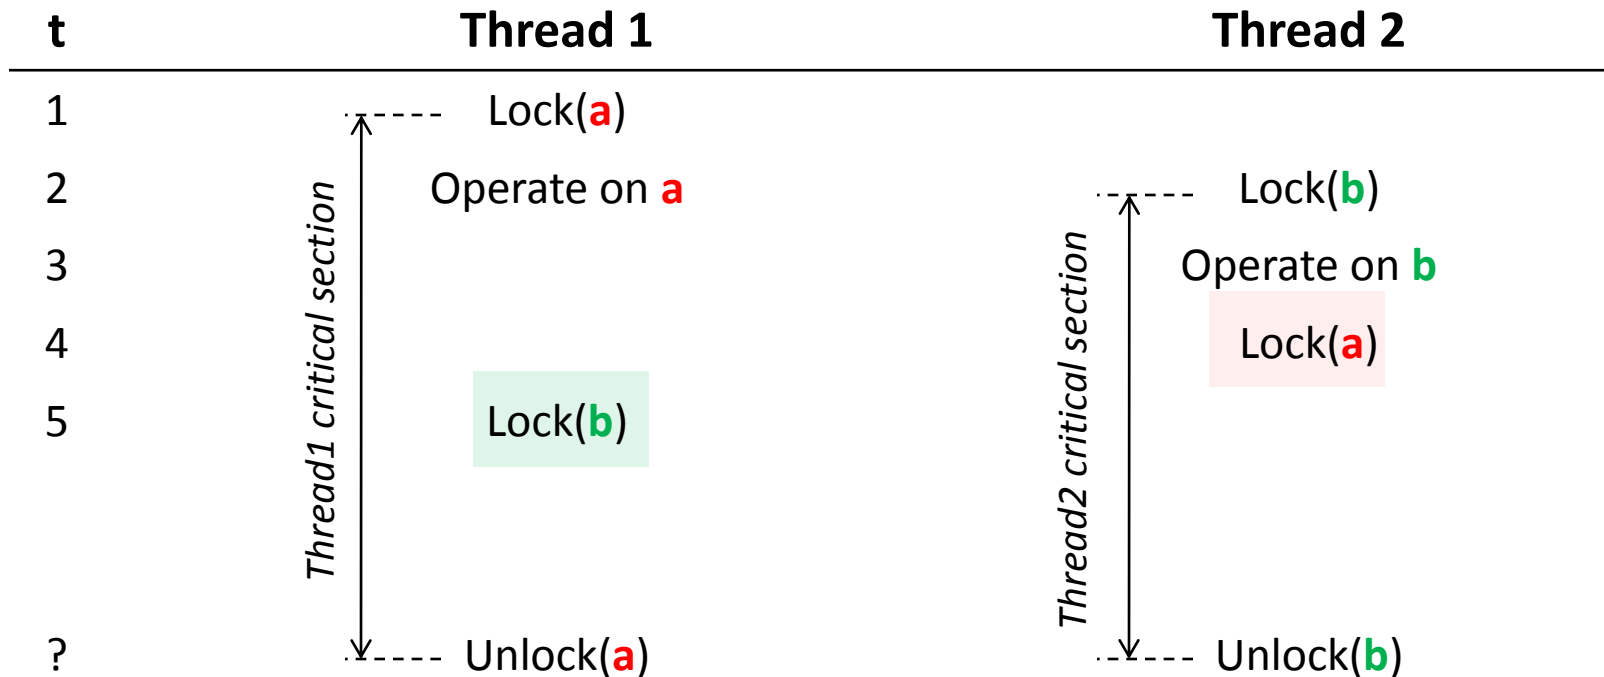
Locking resources

- A typical solution is to lock resources
 - Only one among the executing threads is allowed to access shared resource (**mutual exclusion**)
 - The other threads are blocked and wait until resource is released (unlocked)



Locking issues

- Consider the following example



POSIX locks

- POSIX definition for locks is employed in the mutex concept
- A `mutex_t` variable defines the access port to shared resources

```
#include <pthread.h>

pthread_t my_mutex = PTHREAD_MUTEX_INITIALIZER;
int shared_counter = 0;

/* Executed at generic process Pi */
pthread_mutex_lock(&my_mutex);
Shared_counter++;
pthread_mutex_unlock(&my_mutex);

/* When all done, destroy the lock */
pthread_mutex_destroy(&my_mutex);
```

RLock in POSIX standard


- What it process generic *Pi* locks twice the same lock object?
 - POSIX recursive locks can be acquired an arbitrary number of times by a process

```
#include <pthread.h>
```

```
pthread_t my_mutex;  
pthread_mutexattr_t my_mutex_attr;  
pthread_mutexattr_settype(&my_mutex_attr, PTHREAD_MUTEX_RECUR  
SIVE);  
pthread_mutex_init(& my_mutex_attr);
```

```
pthread_mutex_lock(&my_mutex);  
//...  
pthread_mutex_lock(&my_mutex);  
//...  
pthread_mutex_unlock(&my_mutex);  
pthread_mutex_unlock(&my_mutex);
```

Must be called an equal number
of times as *lock()*



Semaphores

- More advanced
 - Have an internal counter and block if a given number of threads attempted to hold it
- How do they work
 - Counter is decremented at each call to lock (`wait()`)
 - If counter is 0 then block
 - If counter becomes >0 (`post()`) wake up any waiting thread

The producer/consumer problem

- Generally speaking, parallelization leads to a (generalized) producer/consumer problem
 - A producer **P** generates data and send them through ad-hoc interfaces to the consumer(s) **C**
 - A consumer **C** reads incoming data and process them locally
- A producer cannot produce data if the communication via is full
- A consumer cannot read empty buffer

- Example, naive implementation with single semaphore (mutex)

PRODUCER

```
while(1)
{
01    lock(LockObject);
02    write(data, &buffer);
03    release(LockObject);
}
```

CONSUMER

```
while(1)
{
01    lock(LockObject);
02    read(data, &buffer);
03    release(LockObject);
}
```


- Assume shared buffer is an integer
 - What happens in next execution scenario?

PRODUCER

```
while(1)
{
01  lock(LockObject);
02  write(data, &buffer);
03  release(LockObject);
}
```

CONSUMER

```
while(1)
{
01  lock(LockObject);
02  read(data, &buffer);
03  release(LockObject);
}
```



lock: ~~N~~ONE
buffer: ~~E~~MPTY

-
- Previous example shows how a naive solution with a single locking mechanism is not suitable
 - Race condition is **not** solved
 - Generally, the producer/consumer problem cannot be solved using one semaphore
 - We need two semaphores to keep track of free and full buffer events

-
- `wait()` and `post()` are semaphore-related functions, equivalent to previous *lock* and *unlock*

INIT

```
initialize(empty_buffer, 1);  
Initialize(full_buffer, 0);
```

PRODUCER


```
while(1)  
{  
    wait(empty_buffer);  
    write(data, &buffer);  
    post(full_buffer);  
}
```

CONSUMER

```
while(1)  
{  
    wait(full_buffer);  
    read(data, &buffer);  
    post(empty_buffer);  
}
```

Barriers

- Barriers are used to synchronize a set of processes up to a specific point
 - Processes reaching the barrier are allowed to continue execution if **all** processes reached that barrier



Enforcing mutual exclusion


- **Hardware solutions**
 - Disable interrupts, avoiding ISR code to be executed while in the CS and avoiding context switch
 - Not suitable for multiprocessors
 - Test&Set instruction (on shared memory)
 - Spinlock
 - Compare&Swap
 - Hardware synchronizer for MPSoC architectures
- **Software solutions**
 - Locks
 - Mutexes (reentrant mutexes)
 - Semaphores

What's next

- Introduction
- Multi-core architectures
 - Memory hierarchy
 - Interconnect design
- Commercial examples
- Parallel programming issues
 - Forms of parallelism in software (overview)
 - Main issues in parallel programming
- **Wrap-up**

Speedup examples

Source: "Application Performance", RTWH-Aachen



The Multicore Association


- An open membership organization including companies implementing products embracing multicore technology
- To establish an industry-supported set of multicore programming practices and services



Copyright The Multicore Association™

<http://www.multicore-association.org/home.php>

-
- (Some of the) relevant parnters



References

- [1] Culler, David E.; Singh, Jaswinder Pal. “Parallel Computer Architecture: A Hardware/Software Approach”, Morgan Kaufmann
- [2] Hennessy, John L.; Patterson, David A. “Computer Architecture: A Quantitative Approach”, Elsevier
- [3] Herlihy, Maurice; Shavit, Nir. “The Art of Multiprocessor Programming”, Morgan Kaufmann
- [4] Shavit, Nir; Korland, Guy; Cohen, Hila. Slideset and lecture notes for the “*Multicore Programming*” course at Tel Aviv University, Spring 2010. Available at <http://www.cs.tau.ac.il/~multi/?p=slides>
- [5] Nagarajan, Rjagopal. “Multicore technologies and software challenges” Design article. Available at http://www.eetimes.com/design/eda-design/4008860/Multicore-technologies-and-software-challenges?cid=NL_Embedded&Ecosystem=embedded
- [6] Daily, Steve (Intel Co.). “Software Design Issues for Multi-core/Multiprocessor Systems” Design article. Available at http://www.eetimes.com/design/programmable-logic/4006624/Software-Design-Issues-for-Multi-core-Multiprocessor-Systems?cid=NL_Embedded&Ecosystem=embedded

-
- [7] Culler, David E.; Singh, Jaswinder Pal. “Matrix Operation on the GPU”, Slideset and lecture notes available at <http://www.seas.upenn.edu/~cis665/LECTURES/Lecture11.ppt>
- [8] Mitchell, Mark; Oldham, Jeffrey; Samuel, Alex. “Advanced Linux Programming”, Freely available at <http://www.advancedlinuxprogramming.com/downloads.html>


Amdahl's Law

$$Speedup = \frac{1}{(1 - P) + \frac{P}{S}}$$

Speedup of an algorithm due to an improvement with speedup S to the parallel portion P

$$Speedup = \frac{1}{(1 - P) + \frac{P}{N}}$$

Speedup of a program due to an increased number N of processors



Source: Hennessy, J.; Patterson, D.A. "Computer Architecture – A Quantitative Approach", Morgan Kaufman, 2007