

DATE 2010 – Dresden, Germany
Friday Workshops - WP3

Design Space Exploration for Run-Time Resources Management The MULTICUBE View

William FORNACIARI

fornacia@elet.polimi.it

Contributors: Simone Corbetta, Patrick Bellasi

Politecnico di Milano – Dipartimento di Elettronica e Informazione





Introduction

Outline

- **Introduction**
 - Application scenarios & modern requirements
- **Many-core architectures**
 - Technology perspective
 - Benefits
 - MPSoCs
- **Run-Time Resource Management**
 - Needs and value added of a run-time resource manager
 - Dynamic approaches
 - Power budgeting
- **Research directions**
- **The MULTICUBE View**



Introduction



Application scenarios



MOBILE



*Programmable
MPSoC platform*



LIFE-CRITICAL



**PROCESSING
POWER**

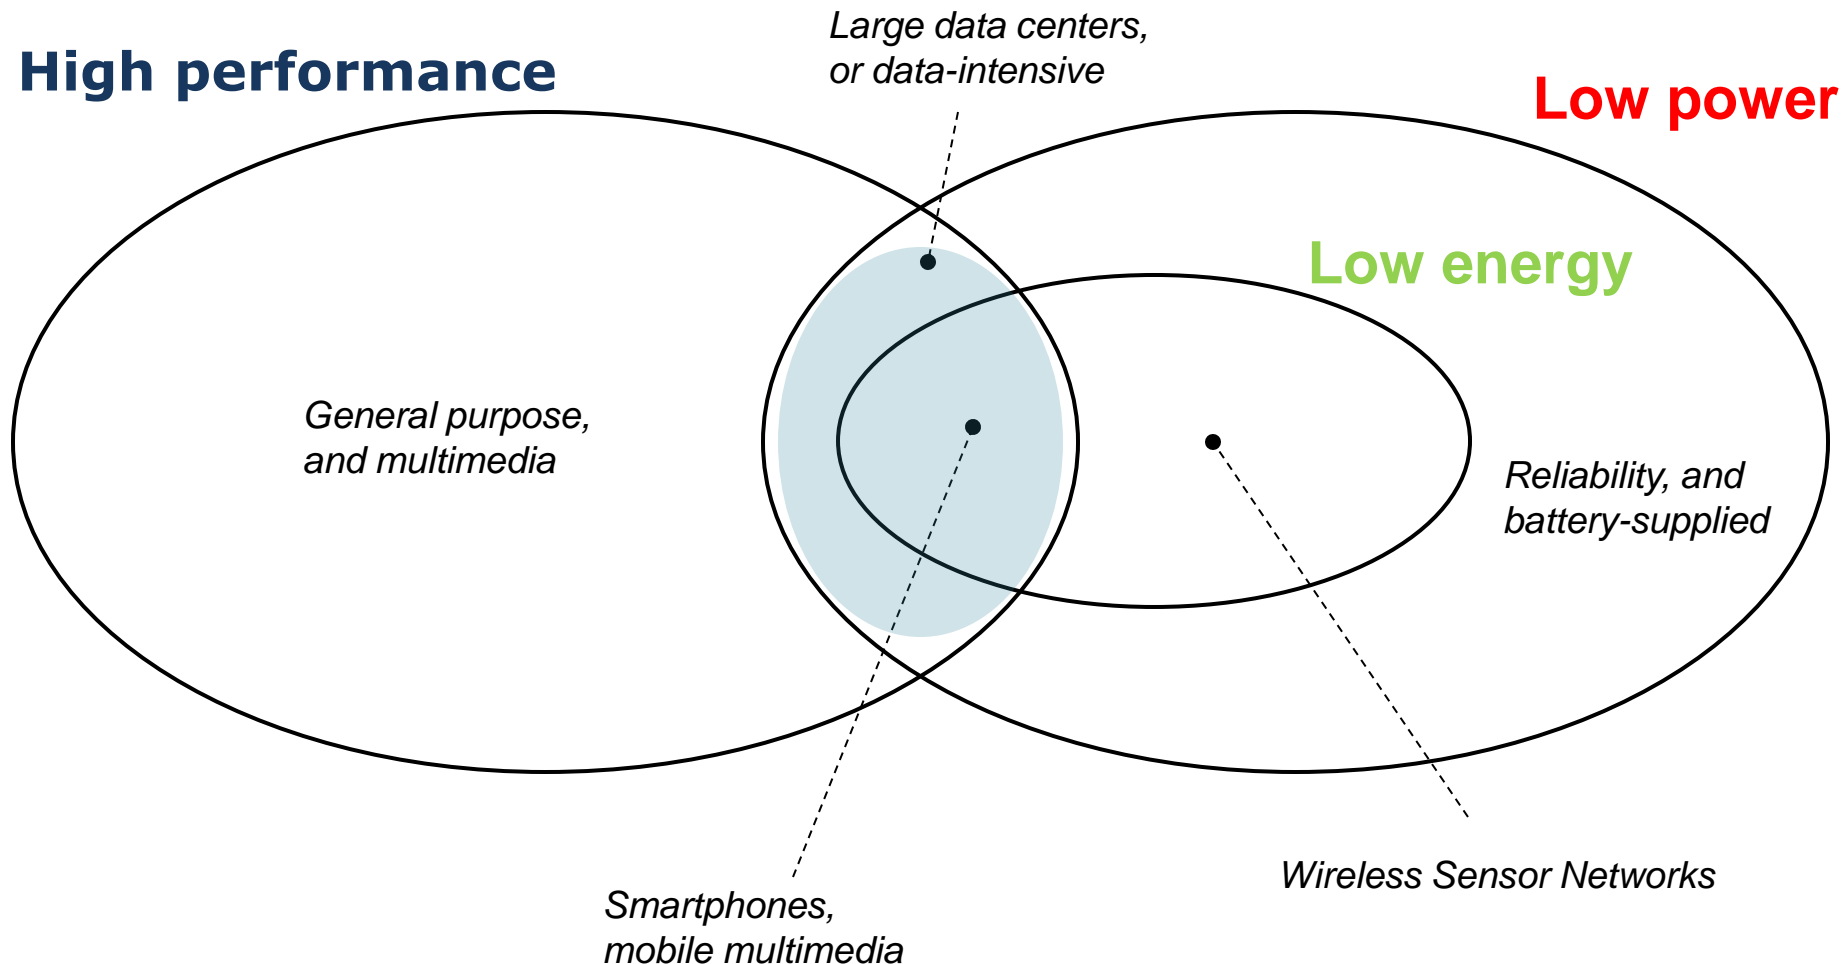


Introduction

Modern power/energy/performance requirements

- **Wide spectrum of application requirements**

High performance






Many-core Architectures

A technology perspective

• SoC architectures evolution



Source: International Technology Roadmap on Semiconductors. 2008 Edition, Design chapter.



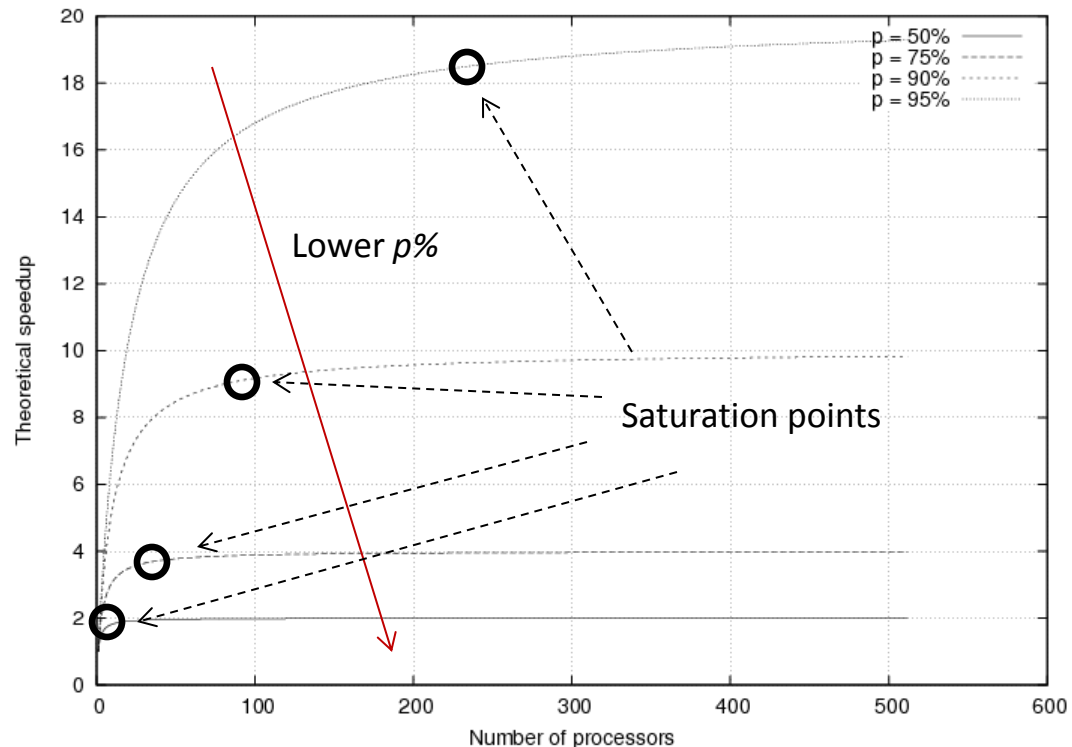
Many-core Architectures

A technology perspective (cont'd)

- **Amdahl's Law limits the speedup due to parallelization**

- Bounded to percentage $(1-p)$ of serial code, and number N of available processors
- Gain saturate soon as p approaches low values

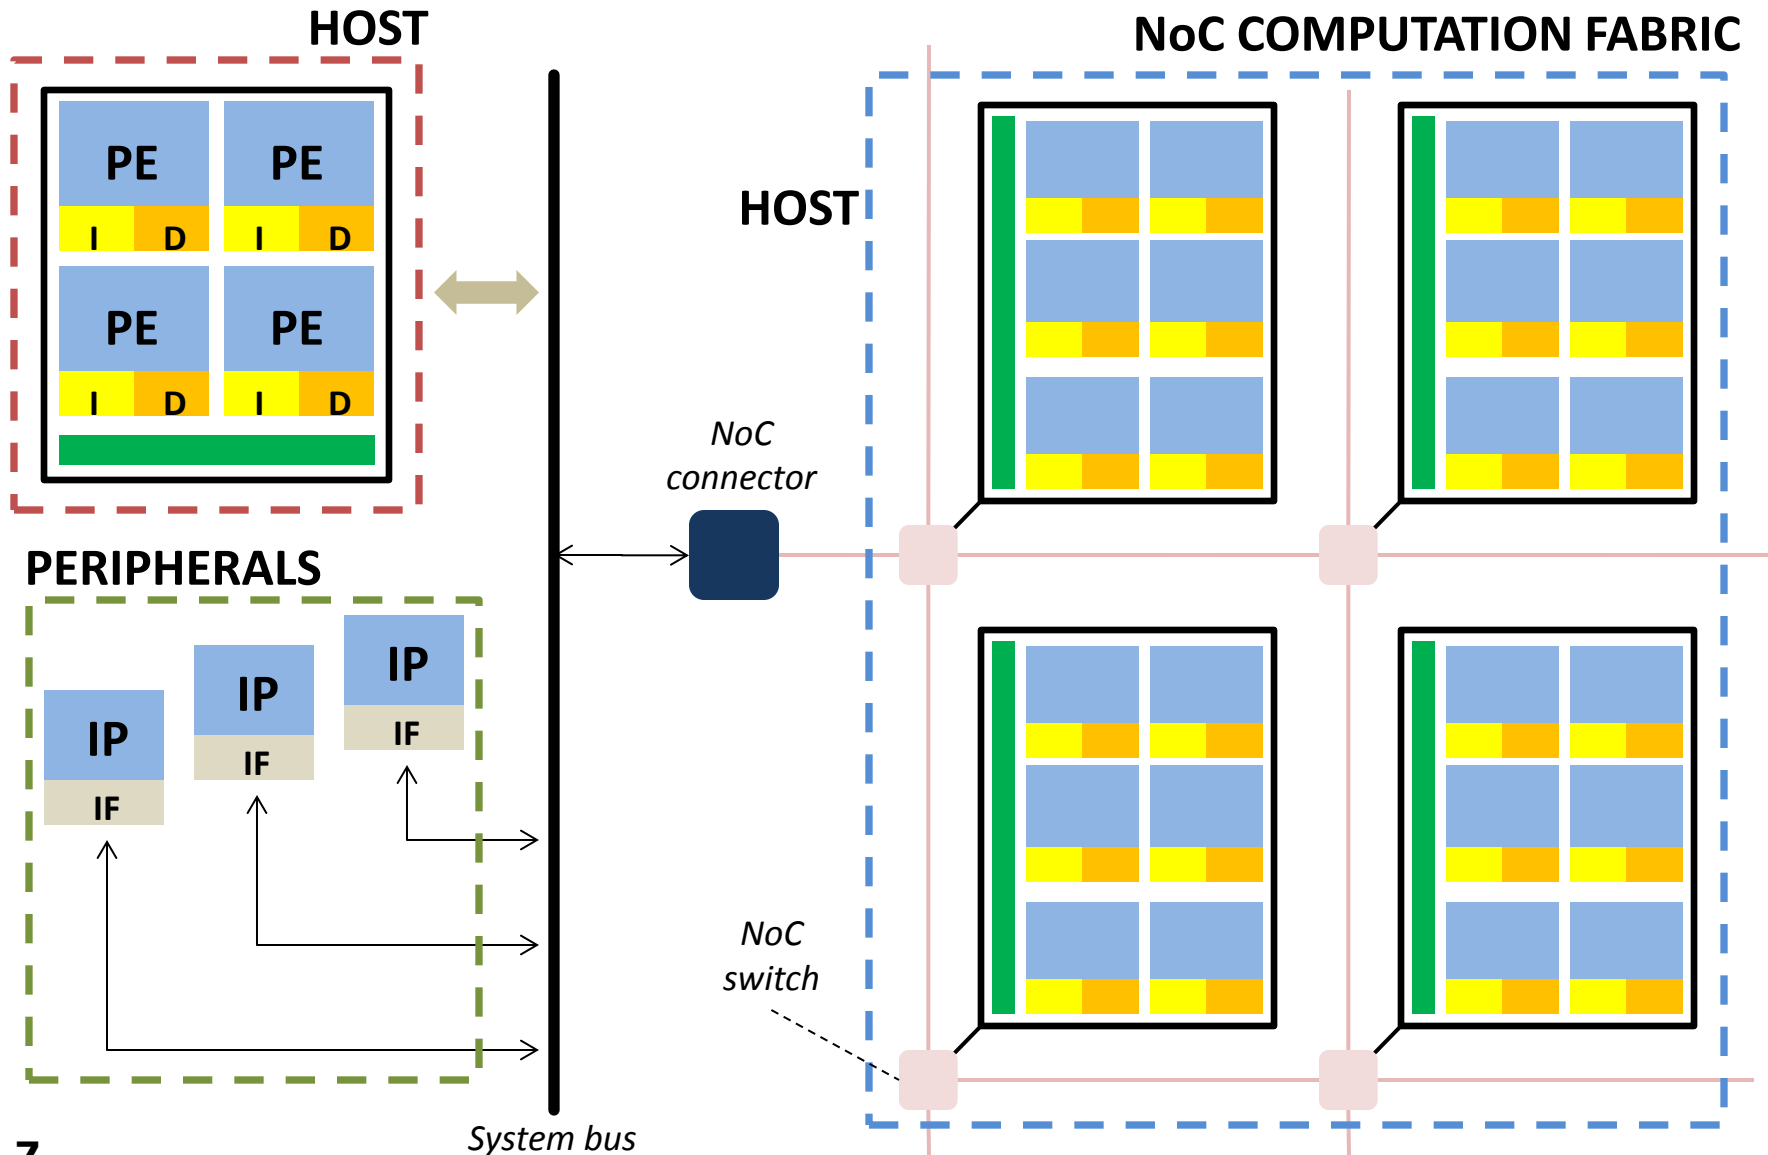
$$Speedup = \frac{1}{(1-p) + \frac{p}{N}}$$





Many-core Architectures

MPSoCs: heterogeneous and Network-on-Chip





Many-core Architectures

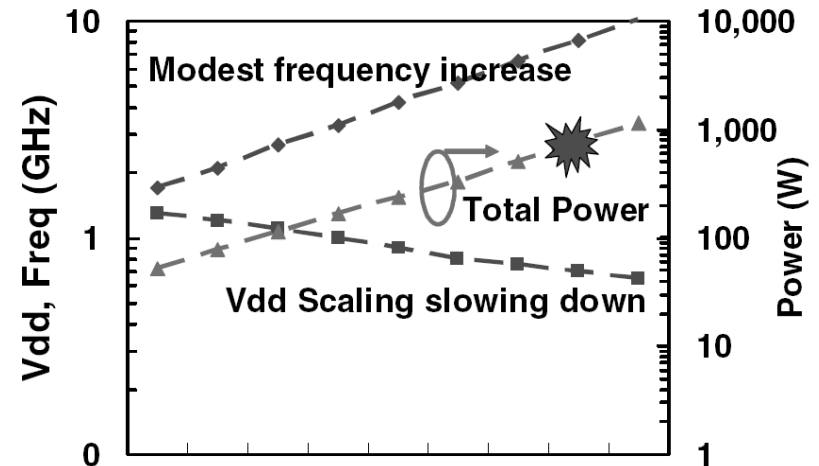
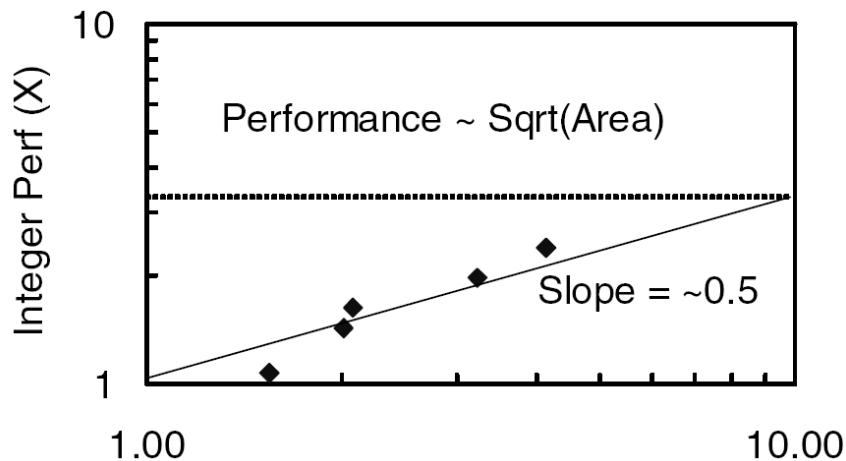
A technology perspective (cont'd)

- **Pollack's Rule**

- (Integer) performance goes with square root
- Doubling complexity leads to 40% of performance gain

- **Power versus integration trend**

- Linear: as soon as systems are more complex, power consumption gets worse



Source: Shekhar Borkar. **Thousand core chips: a technology perspective**. *Inproceedings of the 44th annual Design Automation Conference*, New York, NY, USA. 746-749. 2007.



Many-core Architectures

A technology perspective (cont'd)

- **Needs**

- Need to find suitable architecture design for MPSoCs
 - Extend Amdahl's Law by smart architectural choices
- Need to find suitable approaches to manage on-chip resources
 - Meeting power budget
 - Meeting user expectations and application constraints
 - Exploiting very high-performance

- **MPSoC architectures**

- Resources management is very complex
- Static approaches do not ensure flexibility
 - Dynamic approaches



Many-core Architectures

Typical Goals

- **Power savings**
 - By tuning each core voltage and frequency individually
- **Load balancing**
 - Use major portion of Silicon resources to exploit (not only) parallel operations
- **Lower on-die temperature**
 - Avoid hot-spots, decrease cooling costs and leakage power consumption
- **Reliability**
 - From previous reasons, and for persistence
- **Scalability**
 - Retarget MPSoC platform to different application scenarios



Many-core Architectures

Main issues

- **On-chip resources**
 - Heterogeneous (DSPs, CPUs, GPUs, generic PEs...)
 - Asymmetric w.r.t. performance and power consumption
- **On-chip applications**
 - Multiple scenarios for the same application
 - Multiple applications for the same device (true physical parallelism)
 - Cooperation or Competition
- **On-chip communication is a relevant bottleneck**
 - Collaborative application model must ensure performance, delays and low-power operation



Resources Management

The need for a run-time approach

- **Static approaches**
 - Optimal solution can be found for any application without impact on system performance
 - Design and compile-time lack of flexibility
- **Modern applications exploit dynamic behavior**
 - Static approach no more suitable for run-time requirements
 - Need to manage resources at run-time
 - Upon new application arrival
 - Upon new constraint asserted
 - Upon system fault
 - Upon resources congestion (e.g., NoC links)
 - ...



Resources Management

Benefits of run-time management

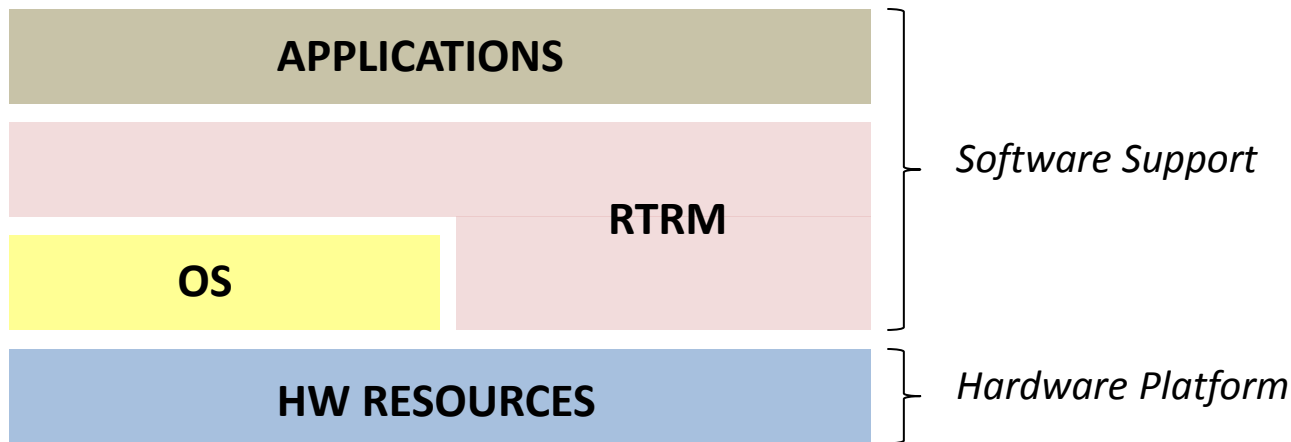
- **Flexibility**
 - Changing requirements and new applications
- **Adaptability**
 - Changing application scenario and requirements
- **Energy efficiency**
 - Power and thermal balancing for energy efficiency
- **Reliability**
 - Provide run-time solutions to application faults
- **Resource utilization**
 - Increase chip area utilization



Resources Management

Run-time manager

- **A run-time manager is required to take decisions and apply choices**
 - A software layer atop the hardware and below the applications
 - Provide applications with *flexible* and *transparent* resource management





Resources Management

Run-time manager components optimization

- **Different aspects of the same system have to be considered while managing resources**
 - Optimization can be performed statically, dynamically or through a smart mix of the two approaches

	Features	Optimization	
		Static	Dynamic
Platform aspects	heterogeneous multicore architecture, multiple PEs, parallel/distributed framework	•	•
Non-functional aspects	real-time, low power and high performance, determinism, responsiveness	•	•
Application aspects	multimedia, adaptive application, QoS, user experience		•

Many problems are interconnected and they can never be solved by looking at each in insulation [....] where do you start? By changing the way we are seeing the world...

[quotes from the Sangiovanni-Vincentelli Keynote @ DATE 2010]


- Each aspect can be mapped to a specific SoC subsystem

Application Aspects

Non-Functional Aspects



SOFTWARE APPLICATIONS



Constraints
assertion



Resources
allocation


HARDWARE PLATFORM



Single
core



Heterogeneous
many-core



Homogeneous
multi-core

Platform Aspects



Resources Management

Design and implementation space

- **Design and implementation space defines the features of the resources manager**
 - *Distribution*
 - Distributed manager
 - Centralized manager
 - *Flexibility*
 - Adaptive
 - Non-adaptive
 - *Implementation aspects*
 - Hardware acceleration
 - Software



Resources Management

Run-time approaches

- **There exist different techniques to manage on-chip resources dynamically**
 - *Task migration*
 - Directly manage executing tasks
 - Load balancing
 - Memory hierarchy management
 - Manage caches
 - Data locality and memory switching power
 - Dynamic thermal (distribution) management
 - Manage tasks, by moving them from overloaded processors
 - Power distribution (focus on static power), avoid hot-spots
 - Dynamic power management
 - Manage power resource directly
 - Dynamic and static power reduction



Task Migration

Main issues

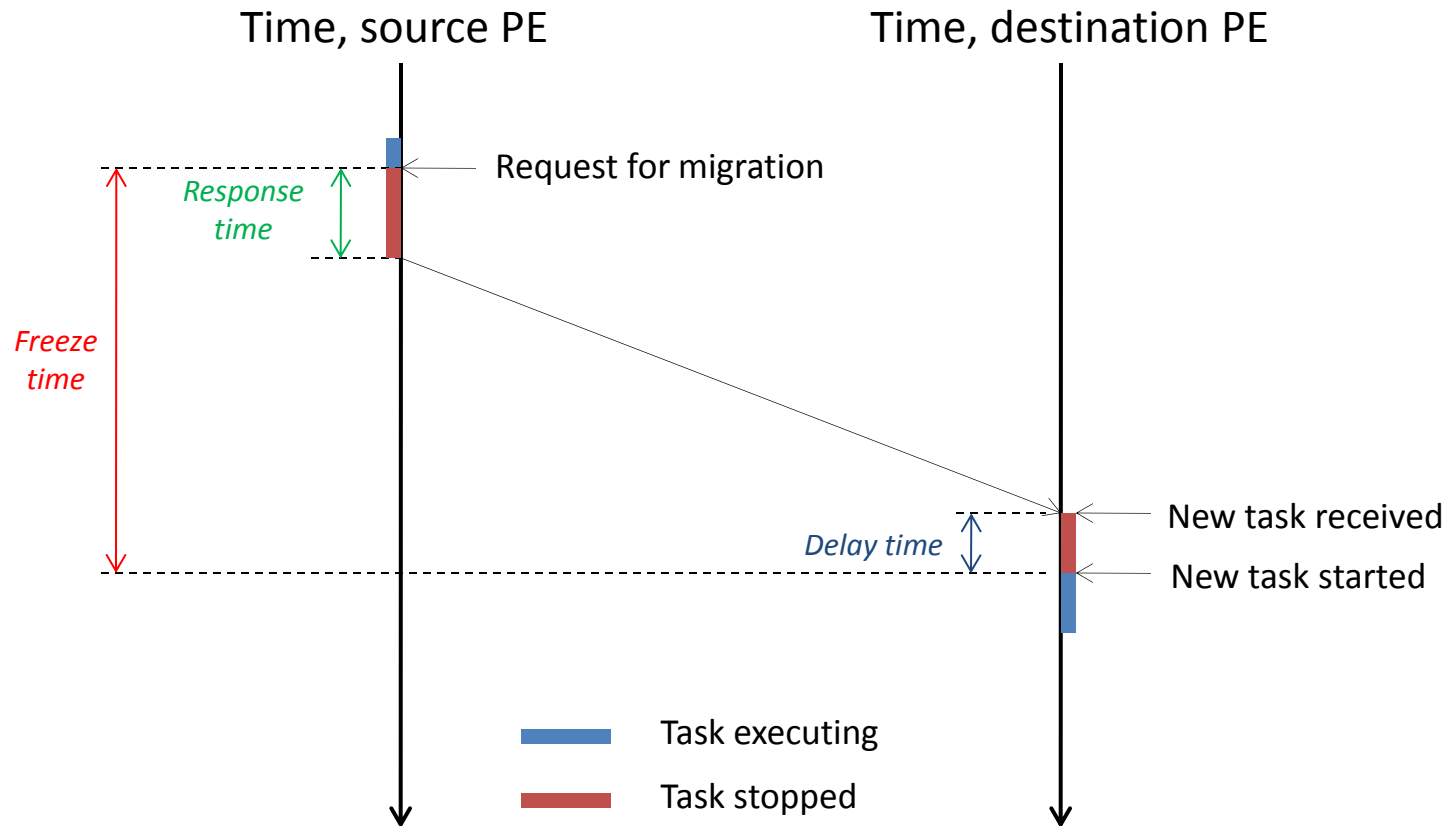
- **Task migration requires communication between source and destination tile**
 - Can be expensive if PEs are physically far each other
 - Executing PE (source and destination) have to communicate with the host
 - Memory hierarchy traversal
 - NoC communication inserts non-null latency
- **Freeze time interval may be too wide**
 - Real-time tasks are not suitable candidates for migration
- **Revenue estimation**
 - How to select *when* and *where* to move a task?



Task Migration

Basic principles

- A task is moved from the source node (a PE) to the destination node (another PE)





Task Migration

Performance measures

- **Response time**

- Time required to initiate migration once the request has been received

- **Freeze time**

- Time interval in which the task is not executing

- **Delay time**

- Delay required at the destination node to insert task in processor queues
- Accounts for destination node internal status upon task arrival and new task insertion



Task Migration

Benefits


- **Why should we take care about task migration?**
 - Load balancing
 - Move tasks from overloaded processor to available processor
 - Access to distributed processing power
 - Power and thermal distribution
 - Higher performance and reliability
 - Meet application requirements
 - Lower power requirements
 - Compensate battery discharging rate on mobile systems
 - Resource locality
 - Remote data access vs migration toward data
 - Fault resilience
 - Detect a fault, and move to an appropriate safer place




Memory Management

Cache architectures templates


- **Different memory organizations for different purposes**
 - Trade-off with respect to power consumption and memory access performance



L1 and L2 shared

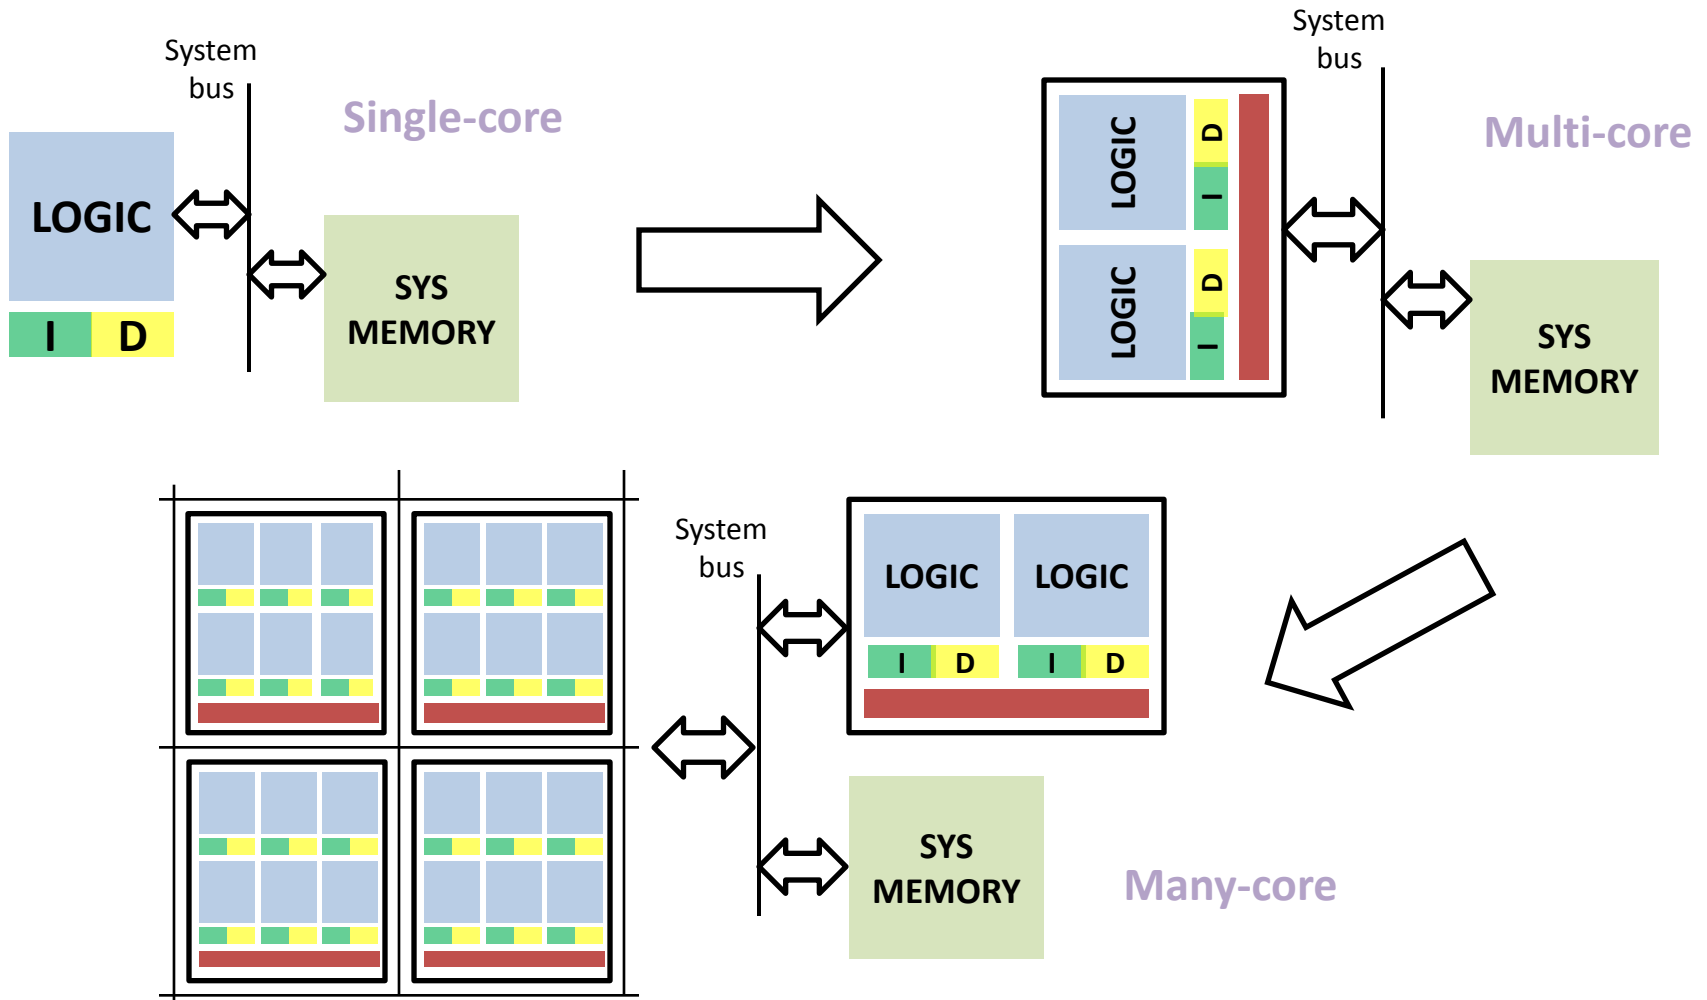


L1 private, L2 shared through bus



L1 shared through crossbar, L2 shared through bus

- The evolution of SoC memory architectures**





Memory Management

Cache architectures

- **Memory hierarchy constitutes a significant portion of the system resources**
 - Energy is spent during memory accesses
- **Active energy is consumed during data accesses**
 - Access bit lines for storing data
- **Standby power is consumed during no-operations**
 - Unwanted bit lines activation
- **Optimization can be done from software implementation guidelines or hierarchy architecture optimization**



Memory Management

Major issues

- **Hierarchy management**
 - Hierarchy depth and complexity
 - Memory consistency
 - Cache coherence
- **Dynamic power management**
 - Operate on data locality
 - Avoid switching memory when unnecessary
- **Static power management**
 - Operate on low-power states
 - Retention voltage, voltage scaling




Thermal Distribution


Dynamic thermal management

- **Thermal distribution**

- Distribute resources so that thermal profile of the integrated circuit is balanced



VS



 HIGH-T
 LOW-T


● Esecuting task



Thermal Distribution

Benefits

- **Thermal distribution spread power sources chip-wide**
 - Leakage phenomena are reduced locally (at each PE)
 - Each core has improved reliability
 - High temperature reduces stability, reliability and performance
 - Lower cooling costs
 - Cooling circuit is much more efficient
- **Thermal distribution is provided at chip-level, not at core-level**



Thermal Distribution

Dynamic management

- **Thermal distribution management can not be addressed a priori (i.e., statically)**
 - There is no knowledge on work load
 - Statistical approaches lack of flexibility
 - What about if new applications arrive?
 - What about deviation from predicted behavior?
- **Need to manage it at run-time**
 - Monitor per-core thermal (and power) profiles
 - Make choices according to
 - Thermal profile
 - Communication between tasks



Power Management

Dynamic power management

- **Basic idea**

- Manage power resource during normal system execution
- Hardware and software integrated approach
 - Hardware provide basic mechanisms, software provides system-wide optimization
- Require both hardware blocks knowledge and software support exploitation
 - OS level, kernel code

- **Directly address the problem from the power view-point**

- RTRM is a more general problem

Source: Patrick Bellasi, Simone Corbetta, and William Fornaciari. **Hardware/Software Codesign for Dynamic Power Management.**



Open research directions

Run time resource management - Summary

- **MPSoCs and application complexity**
 - Energy-efficient resources management becomes challenging
 - Need to provide both static and dynamic support
 - Flexibility and system-wide power/performance optimization
- **Different techniques exist**
 - Not so appropriate for heterogeneous highly distributed and parallel architectures
 - Need to define novel approaches for system-wide power management
- **Future works**
 - Find suitable RTRM solution that is scalable toward future manycore architectures
 - Hierarchical power management
 - Mobile embedded systems application scenarios
 - Industrial feedback and support with real circuitry



MULTICUBE Project

Partners



MULTI-OBJECTIVE DESIGN SPACE EXPLORATION OF MULTI-PROCESSOR SOC ARCHITECTURES FOR EMBEDDED MULTIMEDIA APPLICATIONS

www.multicube.eu

Project Duration: from January 2008 to June 2010



Politecnico di Milano (POLIMI) – Italy (Project Coordinator)



DS2 – Spain



**Università della Svizzera Italiana
(ALaRI) - CH**



IMEC - Belgium



University of Cantabria - Spain



STMicroelectronics - Italy



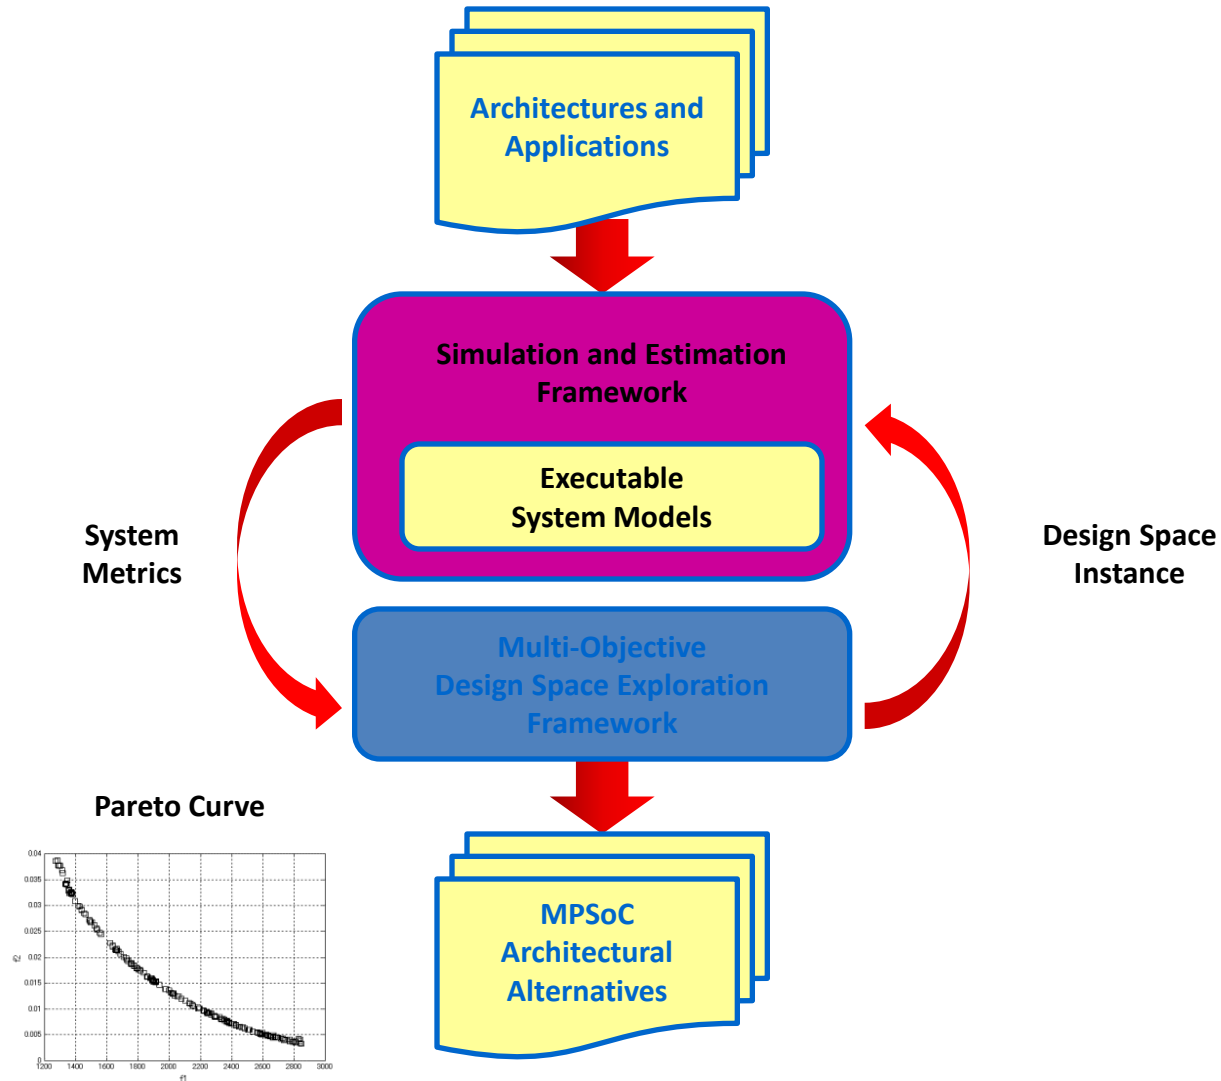
STMicroelectronics - China

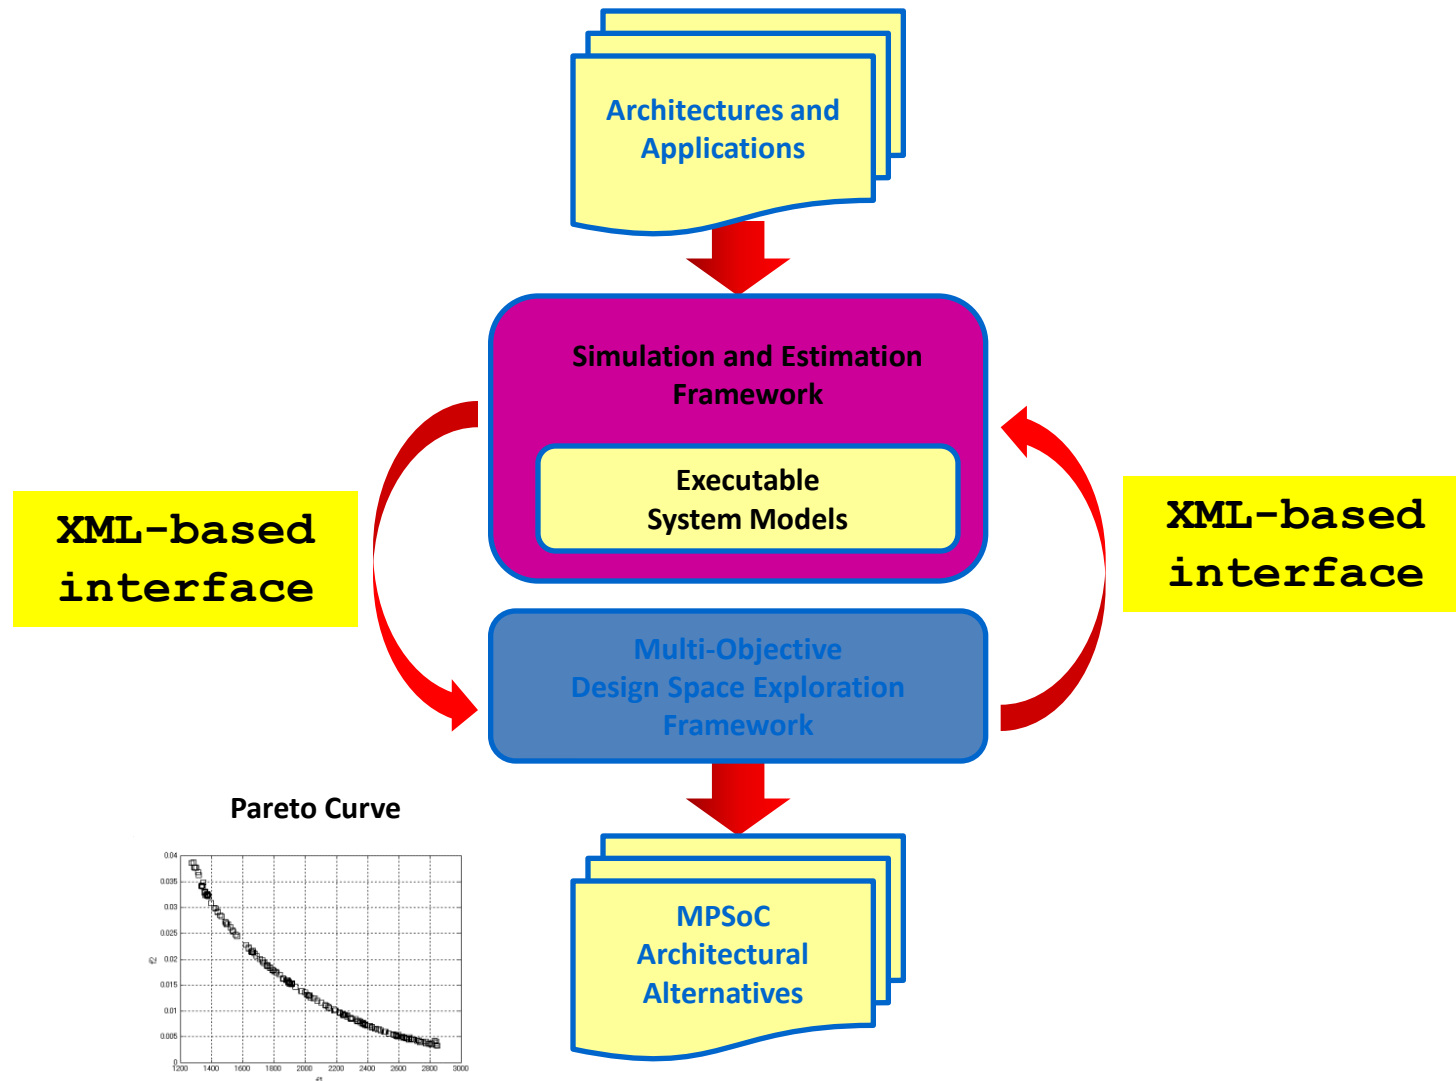


ESTECO - Italy



**Institute of Computing Technology (ICT)
China**







Automatic DSE

Main benefits



- Efficiency of DSE process is improved
 1. Minimizing the numbers of simulations to be executed by using **exploration heuristics** such as state-of-art evolutionary algorithms
 2. Speeding up simulations
 3. Simulating at higher abstraction levels
 4. Defining an **analytical response model** of the system behavior based on a subset of simulations to predict the unknown system response



MULTICUBE Explorer

What is it?



- MULTICUBE Explorer is **not** *"yet another simulator"*, there already many simulators....
- MULTICUBE Explorer is an automation and **acceleration infrastructure** to minimize the numbers of simulations to be executed during the DSE process
- MULTICUBE Explorer is a DSE infrastructure where you can easily **plug-in your own simulator**
- MULTICUBE Explorer can **interface** several simulators at **multiple abstraction levels**



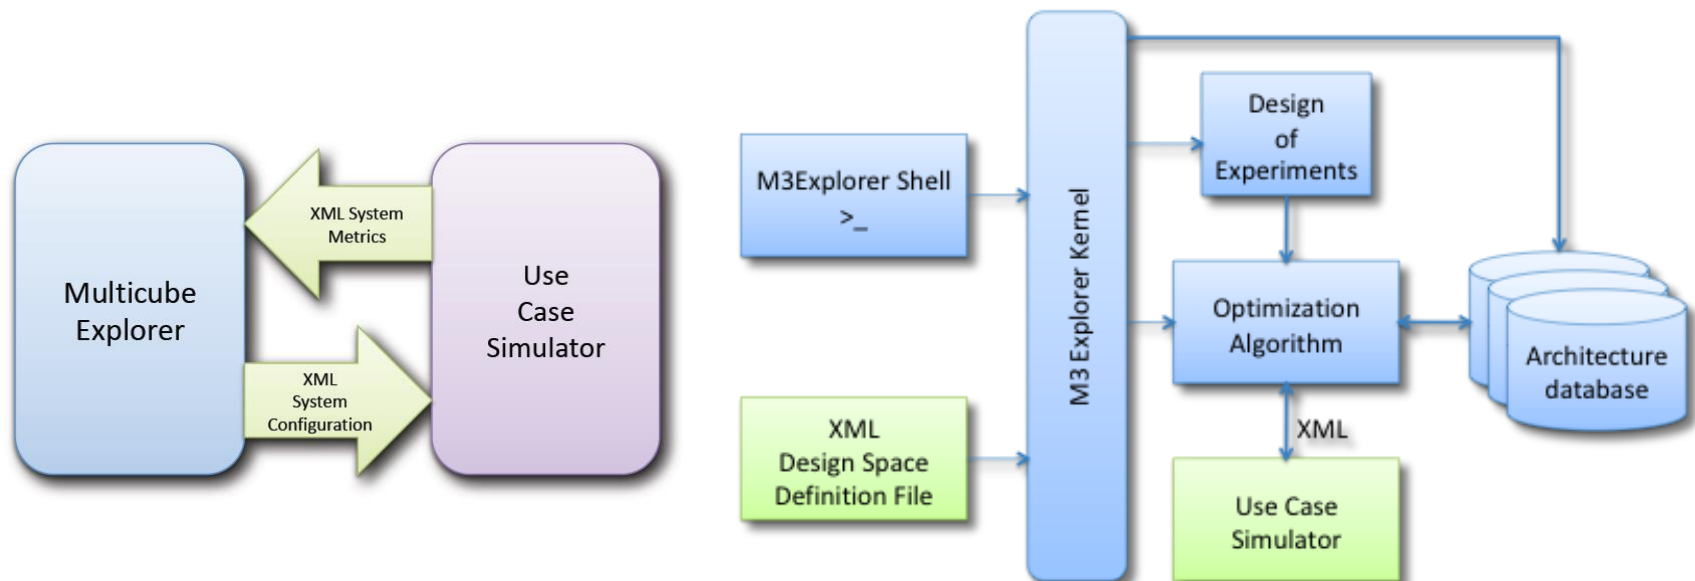
MULTICUBE Explorer

Main Goals



- **Open-source prototype exploration framework (MULTICUBE Explorer):**
The tool enables a fast automatic optimization of parameterized system architectures towards a set of multiple objectives

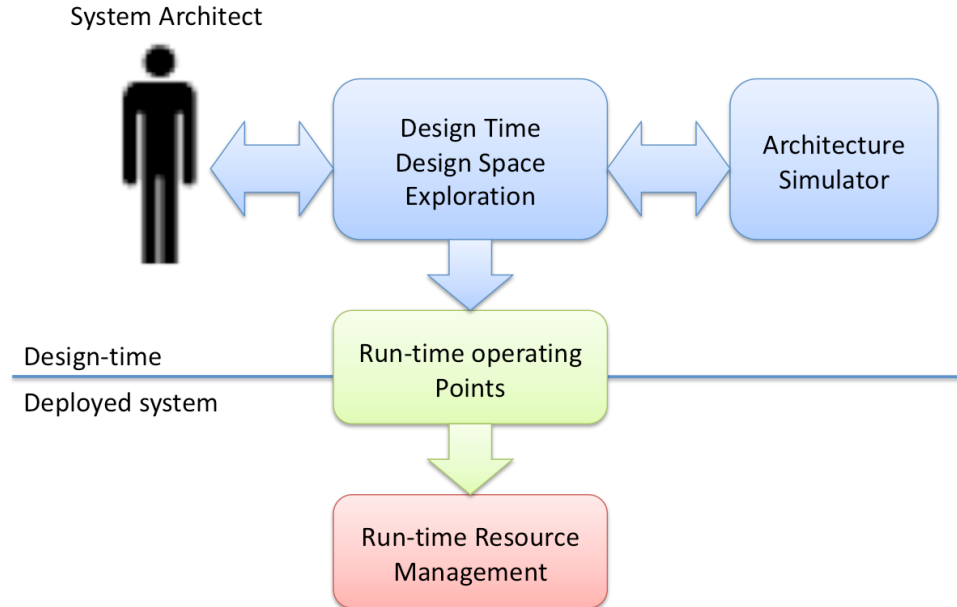
www.multicube.eu





MULTICUBE Explorer


Identification of run-time operating points



- **Design Time DSE** of run-time configurable parameters
 - Pareto optimal Run-time operating Points
- **Run-time Resource Management**
 - Selects dynamically the pre-identified operating points
- **Value added of MULTICUBE Explorer**
 - Fast design space exploration of usage scenarios
 - Pre-computation of run-time manager behavior

- **An example of usage of DSE tool for Run-Time Resource Management**

- An application is profiled off-line
 - Operating points are identified on a Pareto frontier
- Two instances of the same application are run
 - Each application has its own set of requirements
- The number of processors allocated for each application instance is determined at run-time
 - According to resources availability and constraints





From “AN INDUSTRIAL DESIGN SPACE EXPLORATION FRAMEWORK FOR SUPPORTING RUN-TIME RESOURCE MANAGEMENT ON MULTI-CORE SYSTEMS”

DSE integration with RTRM

Evolution of the MULTICUBE approach → 2PARMA


Hierarchical



Requirements
Aggregation

P_g

Run-Time Resource Manager



Retargetability



The End

- **Links to FP7 projects**
 - www.multicube.eu end June 2010
 - www.2parma.eu FP7 Strep
 - complex.offis.de FP7 IP
- **MULTICUBE Explorer download**
 - Follow links reported in www.multicube.eu
- **Info regarding MULTICUBE**
 - A book from Springer will appear in early 2011



THANKS FOR YOUR ATTENTION

William Fornaciari

fornacia@elet.polimi.it