



# Autonomous Agents and Multiagent Systems

*Distributed Constraint Optimization*

Francesco Amigoni

Many of the following slides are taken from:

- the “lecture slides provided by authors”, Chapter 12, available at:  
<http://www.the-mas-book.info/index.html>
- the “Tutorial on Multi-agent Distributed Constrained Optimization” at AAMAS2019, available at:  
<https://www2.isye.gatech.edu/~fferdinando3/cfp/AAMAS19/>

# Chapter 12: Distributed Constraint Handling and Optimization

A. Farinelli<sup>1</sup> M. Vinyals<sup>1</sup> A. Rogers<sup>2</sup> N.R. Jennings<sup>2</sup>

<sup>1</sup>Computer Science Department  
University of Verona, Italy

<sup>2</sup>Agents, Interaction and Complexity Group  
School of Electronics and Computer Science  
University of Southampton, UK

Multi-Agent Systems. MIT Press, 2012 (2nd edition), edited by  
Gerhard Weiss. <http://www.the-mas-book.info>

# Constraints

- Pervade our everyday lives
- Are usually perceived as elements that limit solutions to the problems we face



# Constraints

From a computational point of view, they:

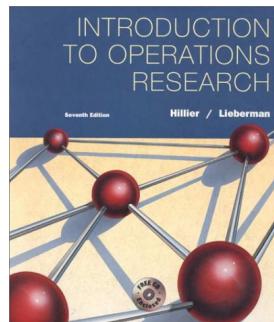
- Reduce the space of possible solutions
- Encode knowledge about the problem at hand
- Are key components for efficiently solving hard problems

# Constraint Processing

*Many different disciplines deal with hard computational problems that can be made tractable by carefully considering the constraints that define the structure of the problem.*



Planning  
Scheduling



Operational  
Research



Automated Reasoning  
Decision Theory

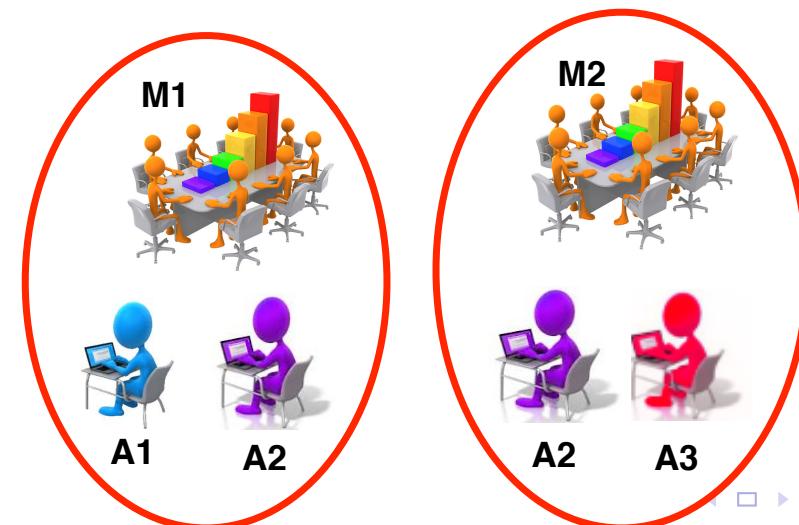


Computer  
Vision

# Constraint Processing in Multi-Agent Systems

Focus on how constraint processing can be used to address optimization problems in **Multi-Agent Systems (MAS)** where:

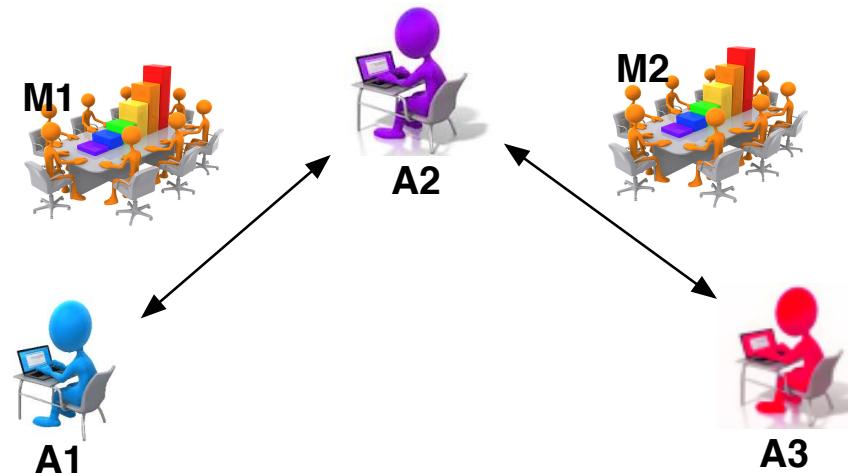
*A set of **agents** must come to some **agreement**, typically via some form of negotiation, about **which action** each agent **should take** in order to jointly obtain the **best solution** for the **whole system**.*



# Distributed Constraint Optimization Problems (DCOPs)

We will consider **Distributed Constraint Optimization Problems (DCOP)** where:

*Each agent negotiates locally with just a subset of other agents (usually called neighbors) that are those that can directly influence his/her behavior.*



# Constraint Networks

A **constraint network**  $\mathcal{N}$  is formally defined as a tuple  $\langle X, D, C \rangle$  where:

- $X = \{x_1, \dots, x_n\}$  is a set of **discrete variables**;
- $D = \{D_1, \dots, D_n\}$  is a set of **variable domains**, which enumerate all possible values of the corresponding variables; and
- $C = \{C_1, \dots, C_m\}$  is a set of **constraints**; where a constraint  $C_i$  is defined on a subset of variables  $S_i \subseteq X$  which comprise the scope of the constraint
  - $r = |S_i|$  is the arity of the constraint
  - Two types: **hard** or **soft**

# Hard constraints

- A **hard constraint**  $C_i^h$  is a relation  $R_i$  that **enumerates** all the **valid joint assignments** of all variables in the scope of the constraint.

$$R_i \subseteq D_{i_1} \times \dots \times D_{i_r}$$

$$D_j = D_k = \{0, 1\}$$

$R_i$	$x_j$	$x_k$
	0	1
	1	0

# Soft constraints

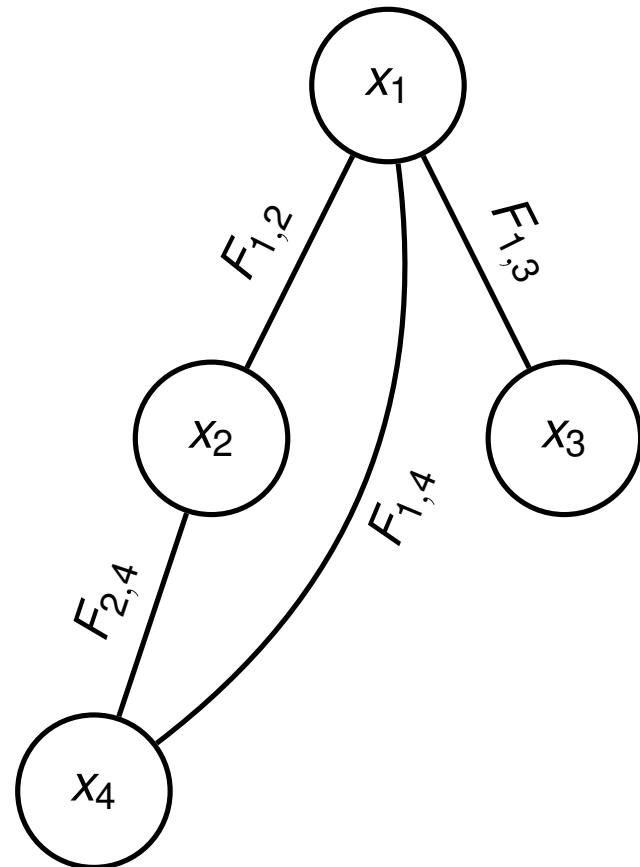
- A **soft constraint**  $C_i^s$  is a function  $F_i$  that **maps** every possible **joint assignment** of all variables in the scope **to a real value**.

$$F_i : D_{i_1} \times \dots \times D_{i_r} \rightarrow \Re$$

$F_i$	$x_j$	$x_k$
2	0	0
0	0	1
0	1	0
1	1	1

# Binary Constraint Networks

- Binary constraint networks are those where:
  - Each constraint (soft or hard) is defined over two variables.
- Every constraint network can be mapped to a binary constraint network
  - requires the addition of variables and constraints
  - may add complexity to the model
- They can be represented by a constraint graph



# Different objectives, different problems

- **Constraint Satisfaction Problem (CSP)**

- Objective: find an assignment for all the variables in the network that satisfies all constraints.

- **Constraint Optimization Problem (COP)**

- Objective: find an assignment for all the variables in the network that satisfies all constraints and optimizes a global function.

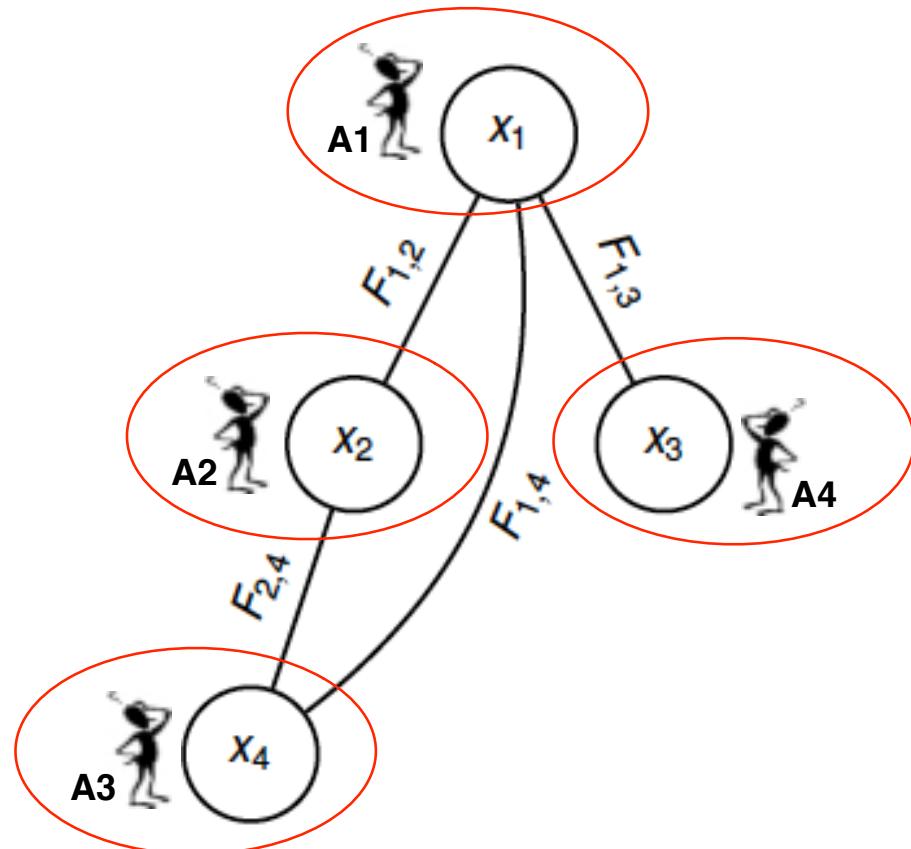
- Global function = aggregation (typically sum) of local functions.

$$F(x) = \sum_i F_i(x_i)$$

# Distributed Constraint Reasoning

When operating in a decentralized context:

- a set of agents control variables
- agents interact to find a solution to the constraint network



# Distributed Constraint Reasoning

Two types of decentralized problems:

- **distributed CSP (DCSP)**
- **distributed COP (DCOP)**

Here, we focus on DCOPs.

# Distributed Constraint Optimization Problem (DCOP)

A DCOP consists of a constraint network  $\mathcal{N} = \langle X, D, C \rangle$  and a set of agents  $A = \{A_1, \dots, A_k\}$  where each agent:

- controls a subset of the variables  $X_i \subseteq X$
- is only aware of constraints that involve variable it controls
- communicates only with its neighbours

# Distributed Constraint Optimization Problem (DCOP)

- Agents are assumed to be fully cooperative
  - Goal: find the assignment that optimizes the global function, not their local local utilities.
- Solving a COP is NP-Hard and DCOP is as hard as COP.

# Motivation

## Why distribute?

- Privacy
- Robustness
- Scalability

# Real World Applications

Many **standard benchmark problems** in computer science can be modeled using the DCOP framework:

- graph coloring

As can many **real world applications**:

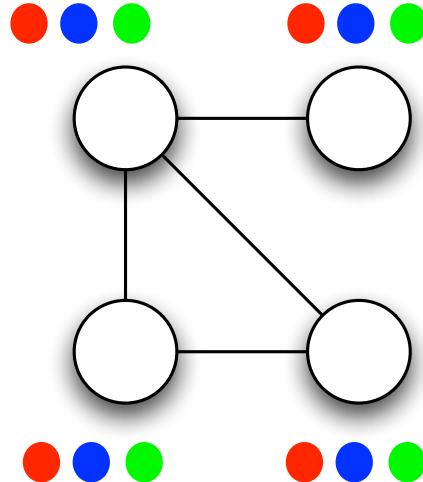
- human-agent organizations (e.g. **meeting scheduling**)
- sensor networks and robotics (e.g. **target tracking**)

# Graph coloring

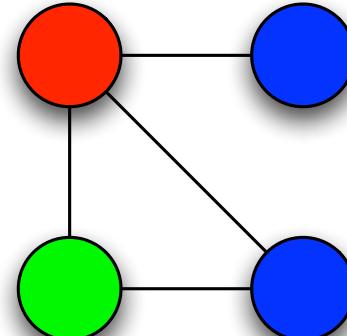
- Popular **benchmark**
- **Simple formulation**
- **Complexity** controlled with **few parameters**:
  - Number of available colors
  - Number of nodes
  - Density ( $\#nodes/\#constraints$ )
- **Many versions** of the problem:
  - CSP, MaxCSP, COP

# Graph coloring - CSP

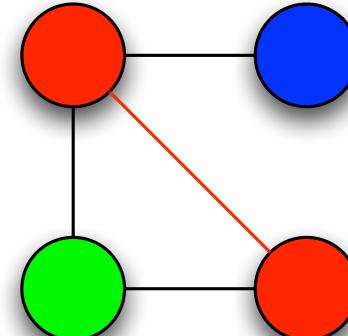
- Nodes can take  $k$  colors
- Any two adjacent nodes should have different colors
  - If it happens this is a conflict



Yes!

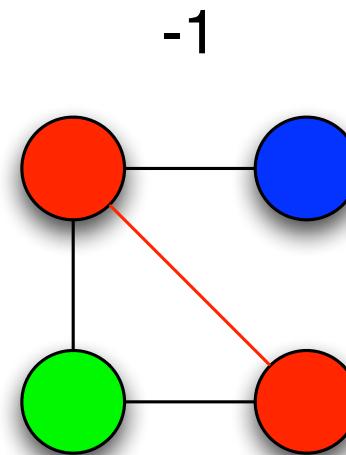
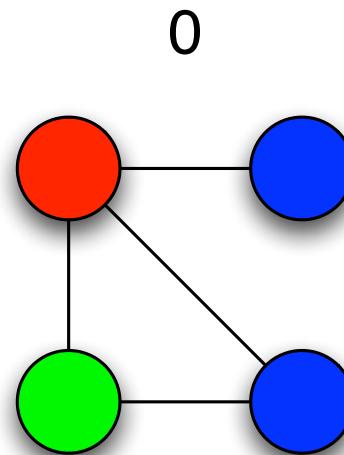
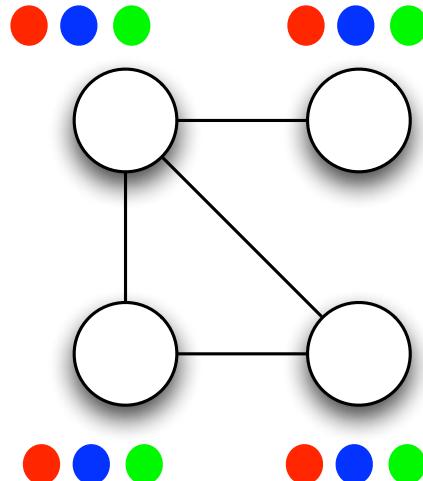


No!



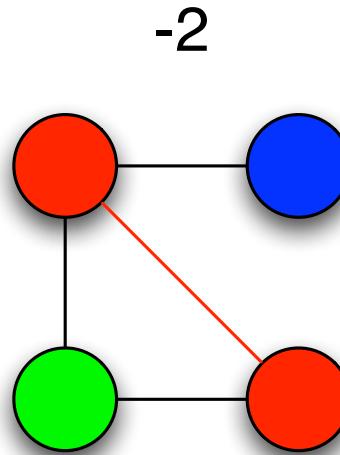
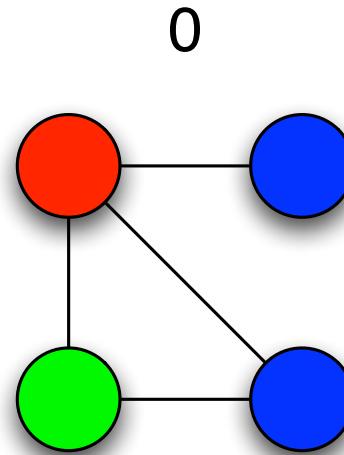
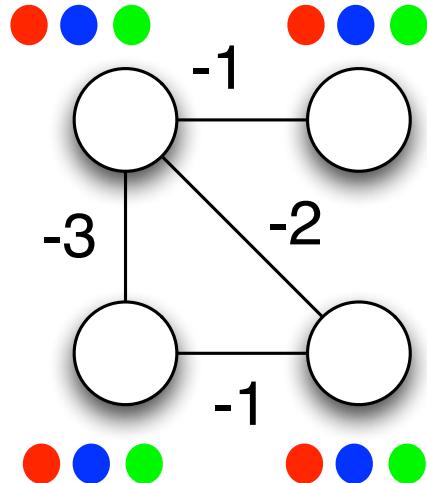
# Graph coloring - Max-CSP

- Minimize the number of conflicts



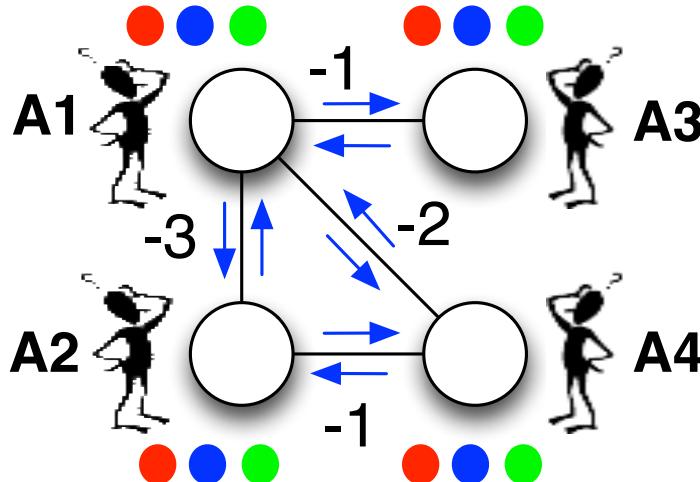
# Graph coloring - COP

- Different **weights** to violated **constraints**
- Preferences for different **colors**



## Graph coloring - DCOP

- Each node:
  - controlled by one agent
- Each agent:
  - Preferences for different colors
  - Communicates with its direct neighbours in the graph



- A1 and A2 exchange preferences and conflicts
- A3 and A4 do not communicate

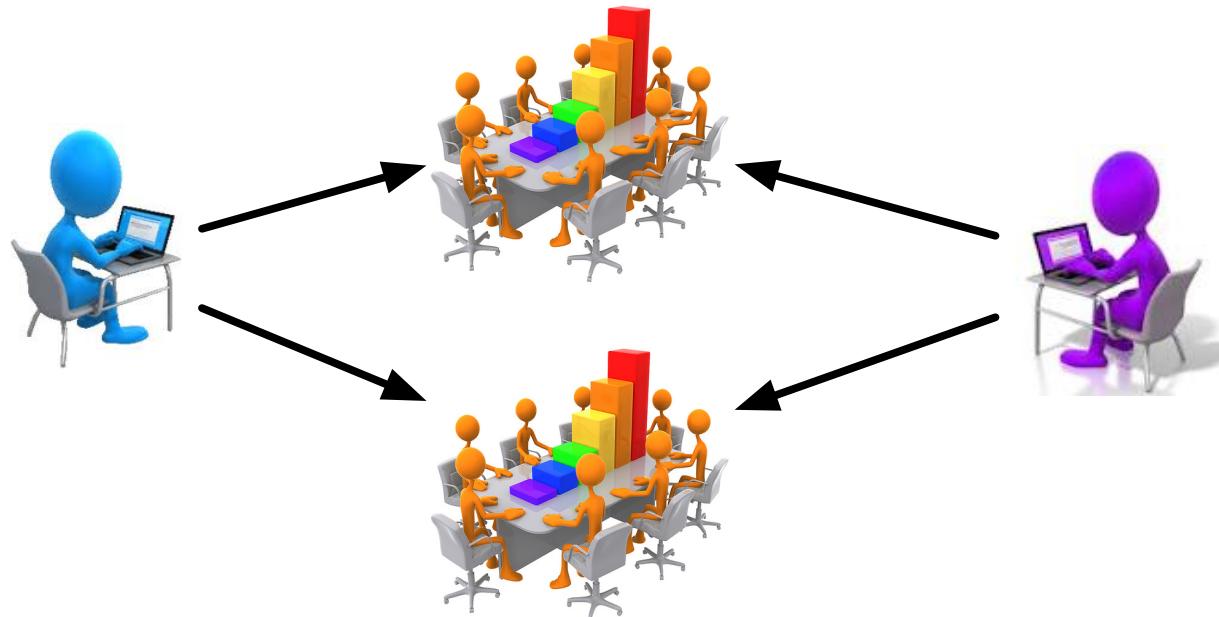
# Meeting Scheduling

## Motivation:

- Privacy
- Robustness
- Scalability

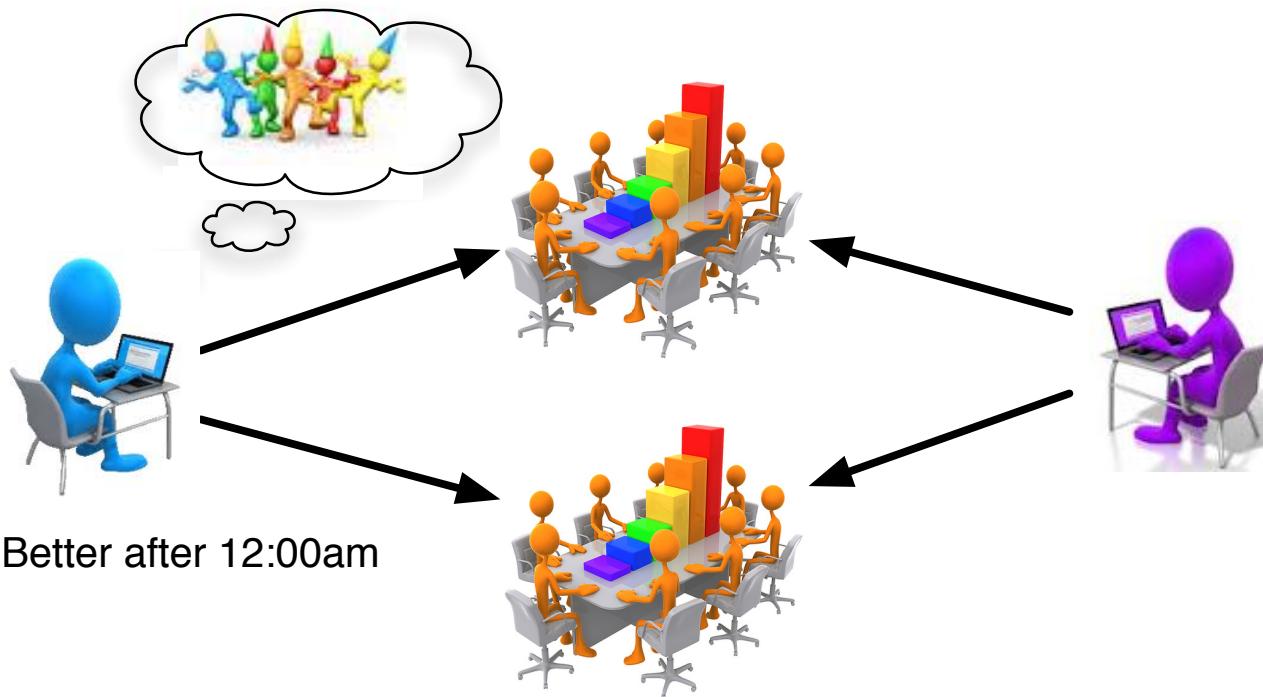
# Meeting Scheduling

*In large organizations many people, possibly working in different departments, are involved in a number of work meetings.*



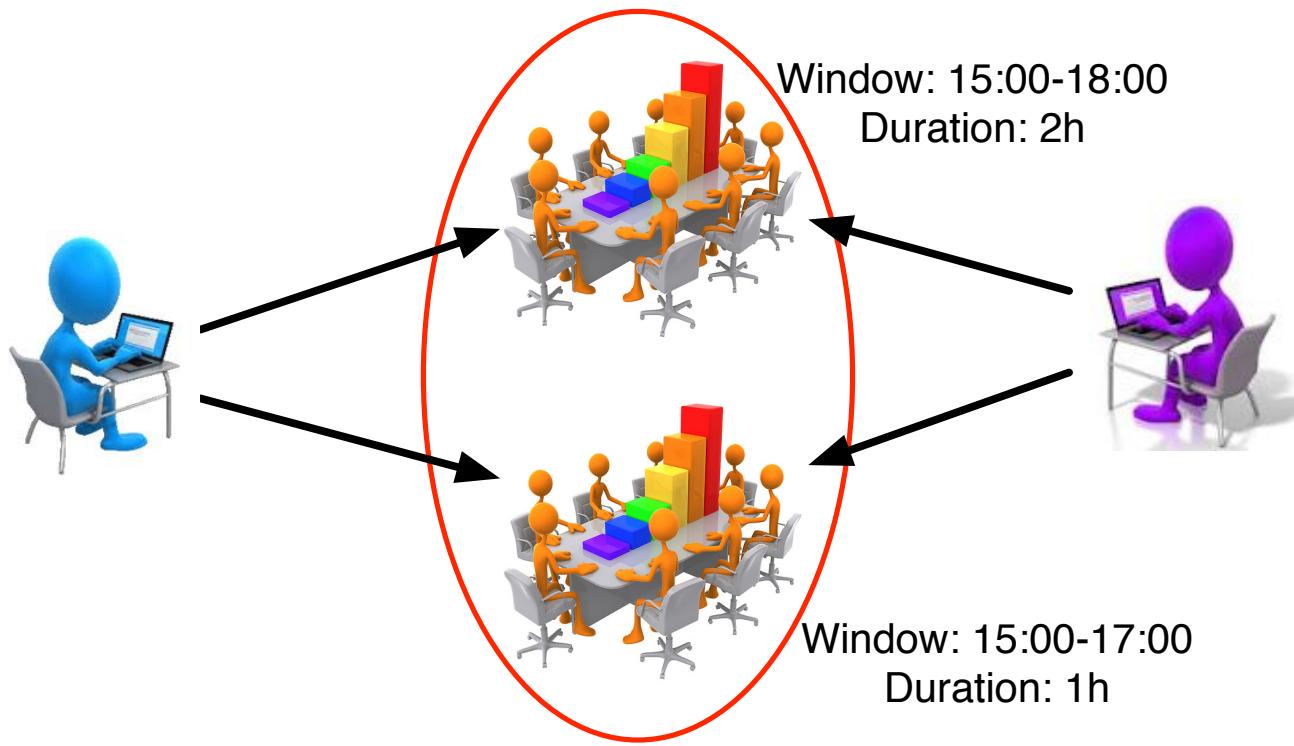
# Meeting Scheduling

*People might have various private preferences on meeting start times*



# Meeting Scheduling

*Two meetings that share a participant cannot overlap*



# DCOP formalization for the meeting scheduling problem

- A set of **agents** representing **participants**
- A set of **variables** representing **meeting starting times** according to a participant.
- Hard Constraints:
  - Starting meeting times across different **agents** are **equal**
  - Meetings for the same agent are **non-overlapping**.
- Soft Constraints:
  - Represent **agent preferences** on meeting starting times.

*Objective:* find a valid schedule for the meeting while maximizing the sum of individuals' preferences.

# Target Tracking

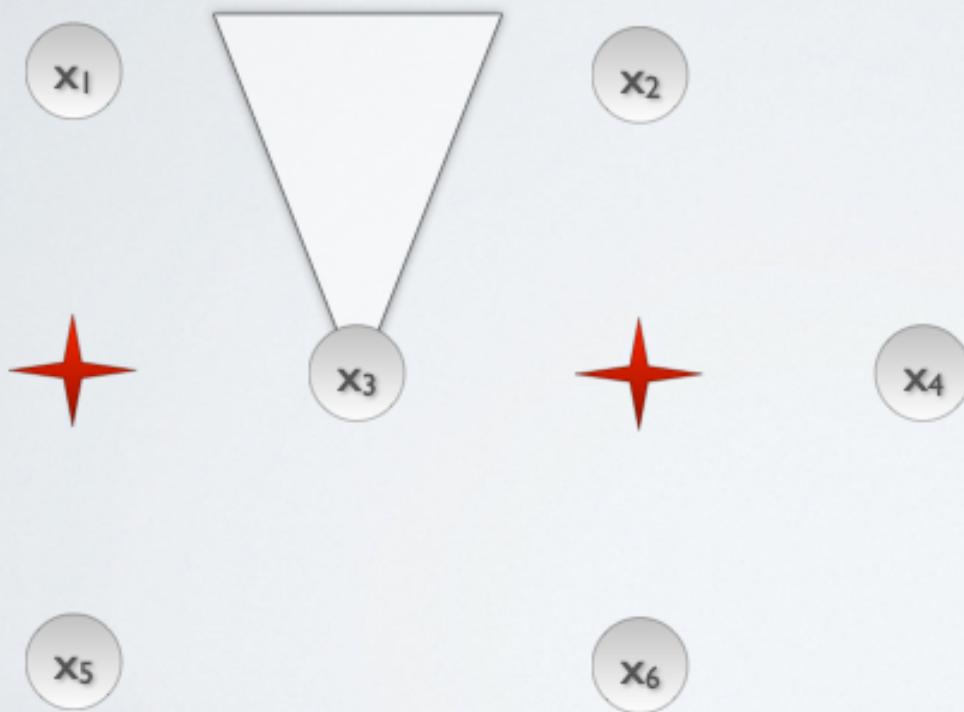
Motivation:

- Privacy
- Robustness
- Scalability

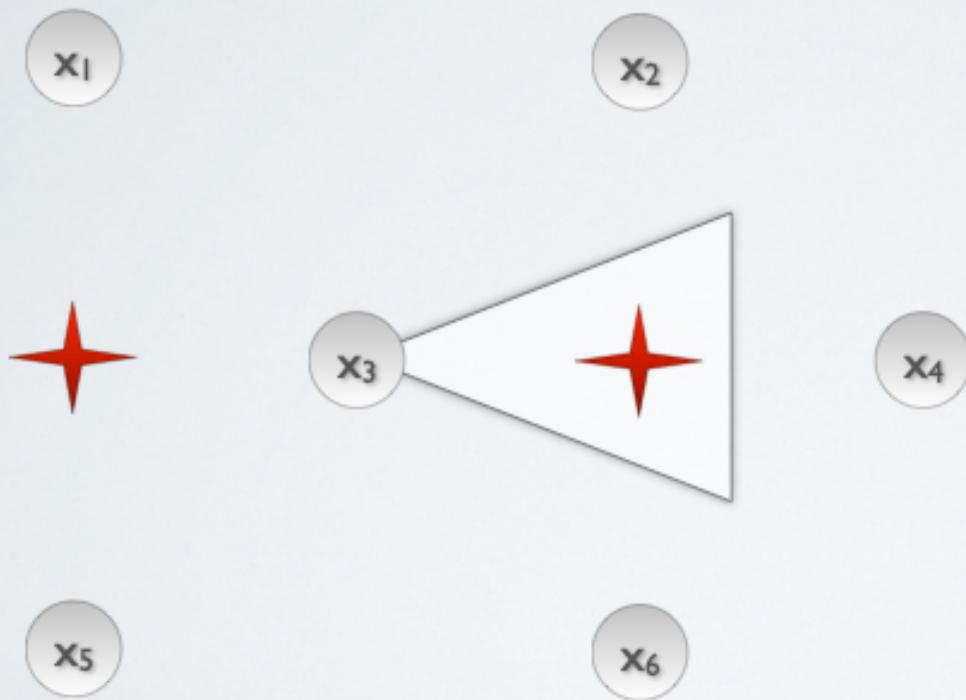
# MOTIVATING DOMAIN: SENSOR NETWORK



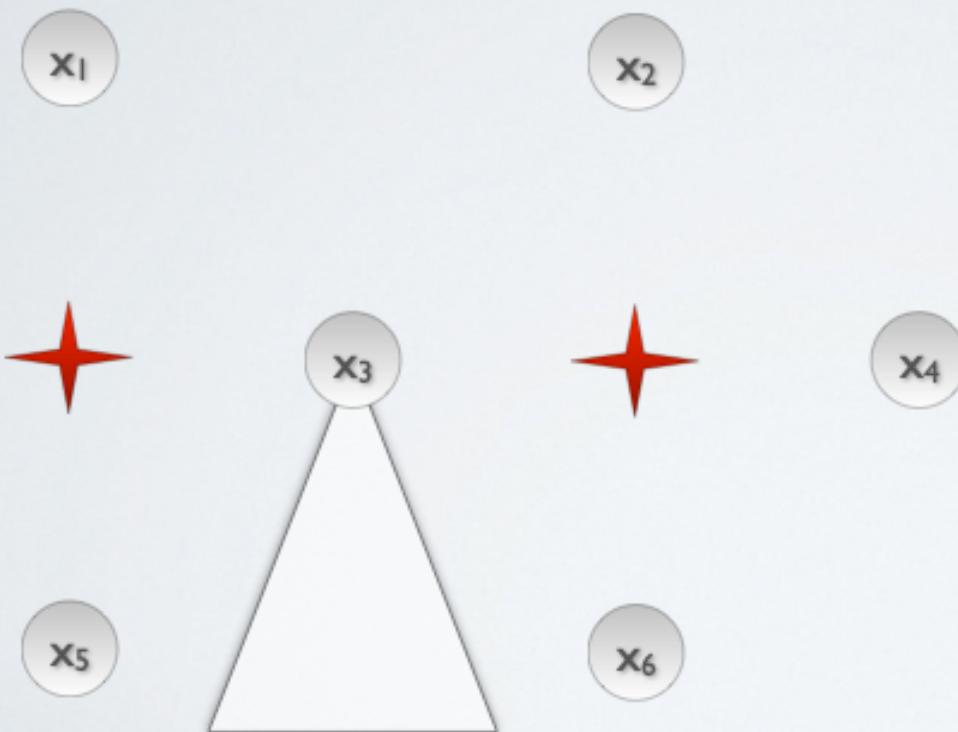
# MOTIVATING DOMAIN: SENSOR NETWORK



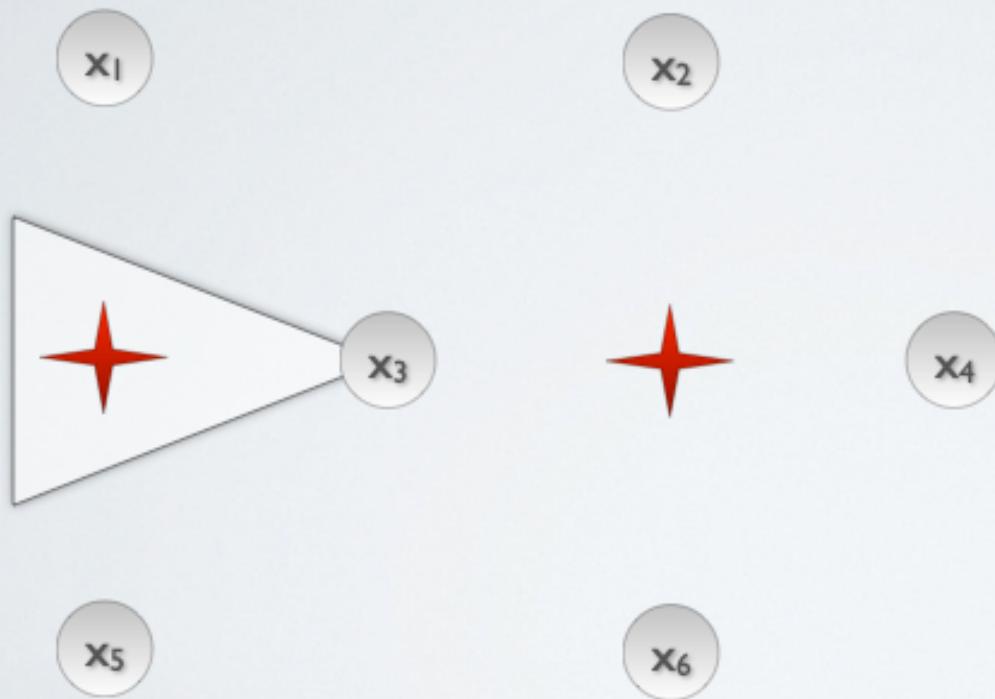
# MOTIVATING DOMAIN: SENSOR NETWORK



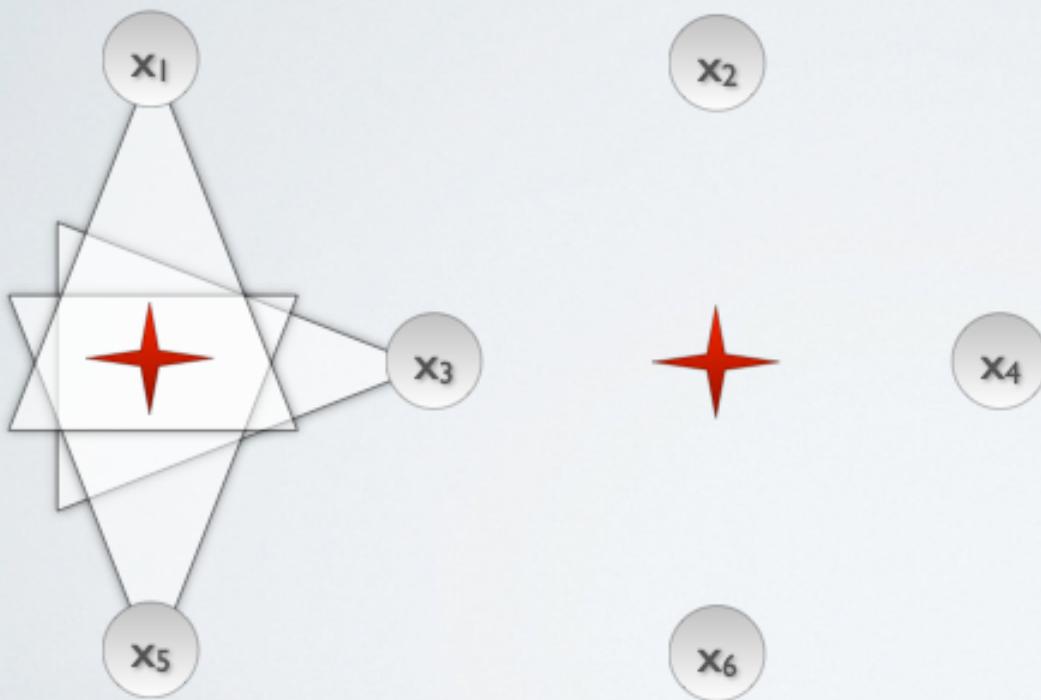
# MOTIVATING DOMAIN: SENSOR NETWORK



# MOTIVATING DOMAIN: SENSOR NETWORK



# MOTIVATING DOMAIN: SENSOR NETWORK

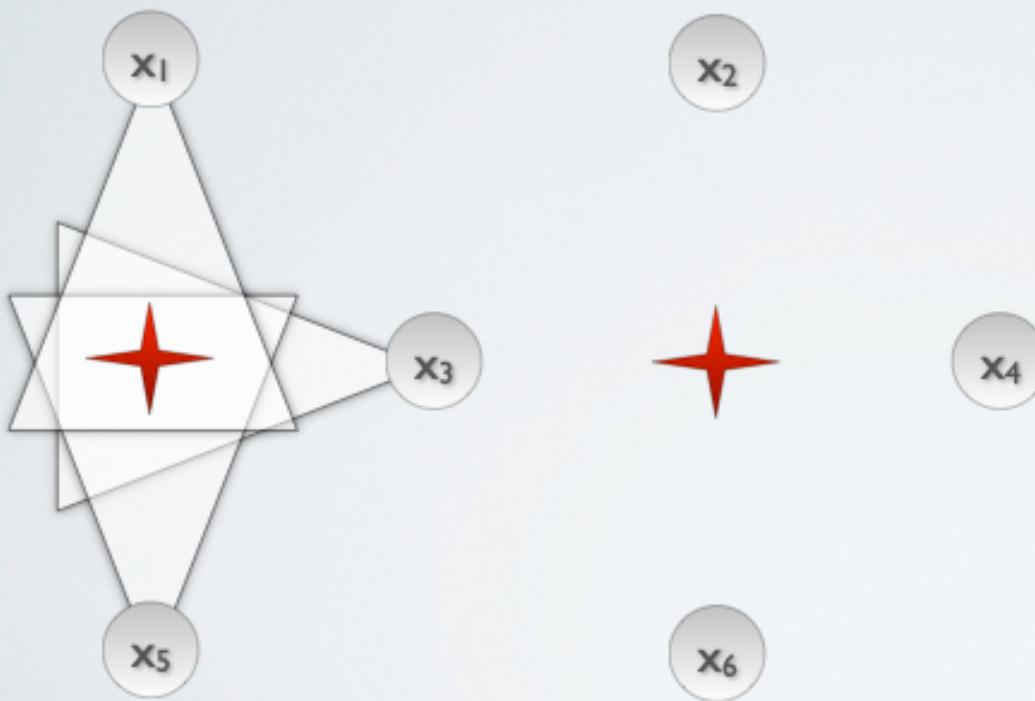


<b><math>x_1</math></b>	<b><math>x_3</math></b>	<b><math>x_5</math></b>	<b>Sat?</b>
N	N	N	X
N	N	E	X
...			X
S	W	N	✓
...			X
W	W	W	X

Model the problem as a  
CSP

# DCOP

## DISTRIBUTED CONSTRAINT OPTIMIZATION

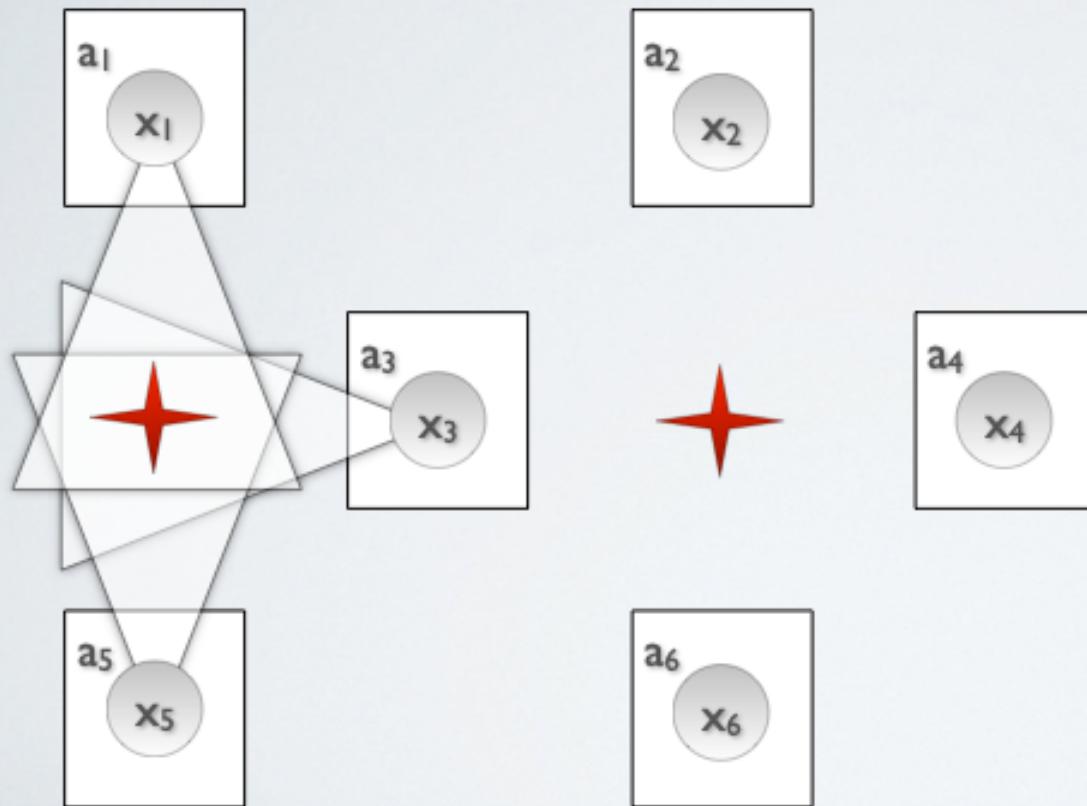


Imagine that each sensor is an autonomous agent.

*How should this problem be modeled and solved in a decentralized manner?*

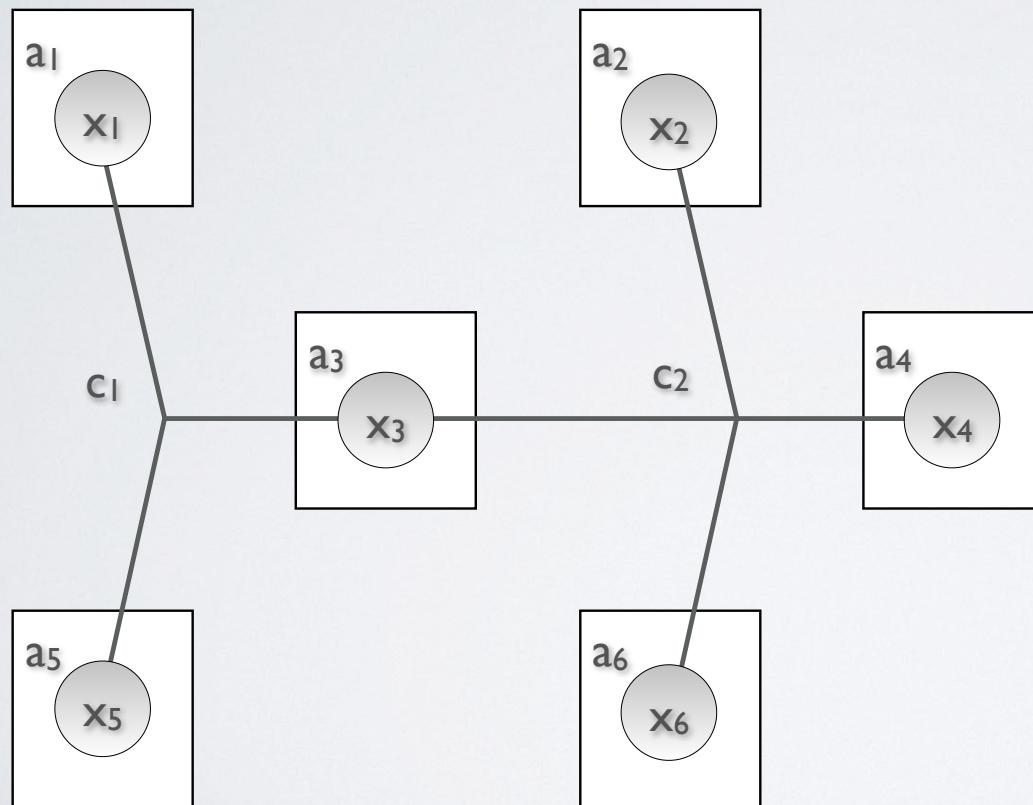
# DCOP

## DISTRIBUTED CONSTRAINT OPTIMIZATION



# DCOP

## DISTRIBUTED CONSTRAINT OPTIMIZATION



# Complete Algorithms

-  Always find an **optimal solution**
-  Exhibit an **exponentially** increasing coordination **overhead**
-  Very **limited scalability** on general problems.

# Complete Algorithms

- **Completely decentralised**
  - Search-based.
    - Synchronous: SyncBB, AND/OR search
    - Asynchronous: ADOPT, NCBB and AFB.
  - Dynamic programming.
- **Partially decentralised**
  - OptAPO

Next, we focus on **completely decentralised algorithms**

# Decentralised Complete Algorithms

## Search-based

- Uses distributed search
- Exchange individual values
- Small messages but  
... exponentially many

Representative: ADOPT [Modi et al., 2005]

## Dynamic programming

- Uses distributed inference
- Exchange constraints
- Few messages but  
... exponentially large

Representative: DPOP [Petcu and Faltings, 2005]

We will discuss DPOP in detail

# DPOP

**DPOP** (Dynamic Programming Optimization Protocol) [Petcu and Faltings, 2005]:

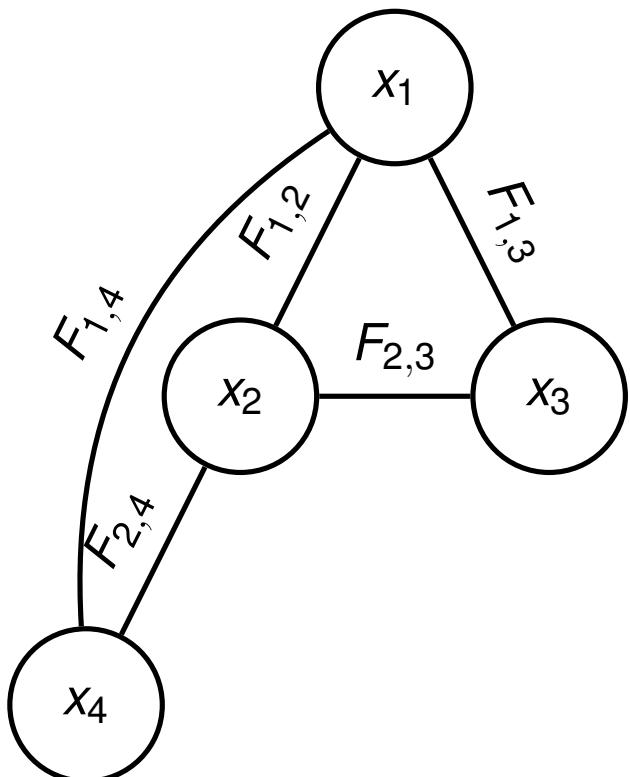
- Based on the dynamic programming paradigm.
- Special case of Bucket Tree Elimination Algorithm (BTE)  
[Dechter, 2003].

# DFS arrangement

The same happens  
for DPOP!

- Before executing ADOPT, **agents** must be arranged in a **depth first search (DFS) tree**.
- **DFS trees** have been frequently used in optimization because they have two interesting properties:
  - Agents in **different branches** of the tree **do not share any constraints**;
  - **Every constraint network admits a DFS tree.**

# DPOP by example

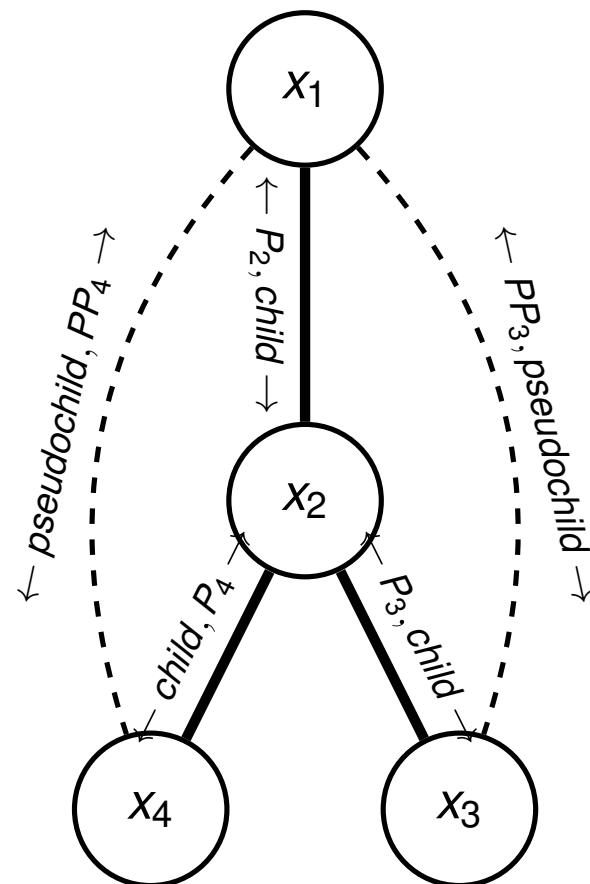


$F_{i,j}$	$x_i$	$x_j$
2	0	0
0	0	1
0	1	0
1	1	1

$\Rightarrow$

## DFS arrangement

# Objective: find assignment with maximal value

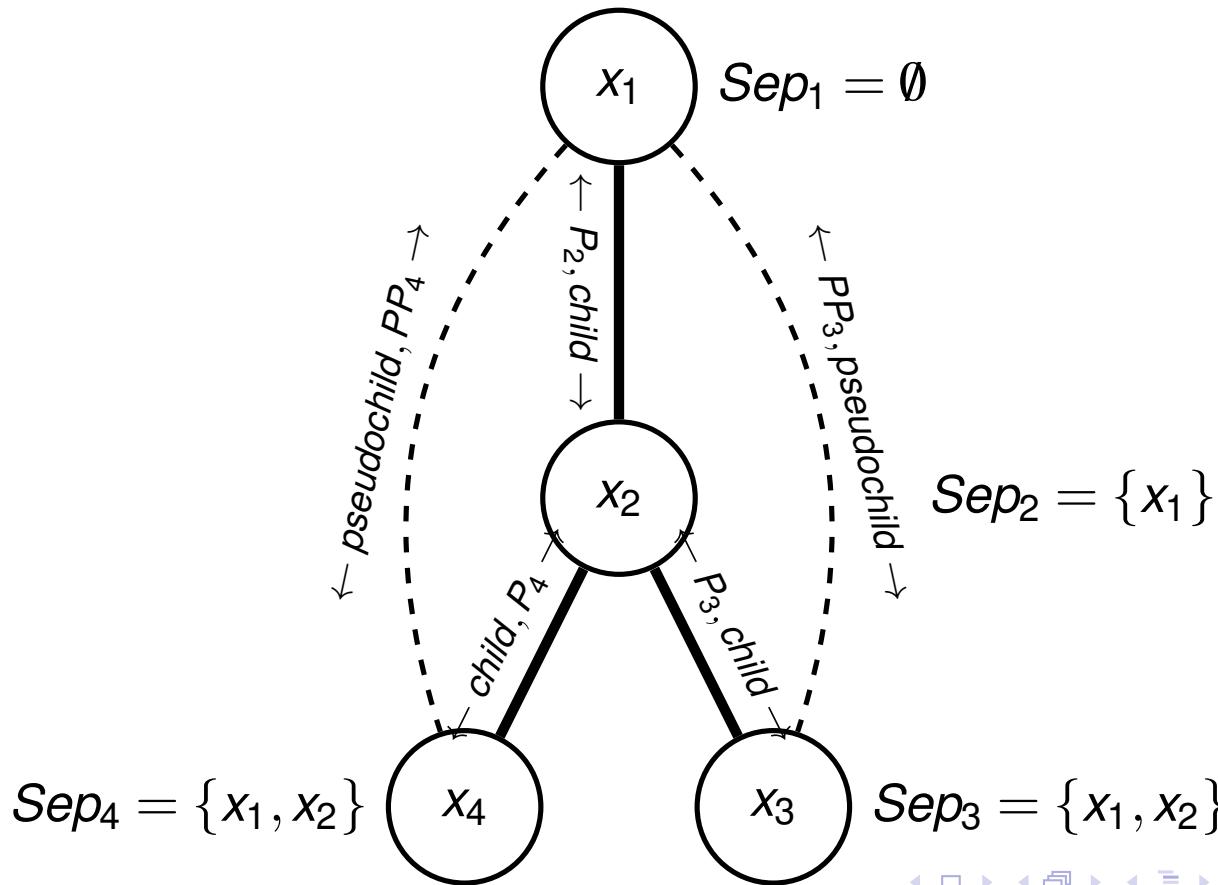


# DPOP phases

Given a DFS tree structure, DPOP runs in **two phases**:

- ***Util* propagation:** agents exchange *util* messages up the tree.
  - Aim: aggregate all info so that root agent can choose optimal value
- ***Value* propagation:** agents exchange **value messages** down the tree.
  - Aim: propagate info so that all agents can make their choice given choices of ancestors

$Sep_i$ : set of agents preceding  $A_i$  in the pseudo-tree order that are connected with  $A_i$  or with a descendant of  $A_i$ .



## *Util* message

The *Util* message  $U_{i \rightarrow j}$  that agent  $A_i$  sends to its parent  $A_j$  can be computed as:

$$U_{i \rightarrow j}(Sep_i) = \max_{x_i} \left( \bigotimes_{A_k \in C_i} U_{k \rightarrow i} \bigotimes \bigotimes_{A_p \in P_i \cup PP_i} F_{i,p} \right)$$

Size exponential  
in  $Sep_i$

All incoming messages  
from children

Shared constraints with  
parents/pseudoparents

The  $\otimes$  operator is a join operator that sums up functions with different but overlapping scores consistently.

# Join operator

$F_{2,4}$	$x_2$	$x_4$
2	0	0
0	0	1
0	1	0
1	1	1

$F_{1,4}$	$x_1$	$x_4$
2	0	0
0	0	1
0	1	0
1	1	1

$F_{1,4} \otimes F_{2,4}$		
	$x_1$	$x_2$
4	0	0
0	0	1
2	0	0
1	0	1
2	1	0
2	1	0
0	1	1
2	1	1

Add

→

Project  
out  $x_4$

$\max_{\{x_4\}}(F_{1,4} \otimes F_{2,4})$	$x_1$	$x_2$	$x_4$
0	0	0	0
0	0	0	1
0	1	0	0
0	1	1	1
1	0	0	0
1	0	0	1
1	1	0	0
1	1	1	1
1	0	0	0
1	0	0	1
1	1	0	0
1	1	1	1

$\max(4,0)$

$\max(2,1)$

$\max(2,2)$

$\max(0,2)$

$\max_{\{x_4\}}(F_{1,4} \otimes F_{2,4})$

$x_1$     $x_2$     $x_4$

0      0      0

0      0      1

0      1      0

0      1      1

1      0      0

1      0      1

1      1      0

1      1      1

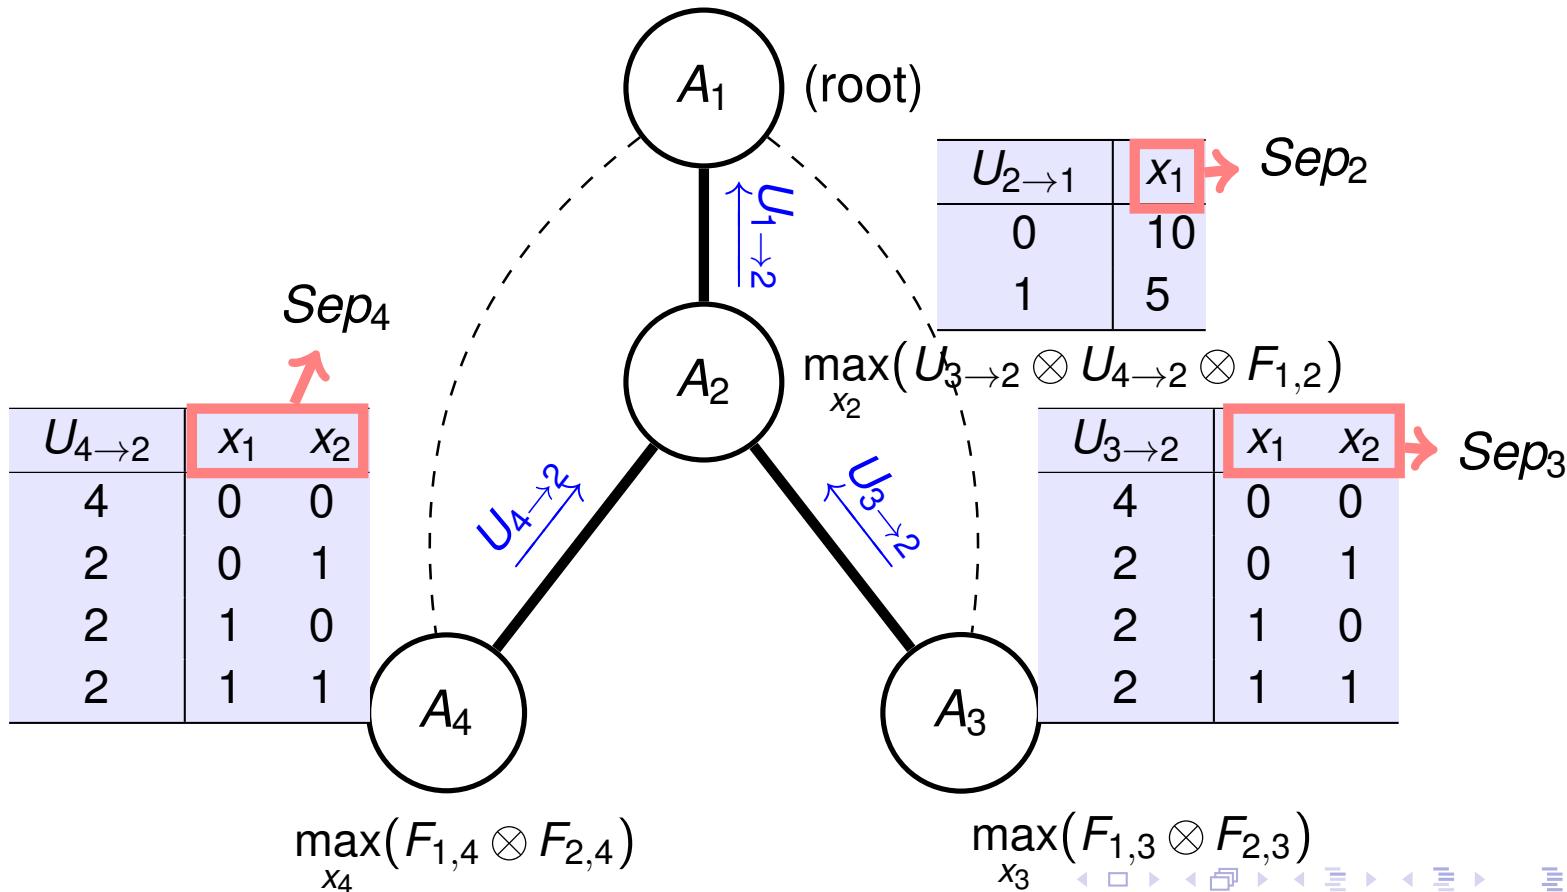
1      0      0

1      0      1

1      1      0

1      1      1

Complexity exponential to the largest  $Sep_i$ .  
 Largest  $Sep_i$  = induced width of the DFS tree ordering used.



## Value message

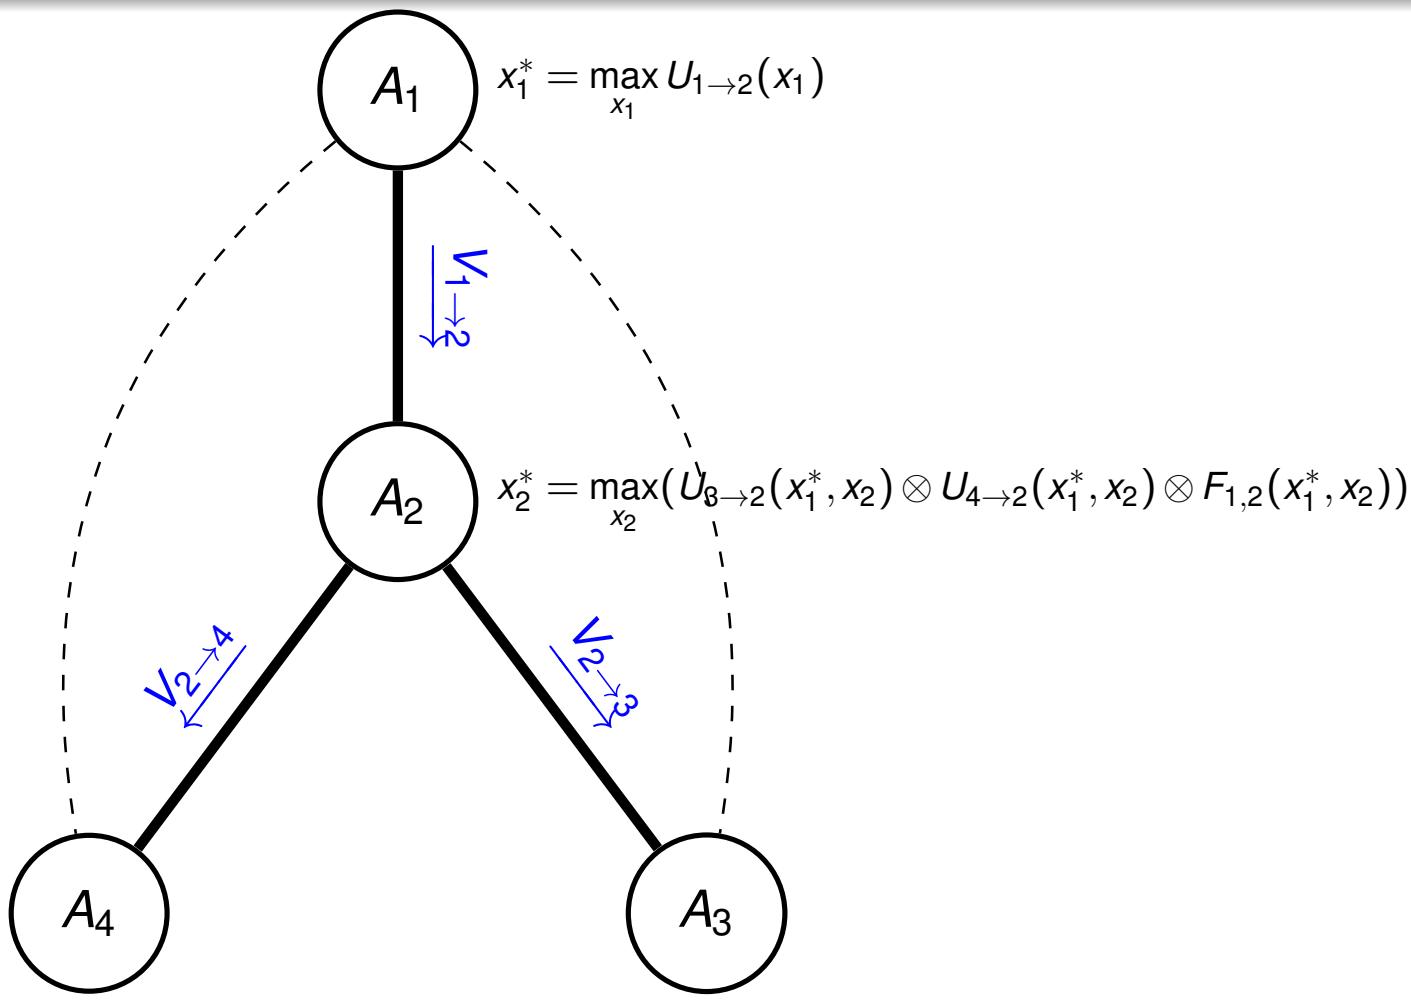
Keeping fixed the value of parent/pseudoparents, finds the value that maximizes the computed cost function in the util phase:

$$x_i^* = \arg \max_{x_i} \left( \sum_{A_j \in C_i} U_{j \rightarrow i}(x_i, x_p^*) + \sum_{A_j \in P_i \cup PP_i} F_{i,j}(x_i, x_j^*) \right)$$

where  $x_p^* = \bigcup_{A_j \in P_i \cup PP_i} \{x_j^*\}$  is the set of optimal values for  $A_i$ 's parent and pseudoparents received from  $A_i$ 's parent.

Propagates this value through children down the tree:

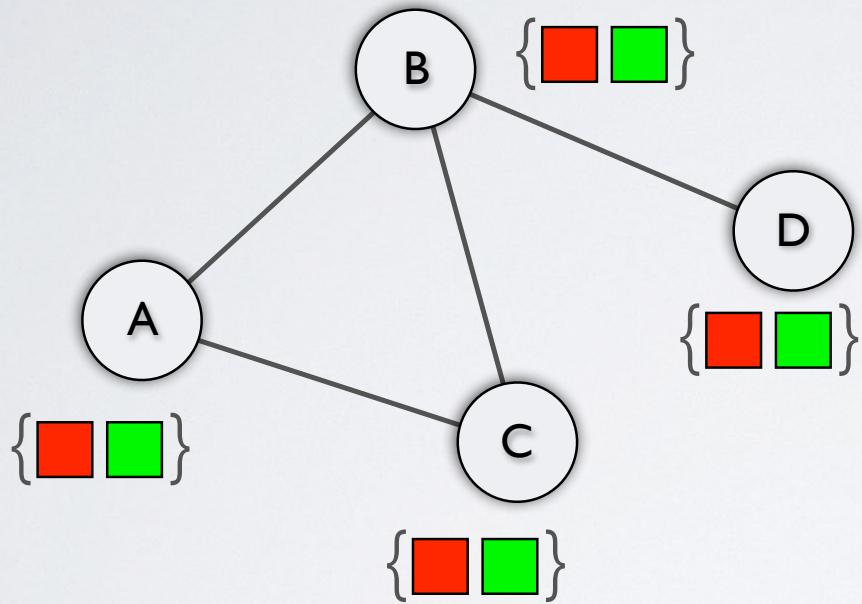
$$V_{i \rightarrow j} = \{x_i = x_i^*\} \cup \bigcup_{x_s \in Sep_i \cap Sep_j} \{x_s = x_s^*\}$$



# DPOP extensions

- **MB-DPOP** [Petcu and Faltings, 2007] trades-off message size against the number of messages.
- **A-DPOP** trades-off message size against solution quality [Petcu and Faltings, 2005(2)].

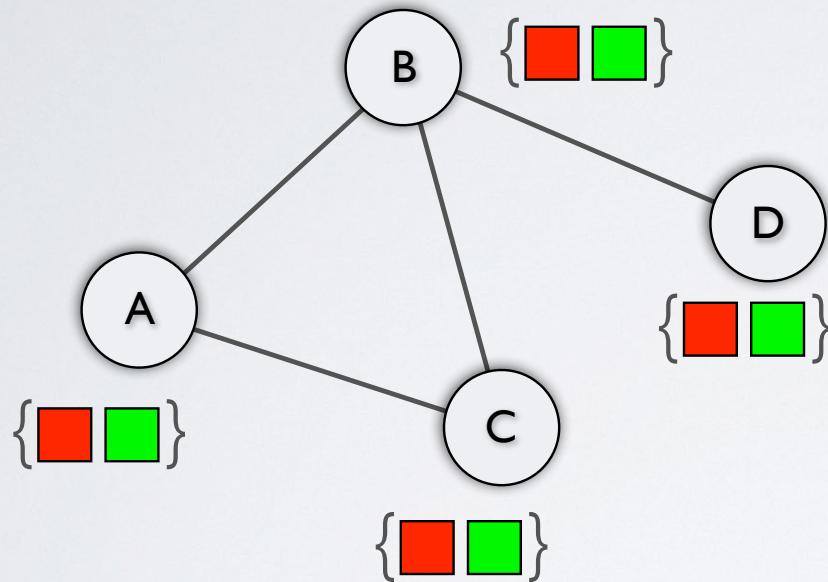
# Example run of DPOP



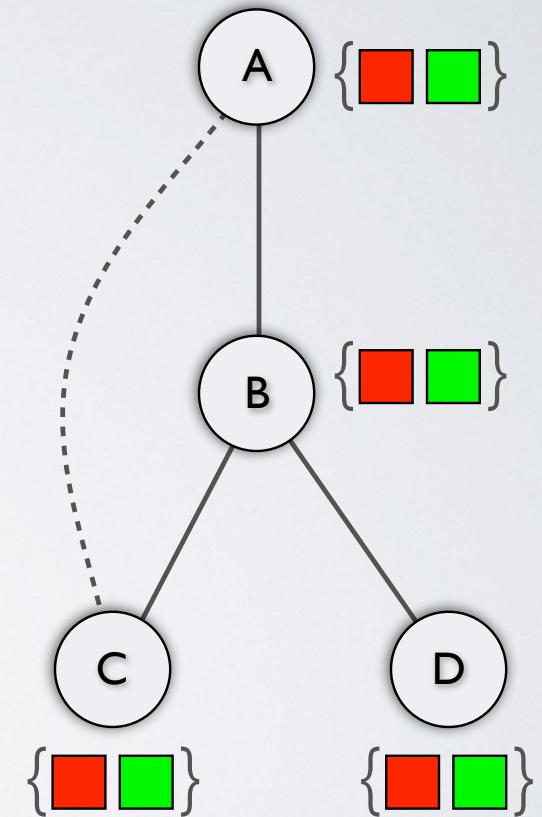
$x_i$	$x_j$	Cost (A,B)	Cost (A,C)	Cost (B,C)	Cost (B,D)
Red	Red	5	5	5	3
Red	Green	8	10	4	8
Green	Red	20	20	3	10
Green	Green	3	3	3	3

Katsutoshi Hirayama, Makoto Yokoo: Distributed Partial Constraint Satisfaction Problem. CP 1997: 222-236

# PSEUDO-TREE



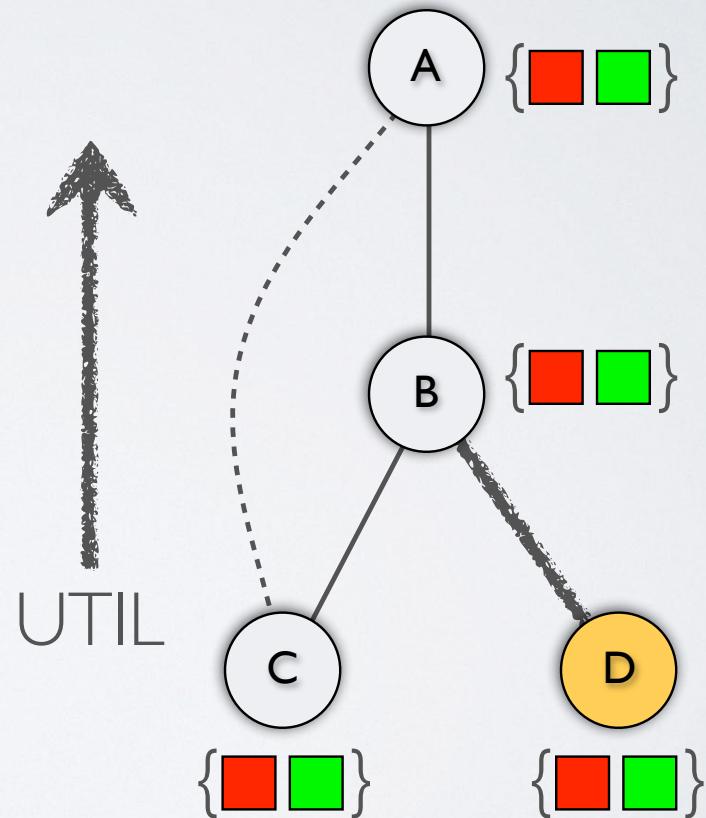
Definition: A *spanning tree* of the constraint graph such that no two nodes in sibling subtrees share a constraint in the constraint graph



# DPOP

B	D	(B,D)
r	r	3
r	g	8
g	r	10
g	g	3

Pseudo-tree Ordering



# DPOP

B	D	(B,D)
r	r	3
r	g	8
g	r	10
g	g	3

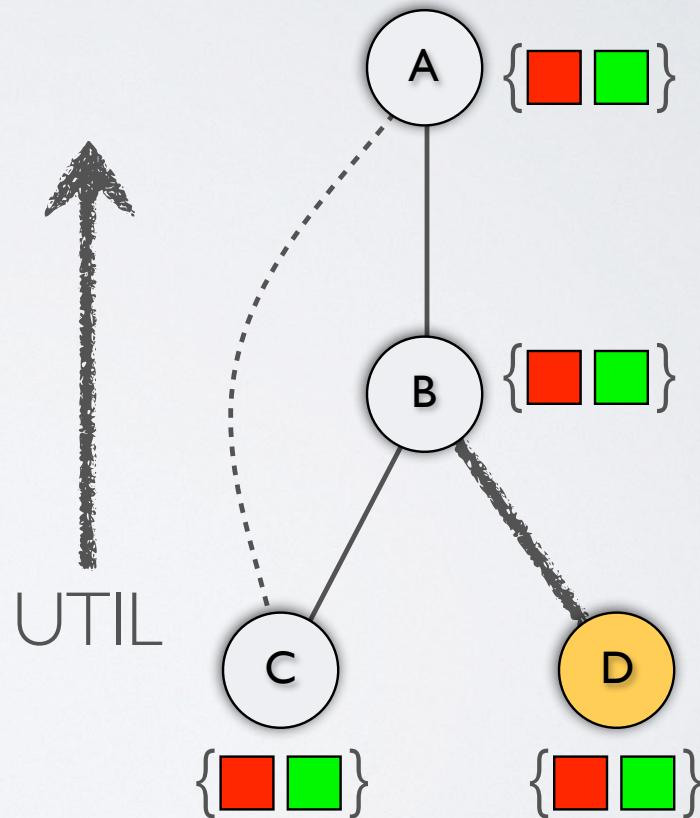
$$\min\{3, 8\} = 3$$

$$\min\{10, 3\} = 3$$

MSG to B

B	cost
r	3
g	3

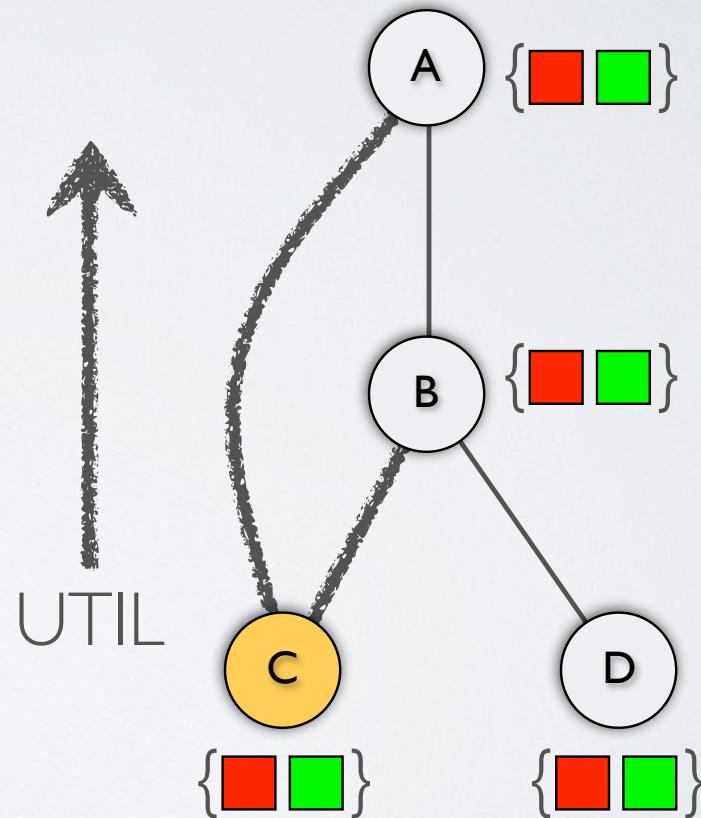
Pseudo-tree Ordering



# DPOP

A	B	C	(B,C)	(A,C)	
r	r	r	5	5	10
r	r	g	4	8	12
r	g	r	3	5	8
r	g	g	3	8	11
g	r	r	5	10	15
g	r	g	4	3	7
g	g	r	3	10	13
g	g	g	3	3	6

Pseudo-tree Ordering



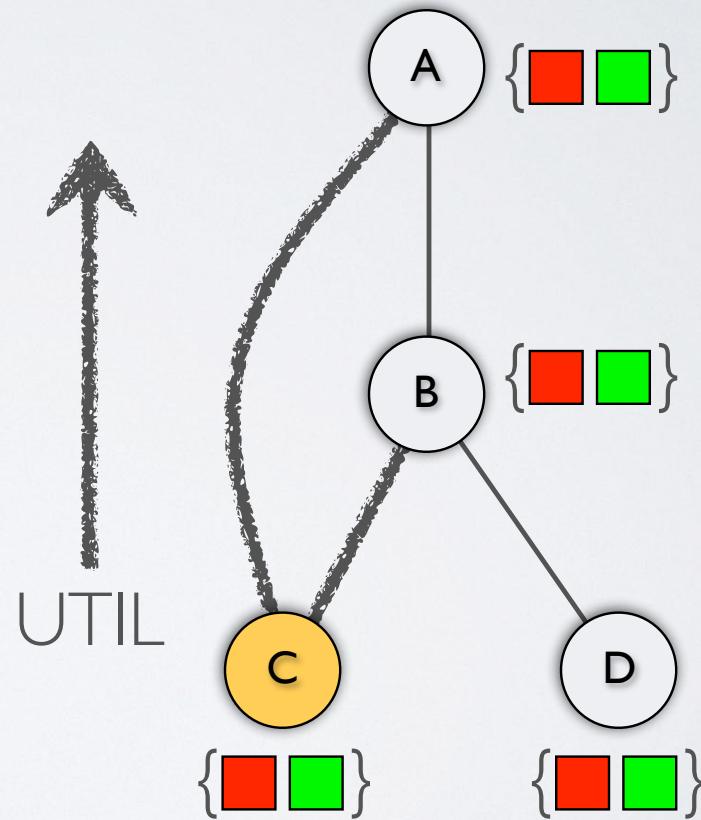
# DPOP

A	B	C	(B,C)	(A,C)	
r	r	r	5	5	10
r	r	g	4	8	12
r	g	r	3	5	8
r	g	g	3	8	11
g	r	r	5	10	15
g	r	g	4	3	7
g	g	r	3	10	13
g	g	g	3	3	6

MSG to B

A	B	
r	r	10
r	g	8
g	r	7
g	g	6

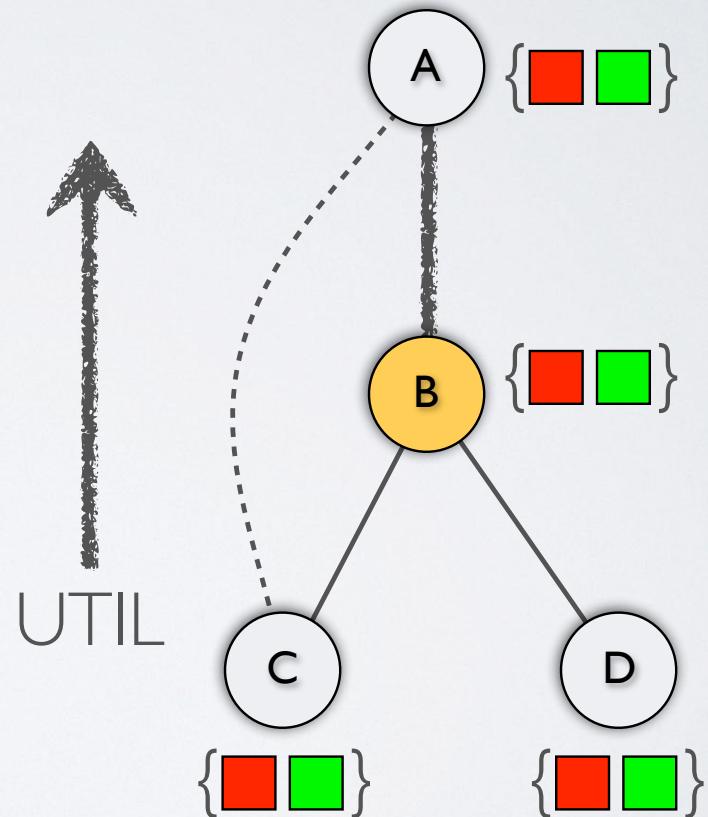
Pseudo-tree Ordering



# DPOP

A	B	(A,B)	Util C	Util D	
r	r	5	10	3	18
r	g	8	8	3	19
g	r	20	7	3	30
g	g	3	6	3	12

Pseudo-tree Ordering



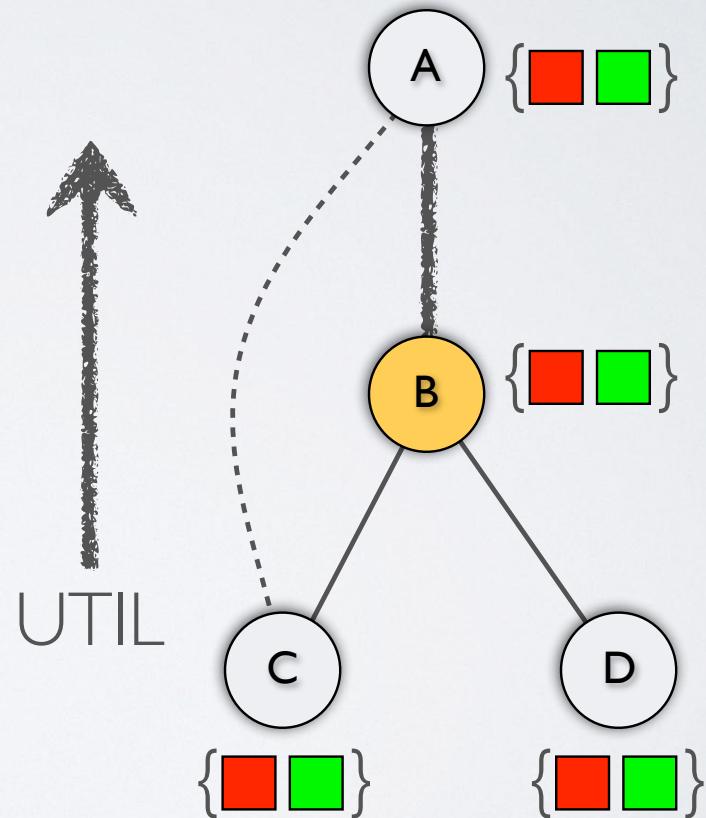
# DPOP

A	B	(A,B)	Util C	Util D	
r	r	5	10	3	18
r	g	8	8	3	19
g	r	20	7	3	30
g	g	3	6	3	12

MSG to A

A	cost
r	18
g	12

Pseudo-tree Ordering

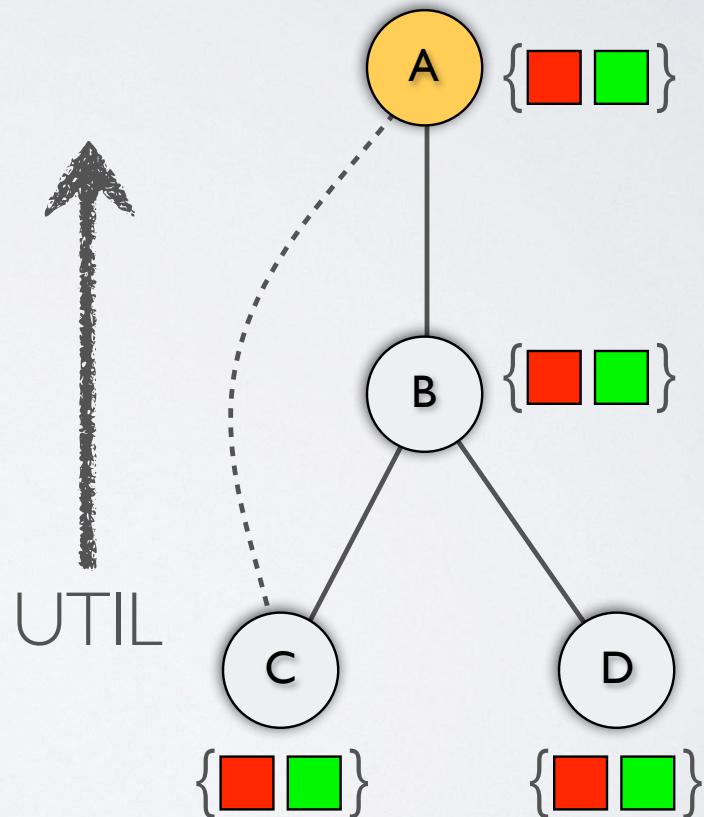


# DPOP

A	cost
r	18
g	12

optimal cost = 12

Pseudo-tree Ordering

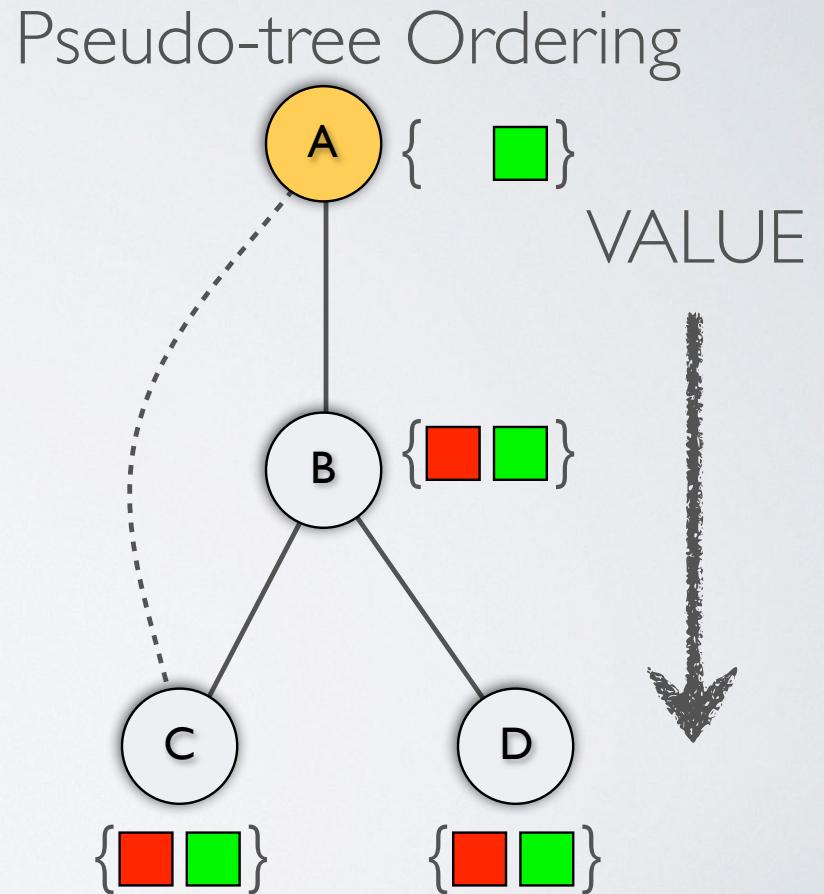


# DPOP

A	cost
r	18
g	12

- Select value for A = 'g'
- Send MSG A = 'g' to agents B and C

Alternative version in which value messages  $x_i = x_i^*$  are sent to children and pseudo-children



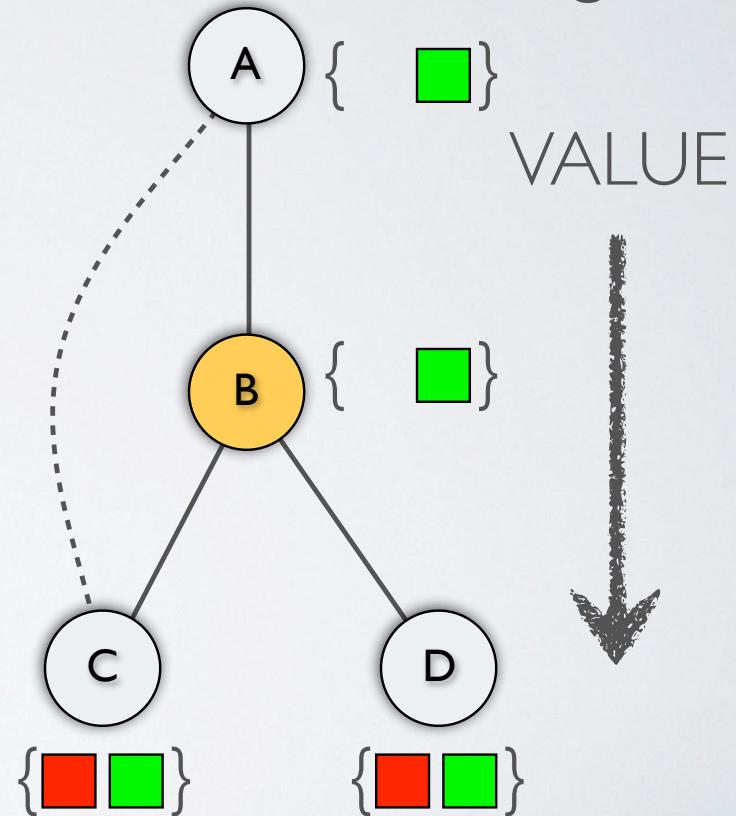
# DPOP



A	B	(A,B)	Util C	Util D	
r	r	5	10	3	18
r	g	8	8	3	19
g	r	20	7	3	30
g	g	3	6	3	12

- Select value for B = 'g'
- Send MSG B = 'g' to agents C and D

Pseudo-tree Ordering

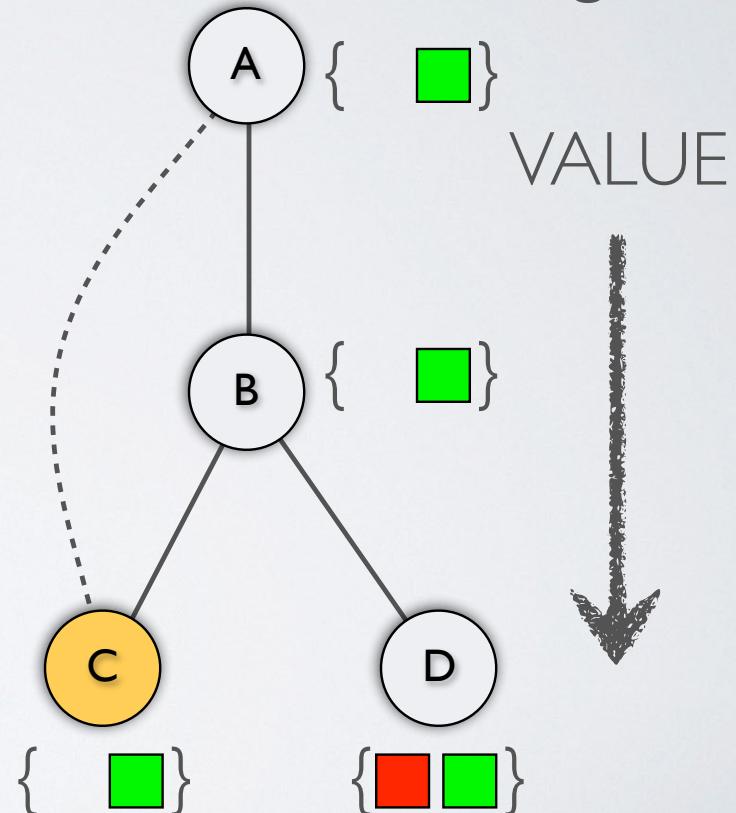


# DPOP



A	B	C	(B,C)	(A,C)	
r	r	r	5	5	10
r	r	g	4	8	12
r	g	r	3	5	8
r	g	g	3	8	11
g	r	r	5	10	15
g	r	g	4	3	7
g	g	r	3	10	13
g	g	g	3	3	6

Pseudo-tree Ordering



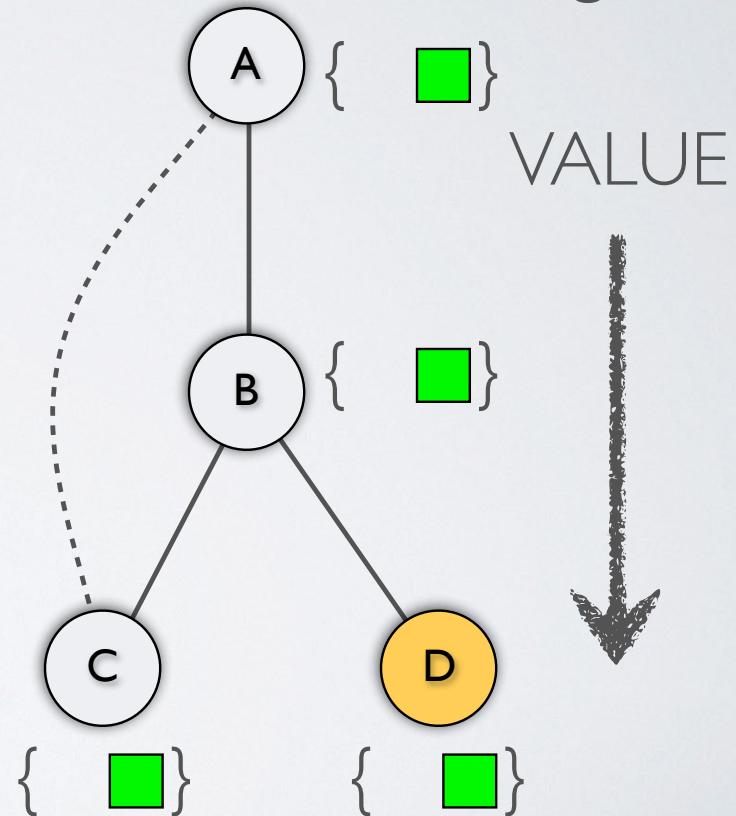
- Select value for C = 'g'

# DPOP



B	D	(B,D)
r	r	3
r	g	8
g	r	10
g	g	3

Pseudo-tree Ordering



- Select value for D = 'g'

# Why Approximate Algorithms

“Very often *optimality* in practical applications is *not achievable*”

## Approximate algorithms

- Sacrifice optimality in favor of computational and communication efficiency
- Well-suited for large scale distributed applications:
  - sensor networks
  - mobile robots

# Centralized Local Greedy approaches

- Start from a random assignment for all the variables
- Do local moves if the new assignment improves the value (local gain)
- Local: changing the value of a small set of variables (in most case just one)
- The search stops when there is no local move that provides a positive gain, i.e., when the process reaches a local maximum.

# Distributed Local Greedy approaches

When operating in a **decentralized context**:

- **Problem:** Out-of-date local knowledge
  - Assumption that other agents do not change their values
  - A greedy local move might be harmful/useless
- **Solution:**
  - Stochasticity on the decision to perform a move (**DSA**)
  - Coordination among neighbours on who is the agent that should move (**MGM**)

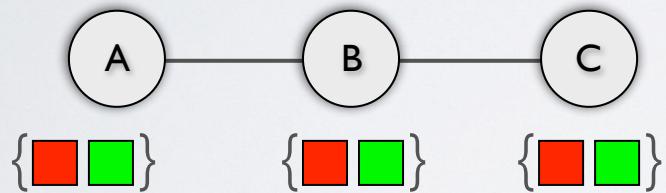
# Distributed Stochastic Algorithm (DSA)

***Activation probability to mitigate issues with parallel executions***

[S. Fitzpatrick and L. Meetrens, 2003]

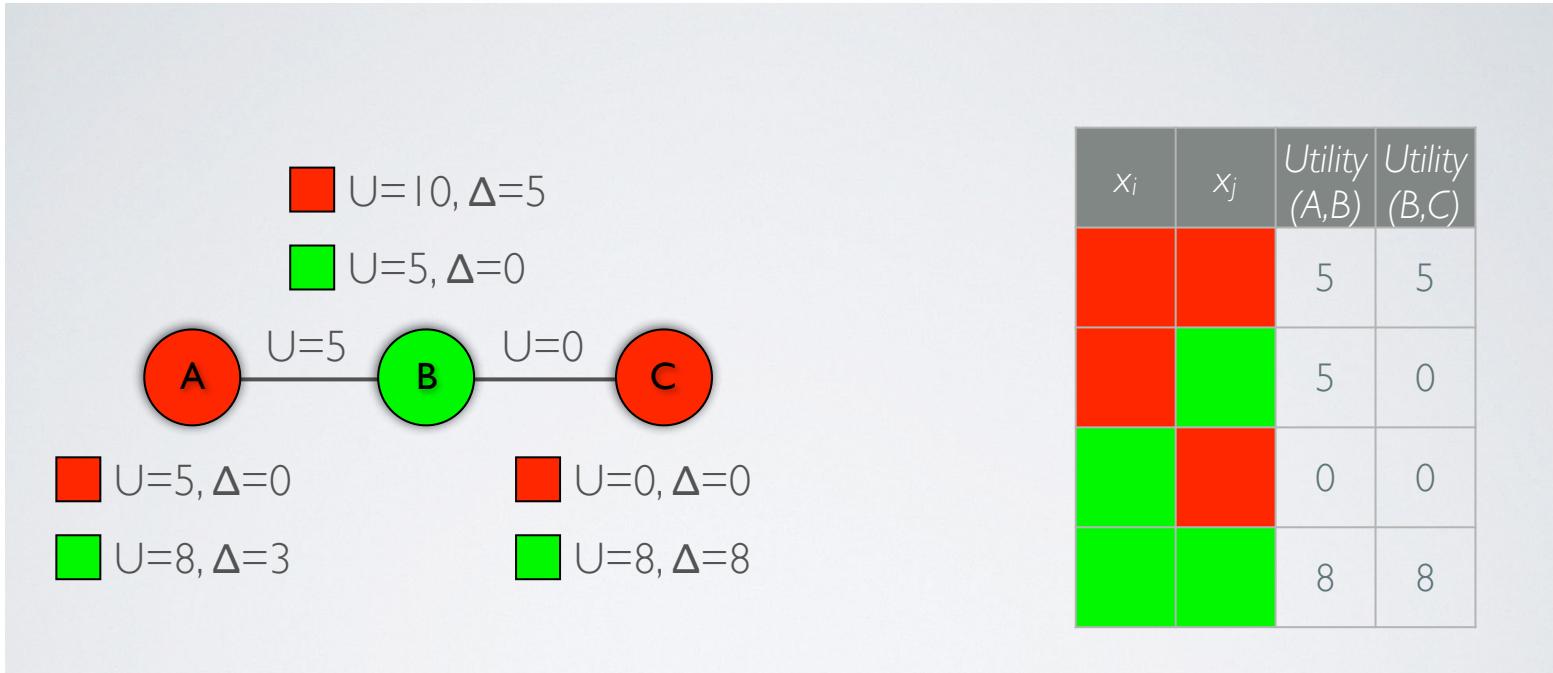
- Initialize agents with a random assignment and communicate values to neighbors
- Each agent:
  - Generates a random number and executes only if it is less than the activation probability
  - When executing choose a value for the variable such that the local gain is maximized
  - Communicate and receive possible variables change to/from neighbors

# Example run of DSA

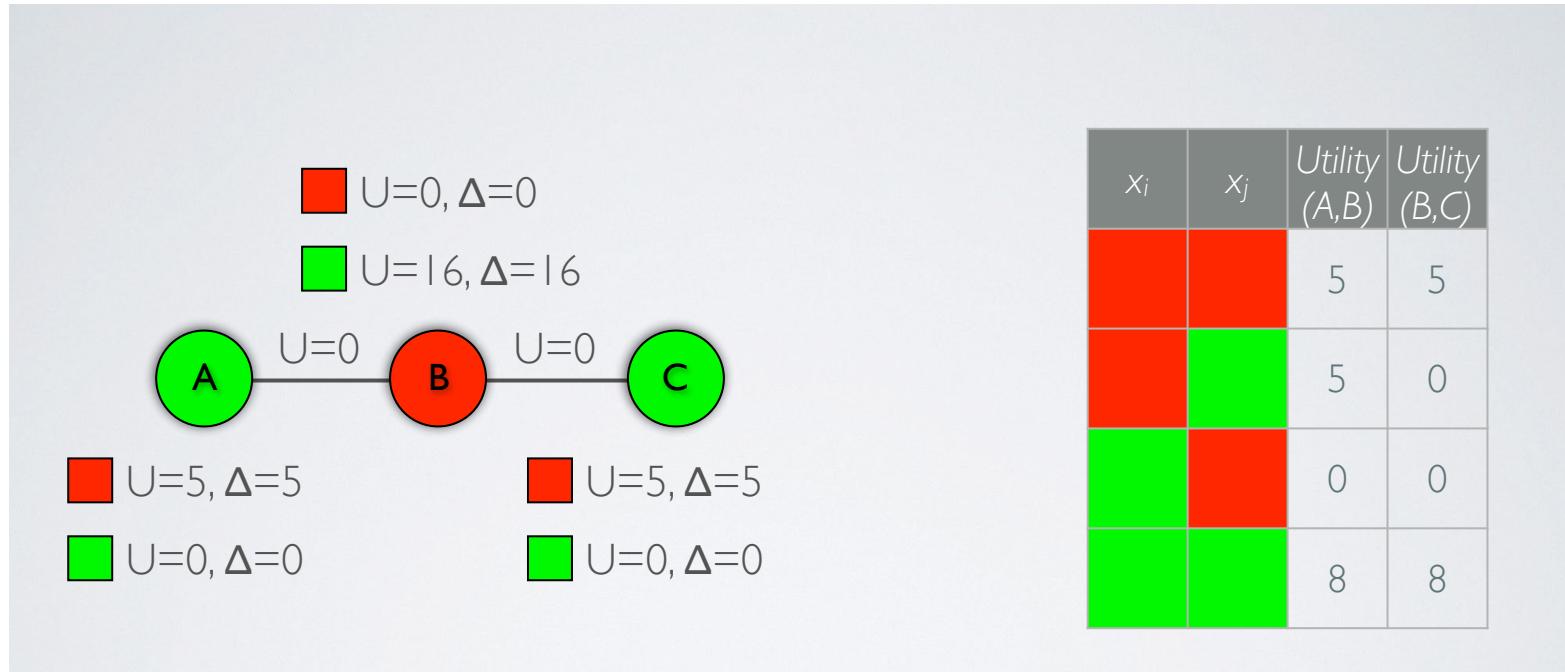


$x_i$	$x_j$	Utility (A,B)	Utility (B,C)
Red	Red	5	5
Red	Green	5	0
Green	Red	0	0
Green	Green	8	8

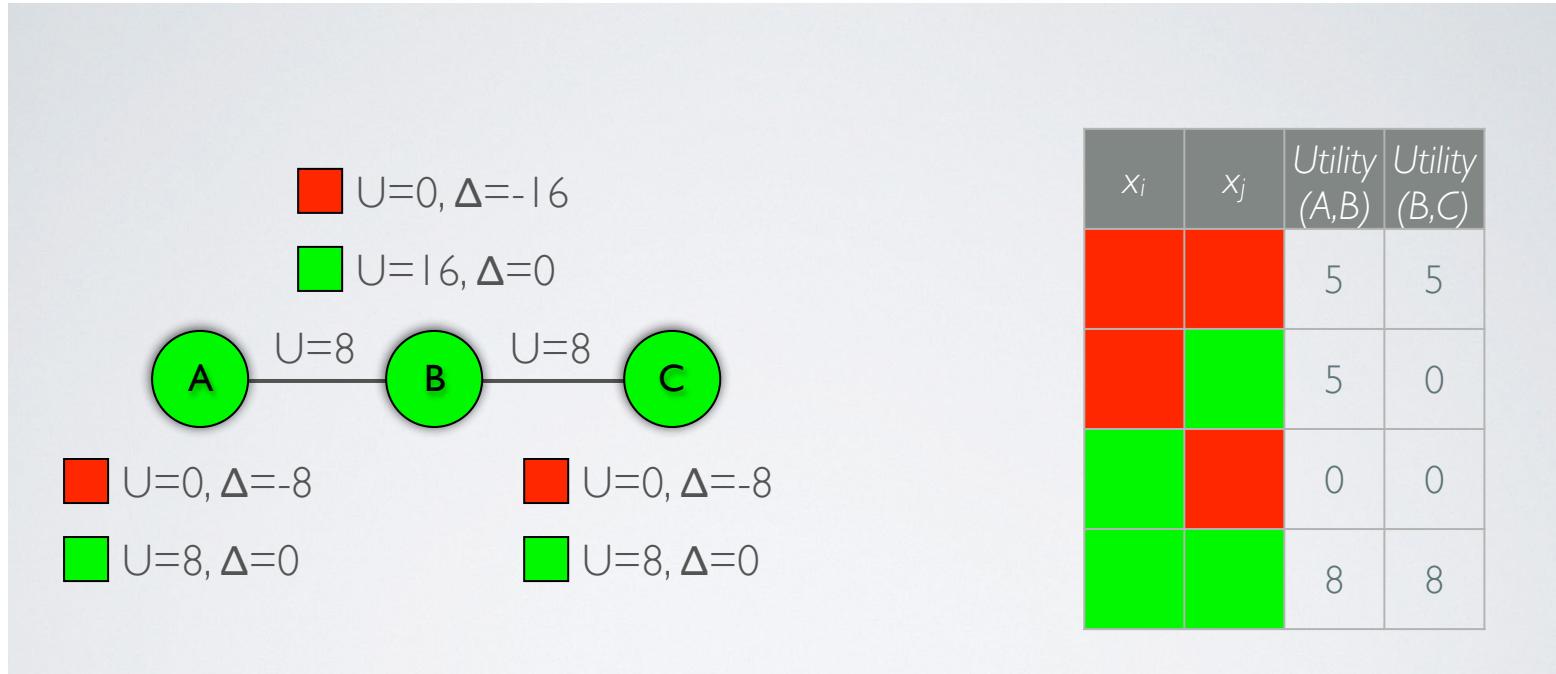
# Example run of DSA



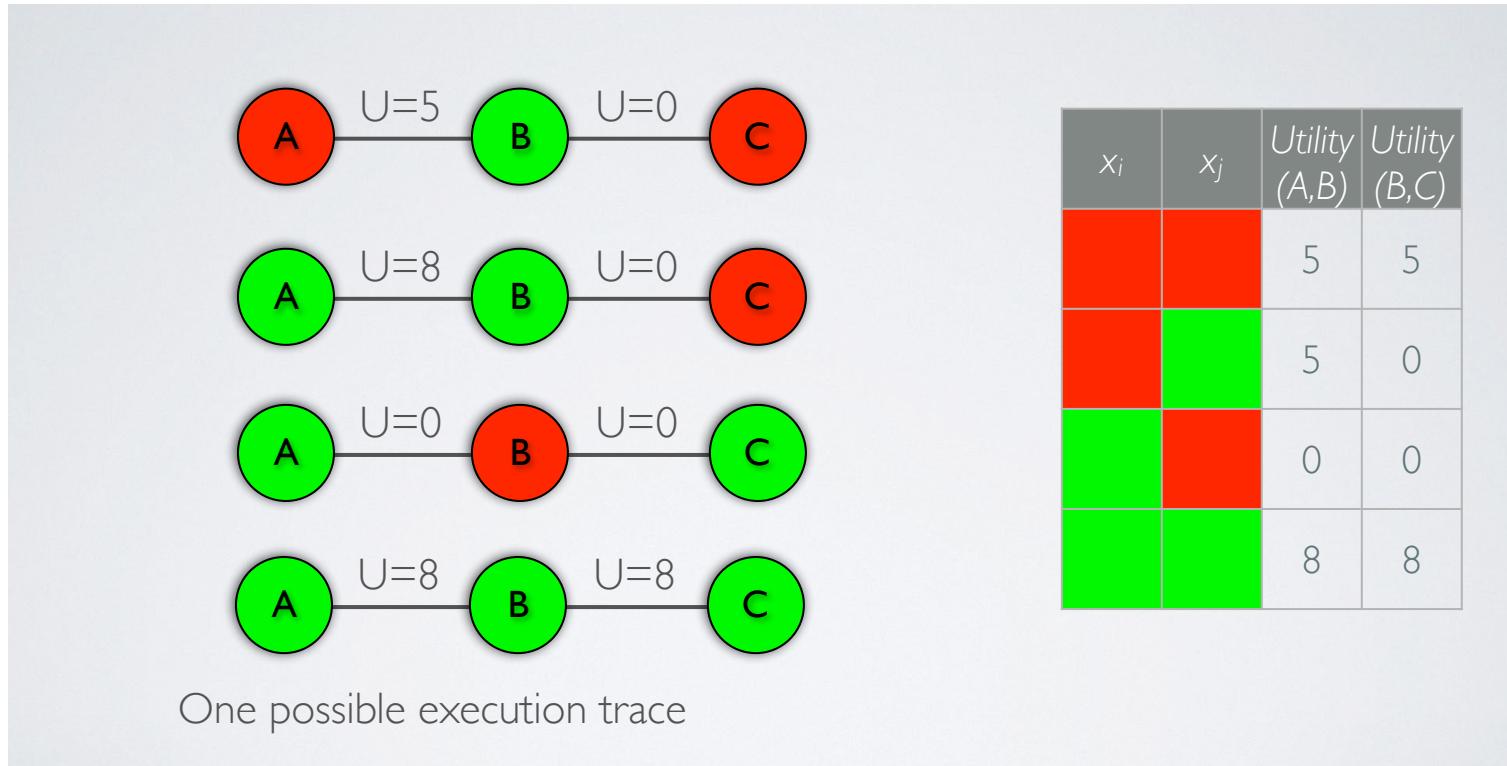
# Example run of DSA



# Example run of DSA



# Example run of DSA



# DSA-1: discussion

-  Extremely low computation/communication
-  Shows an anytime property (not guaranteed)
-  Activation probability:
  - Must be tuned
  - Domain dependent (no general rule)

# Maximum Gain Message (MGM-1)

***Coordination among neighbours to decide which single agent can perform the move.***

[R. T. Maheswaran et al., 2004]

- Initialize agents with a random assignment and communicate values to neighbors
- Each agent:
  - Compute and exchange possible gains
  - Agent with maximum (positive) gain executes
  - Communicate and receive possible variables changes to/from neighbors

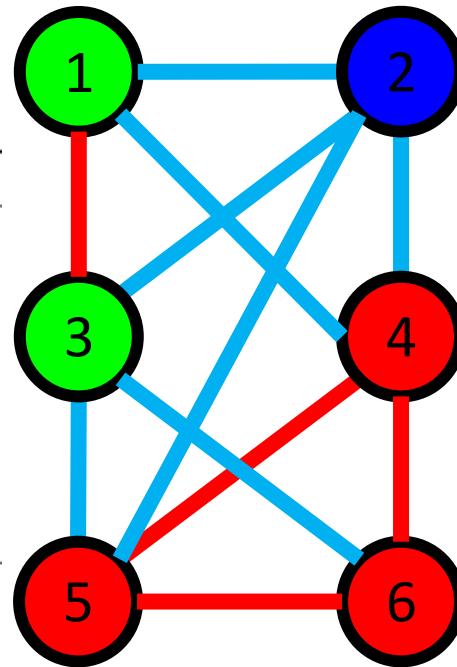
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



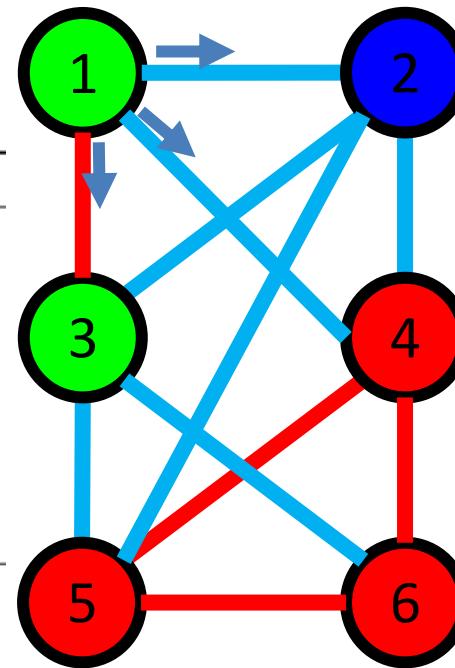
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



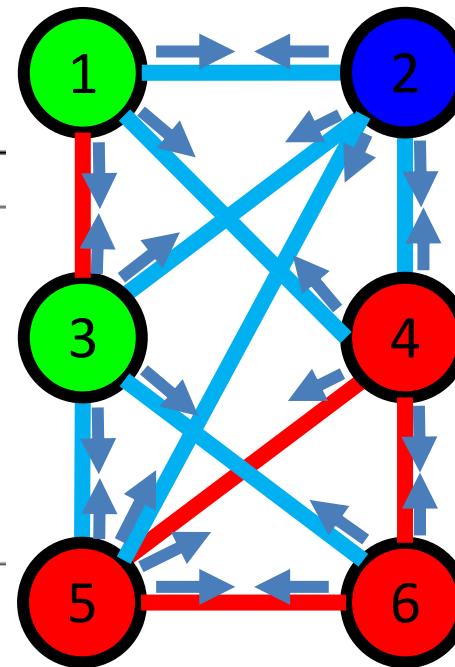
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



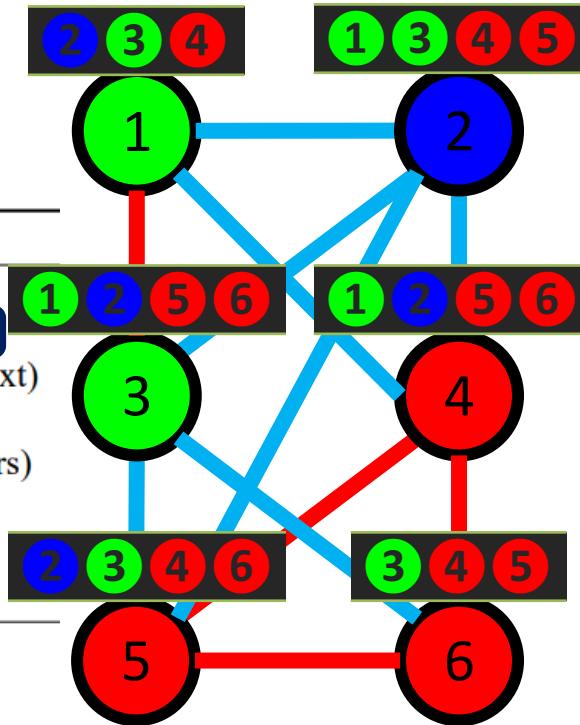
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



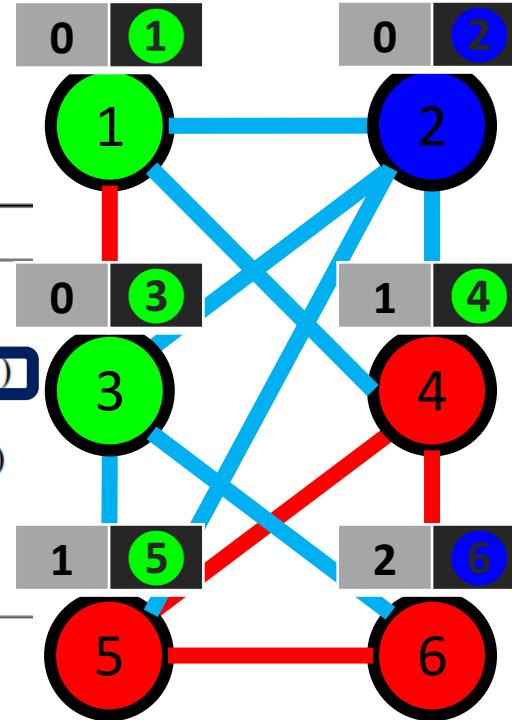
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



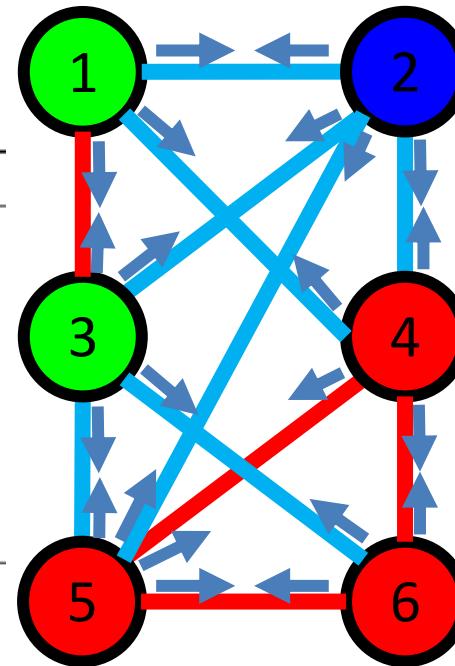
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain) █
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



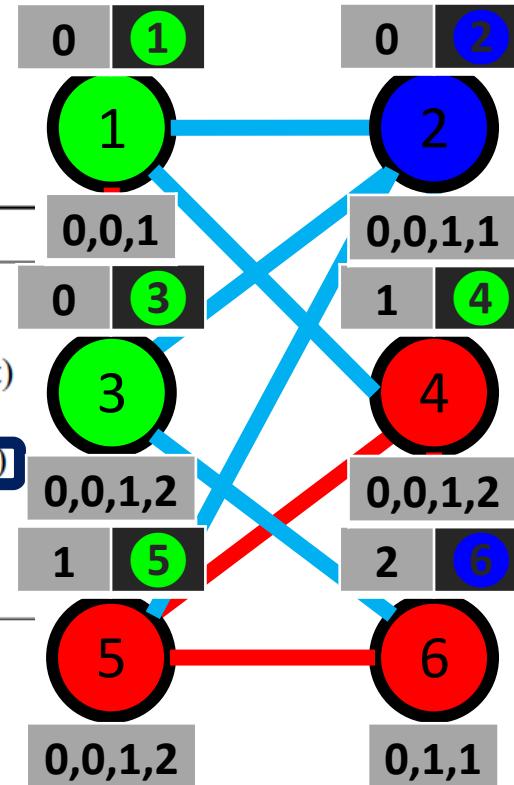
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



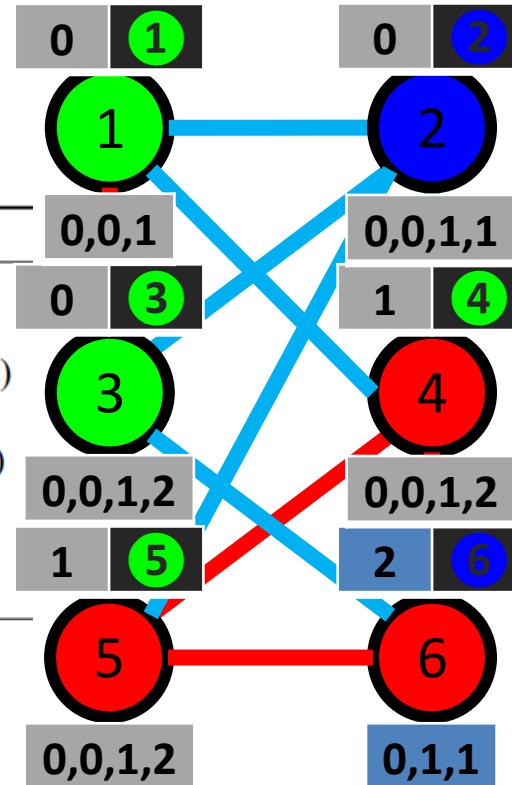
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



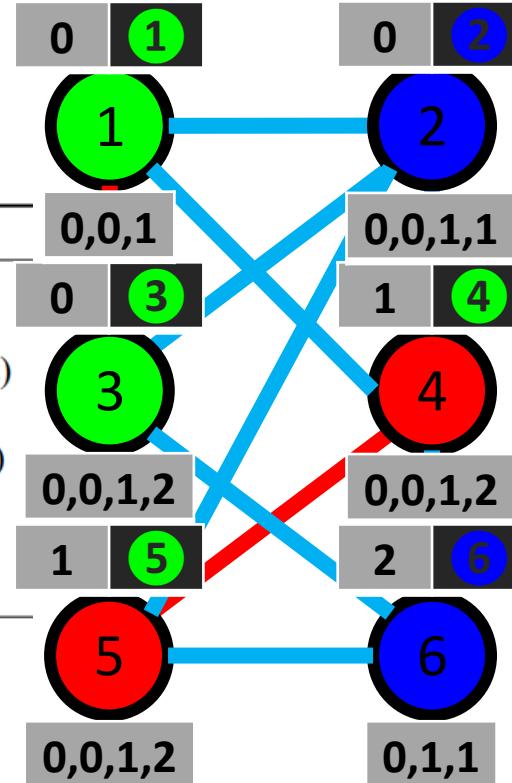
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



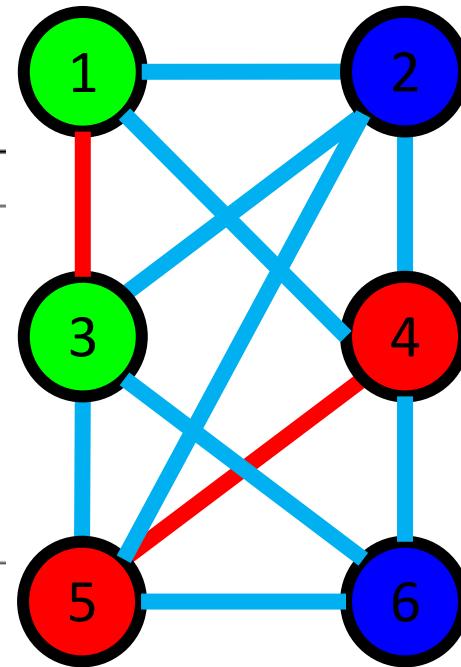
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



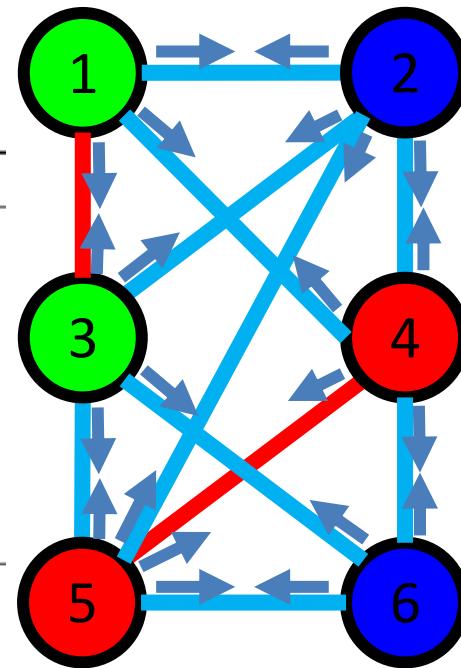
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



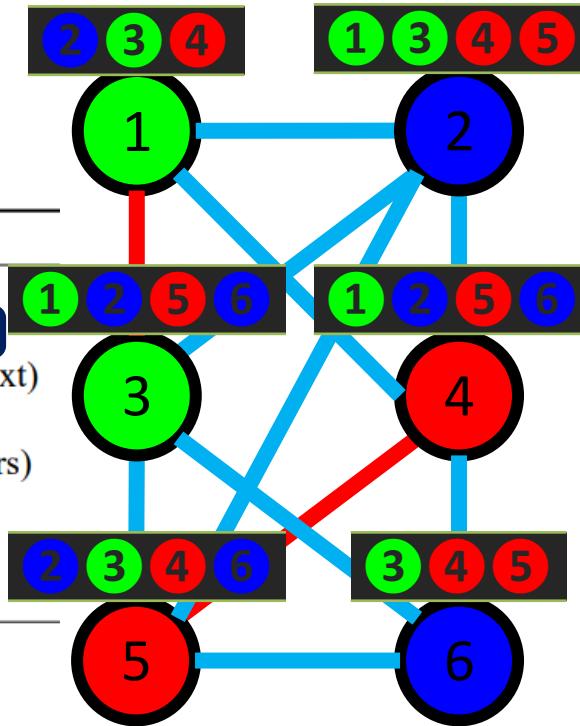
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

---



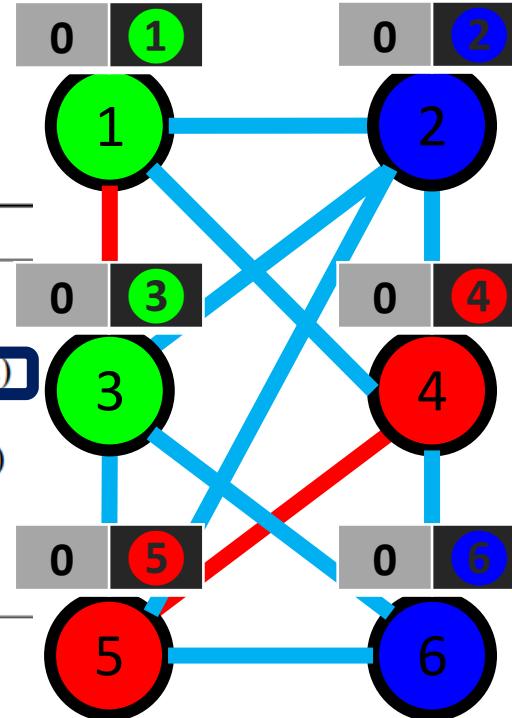
# Example run of MGM-1

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

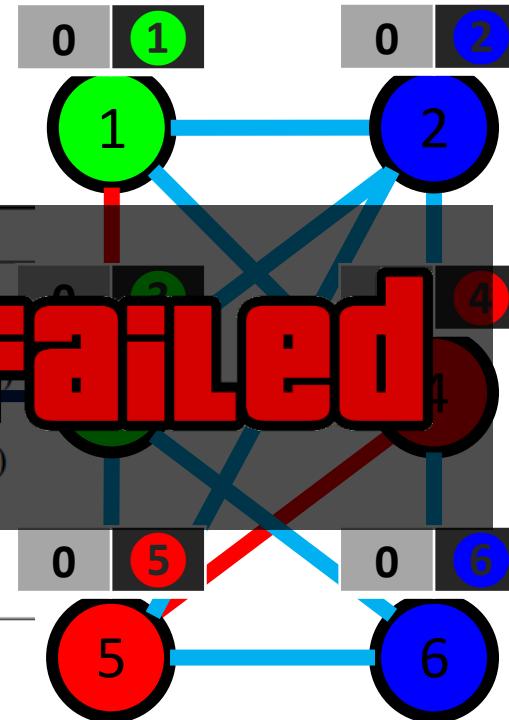
---



# Example run of MGM-1

## Algorithm 1 MGM (allNeighbors, currentValue)

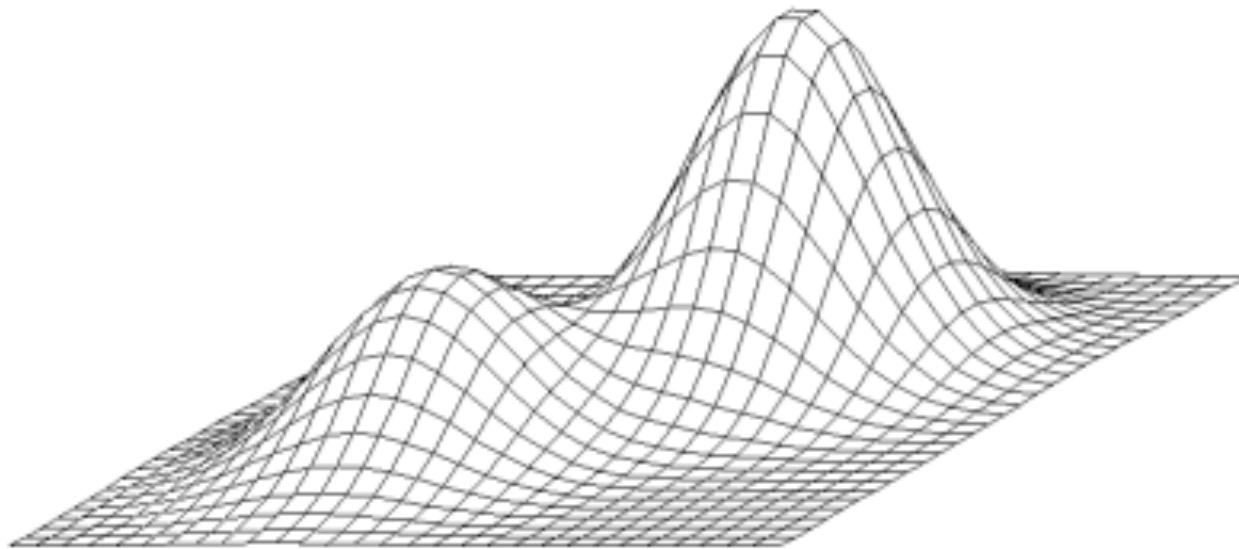
```
1: SendValueMessage(allNeighbors, currentValue)
2: currentGain = 0
3: if gain > max(neighborsGains) then
4:   currentValue = newValue
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```



# Why?

Local maximum vs. global maximum

MGM has reached a local maximum and, changing only a value, cannot escape from it



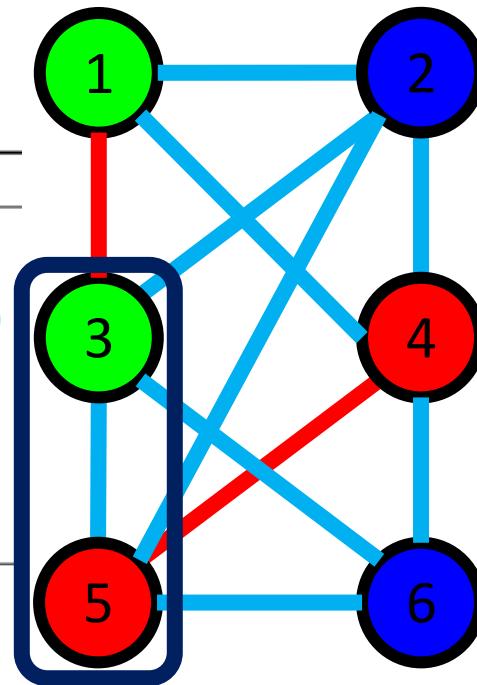
# Local maximum vs. global maximum

**Algorithm 1** MGM (allNeighbors, currentValue)

```

1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if

```



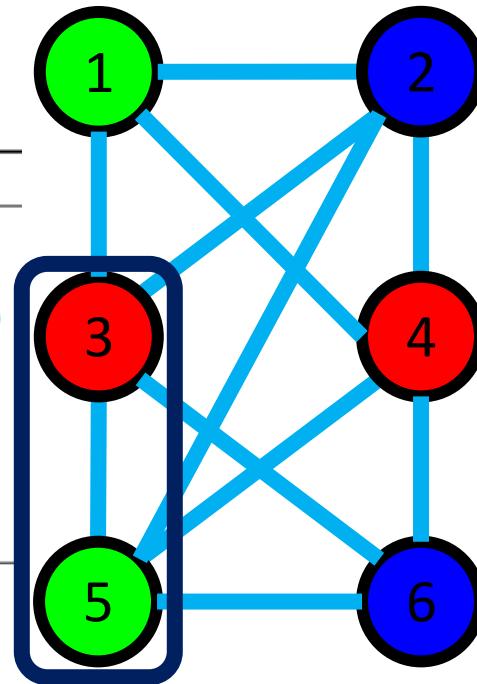
# Local maximum vs. global maximum

---

**Algorithm 1** MGM (allNeighbors, currentValue)

```
1: SendValueMessage(allNeighbors, currentValue)
2: currentContext = GetValueMessages(allNeighbors)
3: [gain,newValue] = BestUnilateralGain(currentContext)
4: SendGainMessage(allNeighbors,gain)
5: neighborGains = ReceiveGainMessages(allNeighbors)
6: if gain > max(neighborGains) then
7:   currentValue = newValue
8: end if
```

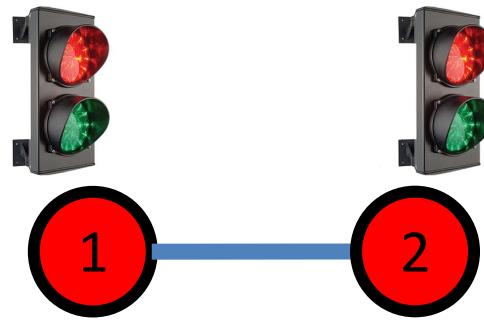
---



# MGM-1: discussion

- ≡ More communication than DSA but still linear
- ≡ Empirically similar to DSA
- thumb up Guaranteed to be anytime
- thumb up Does not require any parameter tuning.

# MGM vs DSA



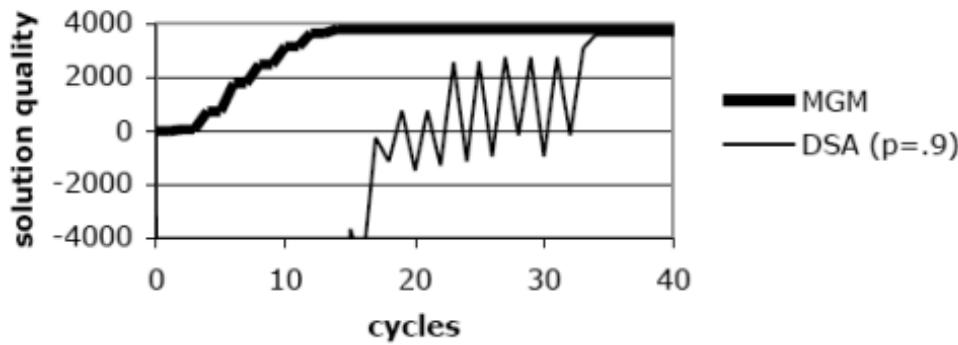
$$\begin{aligned} U(\text{red}, \text{red}) &= 0 \\ U(\text{red}, \text{green}) &= U(\text{green}, \text{red}) = 1 \\ U(\text{green}, \text{green}) &= -1000 \end{aligned}$$

If both agents are allowed to alter their value in the same round, we could end up in the adverse state (**green, green**)

When using DSA, there is always a positive probability for any time horizon that (**green, green**) will be the resulting assignment

# MGM vs DSA

When applying MGM, the global utility strictly increases at each round until convergence

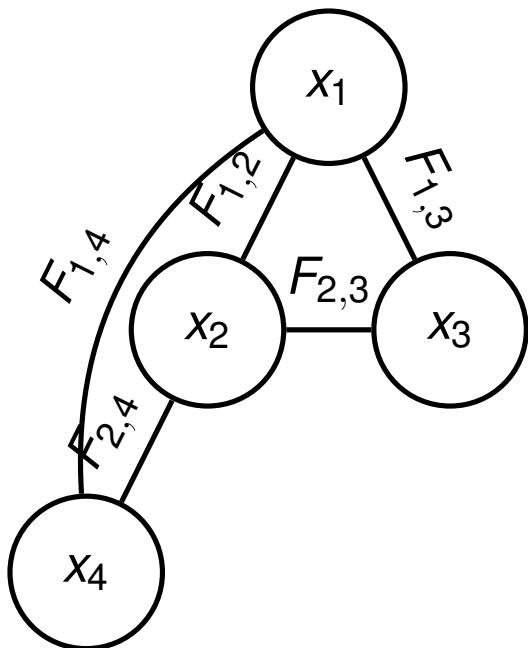


# Decentralised greedy approaches

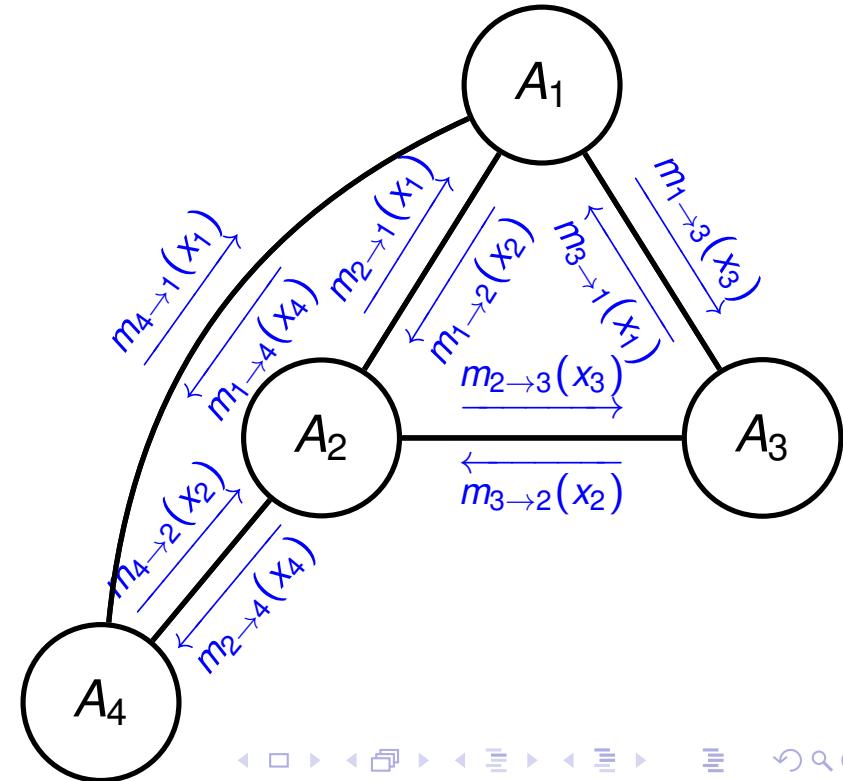
-  Very little memory and computation
-  Anytime behaviours
-  Could result in very bad solutions
  - local maxima arbitrarily far from optimal

# The Max-Sum algorithm

Agents iteratively exchange messages to build a local function that depends only on the variables they control



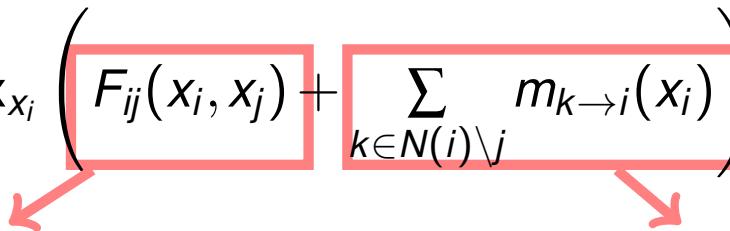
=>



## Max-Sum messages

At each execution step, each agent  $A_i$  sends to each of its neighbors  $A_j$  the message:

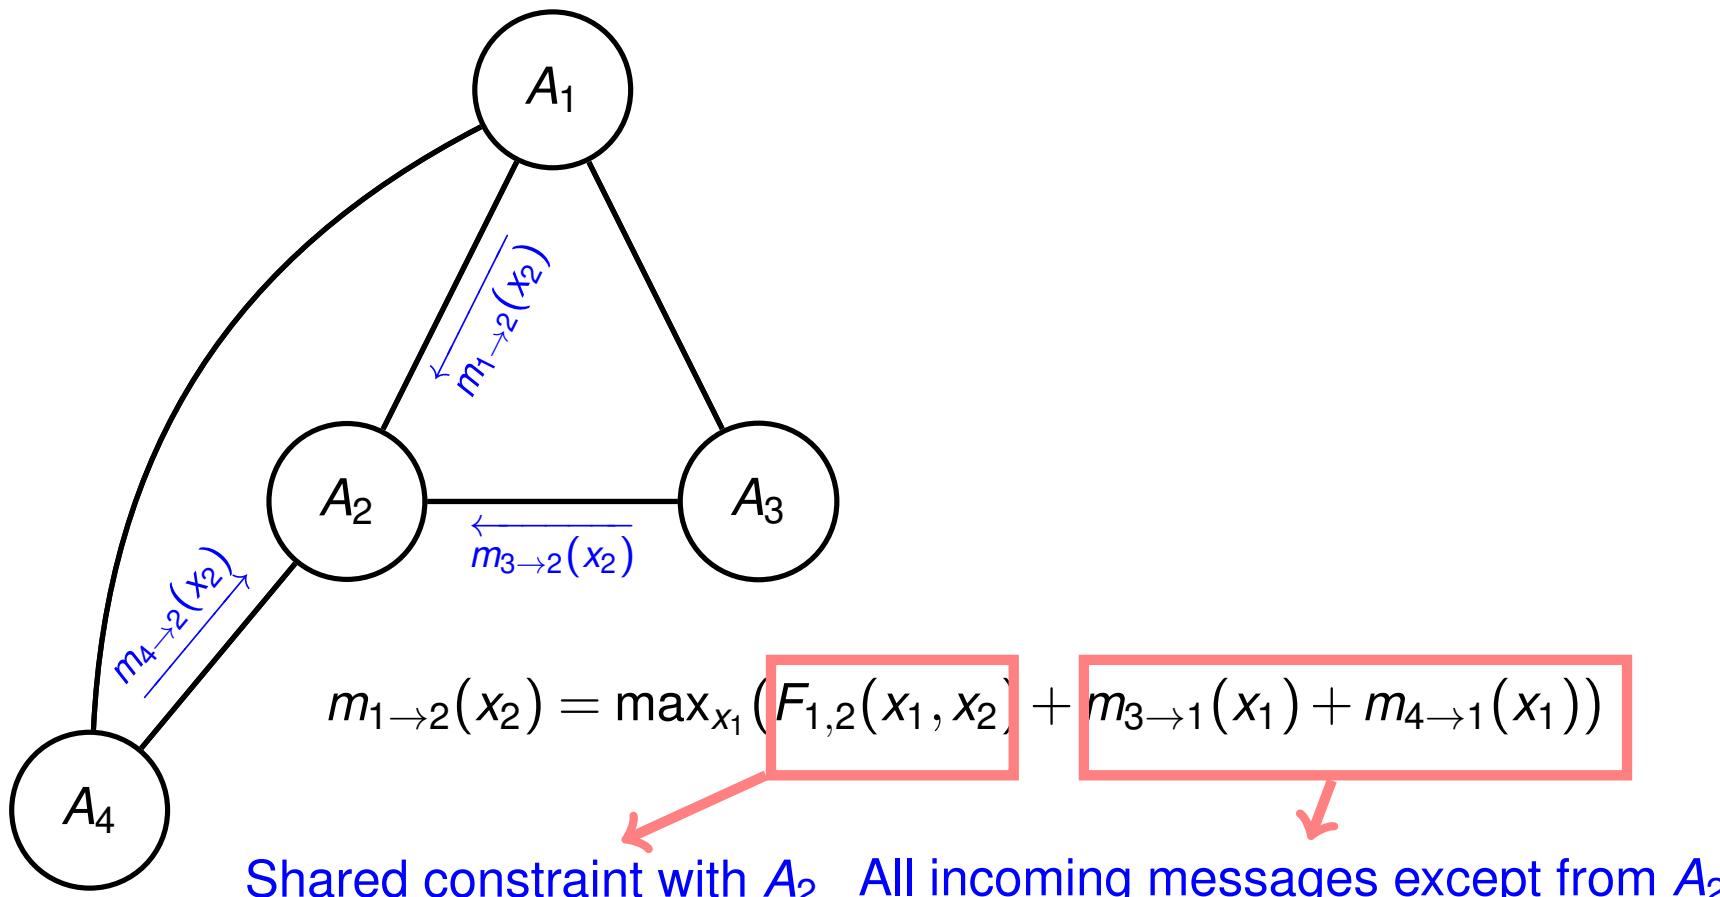
$$m_{i \rightarrow j}(x_j) = \alpha_{ij} + \max_{x_i} \left( F_{ij}(x_i, x_j) + \sum_{k \in N(i) \setminus j} m_{k \rightarrow i}(x_i) \right)$$


  
 Shared constraint with  $A_j$       All incoming messages except from  $A_j$

where:

- $\alpha_{ij}$  is a normalization constant added to all components of the message so that  $\sum_{x_j} m_{i \rightarrow j}(x_j) = 0$
- $N(i)$  is the set of indices for variables that are connected to  $x_i$

## Max-Sum by example



# Max-Sum assignments

At each iteration, each agent  $A_i$ :

- computes its local function as:

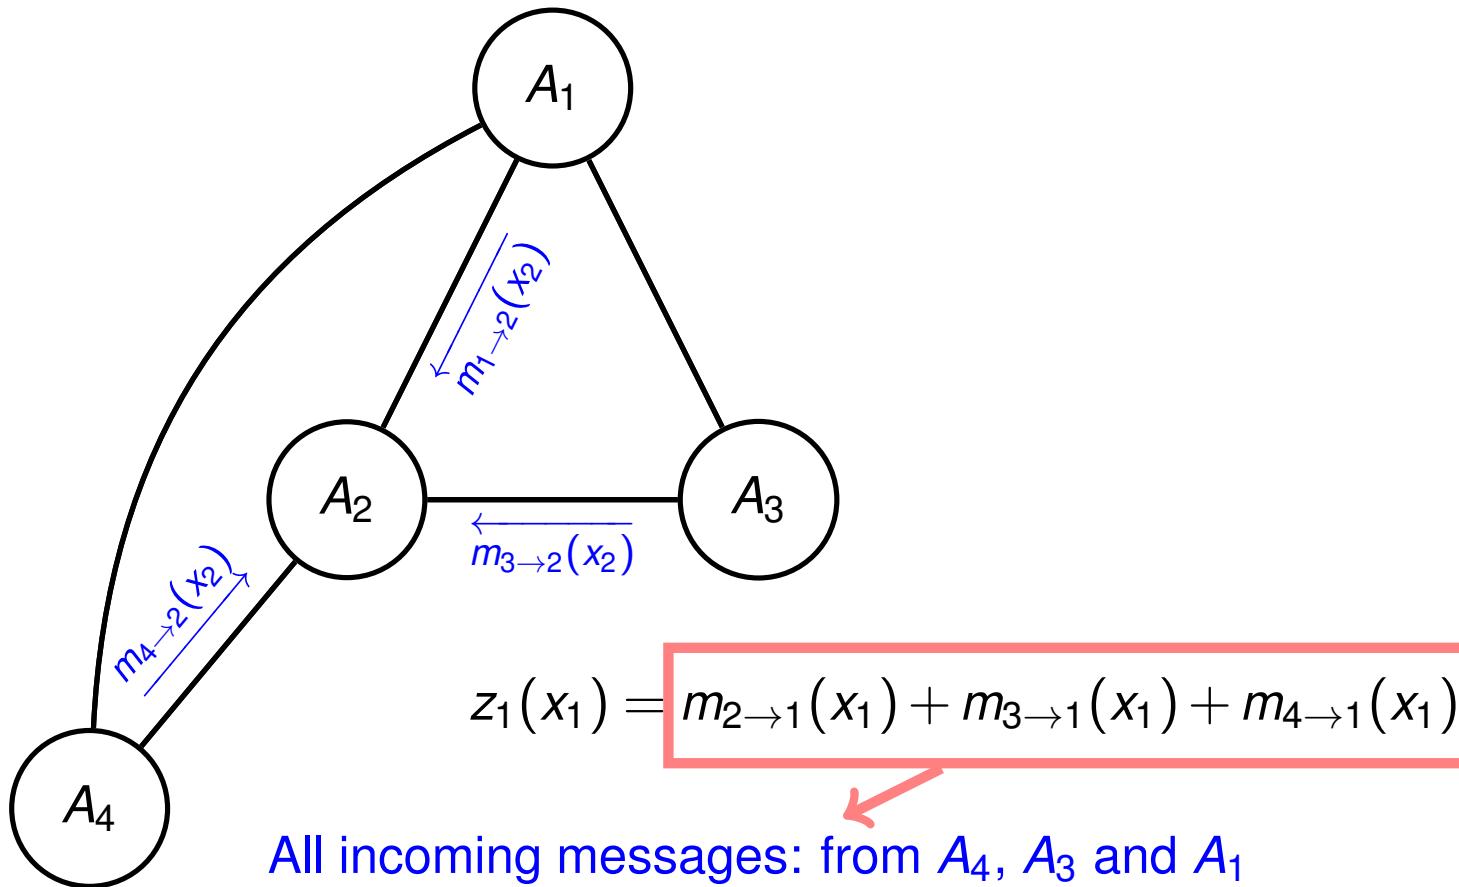
$$z_i(x_i) = \sum_{k \in N(i)} m_{k \rightarrow i}(x_i)$$

All incoming messages

- sets its assignment as the value that maximizes its local function:

$$\tilde{x}_i = \arg \max_{x_i} z_i(x_i)$$

## Max-Sum by example

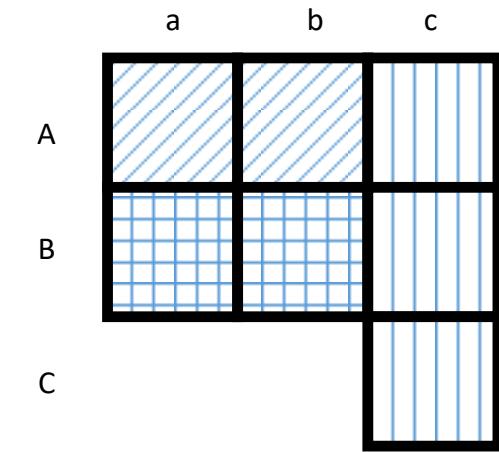


# Example run of Max-Sum

Consider this grid environment, divided in three areas (denoted by patterns)

Each area is controlled by an agent, which has to place a sensor in the area with the constraint that there cannot be two sensors on the same row (A, B, C) nor on the same column (a,b,c)

Agents are {1,2,3}, their variables are the positions of the sensors in their areas {P1, P2, P3}, their domains are {Aa, Ab}, {Ac, Bc, Cc}, and {Ba, Bb}, respectively, and the constraints are:



$F_{12}$	$P_1$	$P_2$
$-\infty$	Aa	Ac
0	Aa	Bc
0	Aa	Cc
$-\infty$	Ab	Ac
0	Ab	Bc
0	Ab	Cc

$F_{13}$	$P_1$	$P_3$
$-\infty$	Aa	Ba
0	Aa	Bb
0	Ab	Ba
$-\infty$	Ab	Bb

$F_{23}$	$P_2$	$P_3$
0	Ac	Ba
0	Ac	Bb
$-\infty$	Bc	Ba
$-\infty$	Bc	Bb
0	Cc	Ba
0	Cc	Bb

# Example run of Max-Sum

Assume normalization constants  $\alpha$  equal to 0

In the first execution round, agent 1 sends the following messages:

$$m_{1 \rightarrow 2}(P_2) = \alpha_{12} + \max_{P_1} F_{12}(P_1, P_2) =$$

P <sub>2</sub>	m <sub>1 → 2</sub>
A <sub>c</sub>	-∞
B <sub>c</sub>	0
C <sub>c</sub>	0

$$m_{1 \rightarrow 3}(P_3) = \alpha_{13} + \max_{P_1} F_{13}(P_1, P_3) =$$

P <sub>3</sub>	m <sub>1 → 3</sub>
B <sub>a</sub>	0
B <sub>b</sub>	0

Similarly, agent 2 sends the following messages:

P <sub>1</sub>	m <sub>2 → 1</sub>
A <sub>a</sub>	0
A <sub>b</sub>	0

P <sub>3</sub>	m <sub>2 → 3</sub>
B <sub>a</sub>	0
B <sub>b</sub>	0

And agent 3 sends the following messages:

P <sub>1</sub>	m <sub>3 → 1</sub>
A <sub>a</sub>	0
A <sub>b</sub>	0

P <sub>2</sub>	m <sub>3 → 2</sub>
A <sub>c</sub>	0
B <sub>c</sub>	-∞
C <sub>c</sub>	0

# Example run of Max-Sum

Agent 1 computes  $z_1(P_1) = m_{2 \rightarrow 1}(P_1) + m_{3 \rightarrow 1}(P_1) =$

P <sub>1</sub>	z <sub>1</sub>
Aa	0
Ab	0

and selects one value that maximizes  $z_1(P_1)$ , namely  $\sim P_1 = \text{Aa}$  (according to the lexicographic order)

Agent 2 does the same: computes  $z_2(P_2) = m_{1 \rightarrow 2}(P_2) + m_{3 \rightarrow 2}(P_2) =$

P <sub>2</sub>	z <sub>2</sub>
Ac	$-\infty$
Bc	$-\infty$
Cc	0

and selects the value  $\sim P_2 = \text{Cc}$

Agent 3 does the same: computes  $z_3(P_3) = m_{1 \rightarrow 3}(P_3) + m_{2 \rightarrow 3}(P_3) =$

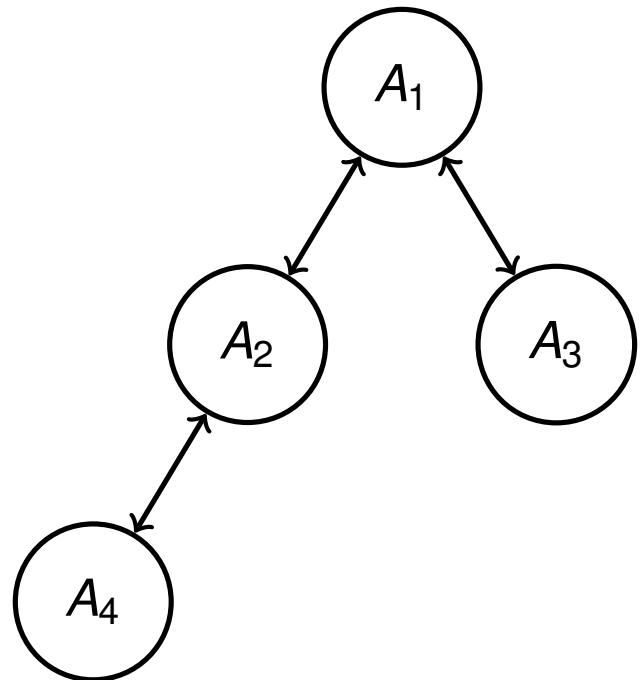
P <sub>3</sub>	z <sub>3</sub>
Ba	0
Bb	0

and selects the value  $\sim P_3 = \text{Ba}$  (according to the lexicographic order)

Then, Max-Sum continues with the next iteration...

# Max-Sum on acyclic graphs

- Optimal on acyclic graphs
  - Different branches are independent
  - $z$  functions provide correct estimations of agents contributions to the global problem
- Convergence guaranteed in a polynomial number of cycles



# Max-Sum on cyclic graphs

On cyclic graphs, limited theoretical results:

- Lack of convergence guarantees
- When converges, it does converge to a neighborhood maximum
- Neighborhood maximum: guaranteed to be greater than all other maxima within a particular region of the search space

# Quality guarantees

So far, **algorithms presented** (DSA-1, MGM-1, Max-Sum) do **not** provide any **guarantee** on the **quality** of their solutions

- Quality highly **dependent** on many factors which **cannot** always be properly **assessed** before deploying the system.
- Particularly **adverse** behaviour on **specific pathological instances**.

Challenge:

- **Quality assessment** on approximate algorithms

# Quality guarantees for approx. techniques

- Key area of research
- Address trade-off between guarantees and computational effort
- Particularly important for:
  - Dynamic settings
  - Severe constrained resources (e.g. embedded devices)
  - Safety critical applications (e.g. search and rescue)

# DCOP APPLICATIONS

- Scheduling Problems
  - Taking DCOP to the Real World: Efficient Complete Solutions for Distributed Multi-Event Scheduling. AAMAS 2004
- Radio Frequency Allocation Problems
  - Improving DPOP with Branch Consistency for Solving Distributed Constraint Optimization Problems. CP 2014
- Sensor Networks
  - Preprocessing techniques for accelerating the DCOP algorithm ADOPT. AAMAS 2005
- Home Automation
  - A Multiagent System Approach to Scheduling Devices in Smart Homes. AAMAS 2017, IJCAI 2016
- Traffic Light Synchronization
  - Evaluating the performance of DCOP algorithms in a real world, dynamic problem. AAMAS 2008
- Disaster Evacuation
  - Disaster Evacuation Support. AAAI 2007; JAIR 2017
- Combinatorial Auction Winner Determination
  - H-DPOP: Using Hard Constraints for Search Space Pruning in DCOP. AAAI 2008

# TRAFFIC FLOW CONTROL

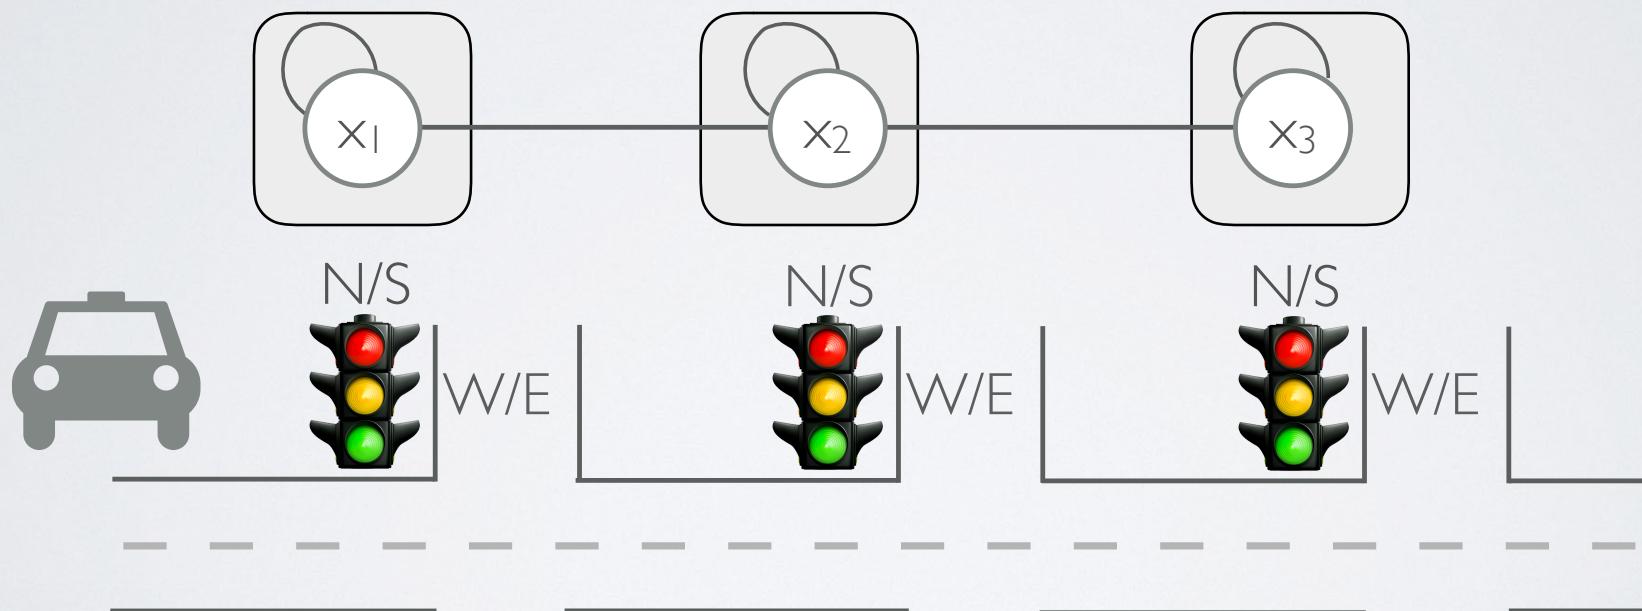
- Given a set of traffic lights in adjacent intersections
- How coordinate them to create green waves?



de Oliveira, D., Bazzan, A. L., & Lesser, V.. Using cooperative mediation to coordinate traffic lights: a case study. AAMAS 2005: 371-378  
Junges, R., & Bazzan, A. L. Evaluating the performance of DCOP algorithms in a real world, dynamic problem. AAMAS 2008: 463-470

# TRAFFIC FLOW CONTROL

- Agents: Each traffic light
- Values: Flow traffic direction



de Oliveira, D., Bazzan, A. L., & Lesser, V.. Using cooperative mediation to coordinate traffic lights: a case study. AAMAS 2005: 371-378  
Junges, R., & Bazzan, A. L. Evaluating the performance of DCOP algorithms in a real world, dynamic problem. AAMAS 2008: 463-470

# TRAFFIC FLOW CONTROL

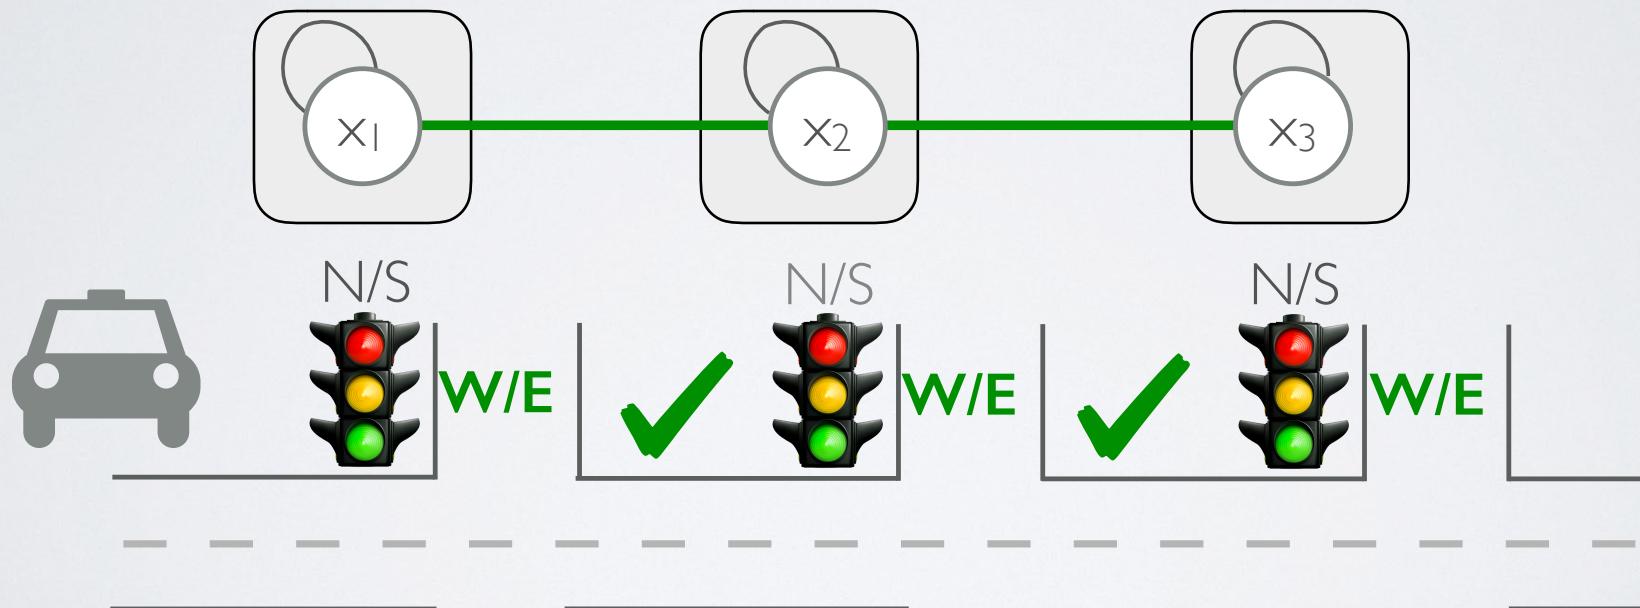
- Agents: Each traffic light
- Values: Flow traffic direction
- Conflict if 2 neighboring signals choose different directions



de Oliveira, D., Bazzan, A. L., & Lesser, V.. Using cooperative mediation to coordinate traffic lights: a case study. AAMAS 2005: 371-378  
Junges, R., & Bazzan, A. L. Evaluating the performance of DCOP algorithms in a real world, dynamic problem. AAMAS 2008: 463-470

# TRAFFIC FLOW CONTROL

- Cost functions model the number of incoming vehicles
- Maximize the traffic flow



de Oliveira, D., Bazzan, A. L., & Lesser, V.. Using cooperative mediation to coordinate traffic lights: a case study. AAMAS 2005: 371-378  
Junges, R., & Bazzan, A. L. Evaluating the performance of DCOP algorithms in a real world, dynamic problem. AAMAS 2008: 463-470

# SMART DEVICES



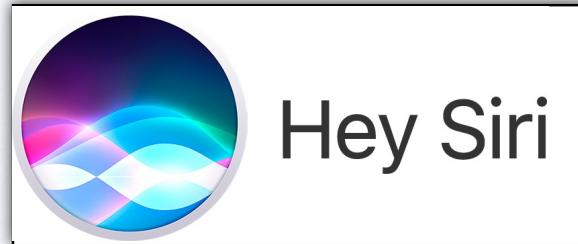
Introducing the  
**SAMSUNG** SmartFridge  
Say goodbye to food gone bad.

The SmartFridge makes it easy to keep track of what food you have and when it gets old.

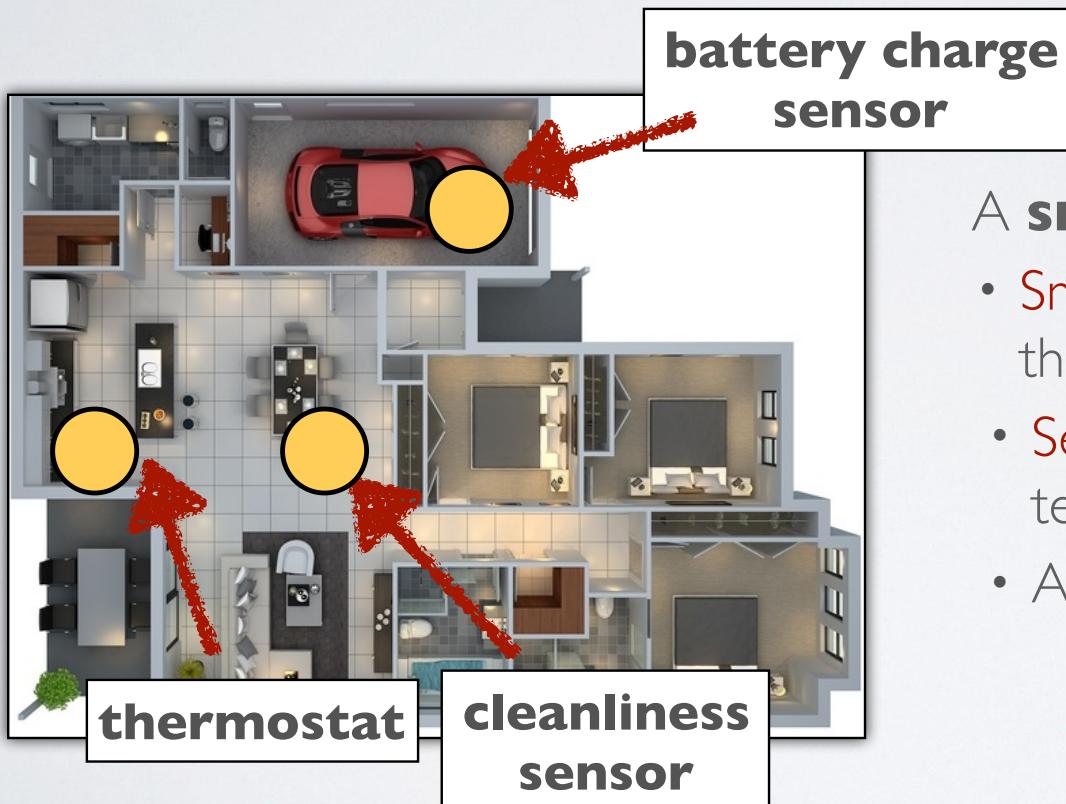
You can add food to the list by:

- Scanning it via the camera,
- Typing it in,
- Dictating it via the microphone.

# HOME ASSISTANTS



# SMART HOMES DEVICE SCHEDULING



A **smart home** has:

- Smart devices (roomba, HVAC) that it can control
- Sensors (cleanliness, temperature)
- A set of locations

Ferdinando Fioretto, William Yeoh, Enrico Pontelli. "A Multiagent System Approach to Scheduling Devices in Smart Homes". AAMAS, 2017.

# SMART HOMES DEVICE SCHEDULING



## Smart device:

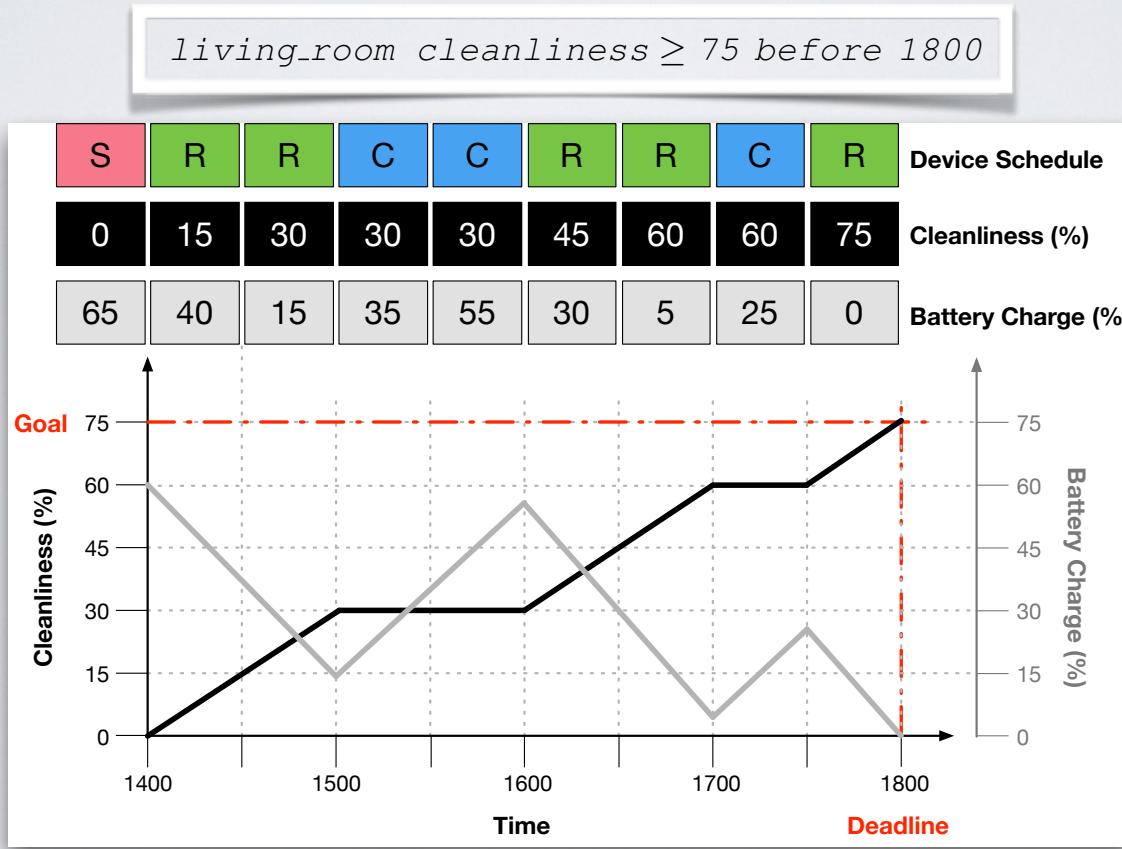
- A set of **actions** it can perform (clean, charge)
- Power consumption associated to each action.

## Scheduling Rules:

*living\_room cleanliness ≥ 75 before 1800*

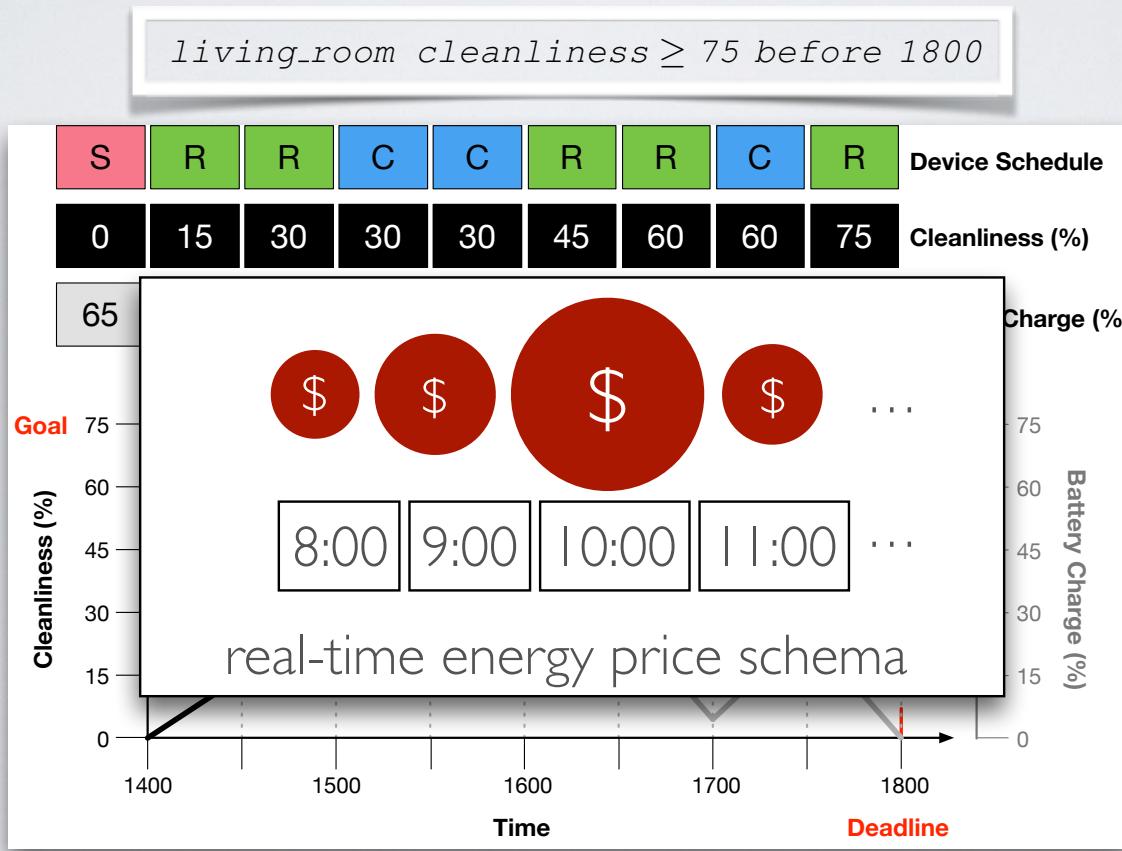
Ferdinando Fioretto, William Yeoh, Enrico Pontelli. "A Multiagent System Approach to Scheduling Devices in Smart Homes". AAMAS, 2017.

# SMART HOMES DEVICE SCHEDULING



Ferdinando Fioretto, William Yeoh, Enrico Pontelli. "A Multiagent System Approach to Scheduling Devices in Smart Homes". AAMAS, 2017.

# SMART HOMES DEVICE SCHEDULING



Ferdinando Fioretto, William Yeoh, Enrico Pontelli. "A Multiagent System Approach to Scheduling Devices in Smart Homes". AAMAS, 2017.

# SMART HOMES DEVICE SCHEDULING

How to schedule smart devices to satisfy the user preferences while

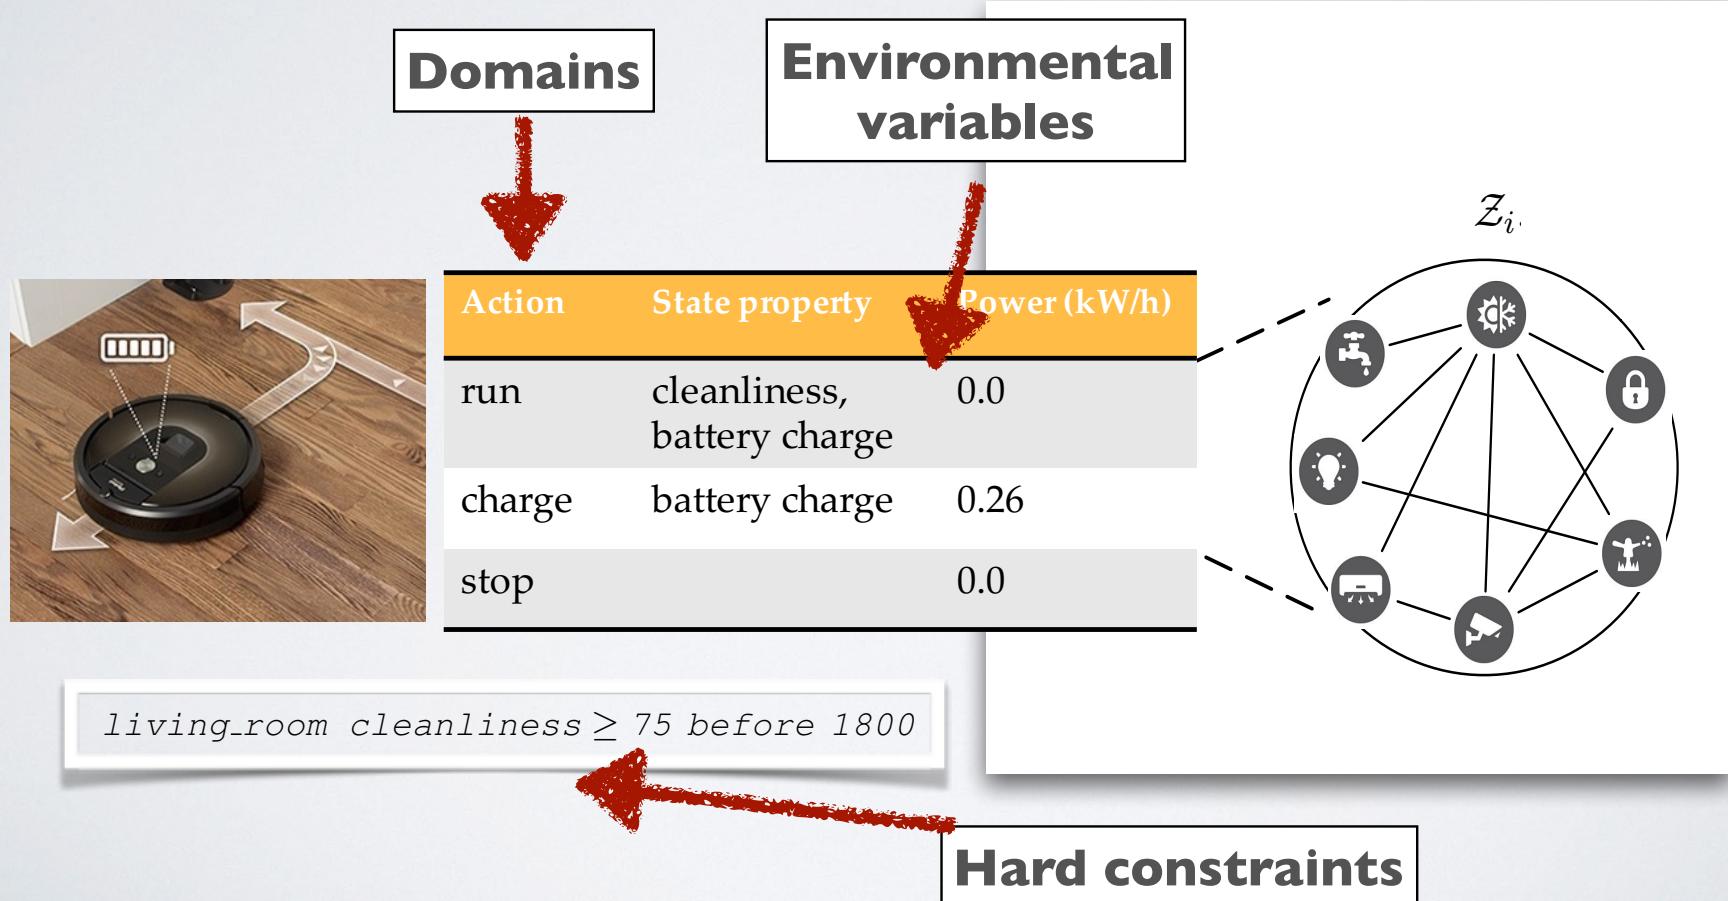
- 1) minimizing energy costs and
- 2) reducing peaks in load demand?

Assumptions: Each home have communication and controllable load capabilities.

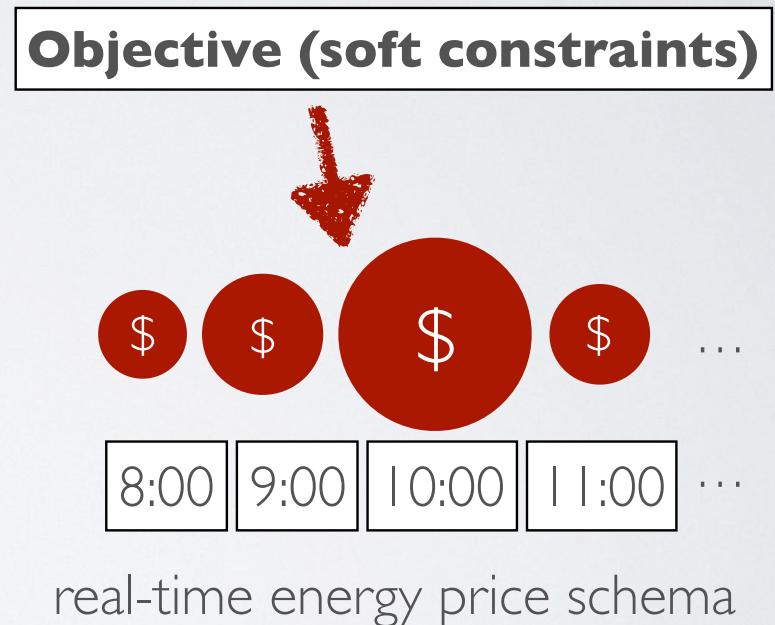
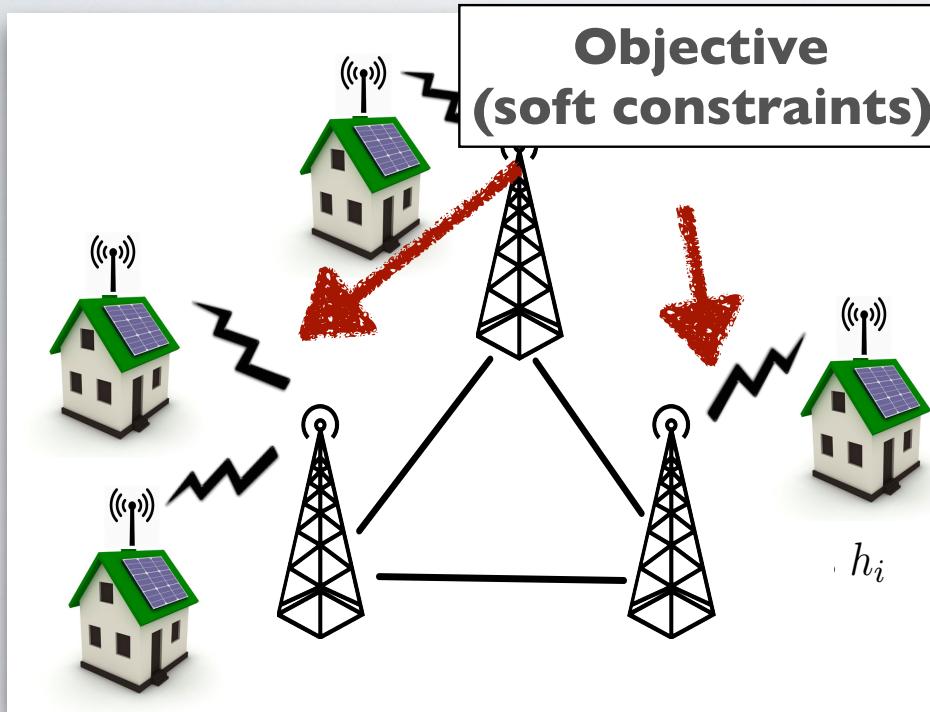


Ferdinando Fioretto, William Yeoh, Enrico Pontelli. "A Multiagent System Approach to Scheduling Devices in Smart Homes". AAMAS, 2017.

# SMART HOMES DEVICE SCHEDULING



# SMART HOMES DEVICE SCHEDULING



Ferdinando Fioretto, William Yeoh, Enrico Pontelli. "A Multiagent System Approach to Scheduling Devices in Smart Homes". AAMAS, 2017.