

Databases 2

1

Transactional Systems


In this lecture

- Introduction to the concept of *transaction*
- *Properties* of a transaction
- Modules of the DBMS responsible of guaranteeing such properties

Definition of Transaction

- An elementary, atomic unit of work performed by an application
- Each transaction is encapsulated within two commands:
 - `begin transaction` (bot)
 - `end transaction` (eot)
- Within a transaction, **one** of the commands below is executed (exactly **once**):
 - `commit-work` (commit)
 - `rollback-work` (abort)
- **Transactional System** (OLTP): a system that supports the execution of transactions on behalf of concurrent applications

Difference between Application and Transaction



1 Transactional Systems

Transactions: an example

```
begin transaction;  
update Account  
    set Balance = Balance + 10 where AccNum = 12202;  
update Account  
    set Balance = Balance - 10 where AccNum = 42177;  
commit-work;  
end transaction;
```

Transactions: an example with alternatives

```
begin transaction;  
update Account  
    set Balance = Balance + 10 where AccNum = 12202;  
update Account  
    set Balance = Balance - 10 where AccNum = 42177;  
select Balance into A from Account  
    where AccNum = 42177;  
if ( A >= 0 ) then commit-work  
                else rollback-work;  
end transaction;
```

Well-formed Transactions


begin transaction

code for data manipulation (reads and writes on DB)

commit-work / **rollback**-work

possibly more code, but no data manipulation

end transaction




ACID Properties of Transactions

- A transaction is a unit of work enjoying the following properties:
 - Atomicity
 - Consistency
 - Isolation
 - Durability

Atomicity

- A transaction is an atomic unit of execution, it can not be divided
 - Either all the operations in the transaction are executed or none is executed



Atomicity


- In the case of the bank transfer, the execution of a single update statement would be disastrous
- The instant in which **COMMIT** is executed marks the atomic and indivisible instant in which the transaction ends successfully:
 - An error before should cause the rollback of the work performed
 - An error afterwards should not alter the effect of the transaction

Atomicity

- The **ROLLBACK** of the work performed can be caused
 - By a ROLLBACK statement
 - By the DBMS, for example for the violation of integrity constraints or for concurrency management
- In case of a rollback, the work performed must be undone, bringing the database to the state it had before the start of the transaction

Atomicity

- Possible behaviours:
 - commit-work: **SUCCESS**
 - rollback-work or fault before commit-work: **UNDO**
 - fault after commit-work: **REDO**



Consistency

- A transaction must satisfy the DB integrity constraints
- Consequence:
 - **if** the initial state S_0 is consistent
 - **then** the final state S_f is also consistent

This is not necessarily true for the intermediate states S_i

Consistency - EXAMPLE

- A table can not contain two rows with the same value for a column declared as a primary key
- The sum of the balances of every couple of accounts must remain the same
 - If $\text{Balance}(a) + \text{Balance}(b) = \text{Costant}$ before the transaction then $\text{Balance}(a) + \text{Balance}(b) = \text{Costant}$ after the transaction
 - This type of constraint can be temporarily violated during the execution of the transaction, but must be satisfied at the end

Isolation

- The execution of a transaction must be independent from the concurrent execution of other transactions
- In particular, the concurrent execution of a number of transaction must produce the same result as the execution of the same transactions in a sequence
- E.g. the concurrent execution of T1 and T2 must produce the same results that can be obtained by executing one of these sequences
 - T1,T2
 - T2,T1

Durability

- The effect of a transaction that has successfully committed will last "**forever**"
 - Independently of any system fault
 - Sounds obvious... but every DBMS manipulates data in *main memory*

Transaction Properties and related mechanisms

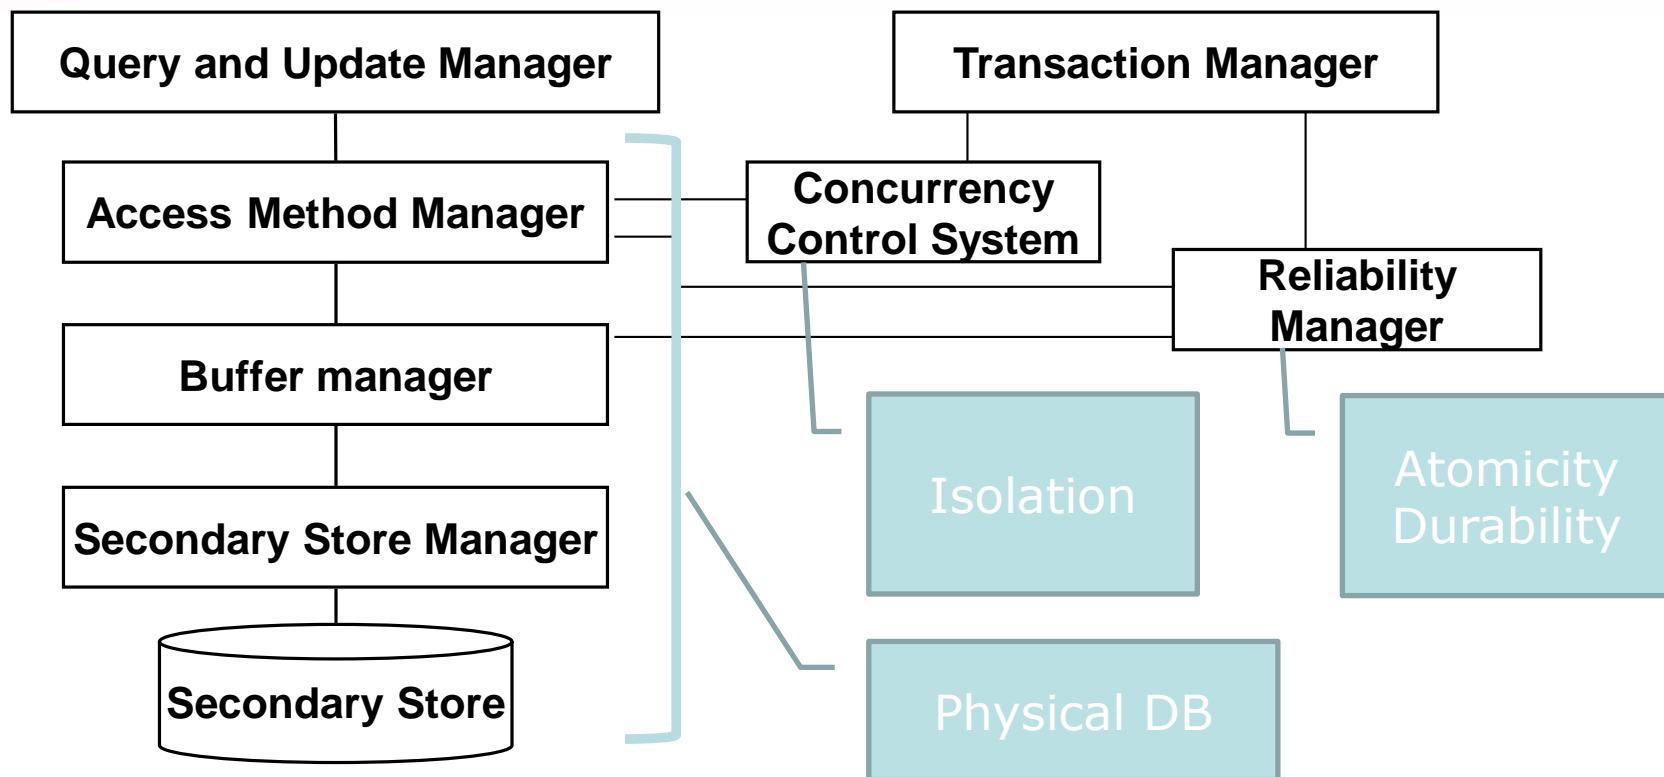
- **A**tomicity
 - Abort-rollback-restart, Commit protocols
- **C**onsistency
 - Integrity checking of the DBMS
- **I**solation
 - Concurrency control
- **D**urability
 - Recovery management

**ACID
properties**

Transactions and DBMS modules

- Atomicity and Durability
 - Reliability Manager
- Isolation
 - Concurrency Control System
- Consistency
 - Integrity Control System at query execution time
(with the support of the DDL compiler)

Logical Architecture of a DBMS



Next Topics...

- Concurrency Control (Isolation)
 - Theory
 - DBMS adopted method
- Reliability Control (Atomicity and Durability)
 - Commit protocols and 2PC
 - Logging and recovery on a single DBMS