

Databases

9

Advanced Databases: NoSQL databases

Database systems 2

Instructor: Sara Comai
sara.comai@polimi.it

DB EVOLUTION

1990-2000

- 1970-72: Relational model (P. Chen)
- 1974-77: Two major RDBMS prototypes
 - Ingres (UCB - Berkeley) – QUEL language → Ingres (later Postgres), MS SQL Server, Sybase
 - System R (IBM) – SEQUEL language → DB2, Oracle, Allbase (HP)
- 1976: Entity-Relationship Model (P. Chen)
- 1980s: SQL standard
- 1990s:
 - Object Oriented Databases (OODB)
 - Object-relational Databases (ORDB)
 - Object-Relational Mapping (ORM) – most popular
- Mid '90: WEB! Client-server DBMSs
 - Internet DB connectors (Front Page, Active Server Pages, JDBC, Java Servlet, Dream Weaver, ColdFusion, Enterprise Java Beans, and Oracle Developer 2000)
- 1998: XML

'90 - New needs

- Complex objects
 - complex aggregations
 - non-numerical information — images, spatial data, temporal series, ...
 - pre-defined and user-defined types (with reuse)
 - Complex operations
 - customized on the type of information — e.g., multimedia or user-defined types
 - long-lived transactions
- Supported by OO programming languages
- BUT: Object-oriented application development and relational data management use two incompatible models
 - IMPEDANCE MISMATCH PROBLEM!

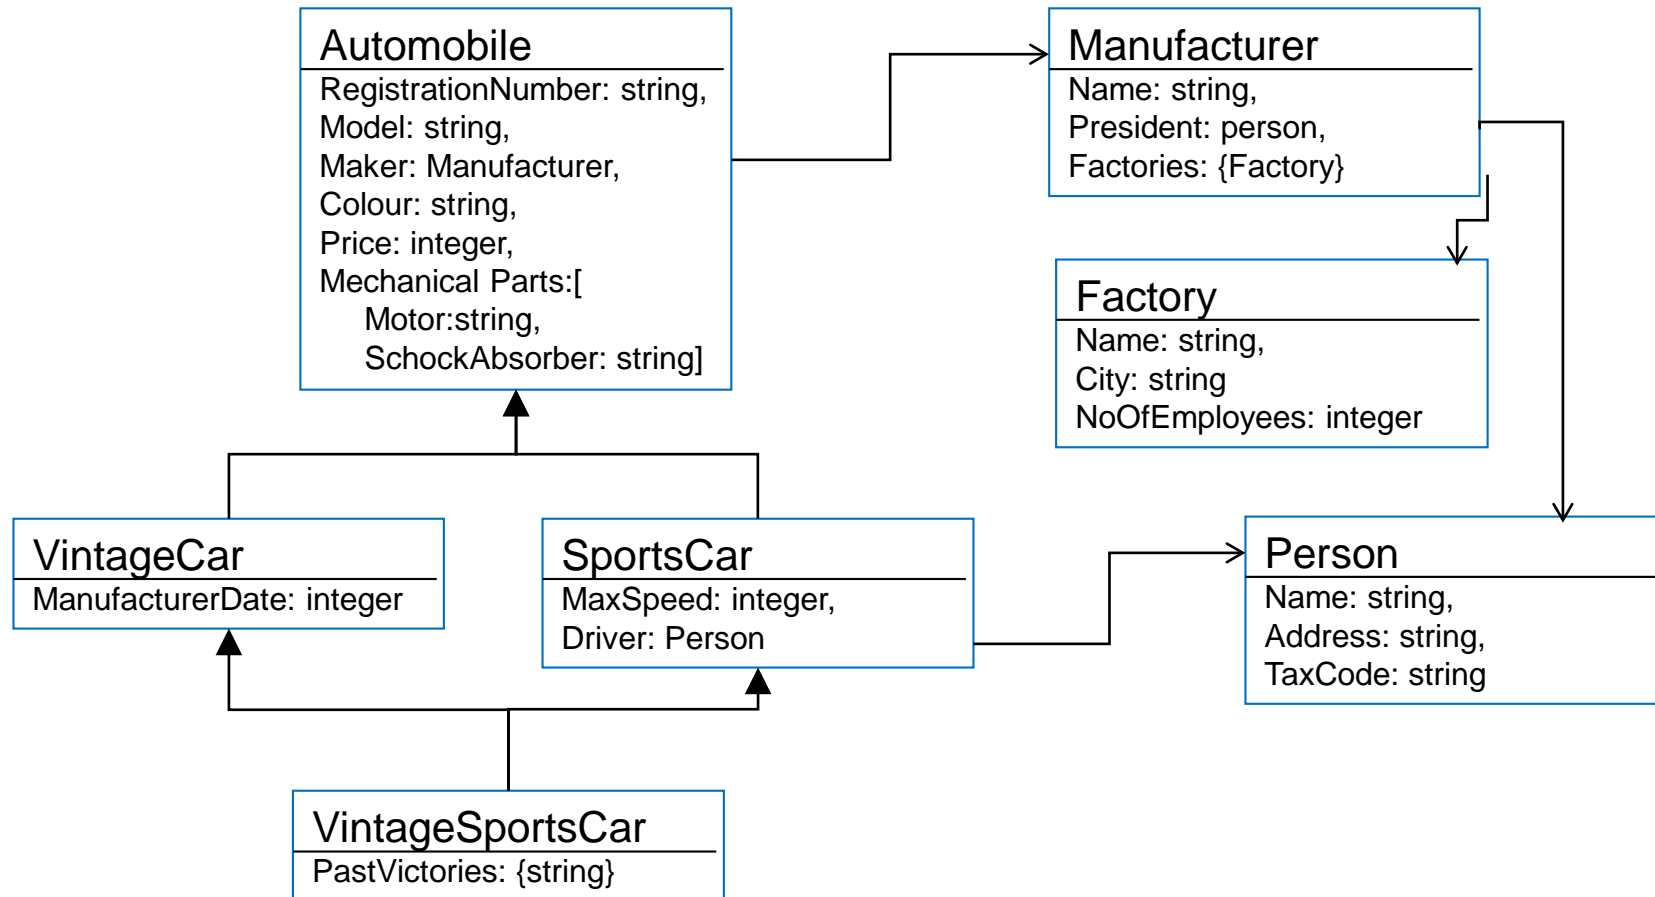
ODBMS Technology

- First generation of ODBMS: it is represented by *persistent* object programming languages
 - offer only some of the services of DBMSs, with no support for queries
- Second generation of ODBMS: manage and store persistent objects → offer a greater number of DBMS services and typically offer support for queries
- Two main ODBMS solutions
 - OODBMS (*Object-Oriented*): a *revolutionary* solution, with respect to RDBMS
 - ORDBMS (*Object-Relational*): an *evolutionary* solution, with respect to RDBMS

OODBMS

- Data represented as objects
 - object identity
 - attributes and methods
 - references, relationships, associations
- Extensible type hierarchy
 - user-defined types, abstract data types
 - single or multiple inheritance
 - overloading, overriding, late binding
- Declarative language for ad hoc purposes (ODL, OQL)
- Binding for object-oriented programming language

OODBMS - Example



OODBMS - Example of ODL schema

```
interface Automobile
{ extent Automobiles key RegistrationNumber}
{ attribute string RegistrationNumber;
  attribute string Model;
  attribute string Colour;
  attribute integer Price
  attribute structure MechanicalParts
  { string Motor,
    string ShockAbsorber};
  relationship <Manufacturer> Maker
    inverse Manufacturer::Builds;
  Automobile init (in string RegistrationNumber_par,
                  in string Model_par,
                  in string Colour_par,
                  in integer Price_par);
  void Increase (in integer Amount) raises(ExcessivePrice); }
```

Complex structure

Bidirectional relationships

Methods

inheritance

```
interface VintageCar: Automobile
{attribute integer ManufacturerDate}
```


OODBMS - OQL by examples

- Find the registration numbers of the vintage cars built at Maranello and driven by Fangio:

```
select x.RegistrationNumber
from x in VintageSportsCar
where x.Driver.Name = "Fangio" and "Maranello" in x.Maker.Factories.Name
```

Path expressions

- Find the names of the manufacturers who sell sports cars at a price higher than 200000; for each of them, list the city and number of employees of the factories:

```
select distinct struct(Name: x.Maker.Name,
                        Fact: (select struct (Cit: y.City,
                                                Emp: y.NoOfEmployees)
                               from y in Factory
                               where y in x.Maker.Factories))
from x in SportsCar
where x.Price > 200000
```

Nested structure

Commercial OODBMS

- Objectivity/DB (C++/Java/Smalltalk/Python/SQL++ interfaces)
 - ObjectStore PSE Pro for Java
 - Versant (C++/Java/Smalltalk)
 - O2 (C++/Java)
 - db4o (Java/.NET)
-
- Language-Integrated Query (LINQ), Microsoft .NET has a similar syntax

ORDBMS

- Relational model extended
 - Nested relations
 - References
 - Sets
 - Row types, abstract types
 - Functions
- Declarative language extended
- Fundamental impedance mismatch remains

Object-Relational databases (ORDBMSs)

- Revision of
 - Data model (non-1NF)
 - Query language
 - SQL-3 (SQL-99 is the official name)

Tuple types

```
create row type PerType(  
    Name varchar(30) not null,  
    Residence varchar(30),  
    TaxCode char(16) primary key)
```

```
create table Person of type PerType  
create table Industrial of type PerType
```

```
create table Manufacturer of type ManufacturerType  
values for ManufID are system generated  
scope for President is Person
```

- Tuples correspond to the objects
- Relations correspond to the classes
- It is possible to use references and to include objects

Queries

- Select the Name of the President of the Maker of Automobiles using the Motor "XV154"

```
select Maker -> President -> Name  
from Automobile  
where MechanicalParts -> Motor = 'XV154'
```

- Extract for each city the set of its Manufacturers

```
select City, set(name)  
from Manufacturer  
group by City
```

Adoption of ODBMS in practice

- ODBMS originally thought of to replace RDBMS because of their better fit with object-oriented programming languages
- But:
 - high switching cost
 - introduction of ORDBMS
 - emergence of object-relational mappers (ORMs)have made RDBMS successfully defend their dominance in the data center for server-side persistence
- Object databases are now established as a complement
 - as embeddable persistence solutions in devices, on clients, in packaged software, in real-time control systems, and to power websites.

DB EVOLUTION

2000-NOW

Big Data EveryWhere!




Social media and networks
(all of us are generating data)



Scientific instruments
(collecting all sorts of data)



Mobile devices
(tracking all objects all the time)



Sensor technology and networks
(measuring all kinds of data)

Needs for ability to manage, analyze, summarize, visualize, and discover knowledge from the collected data in a timely manner and in a scalable fashion




Characteristics of Big Data: 1-Scale (Volume)

- **Data Volume**
 - 44x increase from 2009 to 2020
 - From 0.8 zettabytes to 35zb
- Data volume is increasing exponentially




Characteristics of Big Data: 2-Complexity (Variety)

- Various formats, types, and structures
- Text, numerical, images, audio, video, sequences, time series, social media data, multi-dim arrays, etc...
- Static data vs. streaming data
- A single application can be generating/collecting many types of data



To extract knowledge
→ all these types of data need
to be linked together




Characteristics of Big Data: 3-Speed (Velocity)

- Data is begin generated fast and need to be processed fast
- Online Data Analytics
- Late decisions → missing opportunities
- **Examples**
 - **E-Promotions:** Based on your current location, your purchase history, what you like → send promotions right now for store next to you
 - **Healthcare monitoring:** sensors monitoring your activities and body → any abnormal measurements require immediate reaction

Some Make it 4V's


Volume



Data at Rest

Terabytes to
exabytes of existing
data to process


Velocity



Data in Motion


Streaming data,
milliseconds to
seconds to respond

Variety

Data in Many
Forms

Structured,
unstructured, text,
multimedia


Veracity*



Data in Doubt

Uncertainty due to
data inconsistency
& incompleteness,
ambiguities, latency,
deception, model
approximations

NoSQL

- “NoSQL” = “Not Only SQL”
 - Not every data management/analysis problem is best solved exclusively using a traditional DBMS
- Definition (from <http://nosql-database.org/> )

Next Generation Databases mostly addressing some of the points: being **non-relational**, **distributed**, **open-source** and **horizontally scalable**.

The movement began early 2009 and is growing rapidly. Often more characteristics apply such as: **schema-free**, **easy replication support**, **simple API**, **eventually consistent** / **BASE** (not ACID), a **huge amount of data** and more.

NoSQL taxonomy

Key-Value Stores



Document Stores




Graph/Triple Stores



Object Stores



Column-Family Stores



What do they miss compared to RDBMs?

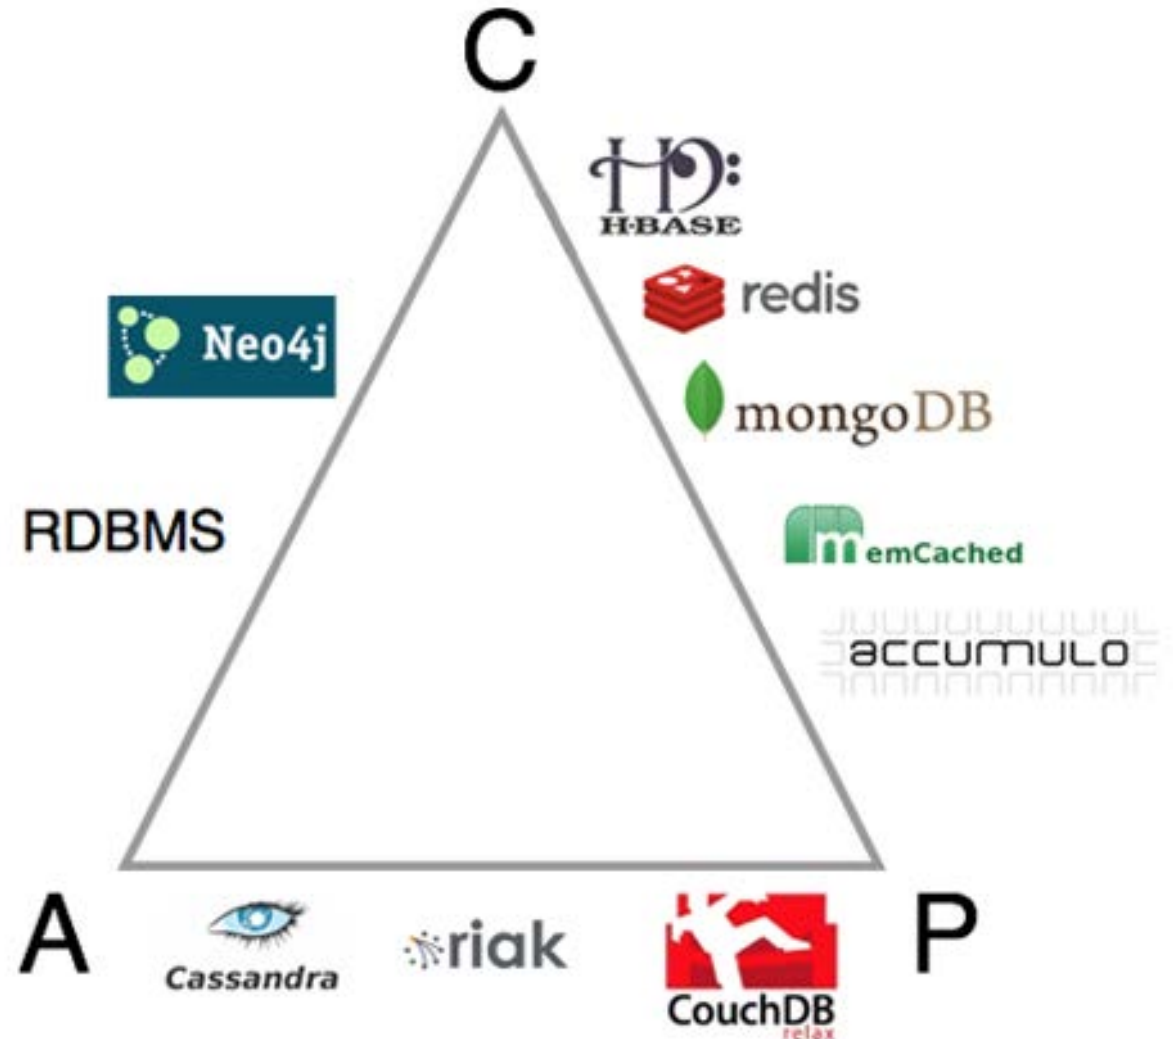
In general, they do not (fully) support:

- Joins
- Group by
- Order by
- ACID transactions

Brewer's CAP Theorem

- Consistency: Every read receives the most recent write or an error
- Availability: Every request receives a (non-error) response – without the guarantee that it contains the most recent write
- Partition tolerance: The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes

Theorem: for any system sharing data it is impossible to guarantee simultaneously all of these three properties



Eventual Consistency

- BASE (**B**asically **A**vailable, **S**oft state, **E**ventual consistency) properties, as opposed to ACID
 - **S**oft state: copies of a data item may be inconsistent
 - **E**ventually Consistent – copies becomes consistent at some later time if there are no more updates to that data item
 - **B**asically **A**vailable – possibilities of faults but not a fault of the whole system

ACID	BASE
Strong consistency Isolation Focus on "commit" Nested transactions Availability? Conservative (pessimistic) Difficult evolution (e.g. schema)	Weak consistency – stale data OK Availability first Best effort Approximate answers OK Aggressive (optimistic) Simpler! Faster Easier evolution

New RDBMs

- To increase **scalability**:

- Oracle RAC
- MySQL Cluster

Provide partitioning of load over multiple nodes

Other options: ScaleDB, VoltDB, NimbusDB, Clustrix

→ also called **NewSQL**

They support the relational data model and SQL, and guarantee ACID properties.

Then, which solutions should we choose?


- RDBMs if performances can be achieved with a **single server** or possibly with indexed caching front-end (there are tools available with the different RDBMS solutions)
- If you already use SQL and need some **scalability**: new RDBMs for scalability
- Key-value stores if you need to **partition** data across different RAMs in a LAN
- If you need to **scale to 100X** the single server → document or column stores

Ranking from <http://db-engines.com/en/ranking>

343 systems in ranking, August 2018

Aug 2018	Rank		DBMS	Database Model	Score		
	Jul 2018	Aug 2017			Aug 2018	Jul 2018	Aug 2017
1.	1.	1.	Oracle +	Relational DBMS	1312.02	+34.24	-55.85
2.	2.	2.	MySQL +	Relational DBMS	1206.81	+10.74	-133.49
3.	3.	3.	Microsoft SQL Server +	Relational DBMS	1072.65	+19.24	-152.82
4.	4.	4.	PostgreSQL +	Relational DBMS	417.50	+11.69	+47.74
5.	5.	5.	MongoDB +	Document store	350.98	+0.65	+20.48
6.	6.	6.	DB2 +	Relational DBMS	181.84	-4.36	-15.62
7.	7.	↑ 9.	Redis +	Key-value store	138.58	-1.34	+16.68
8.	8.	↑ 10.	Elasticsearch +	Search engine	138.12	+1.90	+20.47
9.	9.	↓ 7.	Microsoft Access	Relational DBMS	129.10	-3.48	+2.07
10.	10.	↓ 8.	Cassandra +	Wide column store	119.58	-1.48	-7.14
11.	11.	11.	SQLite +	Relational DBMS	113.73	-1.55	+2.88
12.	12.	12.	Teradata +	Relational DBMS	77.41	-0.82	-1.83


Number of systems per category, August 2018



Popularity of the models

<http://db-engines.com/en/ranking>


Ranking scores per category in percent, August 2018



Popularity of the models

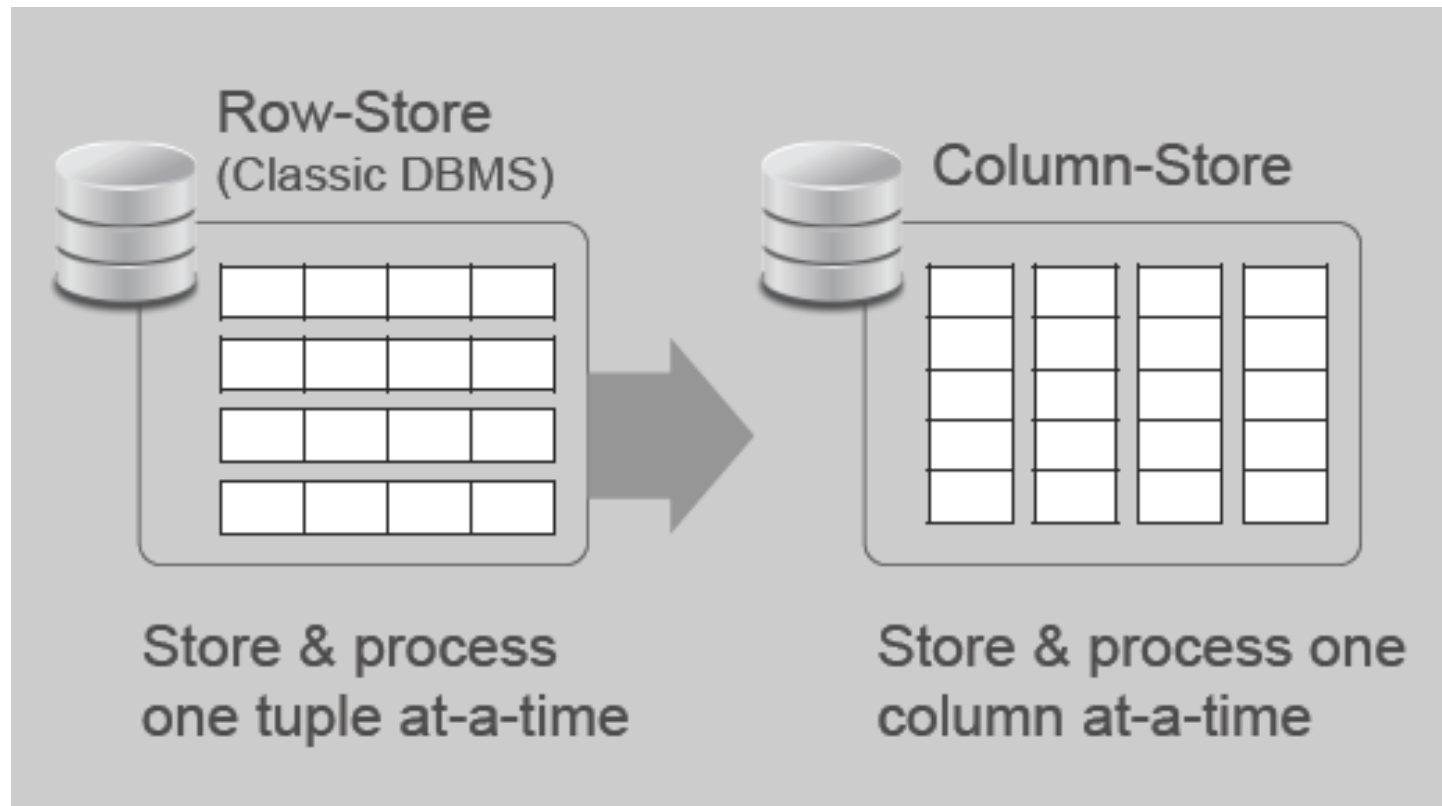
Trend of the last 24 months

<http://db-engines.com/en/ranking>




COLUMN-STORES

Column-store: an overview



Tuple Insertion:  Fast

 Requires multiple seeks

Queries:  May access more data than needed


 Read only relevant data

Perfect for analytics!

Compression

- Run-length encoding



Replace list of identical values by pair (value, count)



⊕ Enables query processing on compressed data directly

⊕ Good for sorted columns

e.g., Select persons in CA



Late tuple materialization

e.g., `SELECT Name FROM Person WHERE City="SD" AND State="CA"`

Π_{Name}

$\sigma_{\text{City} = \text{"SD"} \wedge \text{State} = \text{"CA"}}$

Create tuples

Name	Phone	City	State
Flintstone	858	SD	CA
Barney	619	SD	CA
Simpson	415	SF	CA
Bond	212	NY	NY

+ Reads relevant columns only

Late tuple materialization

e.g., SELECT Name FROM Person WHERE City="SD" AND State="CA"

Extract

Name

σ

City = "SD"

σ

State = "CA"

1

1

0

0

1

1

1

0

Name

Flintstone

Barney

1

1

0

0

Position lists

Name	Phone	City	State
Flintstone	858	SD	CA
Barney	619	SD	CA
Simpson	415	SF	CA
Bond	212	NY	NY


KEY-VALUE STORES

Key-Value Stores

- Extremely simple interface
 - Data model: (key, value) pairs
 - Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)

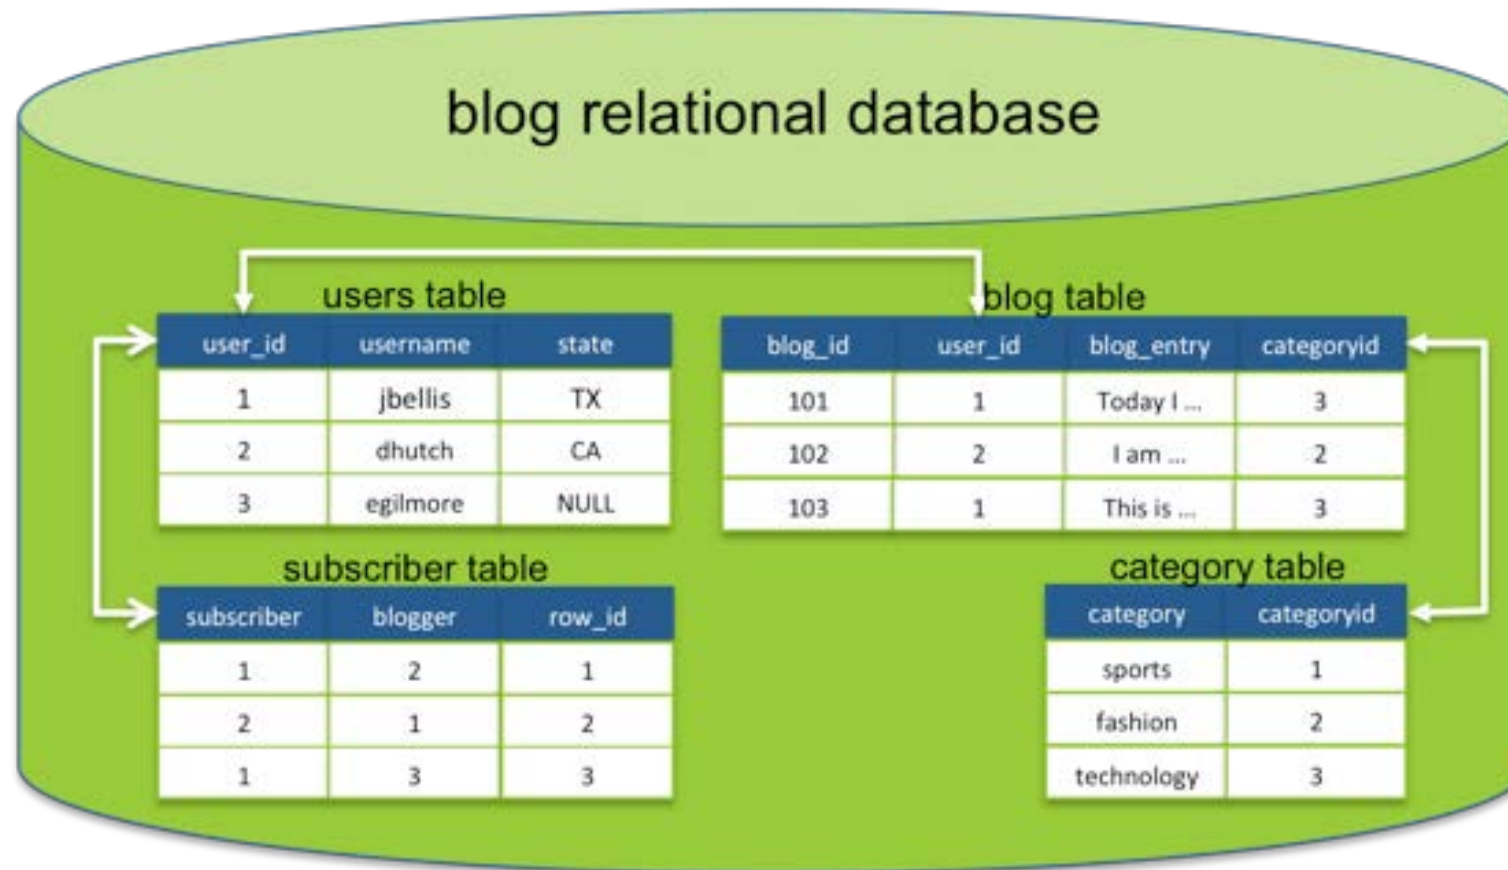
Key-Value Stores

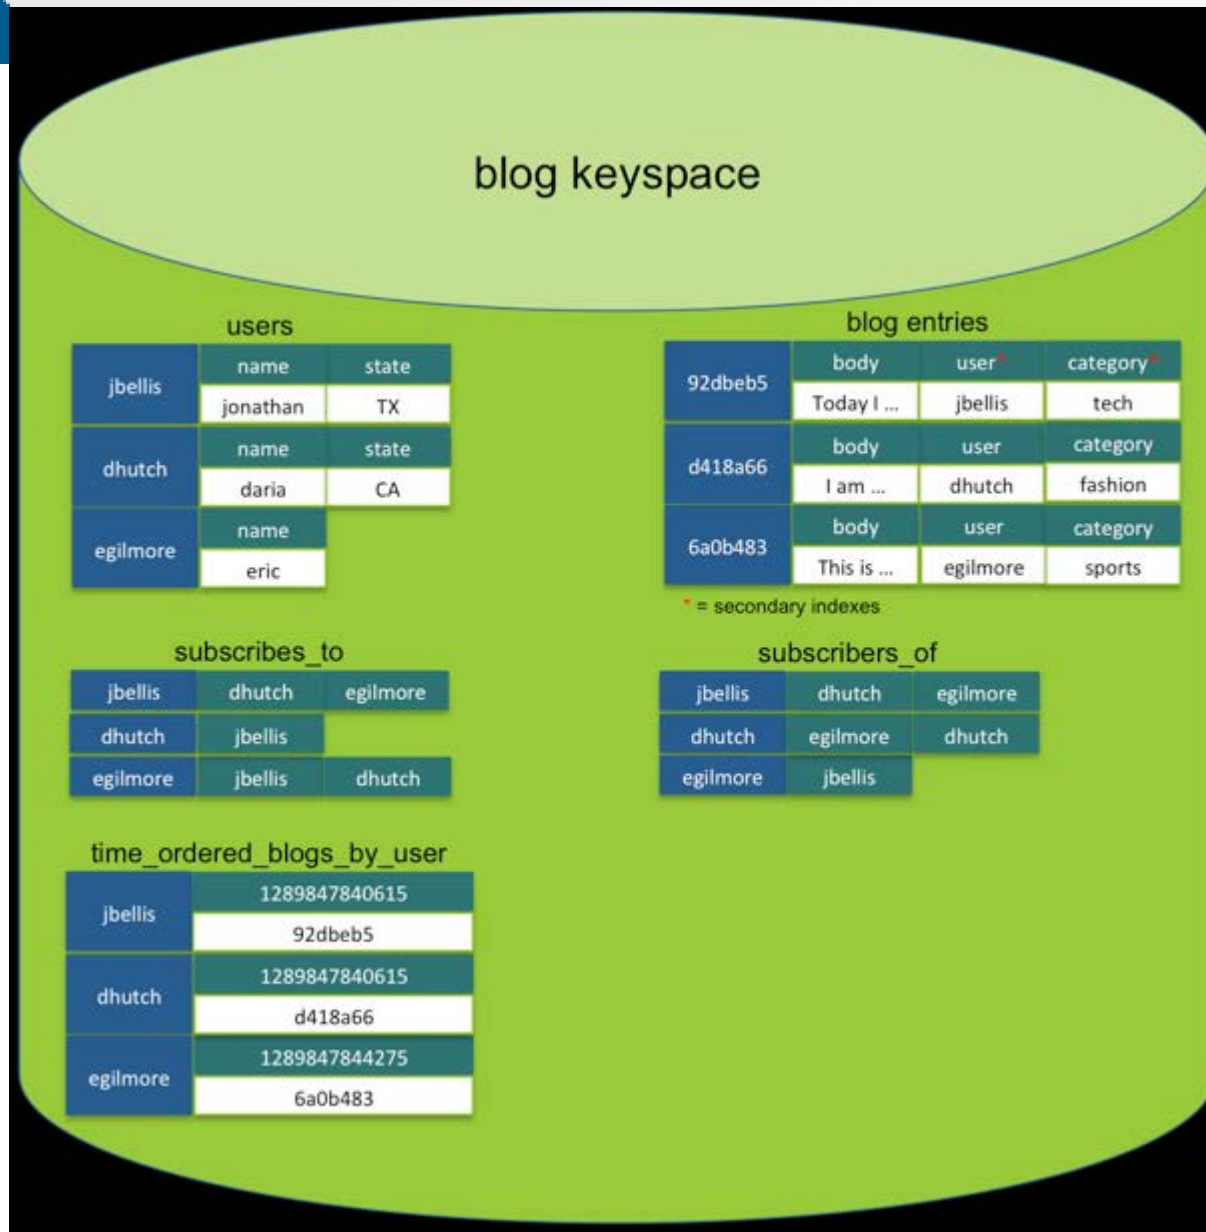
- Primary version of a DB!
- Like associative arrays
- Only for performance reasons



Example (Cassandra data model)

- Consider the relational DB





- One or more *column family* objects, analogous to tables
- Identified by an application-supplied row *key*
- Each row in a column family is *not* required to have the same set of columns
- There are no formal foreign keys and joining column families at query time is not supported

Cassandra Query Language

Examples:

- `INSERT INTO users (KEY, name, state)
VALUES ('jbellis', 'jonathan', 'TX');`
- `Select * from users where KEY='jbellis';`

DOCUMENT STORES

Document Stores

- Like Key-Value Stores except that “value” is a document
 - Data model: (key, document) pairs
 - Document: JSON, XML, other semistructured formats
 - Basic operations: Insert(key,document), Fetch(key), Update(key), Delete(key)
 - Also Fetch based on document contents

MongoDB example

A document representing a wiki article may look like the following in JSON notation:

```
{ title : " MongoDB " ,  
  last_editor : "172.5.123.91" ,  
  last_modified : new Date ("9/23/2010") ,  
  body : " MongoDB is a ..." ,  
  categories : [" Database " , " NoSQL " , " Document Database " ] ,  
  reviewed : false }
```

Insertion into a MongoDB collection:

```
db . < collection > . insert ( { title : " MongoDB " , last_editor : ... } ) ;
```

Retrieval examples:

```
db . < collection > . find ( { title : " MongoDB " } ) ;  
db . < collection > . find ( { categories : { $in : [" NoSQL " , " Document Databases " ] } } ) ;  
db . < collection > . find ( { categories : { $all : [" NoSQL " , " Document Databases " ] } } ) ;
```

GRAPH DATABASES

Graph Database Systems

- Data model: nodes and edges
- Nodes may have properties (including ID)
- Edges may have labels or roles
- Interfaces and query languages vary
- Example systems
 - Neo4j, FlockDB, Pregel, ...
- Aim to manage:
 - Multi-valued and/or complex attributes (Violations of the 1NF)
 - Dynamic data → accommodate changes
 - Data and query in a uniform way (to facilitate learning and reasoning)

Database Representation

- Sailors(sid:integer, sname:char(10), rating: integer, age:real)
- Boats(bid:integer, bname:char(10), color:char(10))
- Reserve(sid:integer, bid:integer, day:date)

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

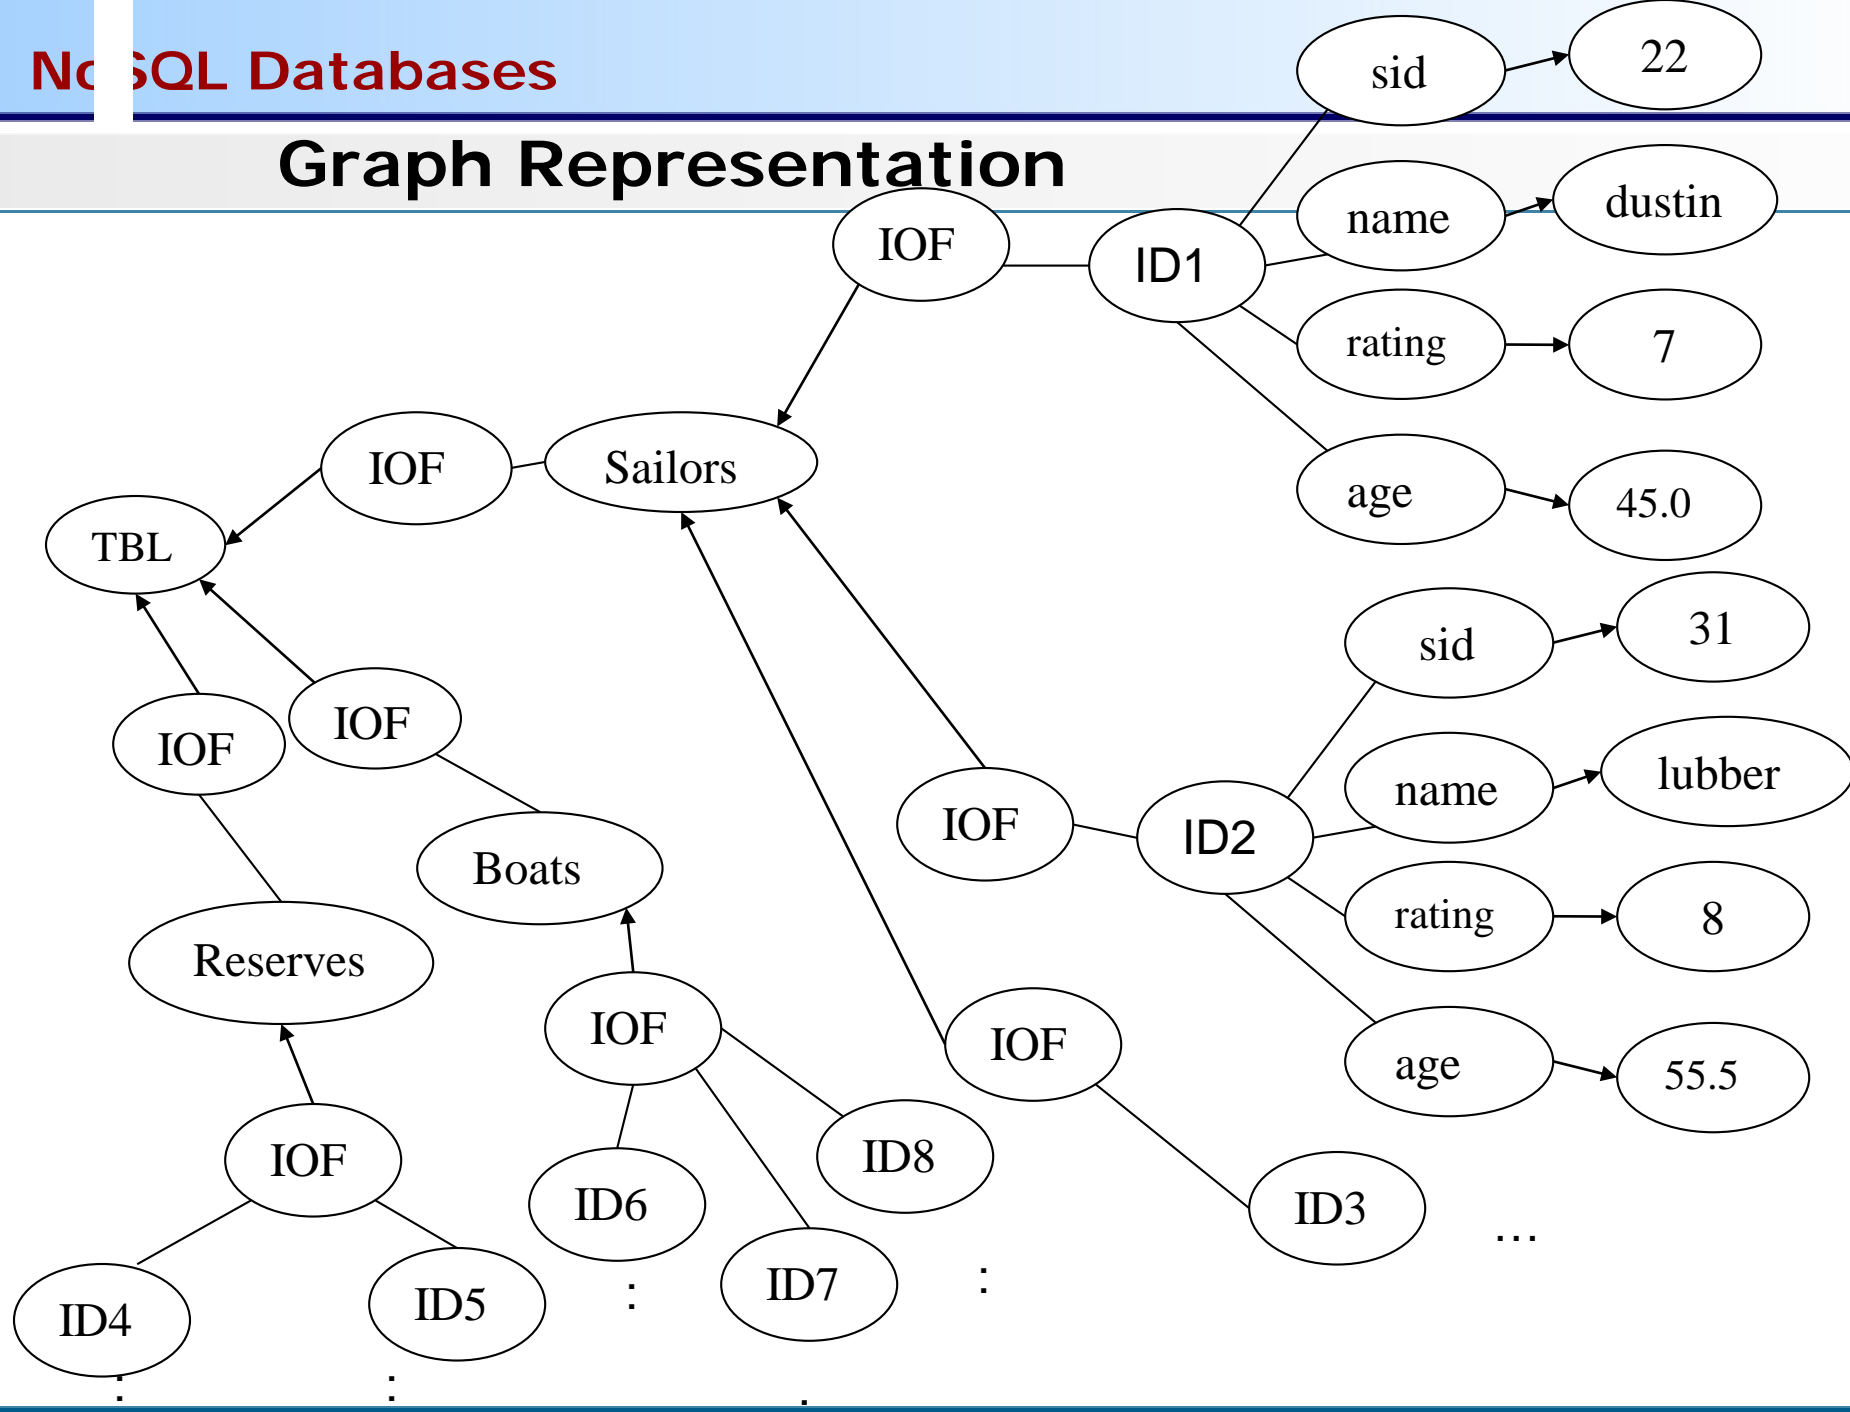
Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96


Boats

<u>bid</u>	bname	color
101	Interlake	red
102	Clipper	green
103	Marine	red

Graph Representation



Data Representation in the GDB DDL



- ID: (Name1=Val1,...,NameN=ValN)

Examples:

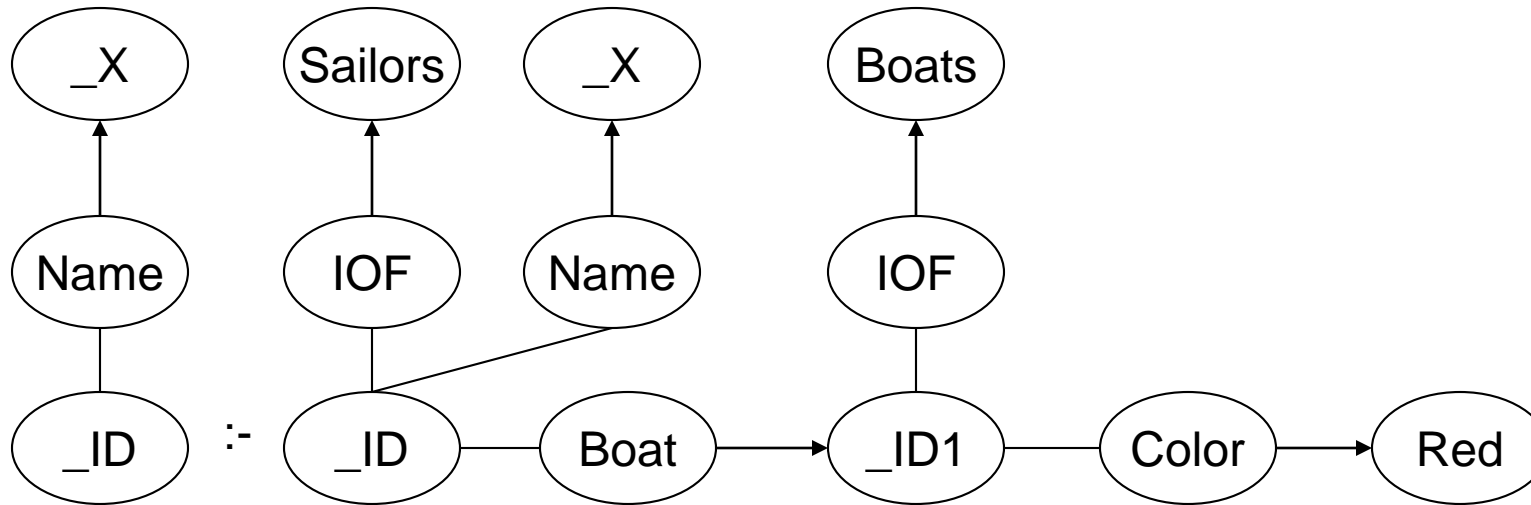
ID1: (sid=22, name="Dustin", rating=7, age=45.0)

ID4: (sailor=ID1, day="10/10/96", boat=ID6)

ID6: (bid=101, bname="Interlake", color="red")

[DML-QL] Writing simple queries:

- The names of all sailors who have reserved a red boat



$_ID:(Name = _X) :-$

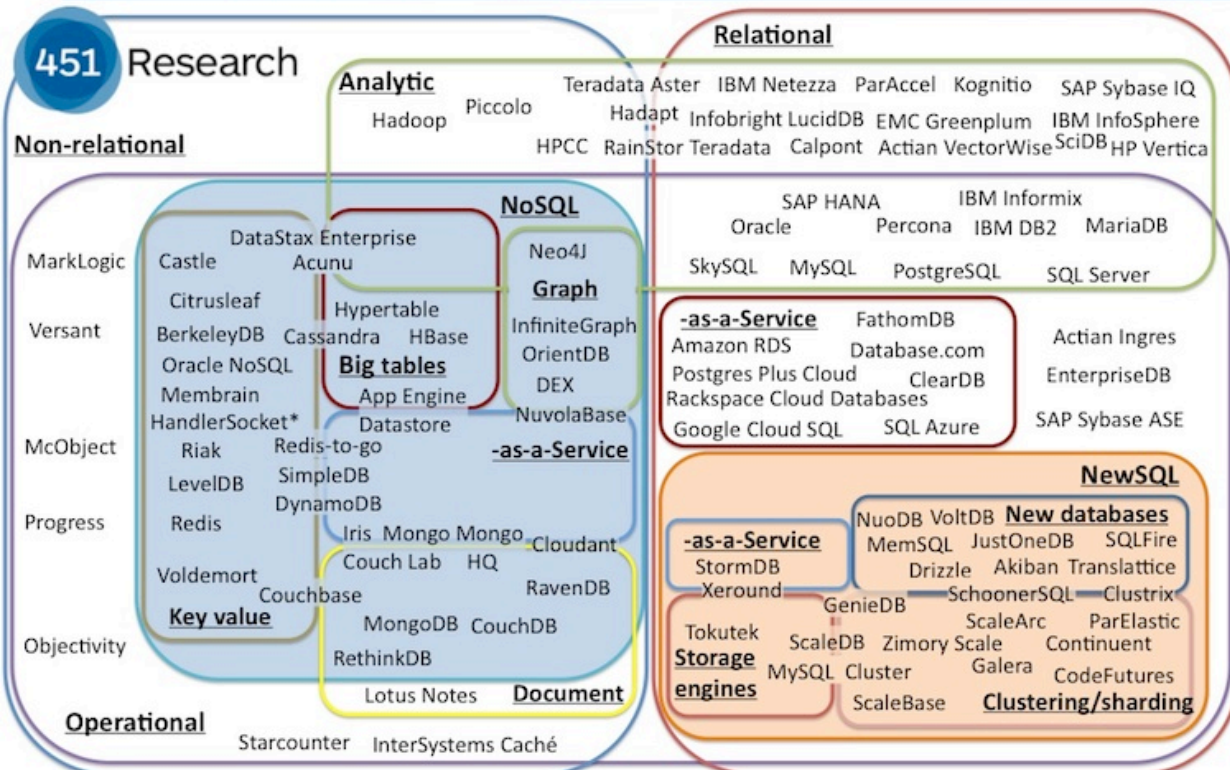
$_ID:(IOF = Sailors, Boat = _ID1, Name = _X),$

$_ID1:(IOF = Boats, Color = Red).$

The query is matched against the data graph



The evolving database landscape



© 2012 by The 451 Group. All rights reserved

