





Machine Learning

Linear Models for Regression



Marcello Restelli

March 10, 2020




Outline

- 1 Linear Regression
- 2 Minimizing Least Squares
- 3 Regularization
- 4 Bayesian Linear Regression

Regression Problems


- The **goal** of regression is to learn a **mapping** from input \mathbf{x} to a continuous output t



- Examples
 - Predict stock market price
 - Predict age of a web user
 - Predict effect of an actuation in robotics
 - Predict the value of a house
 - Predict the temperature in a building

Linear Models

- Many real processes can be **approximated** with linear models
- Linear regression often appears as a **module** of larger systems
- Linear problems can be solved **analytically**
- Linear prediction provides an introduction to many of the **core concepts** of machine learning
- Augmented with kernels, it can model **non-linear** relationships





Linear Function

- Linear function in the **parameters** \mathbf{w} :

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{D-1} w_j x_j = \mathbf{w}^T \mathbf{x}$$

- $\mathbf{x} = (1, x_1, \dots, x_{D-1})$
- w_0 is the offset



Loss Functions for Regression

- We need to quantify what it means to do **well or poorly** on a task
- We need to define a **loss** (error) function: $L(t, y(\mathbf{x}))$
- The **average**, or expected, loss is given by:


$$\mathbb{E}[L] = \int \int L(t, y(\mathbf{x})) p(\mathbf{x}, t) d\mathbf{x} dt$$

- A common choice is the **squared loss function**

$$\mathbb{E}[L] = \int \int (t - y(\mathbf{x}))^2 p(\mathbf{x}, t) d\mathbf{x} dt$$

- The optimal solution (if we assume a completely flexible function) is the **conditional average**:

$$y(\mathbf{x}) = \int t p(t|\mathbf{x}) dt = \mathbb{E}[t|\mathbf{x}]$$



Other Loss Functions

- Simple generalization of the squared loss, called the **Minkowski** loss:

$$\mathbb{E}[L] = \int \int |t - y(\mathbf{x})|^q p(\mathbf{x}, t) d\mathbf{x} dt$$

- The minimum of $\mathbb{E}[L]$ is given by:
 - the conditional mean for $q = 2$
 - the conditional median for $q = 1$
 - the conditional mode for $q \rightarrow 0$

Basis Functions

- To consider non-linear functions we can use non-linear **basis function**:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$


- $\boldsymbol{\phi}(\mathbf{x}) = (1, \phi_1(\mathbf{x}), \dots, \phi_{M-1}(\mathbf{x}))^T$

- Examples:

- Polynomial: $\phi_j(x) = x^j$
- Gaussian:

$$\phi_j(x) = \exp\left(-\frac{(x-\mu_j)^2}{2\sigma^2}\right)$$
- Sigmoidal:

$$\phi_j(x) = \frac{1}{1+\exp\left(\frac{\mu_j-x}{\sigma}\right)}$$



Basis Functions


- To consider non-linear functions we can use non-linear **basis function**:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

- $\boldsymbol{\phi}(\mathbf{x}) = (1, \phi_1(\mathbf{x}), \dots, \phi_{M-1}(\mathbf{x}))^T$

- Examples:

- Polynomial: $\phi_j(x) = x^j$
- Gaussian: $\phi_j(x) = \exp\left(-\frac{(x-\mu_j)^2}{2\sigma^2}\right)$
- Sigmoidal: $\phi_j(x) = \frac{1}{1+\exp\left(\frac{\mu_j-x}{\sigma}\right)}$



Basis Functions

- To consider non-linear functions we can use non-linear **basis function**:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$


- $\boldsymbol{\phi}(\mathbf{x}) = (1, \phi_1(\mathbf{x}), \dots, \phi_{M-1}(\mathbf{x}))^T$

- Examples:

- Polynomial: $\phi_j(x) = x^j$
- Gaussian:





$$\phi_j(x) = \exp\left(-\frac{(x-\mu_j)^2}{2\sigma^2}\right)$$
- Sigmoidal:

$$\phi_j(x) = \frac{1}{1+\exp\left(\frac{\mu_j-x}{\sigma}\right)}$$



Basis Functions

Example



Discriminative vs Generative

- **Generative approach:**

- Model the **joint density**: $p(\mathbf{x}, t) = p(\mathbf{x}|t)p(t)$
- Infer **conditional density**: $p(t|\mathbf{x}) = \frac{p(\mathbf{x}, t)}{p(\mathbf{x})}$
- Marginalize to find **conditional mean**: $\mathbb{E}[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt$

- **Discriminative approach:**

- Model **conditional density** $p(t|\mathbf{x})$
- Marginalize to find **conditional mean**: $\mathbb{E}[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt$

- **Direct approach**

- Find a regression function $y(\mathbf{x})$ directly from the training data



Minimizing Least Squares

- Given a data set with N samples, let us consider the following **error (loss) function**

$$L(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

- This is (half) the **residual sum of squares (RSS)**, a.k.a. **sum of squared errors (SSE)**
- It can also be written as the sum of the ℓ_2 -norm of the vector of **residual errors**

$$RSS(\mathbf{w}) = \|\epsilon\|_2^2 = \sum_{i=1}^N \epsilon_i^2$$



Ordinary Least Squares

Closed-Form Optimization

- Let's write RSS in matrix form with $\Phi = (\phi(\mathbf{x}_1), \dots, \phi(\mathbf{x}_N))^T$ and $\mathbf{t} = (t_1, \dots, t_N)^T$

$$L(\mathbf{w}) = \frac{1}{2}RSS(\mathbf{w}) = \frac{1}{2}(\mathbf{t} - \Phi\mathbf{w})^T (\mathbf{t} - \Phi\mathbf{w})$$

- Compute first and second derivative

$$\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} = -\Phi^T (\mathbf{t} - \Phi\mathbf{w}); \quad \frac{\partial^2 L(\mathbf{w})}{\partial \mathbf{w} \partial \mathbf{w}^T} = \Phi^T \Phi$$

- Assuming $\Phi^T \Phi$ in **nonsingular**

$$\hat{\mathbf{w}}_{OLS} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- Complexity $O(NM^2 + M^3)$
 - Cholesky: $M^3 + NM^2/2$
 - QR: NM^2

Gradient Optimization

- Closed-form solution is not practical with **big data**
- We can use **sequential (online)** updates
- **Stochastic gradient descent**
 - If the loss function can be expressed as a **sum over samples** ($L(\mathbf{x}) = \sum_n L(x_n)$)

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha^{(k)} \nabla L(x_n)$$

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha^{(k)} \left(\mathbf{w}^{(k)T} \phi(\mathbf{x}_n) - t_n \right) \phi(\mathbf{x}_n)$$

- where k is the iteration and α is a **learning rate**
- For **convergence** the learning rate has to satisfy

$$\sum_{k=0}^{\infty} \alpha^{(k)} = +\infty$$

$$\sum_{k=0}^{\infty} \alpha^{(k)^2} < +\infty$$

Geometric Interpretation

- \mathbf{t} is an N -dimensional vector
- Let's denote with φ_j the j -th column of Φ
- Define $\hat{\mathbf{t}}$ the N -dimensional vector, whose n -th element is $y(\mathbf{x}_n, \mathbf{w})$
- $\hat{\mathbf{t}}$ is a linear combination of $\varphi_1, \dots, \varphi_M$
- so $\hat{\mathbf{t}}$ lies in an M -subspace \mathcal{S}
- Since $\hat{\mathbf{t}}$ minimizes the SSE with respect to \mathbf{t} , it represents the projections of \mathbf{t} onto the subspace \mathcal{S}


$$\hat{\mathbf{t}} = \Phi \hat{\mathbf{w}} = \Phi (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- $H = \Phi (\Phi^T \Phi)^{-1} \Phi^T$ is called the **hat matrix**

Geometric Example

- Assume $N = 3$ and $M = D = 2$

$$\Phi = \mathbf{X} = \begin{pmatrix} 1 & 2 \\ 1 & -2 \\ 1 & 2 \end{pmatrix} \quad \mathbf{t} = \begin{pmatrix} 5 \\ 1 \\ 2 \end{pmatrix} \quad \hat{\mathbf{t}} = \begin{pmatrix} 3.5 \\ 1 \\ 3.5 \end{pmatrix}$$



Maximum Likelihood (ML)

- The output variable t can be modeled as a deterministic function y of the input \mathbf{x} and a random noise ϵ

$$t = f(\mathbf{x}) + \epsilon$$

- We want to approximate $f(\mathbf{x})$ with $y(\mathbf{x}, \mathbf{w})$
- We assume $\epsilon \sim \mathcal{N}(0, \sigma^2)$
- Given N samples, with inputs $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ and outputs $\mathbf{t} = (t_1, \dots, t_N)^T$, the likelihood function is

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \sigma^2)$$

Maximum Likelihood (ML)

- Assuming the samples to be **independent and identically distributed (iid)**, we consider the log-likelihood:

$$\begin{aligned}\ell(\mathbf{w}) &= \ln p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \sigma^2) = \sum_{n=1}^N \ln p(t_n|\mathbf{x}_n, \mathbf{w}, \sigma^2) \\ &= -\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} RSS(\mathbf{w})\end{aligned}$$

- To find the maximum likelihood, we equal the gradient to zero

$$\begin{aligned}\nabla \ell(\mathbf{w}) &= \sum_{n=1}^N t_n \phi(\mathbf{x}_n)^T - \mathbf{w}^T \left(\sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T \right) = 0 \\ \mathbf{w}_{ML} &= (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}\end{aligned}$$

Variance of the Parameters

- We assume
 - the observation t_i are **uncorrelated** and have **constant variance** σ^2
 - the x_i are fixed (non random)
- The variance-covariance matrix of the least-squares estimates is

$$\text{Var}(\hat{\mathbf{w}}_{OLS}) = (\Phi^T \Phi)^{-1} \sigma^2$$

- Usually, the variance σ^2 is estimated by

$$\hat{\sigma}^2 = \frac{1}{N - M} \sum_{n=1}^N (t_n - \hat{\mathbf{w}}^T \phi(\mathbf{x}_n))^2$$

- Assuming that the model is linear in the features $\phi_1(), \dots, \phi_M()$ and that the noise is additive and Gaussian

$$\hat{\mathbf{w}} \sim \mathcal{N}(\mathbf{w}, (\Phi^T \Phi)^{-1} \sigma^2) \quad (N - M) \hat{\sigma}^2 \sim \sigma^2 \chi_{N-M}^2$$

- Such properties can be used to form **test hypothesis** and **confidence intervals**

Gauss-Markov Theorem

Theorem (Gauss-Markov)

*The least squares estimate of \mathbf{w} has the **smallest variance** among all linear **unbiased** estimates.*

- It follows that least squares estimator has the **lowest MSE** of all linear estimator with **no bias**
- However, there may exist a **biased** estimator with **smaller MSE**

Multiple Outputs

- Let now consider the case of **multiple outputs**
- We could use a **different** set of basis functions for each output, thus having **independent** regression problems
- Usually, a **single** set of basis functions is considered


$$\hat{\mathbf{W}}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{T}$$

- For each output t_k we have


$$\hat{\mathbf{w}}_k = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}_k$$

- where \mathbf{t}_k is an N -dimensional column vector
- The solution **decouples** between the different outputs
- The pseudo-inverse $\Phi^\dagger = (\Phi^T \Phi)^{-1} \Phi^T$ needs to be computed only **once**


Increasing Model Complexity: Quadratic Function




Increasing Model Complexity: Sinusoidal Function




Increasing Model Complexity: Sinusoidal Function



Increasing Model Complexity: Sinusoidal Function



Increasing Model Complexity: Sinusoidal Function




Under-fitting vs Over-Fitting

- With low-order polynomials we have **under-fitting**
- With high-order polynomials we get excellent fit over the training data, but a poor representation of the true function: **over-fitting**
- We want to have good **generalization**



- We use a **test set** of 100 samples to evaluate generalization

- $$E_{RMS} = \sqrt{\frac{2 * RSS(\hat{\mathbf{w}})}{N}}$$



How to Avoid Over-fitting?

- This is the problem of **model selection** (we will see this later)
- What happens when the number of training samples **increases**?



How to Avoid Over-fitting?

What happens to the parameters when the model gets more complex?

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
\hat{w}_0	0.19	0.82	0.31	0.35
\hat{w}_1		-1.27	7.99	232.37
\hat{w}_2			-25.43	-5321.83
\hat{w}_3				48568.31
\hat{w}_4				-231639.30
\hat{w}_5				640042.26
\hat{w}_6				-1061800.52
\hat{w}_7				1042400.18
\hat{w}_8				-557682.99
\hat{w}_9				125201.43

Ridge Regression

- One way to reduce the MSE is to change the **loss function** as follows

$$L(\mathbf{w}) = L_D(\mathbf{w}) + \lambda L_W(\mathbf{w})$$

- $L_D(\mathbf{w})$: error on data (e.g., RSS)
- $L_W(\mathbf{w})$: model complexity
- By taking $L_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} = \frac{1}{2} \|\mathbf{w}\|_2^2$ we get

$$L(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (t_i - \mathbf{w}^T \phi(\mathbf{x}_i))^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

- It is called **ridge regression** (or weight decay)
- It is a **regularization** (or parameter shrinkage) method
- The loss function is still quadratic in \mathbf{w} :


$$\hat{\mathbf{w}}_{ridge} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

Ridge Regression

Quadratic Example

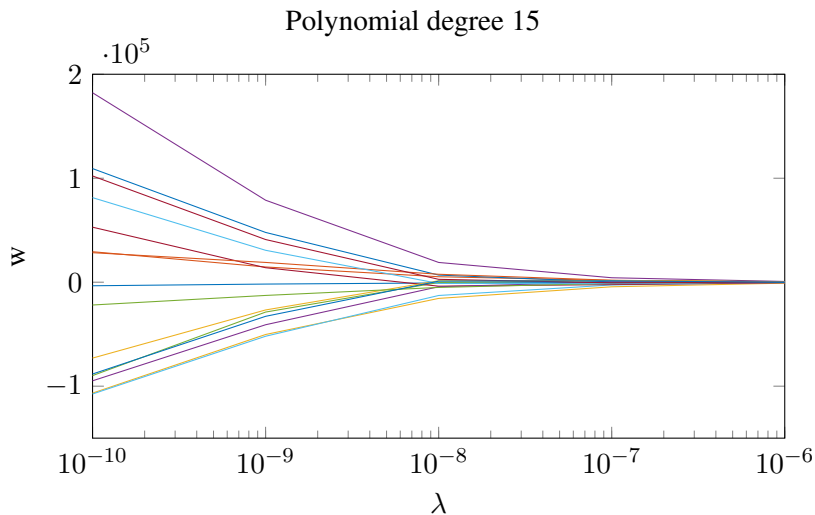
Polynomial degree 15

$$\|\mathbf{w}\|_2 = 2.74\text{e}+6$$




Ridge Regression

Quadratic Example: Weights



Ridge Regression

Sinusoidal Example



Lasso


- Another popular regularization method is **lasso**

$$L(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (t_i - \mathbf{w}^T \phi(\mathbf{x}_i))^2 + \frac{\lambda}{2} \|\mathbf{w}\|_1$$


- where $\|\mathbf{w}\|_1 = \sum_{j=1}^M |w_j|$
- Differently from ridge, lasso is **nonlinear** in the t_i and **no closed-form** solution exists (quadratic programming problem)
- Nonetheless, it has the advantage of making some weights equal to **zero** for values of λ sufficiently large
- Lasso yields **sparse** models

Lasso vs Ridge Regression

Ridge



Lasso



Lasso tends to generate **sparser solutions** than quadratic regularizer

Bayesian Approach

- We formulate our knowledge about the world in a **probabilistic way**
 - We define the **model** that expresses our knowledge **qualitatively**
 - Our model will have some **unknown parameters**
 - We capture our assumptions about unknown parameters by specifying the **prior distribution** over those parameters before seeing the data
- We **observe the data**
- We compute the **posterior probability distribution** for the parameters, given observed data
- We use the posterior distribution to:
 - **Make predictions** by averaging over the posterior distribution
 - **Examine/Account for uncertainty** in the parameter values
 - **Make decisions** by minimizing expected posterior loss

Posterior Distribution

- The **posterior** distribution for the model parameters can be found by combining the **prior** with the **likelihood** for the parameters given data
- This is accomplished using **Bayes' Rule**:

$$P(\text{parameters}|\text{data}) = \frac{P(\text{data}|\text{parameters})P(\text{parameters})}{P(\text{data})}$$

$$p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w})P(\mathbf{w})}{P(\mathcal{D})}$$

where

- $p(\mathbf{w}|\mathcal{D})$ is the **posterior** probability of parameters \mathbf{w} given training data \mathcal{D}
- $p(\mathcal{D}|\mathbf{w})$ is the probability (**likelihood**) of observing \mathcal{D} given \mathbf{w}
- $P(\mathbf{w})$ is the **prior** probability over the parameters
- $P(\mathcal{D})$ is the marginal likelihood (**normalizing constant**):

$$P(\mathcal{D}) = \int p(\mathcal{D}|\mathbf{w})P(\mathbf{w})d\mathbf{w}$$
- Stating Bayes' rule in words: posterior \propto likelihood \times prior
- We want the most probable value of \mathbf{w} given the data: **maximum a posteriori** (MAP). It is the **mode** of the posterior.

Bayesian Linear Regression

- Another approach to avoid the **over-fitting** problem of ML is to use **Bayesian** linear regression
- In the Bayesian approach the parameters of the model are considered as drawn from some **distribution**
- Assuming Gaussian likelihood model, the **conjugate prior** is Gaussian too

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{w}_0, \mathbf{S}_0)$$

- Given the data \mathcal{D} , the **posterior** is still Gaussian

$$p(\mathbf{w}|\mathbf{t}, \Phi, \sigma^2) \propto \mathcal{N}(\mathbf{w}|\mathbf{w}_0, \mathbf{S}_0) \mathcal{N}(\mathbf{t}|\Phi\mathbf{w}, \sigma^2\mathbf{I}_N) = \mathcal{N}(\mathbf{w}|\mathbf{w}_N, \mathbf{S}_N)$$

$$\mathbf{w}_N = \mathbf{S}_N \left(\mathbf{S}_0^{-1} \mathbf{w}_0 + \frac{\Phi^T \mathbf{t}}{\sigma^2} \right)$$

$$\mathbf{S}_N^{-1} = \mathbf{S}_0^{-1} + \frac{\Phi^T \Phi}{\sigma^2}$$



- For **sequential** data, the posterior acts as prior for the next iteration

Relation to Ridge Regression

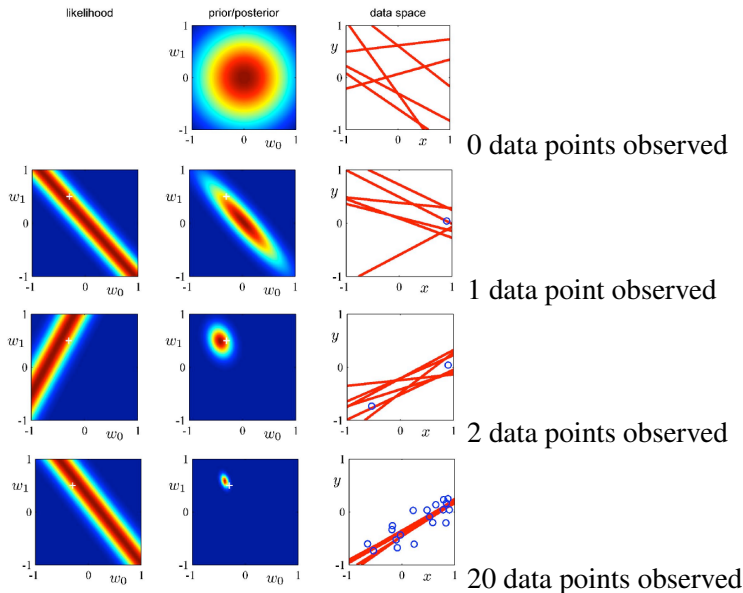
- In Gaussian distributions the **mode** coincides with the **mean**
- It follows that \mathbf{w}_N is the **MAP estimator** (Maximum a posteriori)
- If the prior has infinite variance, \mathbf{w}_N reduces to the **ML estimator**
- If $\mathbf{w}_0 = 0$ and $\mathbf{S}_0 = \tau^2 \mathbf{I}$, then \mathbf{w}_N reduces to the **ridge estimate**, where $\lambda = \frac{\sigma^2}{\tau^2}$

1D Example

- Data generated from: $t(x) = -0.3 + 0.5x + \epsilon$, where $\epsilon \sim \mathcal{N}(0, 0.04)$
- x values taken uniformly from $[-1, 1]$
- Model: $y(x, \mathbf{w}) = w_0 + w_1x$
- We assume to know $\sigma^2 = 0.04$ and $\tau^2 = 0.5$



1D Example



Predictive Distribution

- We are interested in the **posterior predictive distribution**

$$\begin{aligned}
 p(t|\mathbf{x}, \mathcal{D}, \sigma^2) &= \int \mathcal{N}(t|\mathbf{w}^T \phi(\mathbf{x}), \sigma^2) \mathcal{N}(\mathbf{w}|\mathbf{w}_N, \mathbf{S}_N) d\mathbf{w} \\
 &= \mathcal{N}(t|\mathbf{w}_N^T \phi(\mathbf{x}), \sigma_N^2(\mathbf{x})) \\
 \sigma_N^2(\mathbf{x}) &= \sigma^2 + \phi(\mathbf{x})^T \mathbf{S}_N \phi(\mathbf{x})
 \end{aligned}$$

where


$$\sigma_N^2(\mathbf{x}) = \underbrace{\sigma^2}_{\text{noise in the target values}} + \underbrace{\phi(\mathbf{x})^T \mathbf{S}_N \phi(\mathbf{x})}_{\text{Uncertainty associated with parameter values}}$$

- In the limit, as $N \rightarrow \infty$, **the second term goes to zero**
- The variance of the predictive distribution arises only from the additive noise governed by parameter σ




Example

Sinusoidal dataset, 9 Gaussian basis functions

Predictive distribution



Samples from the posterior



Modeling Challenges

- The first challenge is in **specifying suitable model** and **suitable prior distributions**
 - A suitable model should admit all the possibilities that thought to be at all **likely**
 - A suitable prior should **avoid** giving zero or very small probabilities to **possible events**, but should also avoid spreading out the probability over all possibilities
- To avoid uninformative priors, we may need to model **dependencies** between parameters
- One strategy is to introduce **latent variables** into the model and **hyperparameters** into the prior
- Both of these represent the ways of modeling dependencies in a **tractable way**

Computational Challenges

The other big challenge is **computing the posterior distribution**. There are several approaches:

- **Analytical integration:** If we use “**conjugate priors**”, the posterior distribution can be computed **analytically**. Only works for **simple models**
- **Gaussian (Laplace) approximation:** Approximate the posterior distribution with a Gaussian. Works well when there a lot of data compared to the model complexity
- **Monte Carlo integration:** Once we have a sample from the posterior distribution, we can do many things. Currently, the common approach is Markov Chain Monte Carlo (MCMC), that consists in simulating a Markov chain that converges to the posterior distribution
- **Variational approximation:** A cleverer way of approximating the posterior. It is usually faster than MCMC, but it is less general

Pros and Cons of Fixed Basis Functions

- **Advantages**

- Closed-form solution
- Tractable Bayesian treatment
- Arbitrary non-linearity with the proper basis functions

- **Limitations**

- Basis functions are chosen independently from the training set
- Curse of dimensionality