

Machine Learning

Solving Markov Decision Processes




Marcello Restelli

May 28, 2017

0.51	0.72	0.84	1.00
0.27		0.95	-1.00
0.00	0.22	0.37	0.13

VALUES AFTER 5 ITERATIONS



Outline

- 1 Policy Search
- 2 Dynamic Programming
 - Policy Iteration
 - Value Iteration
- 3 Linear Programming

Brute Force

- Solving an MDP means finding an **optimal policy**
- A **naïve** approach consists of
 - **enumerating** all the deterministic Markov policies
 - **evaluate** each policy
 - **return** the best one
- The number of policies is **exponential**: $|\mathcal{A}|^{|\mathcal{S}|}$
- Need a **more intelligent search** for best policies
 - **restrict the search** to a subset of the possible policies
 - using **stochastic optimization** algorithms

What is Dynamic Programming?

- **Dynamic:** sequential or temporal component to the problem
- **Programming:** optimizing a “program”, i.e., a policy
 - c.f. linear programming
- A method for solving **complex** problems
- By breaking them down into **subproblems**
 - **Solve** the subproblems
 - **Combine** solutions to subproblems

Requirements for Dynamic Programming

- Dynamic Programming is a **very general** solution method for problems which have **two properties**:
 - **Optimal substructure**
 - **Principle of optimality** applies
 - Optimal solution can be decomposed into **subproblems**
 - **Overlapping subproblems**
 - Subproblems **recur** many times
 - Solutions can be **cached** and **reused**
- Markov decision processes satisfy both properties
 - **Bellman equation** gives recursive decomposition
 - **Value function** stores and reuses solutions

Planning by Dynamic Programming

- Dynamic Programming assumes **full knowledge** of the MDP
- It is used for **planning** in an MDP
- **Prediction**
 - Input: MDP $\langle \mathcal{S}, \mathcal{A}, P, R, \gamma, \mu \rangle$ and policy π (i.e., MRP $\langle \mathcal{S}, P^\pi, R^\pi, \gamma, \mu \rangle$)
 - Output: value function V^π
- **Control**
 - Input: MDP $\langle \mathcal{S}, \mathcal{A}, P, R, \gamma, \mu \rangle$
 - Output: value function V^* and optimal policy π^*

Other Applications of Dynamic Programming

Dynamic Programming is used to solve many other problems:

- Scheduling algorithms
- String algorithms (e.g., sequence alignment)
- Graph algorithms (e.g., shortest path algorithms)
- Graphical models (e.g., Viterbi algorithm)
- Bioinformatics (e.g., lattice models)

Finite–Horizon Dynamic Programming

- **Principle of optimality:** the tail of an optimal policy is optimal for the “tail” problem
- **Backward induction**
 - **Backward recursion**

$$V_k^*(s) = \max_{a \in \mathcal{A}_k} \left\{ R_k(s, a) + \sum_{s' \in \mathcal{S}_{k+1}} P_k(s'|s, a) V_{k+1}^*(s') \right\}, \quad k = N-1, \dots, 0$$

- **Optimal policy**

$$\pi_k^*(s) \in \arg \max_{a \in \mathcal{A}_k} \left\{ R_k(s, a) + \sum_{s' \in \mathcal{S}_{k+1}} P_k(s'|s, a) V_{k+1}^*(s') \right\}, \quad k = 0, \dots, N-1$$

- **Cost:** $N|\mathcal{S}||\mathcal{A}|$ vs $|\mathcal{A}|^{N|\mathcal{S}|}$ of brute force policy search
- From now on, we will consider **infinite–horizon discounted** MDPs

Policy Evaluation

- For a **given policy** π compute the **state-value function** V^π
- Recall
 - State-value function for policy π :

$$V^\pi(s) = \mathbb{E} \left\{ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right\}$$

- **Bellman equation** for V^π :

$$V^\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^\pi(s') \right]$$

- A **system** of $|\mathcal{S}|$ simultaneous **linear equations**
- Solution in **matrix** notation (complexity $O(n^3)$):

$$V^\pi = (I - \gamma P^\pi)^{-1} R^\pi$$

Iterative Policy Evaluation


- Iterative application of Bellman expectation backup
- $V_0 \rightarrow V_1 \rightarrow \dots \rightarrow V_k \rightarrow V_{k+1} \rightarrow \dots \rightarrow V^\pi$
- A **full policy-evaluation backup**:

$$V_{k+1}(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a|s) \left[R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k(s') \right]$$

- A **sweep** consists of applying a backup operation to each state
- Using **synchronous** backups
 - At each iteration $k + 1$
 - For all states $s \in \mathcal{S}$
 - Update $V_{k+1}(s)$ from $V_k(s')$

Example

Small Gridworld



actions

	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$
on all transitions

- **Undiscounted episodic MDP**

- $\gamma = 1$
- All episodes terminate in **absorbing** terminal state


- **Transient** states $1, \dots, 14$

- One **terminal** state (shown twice as shaded squares)

- Actions that would take agent off the grid leave state **unchanged**

- Reward is -1 until the terminal state is reached


Policy Evaluation in Small Gridworld




Policy Evaluation in Small Gridworld

 $k = 3$


0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0


 $k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0


 $k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



optimal
policy

Policy Improvement

- Consider a **deterministic policy** π
- For a given state s , would it **better** to do an action $a \neq \pi(s)$?
- We can **improve** the policy by acting greedily

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} Q^\pi(s, a)$$

- This improves the value from **any** state s over one step

$$Q^\pi(s, \pi'(s)) = \max_{a \in \mathcal{A}} Q^\pi(s, a) \geq Q^\pi(s, \pi(s)) = V^\pi(s)$$

Policy Improvement Theorem

Theorem

Let π and π' be any pair of deterministic policies such that

$$Q^\pi(s, \pi'(s)) \geq V^\pi(s) \quad , \quad \forall s \in \mathcal{S}$$

Then the policy π' must be as good as, or better than π

$$V^{\pi'}(s) \geq V^\pi(s) \quad , \quad s \in \mathcal{S}$$

Proof.

$$\begin{aligned} V^\pi(s) &\leq Q^\pi(s, \pi'(s)) = \mathbb{E}_{\pi'} [r_{t+1} + \gamma V^\pi(s_{t+1}) | s_t = s] \\ &\leq \mathbb{E}_{\pi'} [r_{t+1} + \gamma Q^\pi(s_{t+1}, \pi'(s_{t+1})) | s_t = s] \\ &\leq \mathbb{E}_{\pi'} [r_{t+1} + \gamma r_{t+2} + \gamma^2 Q^\pi(s_{t+2}, \pi'(s_{t+2})) | s_t = s] \\ &\leq \mathbb{E}_{\pi'} [r_{t+1} + \gamma r_{t+2} + \dots | s_t = s] = V^{\pi'}(s) \end{aligned}$$



Policy Iteration

- What if improvements **stops** ($V^{\pi'} = V^{\pi}$)?

$$Q^{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} Q^{\pi}(s, a) = Q^{\pi}(s, \pi(s)) = V^{\pi}(s)$$


- But this is the **Bellman optimality equation**
- Therefore $V^{\pi}(s) = V^{\pi'}(s) = V^*(s)$ for all $s \in \mathcal{S}$
- So π is an **optimal** policy!

$$\pi_0 \rightarrow V^{\pi_0} \rightarrow \pi_1 \rightarrow V^{\pi_1} \rightarrow \dots \rightarrow \pi^* \rightarrow V^* \rightarrow \pi^*$$


Modified Policy Iteration

- Does policy evaluation **need to converge** to V^π ?
- Or should we introduce a **stopping condition**
 - e.g., ϵ -convergence of value function
- Or simply **stop after k iterations** of iterative policy evaluation?
- For example, in the small gridworld $k = 3$ was sufficient to achieve optimal policy
- Why not update policy **every iteration**? i.e. stop after $k = 1$

Generalized Policy Iteration



- **Policy evaluation:** Estimate V^π
 - e.g., Iterative policy evaluation
- **Policy improvement:** Generate $\pi' \geq \pi$
 - e.g., Greedy policy improvement



Value Iteration

- **Problem:** find optimal policy π
- **Solution:** iterative application of Bellman optimality backup
- $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V^*$
- Using **synchronous backups**
 - At each iteration $k + 1$
 - For all states $s \in \mathcal{S}$
 - Update $V_{k+1}(s)$ from $V_k(s')$
- Unlike policy iteration there is **no explicit policy**
- **Intermediate** value functions **may not correspond** to any policy

Value Iteration demo:

<http://www.cs.ubc.ca/~poole/demos/mdp/vi.html>

Convergence and Contractions

Define the max-norm: $\|V\|_\infty = \max_s |V(s)|$

Theorem

Value Iteration converges to the optimal state-value function $\lim_{k \rightarrow \infty} V_k = V^$*

Proof.

$$\|V_{k+1} - V^*\|_\infty = \|T^*V_k - T^*V^*\|_\infty \leq \gamma \|V_k - V^*\|_\infty \leq \dots \leq \gamma^{k+1} \|V_0 - V^*\|_\infty \rightarrow 0$$



Theorem

$$\|V_{i+1} - V_i\|_\infty < \epsilon \Rightarrow \|V_{i+1} - V^*\|_\infty < \frac{2\epsilon\gamma}{1-\gamma}$$

Synchronous Dynamic Programming Algorithms

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Policy Evaluation (Iterative)
Control	Bellman Expectation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- Algorithms are based on **state-value function** $V^\pi(s)$ or $V^*(s)$
- Complexity $O(mn^2)$ **per iteration**, for m actions and n states
- Could also apply to **action-value function** $Q^\pi(s, a)$ or $Q^*(s, a)$
- Complexity $O(m^2n^2)$ **per iteration**

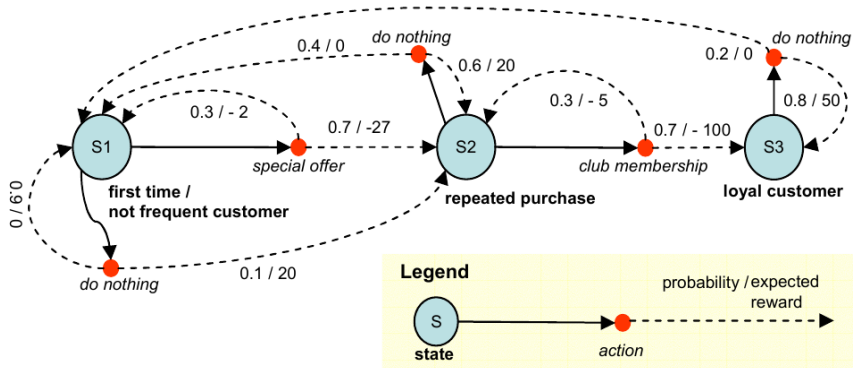
Efficiency of DP

- To find optimal policy is **polynomial** in the number of states...
- **but**, the number of states is often astronomical, e.g., often growing **exponentially** with the number of state variables: **curse of dimensionality**
- In practice, classical DP can be applied to problems with a few millions states
- **Asynchronous DP** can be applied to larger problems, and appropriate for parallel computation
- It is surprisingly **easy** to come up with MDPs for which methods are not practical

Complexity of DP

- DP methods are **polynomial time** algorithms for **fixed-discounted** MDPs
- **Value Iteration:** $O(|S|^2|\mathcal{A}|)$ for each iteration
- **Policy Iteration:** Cost of policy evaluation + Cost of policy iteration
 - Policy evaluation:
 - Linear system of equations: $O(|S|^3)$ or $O(|S|^{2.373})$
 - Iterative: $O\left(|S|^2 \frac{\log(\frac{1}{\epsilon})}{\log(\frac{1}{\gamma})}\right)$
 - Policy improvement: recently proven to be $O\left(\frac{|\mathcal{A}|}{1-\gamma} \log\left(\frac{|S|}{1-\gamma}\right)\right)$
- **Each iteration** of PI is computationally **more expensive** than each iteration of VI
- PI typically requires fewer iterations to converge than VI
- **Exponentially faster** than any **direct policy search**
- Number of states often **grows exponentially** with the number of state variables

Exercise



Infinite Horizon Linear Programming

- Recall, at value iteration convergence we have

$$\forall s \in \mathcal{S} : \quad V^*(s) = \max_{a \in \mathcal{A}} \left\{ R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right\}$$

- LP formulation to find V^* :

$$\begin{aligned} \min_V \quad & \sum_{s \in \mathcal{S}} \mu(s) V(s) \\ \text{s. t.} \quad & V(s) \geq R(s, a) + \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s \in \mathcal{S}, \forall a \in \mathcal{A} \end{aligned}$$

- $|\mathcal{S}|$ variables
- $|\mathcal{S}| |\mathcal{A}|$ constraints

Theorem

V^* is the solution of the above LP.

Theorem Proof

Let T^* be the **optimal Bellman operator**, then the LP can be written as:

$$\begin{array}{ll} \min_V & \mu^\top V \\ \text{s. t.} & V \geq T^*(V) \end{array}$$

- **Monotonicity property:** if $U \geq V$ then $T^*(U) \geq T^*(V)$.
- Hence, if $V \geq T^*(V)$ then $T^*(V) \geq T^*(T^*(V))$, and by **repeated application**, $V \geq T^*(V) \geq T^{*2}(V) \geq T^{*3}(V) \geq \dots \geq T^{*\infty}(V) = V^*$
- Any **feasible solution** to the LP must satisfy $V \geq T^*(V)$, and hence must satisfy $V \geq V^*$
- Hence, assuming all entries μ are positive, V^* is the **optimal solution** to the LP

Dual Linear Program

$$\begin{aligned}
 & \max_{\lambda} \quad \sum_{s \in \mathcal{S}} \sum_{a \in \mathcal{A}} \lambda(s, a) R(s, a) \\
 \text{s. t.} \quad & \sum_{a' \in \mathcal{A}} \lambda(s', a') = \mu(s) + \gamma \sum_{s \in \mathcal{S}} \sum_{a \in \mathcal{A}} \lambda(s, a) P(s' | s, a), \quad \forall s' \in \mathcal{S} \\
 & \lambda(s, a) \geq 0, \quad \forall s \in \mathcal{S}, a \in \mathcal{A}
 \end{aligned}$$

• Interpretation

- $\lambda(s, a) = \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a)$
- Equation 2: ensures λ has the above meaning
- Equation 1: maximize expected discounted sum of rewards

- **Optimal policy:** $\pi^*(s) = \arg \max_a \lambda(s, a)$

Complexity of LP

- LP **worst-case** convergence guarantees are better than those of DP methods
- LP methods become **impractical** at a much smaller number of states than DP methods do